

Building a Enterprise Eventing Platform

Bryan Zelle

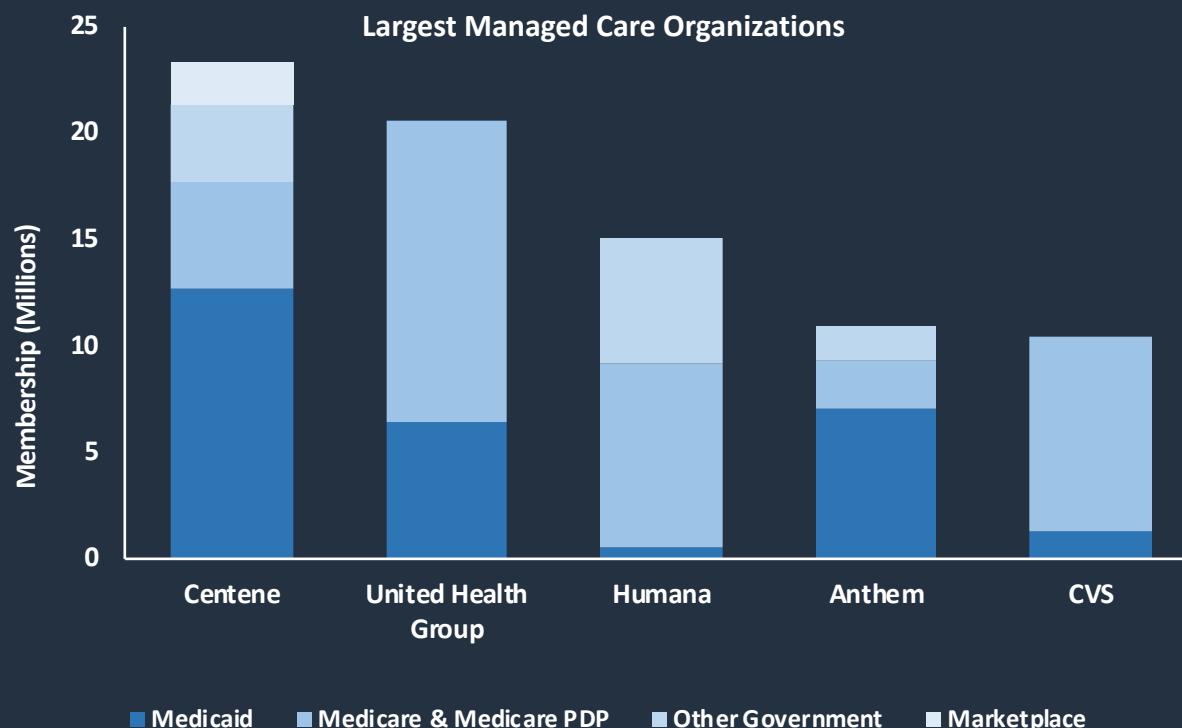
Centene Introduction

Industry:

Largest Medicaid and Medicare Managed Care Provider

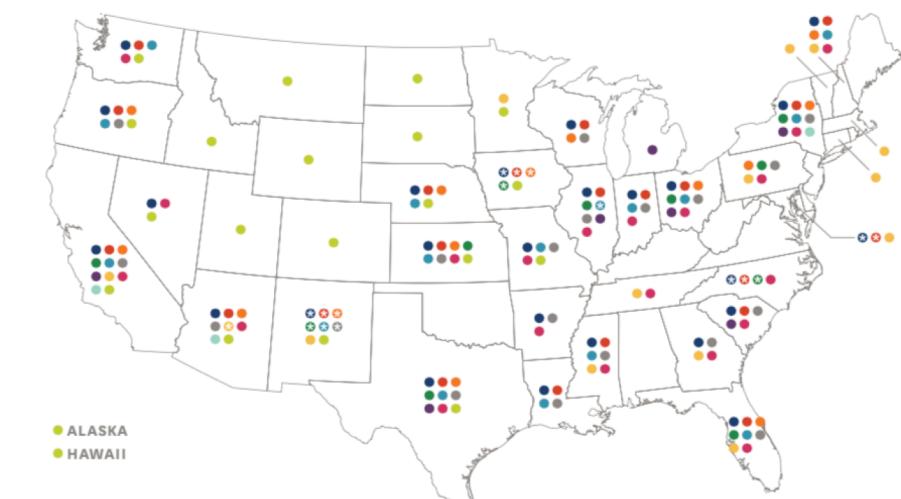
Mission Statement:

Transforming the health of the community, one person at a time



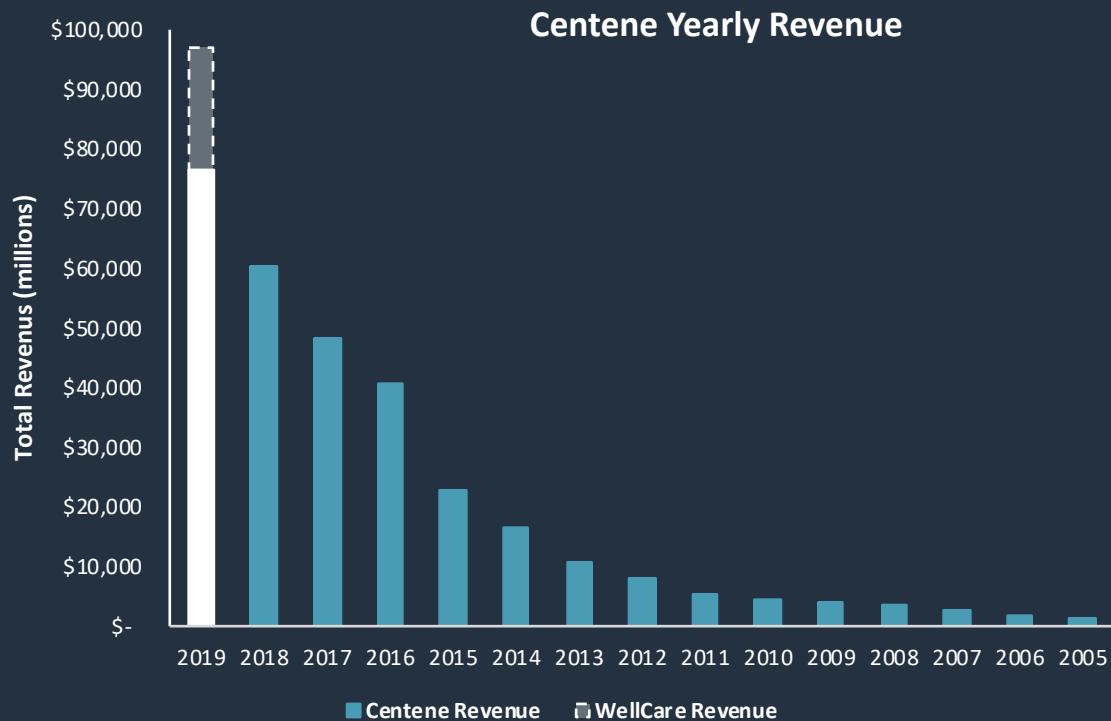
Membership Composition:

Medicaid:	12,700,000	30 States
Medicare (Part D):	4,000,000	50 States
Marketplace:	2,000,000	21 States
Medicare:	1,000,000	28 States
Other:	3,700,000	33 States
Total:	23,400,000	50 States



Summary of Centene's Key Challenges in one word...

Growth



By the numbers:

\$4.1 Billion Revenue to \$96.9 Billion in 10 Years

\$80.4 Billion in growth in past 5 years

\$48.6 Billion in growth in past 2½ years

Cause of the growth...

Mergers & Acquisitions



Centene Growth Outlook

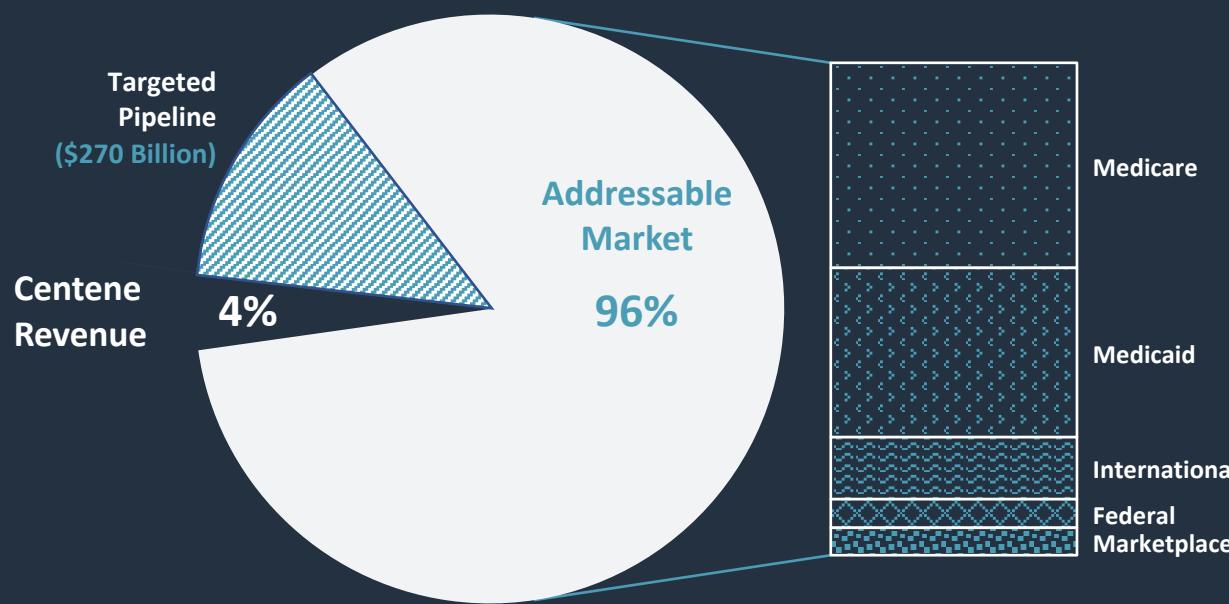
Additional Growth Opportunities

Addressable Market

\$2,000,000,000,000 +

Centene Revenue

\$97,000,000,000 +



\$860 B
40%

\$710 B
33%

\$260 B
12%

\$120B
6%

\$115 B
5%

Federal Medicare

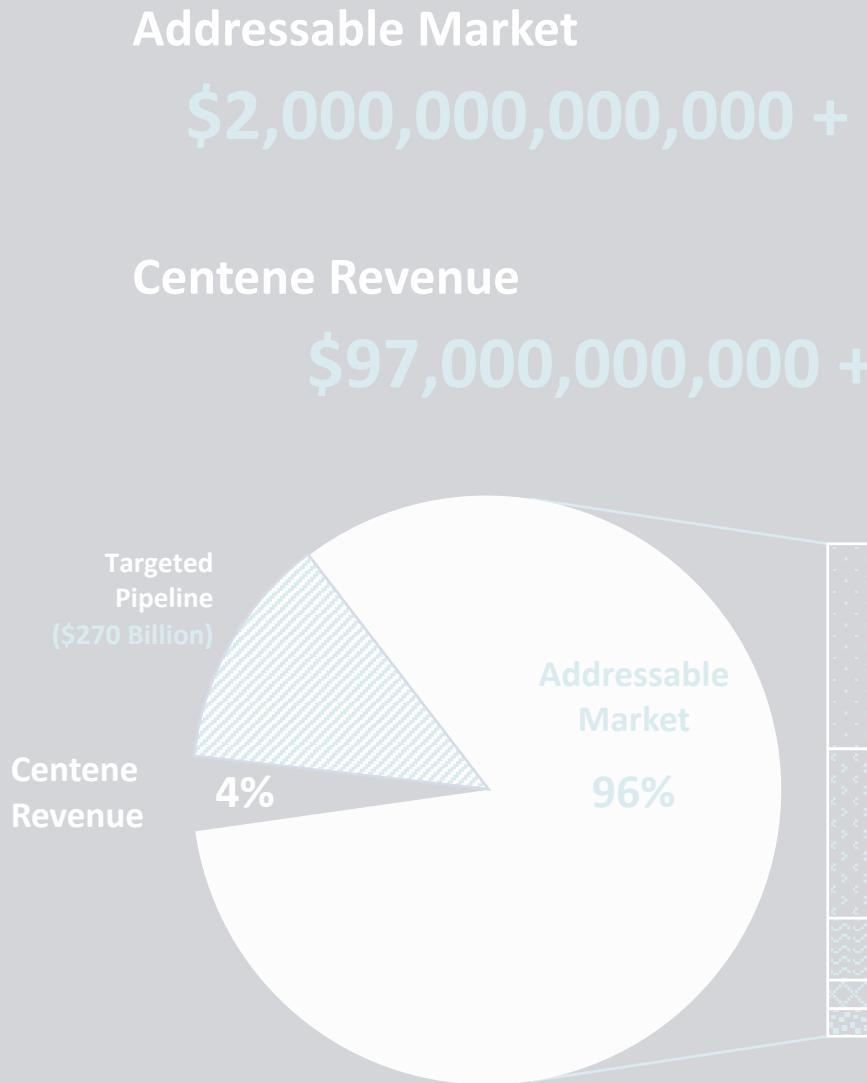
State Medicaid

International Market

Federal Services

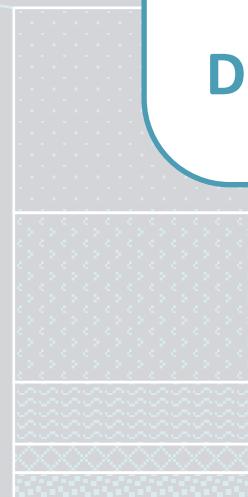
Health Insurance Marketplace

Centene Growth Outlook



Mergers & Acquisitions

Data Integration & Data Migration



Additional Growth Opportunities

Federal Medicare

State Medicaid

International Market

Federal Services

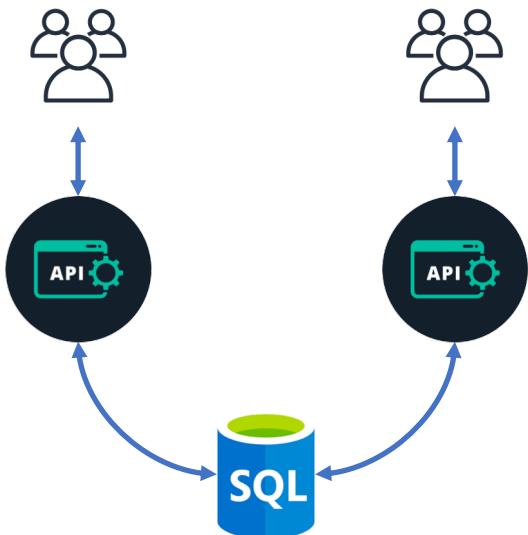
Health Insurance Marketplace

Data Integration & Data Migration

1

Shared Database

- Application Refactor
- Direct Schema Coupling
- Scaling Challenges
- Single Point of Failure

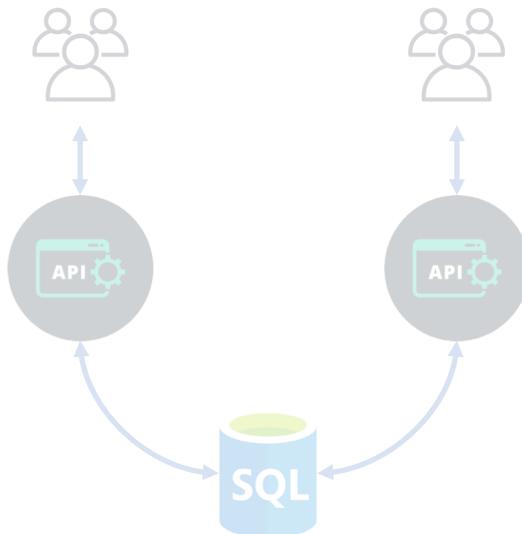


Data Integration & Data Migration



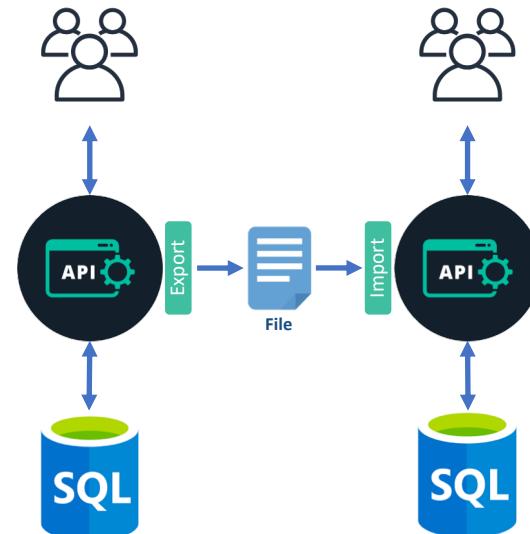
Shared Database

- Application Refactor
- Direct Schema Coupling
- Scaling Challenges
- Single Point of Failure



File Transfer (Batch ETL)

- Latent Data
- Direct Database Load
- Consistency Challenges

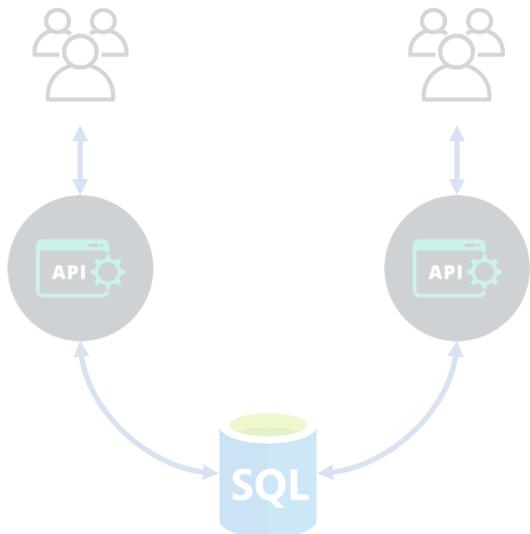


Data Integration & Data Migration



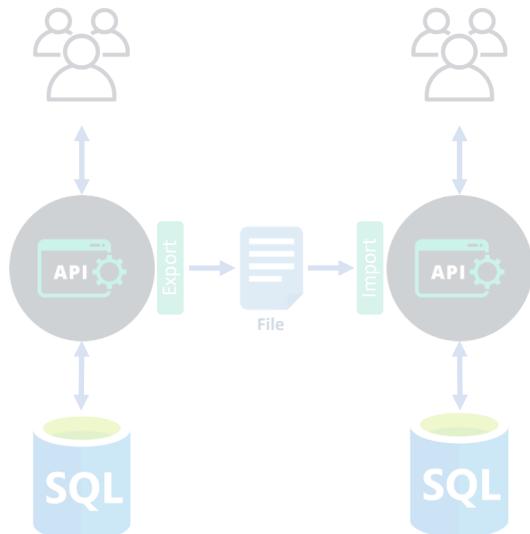
Shared Database

- Application Refactor
- Direct Schema Coupling
- Scaling Challenges
- Single Point of Failure



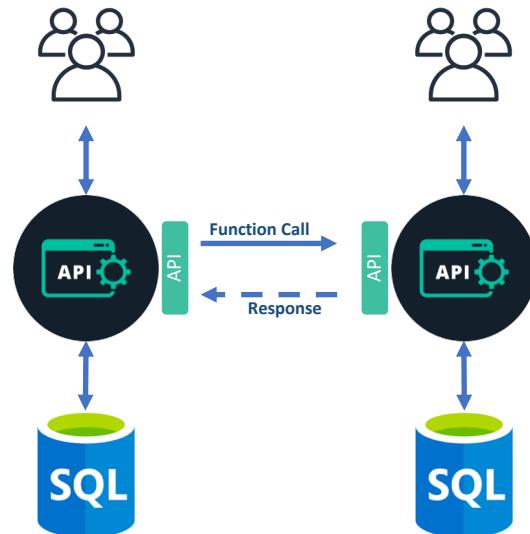
File Transfer (Batch ETL)

- Latent Data
- Direct Database Load
- Consistency Challenges



Remote Procedure Invocation

- Direct Coupling
- Application Refactor
- Availability Concerns
- Scaling Concerns

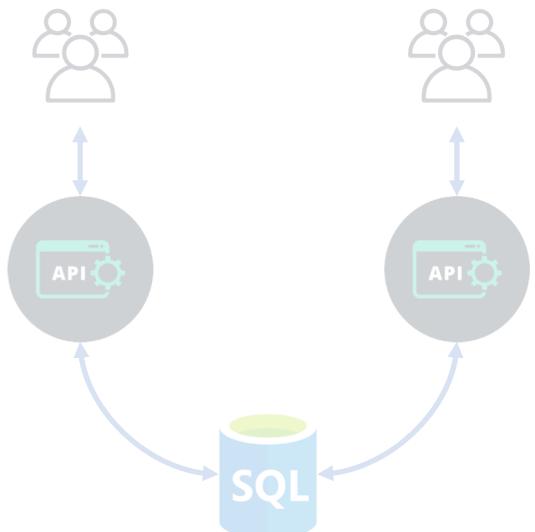


Data Integration & Data Migration



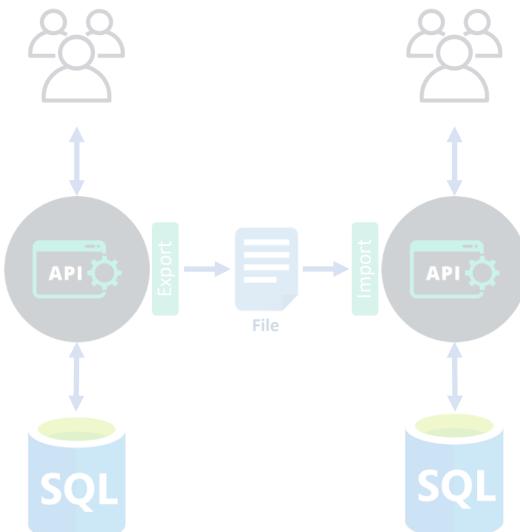
Shared Database

- Application Refactor
- Direct Schema Coupling
- Scaling Challenges
- Single Point of Failure



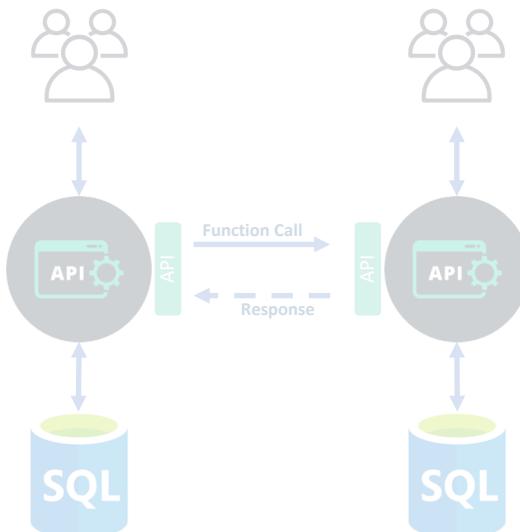
File Transfer (Batch ETL)

- Latent Data
- Direct Database Load
- Consistency Challenges



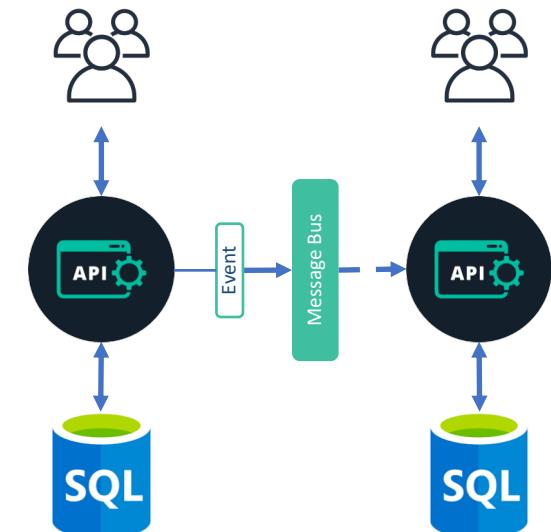
Remote Procedure Invocation

- Direct Coupling
- Application Refactor
- Availability Concerns
- Scaling Concerns



Pub / Sub Messaging (Streaming ETL)

- Loosely Coupled
- No Application Refactor
- Highly Available
- Highly Scalable
- Real-Time Data

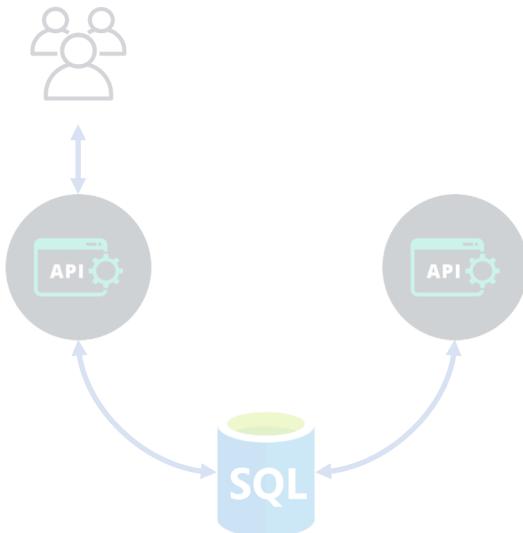


Data Integration & Data Migration



Shared Database

- Application Refactor
- Direct Schema Coupling
- Scaling Challenges
- Single Point of Failure



What is a Event?

Definition: “A significant change in state”

- Statement of fact (immutable)
- Expects no response (or call to action)
- Has a defined “timepoint”

Persistence

- Stateless: Notification Event
- Stateful: Event-Carried State Transfer

Synthesized / Composite Events

- Combination of Events

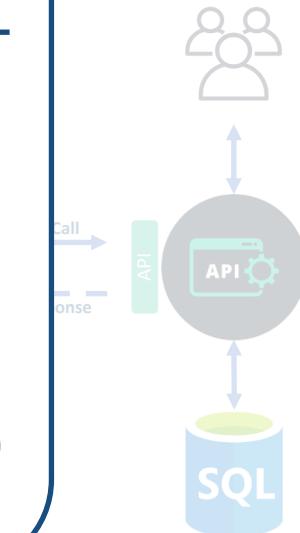


- Absence of an Event



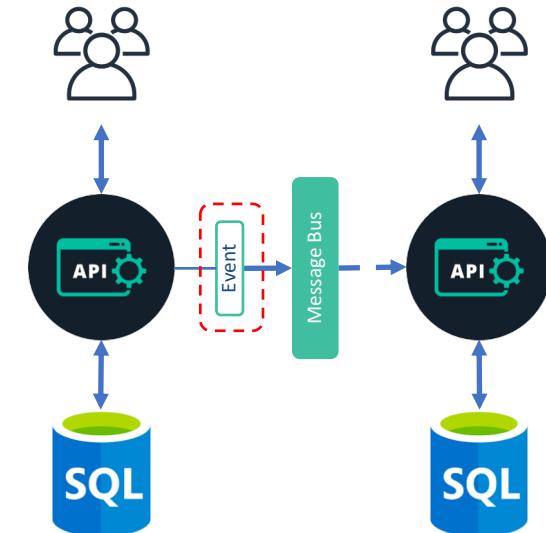
Procedure Migration

Refactor Concerns



Pub / Sub Messaging (Streaming ETL)

- Loosely Coupled
- No Application Refactor*
- Highly Availability
- Highly Scalable
- Real-Time Data

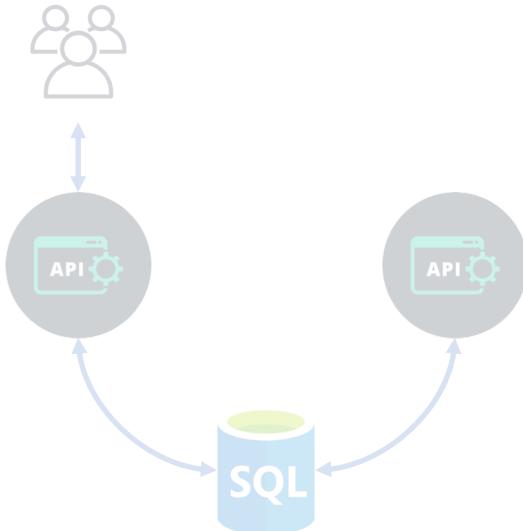


Data Integration & Data Migration



Shared Database

- Application Refactor
- Direct Schema Coupling
- Scaling Challenges
- Single Point of Failure



What is a Event?

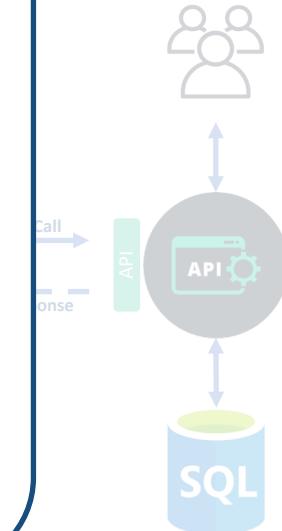
Definition: “A significant change in state”

- Statement of fact (immutable)
- Expects no response (no call to action)
- Has a defined “timepoint”

How do we
publish / consume
meaningful
events?

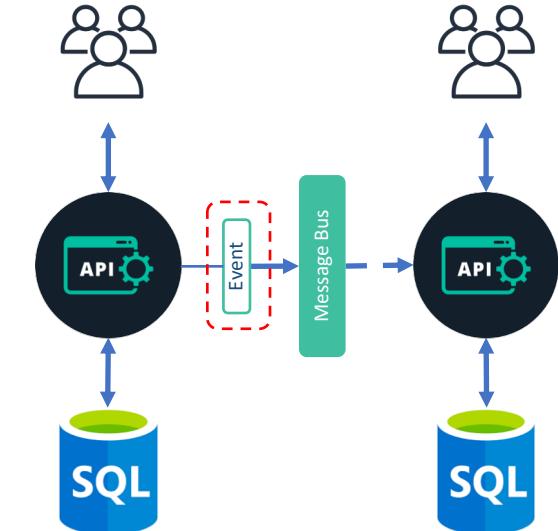


Procedure
ation
ing
Refactor
Concerns
terns



**Pub / Sub Messaging
(Streaming ETL)**

- Loosely Coupled
- No Application Refactor*
- Highly Availability
- Highly Scalable
- Real-Time Data



Change Data Capture (CDC)

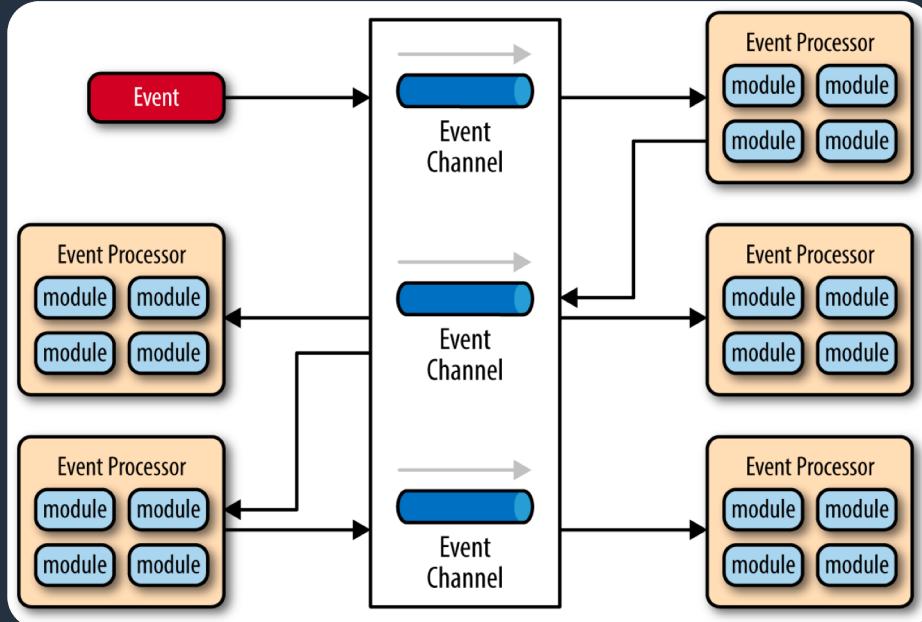
Broker Topology

Subscribe to Event Channels (Topics)

Self-Defined Event Routing

Partial Coupling of Event Channels

Reduced Complexity at cost of reduced coordinating of event execution



Advantages:

- Mature 3rd Party Products / Tooling
- Limited Database Load
- Fast Implementation
- No refactoring of source system

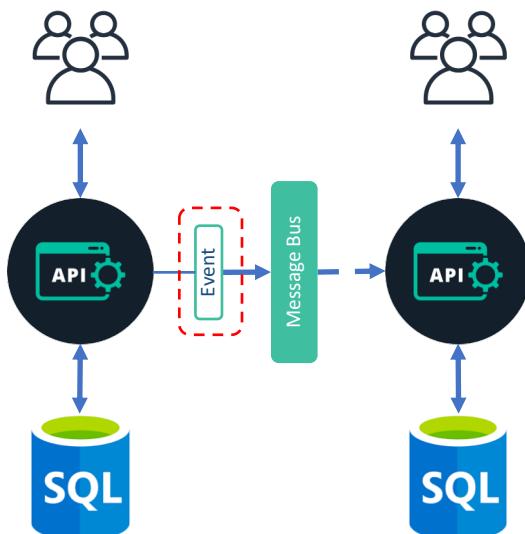
Disadvantages:

- No consistent data structure
- No data governance
- Direct technology coupling



Pub / Sub Messaging (Streaming ETL)

- Loosely Coupled
- No Application Refactor*
- Highly Available
- Highly Scalable
- Real-Time Data



Change Data Capture (CDC)

Transaction
Metadata

```
{  
    "INFA_SEQUENCE": {"string": "2,PWX_GENERIC,1,,2,3,C7084816514A5D260"},  
    "\\"DTL__CAPXUSER\\": {"string": "USER1"},  
    "\\"DTL__CAPXTIMESTAMP\\": {"string": "20180305131540000000000000"},  
    "\\"INFA_OP_TYPE\\": {"string": "UPDATE_EVENT"},  
    "\\"INFA_TABLE_NAME\\": {"string": "d8amisou6p.MEMBER_CONTACT"}  
}
```

Event
Body

```
,\\"MEMBER_PCP\\": {"string": "Dr. Bryan Zelle"},  
,\\"MEMBER_PCP_Present\\": true,  
,\\"MEMBER_PCP_BeforeImage\\": {"string": "Dr. John Smith"},  
,\\"MEMBER_PCP_BeforeImage_Present\\": true  
}
```

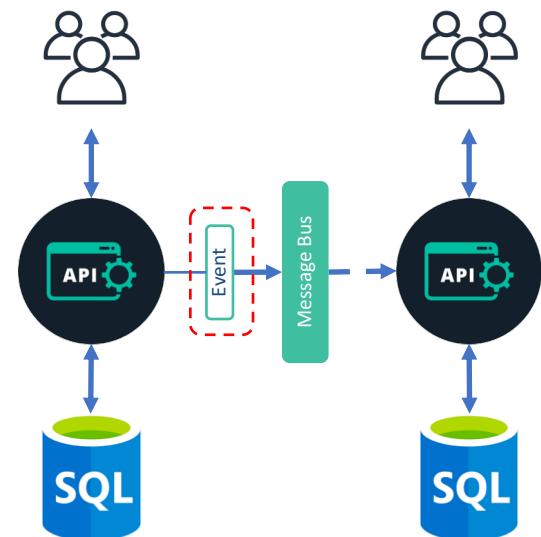
What Event
information are we
capturing?

- Who - Who changed the data ?*
- What - What data changed ?
- When - When the data changed ?
- Where - Where was the data changed ?



Pub / Sub Messaging (Streaming ETL)

- Loosely Coupled
- No Application Refactor*
- Highly Available
- Highly Scalable
- Real-Time Data



Mediated (Orchestrated) Eventing

Mediator Topology

Mediator transfers events to assigned event channel (Topic)

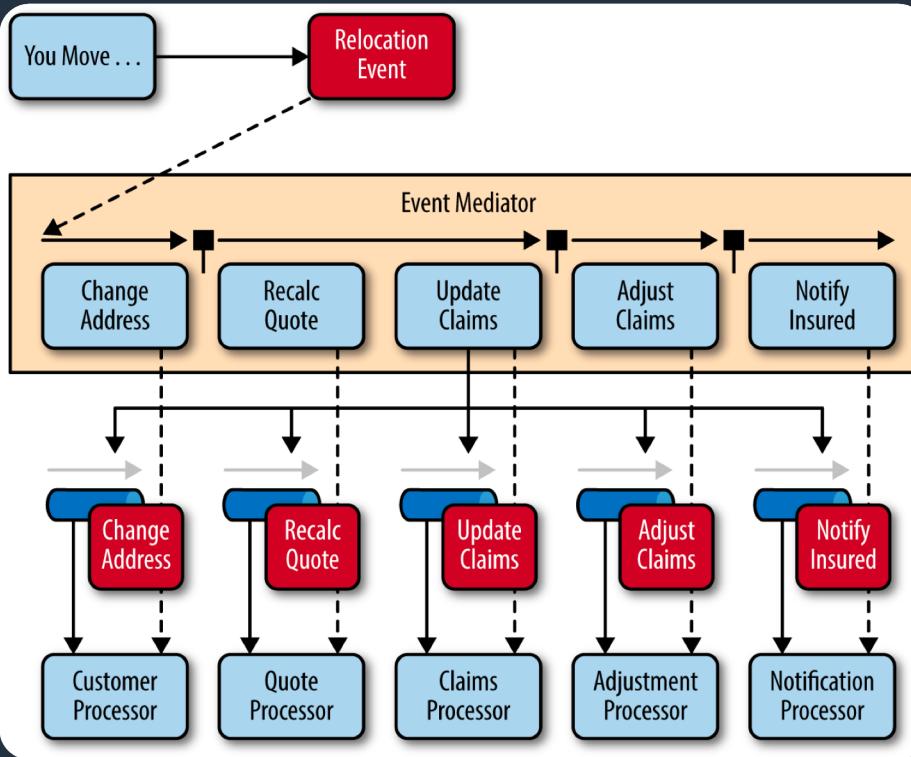
Centrally Coordinated Event Routing

Complete Decoupling of Event Channels

Increased Complexity at cost of increased coordination of event execution

Advantages:

- Consistent / Common Framework
- Enforce Data governance
- Economy of Scale Advantage
- Technology abstraction / decoupling



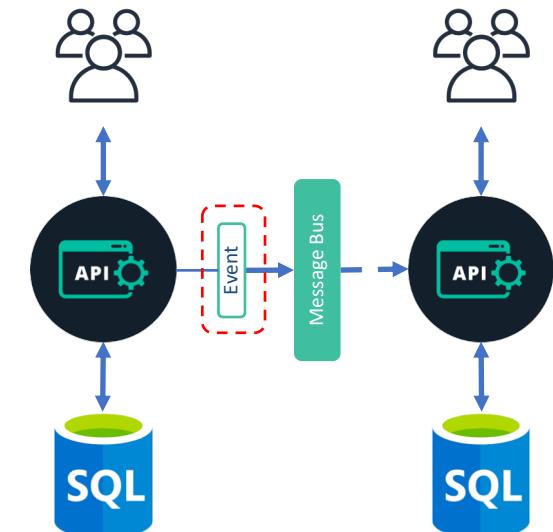
Disadvantages:

- External bottleneck (Mediator Owner)
- Single Point of Failure
- Duplicative data storage



Pub / Sub Messaging (Streaming ETL)

- Loosely Coupled
- No Application Refactor*
- Highly Available
- Highly Scalable
- Real-Time Data



Mediated (Orchestrated) Eventing

Transaction Metadata

Event Body

Example Event Payload (JSON via REST)

```
"Metadata" : {  
    "Transaction ID" : "C7084816514A5D260",  
    "User ID" : "USER1",  
    "Time Stamp" : "201803051315400000000000",  
    "Transaction Type" : "UPDATE",  
    "Source System" : "d8amisou6p.MEMBER_CONTACT" },  
    "Event Body" :{  
        "Event Type" : "Member-PCP-Change",  
        "Previous Value" : "Dr. John Smith",  
        "Updated Value" : "Dr. Bryan Zelle",  
        "Event Source" : "Inbound-Member-Call",  
        "Caller Information" :{  
            "Name" : "Jane Doe",  
            "Inbound Number" : "1-614-847-0982",  
            "Call Resolution Status" : "5 - Highly Satisfied",  
            "First Call Resolution" : "Success",  
            "Internal Representative" : "CN-10238381",*  
            "Call Duration (Minutes)" : "8:19" }  
    }
```

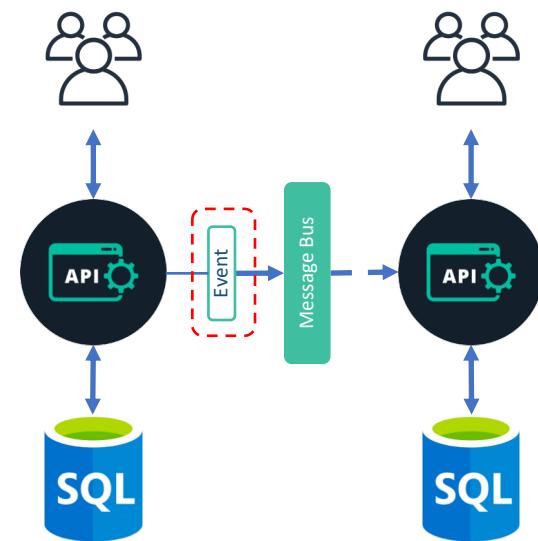
What Event information are we capturing?

- Who** - Who changed the data ?*
- What** - What data changed ?
- When** - When the data changed ?
- Where** - Where was the data changed ?
- Why** - Why was the data changed ?



Pub / Sub Messaging (Streaming ETL)

- Loosely Coupled
- No Application Refactor*
- Highly Available
- Highly Scalable
- Real-Time Data



Why is the “Why” so important?

```
“Event Body” : {  
    “Event Type” : “Member-PCP-Change”,  
    “Previous Value” : “Dr. John Smith”,  
    “Updated Value” : “Dr. Bryan Zelle”,  
    “Event Source” : “Inbound-Member-Call”,  
    “Caller Information” : {  
        “Name” : “Jane Doe”,  
        “Inbound Number” : “1-614-847-0982”,  
        “Internal Representative” : “CN-10238381”,  
        “Call Resolution Status” : “5 - Highly Satisfied”,  
        “First Call Resolution” : “Success”,  
        “Call Duration (Minutes)” : “8:19”  
    }  
}
```

New Events can be Created or Derived

If two separate “Member-PCP-Change” events happen within 2 weeks window -> Create event for Case Manager to Review

If Inbound Number not currently associated with Member -> Create event to add it to the Member’s Profile

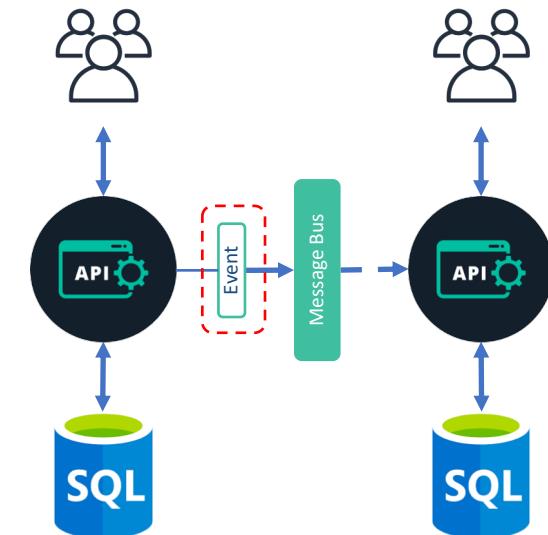
If Resolution Status ≥ 4 -> Create event to assign call rep to member for call-back

If first call resolution \neq “Success” and call duration > 15 min -> Create event to escalate call to Supervisor for Audit



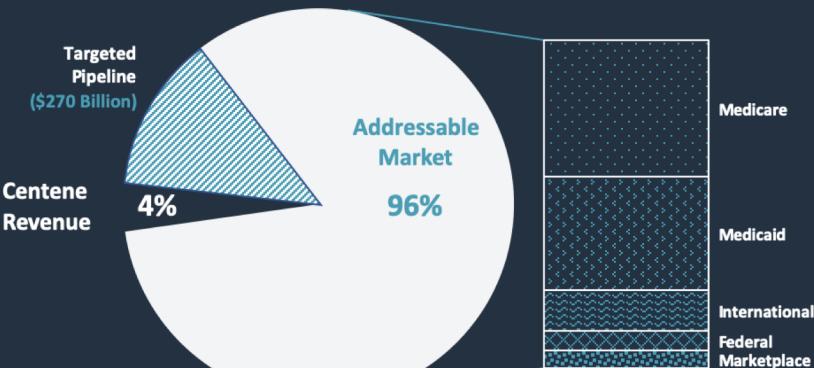
Pub / Sub Messaging (Streaming ETL)

- Loosely Coupled
- No Application Refactor*
- Highly Available
- Highly Scalable
- Real-Time Data



1

Centene's Core Challenge is **Growth** cause by **Mergers & Acquisitions**; causing us to reevaluate our Enterprise **Data Integration** and **Data Migration** Strategies...

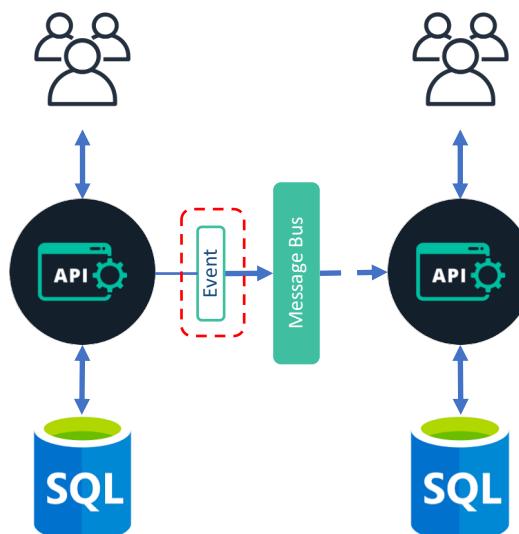


Recap

2

Async Pub / Sub Eventing through Kafka provides us valuable capabilities:

- Highly Scalable
- High Autonomy / Decoupling
- High Availability & Data Resiliency
- Real Time Data Transfer
- Complex Stream Processing



3

Leveraging a **Mediator Topology** enables the creation of **meaningful** events; which provide insight into **why** things are happening, so we can **react** to them in real time...

```
    "Metadata": {  
        "Transaction ID": "C7084816514A5D260",  
        "User ID": "USER1",  
        "Time Stamp": "201803051315400000000000",  
        "Transaction Type": "UPDATE",  
        "Source System": "d8amisou6p.MEMBER_CONTACT" },  
        "Event Body": {  
            "Event Type": "Member-PCP-Change",  
            "Previous Value": "Dr. John Smith",  
            "Updated Value": "Dr. Bryan Zelle",  
            "Event Source": "Inbound-Member-Call",  
            "Caller Information": {  
                "Name": "Jane Doe",  
                "Inbound Number": "1-614-847-0982",  
                "Call Resolution Status": "5 - Highly Satisfied",  
                "First Call Resolution": "Success",  
                "Internal Representative": "CN-10238381",  
                "Call Duration (Minutes)": "8:19" }  
    }
```

Common Core Architecture:

- 1) Event Source
- 2) Event Intake
- 3) Event Channel
- 4) Event Router
- 5) Event Subscription
- 6) Event Destination

Event
Bridge



Mule
ESB



Event
Grid



Apache
Camel



Mediator Alternatives?

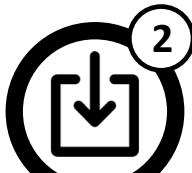
Knative
Eventing



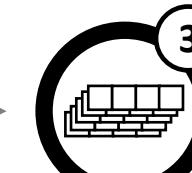
Generic Event Mediator



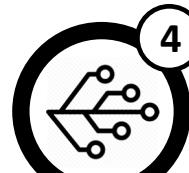
Event
Source



Event
Intake



Event
Channel



Event
Router

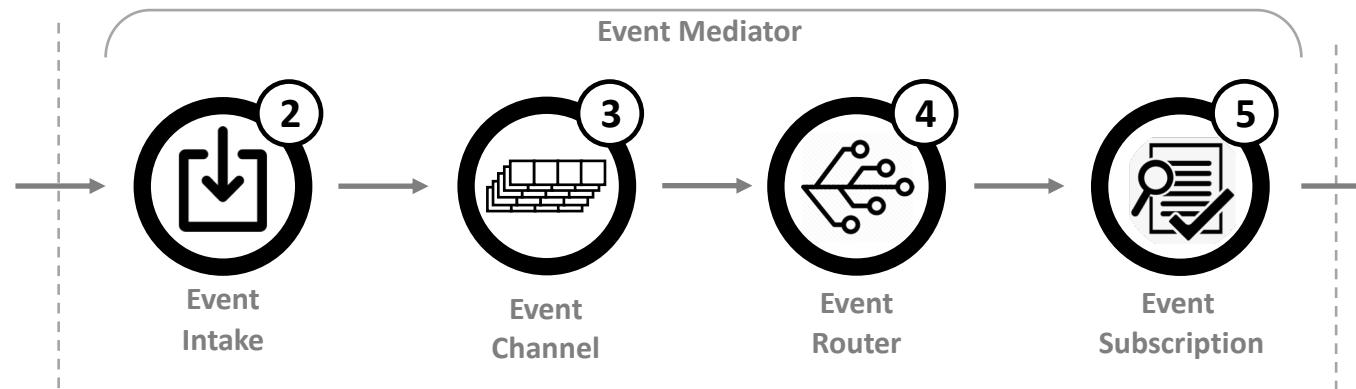


Event
Subscription



Event
Destination

Event Mediator



Required Features & Functionality



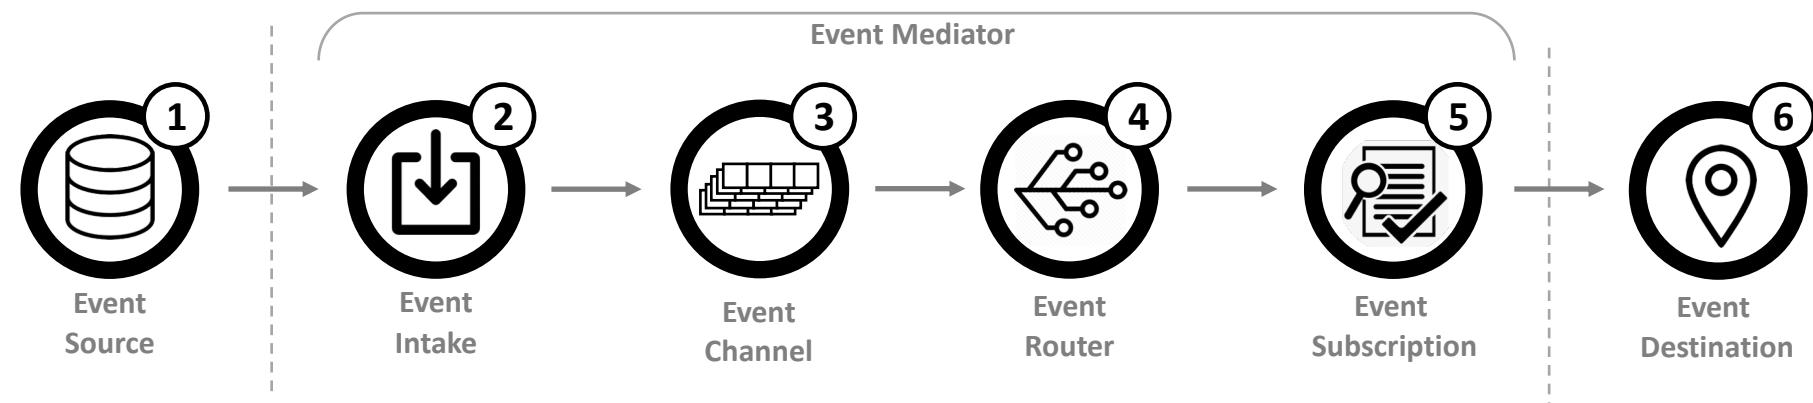
Design Criteria

- 1) AVRO Event Serialization
- 2) JSON Validation of Event Body
- 3) Centralized Event Registry
- 4) Distributed Tracing of Events
- 5) Sensitive Data Redaction
- 6) Turn / Key Self-Service
- 7) Cloud Agnostic
- 8) Permanent Event Storage
- 9) Flexible Ingestion Intake
- 10) Pre-built Monitoring / Dashboards
- 11) Synthetic Events

Business Value

Reduced Message Size -> Reduced Storage Cost in Cloud
Data Validation -> Clean Data
Easily Find Events -> Prevents Event Duplication & Increases Adoption
Tracing -> Provides Event Lineage and Auditability
Data Restriction -> Protects HIPPA data (including PHI/PII)
Automated Configuration -> Reduced manual administrative burden
Multi-Cloud Strategy -> No Reliance on Single Cloud Provider
Event Persistence -> DR Strategy + Event Sourcing / Hydration
Legacy Systems Limitations -> Offer REST, gRPC, SOAP Interfaces & API's
Universal Metrics -> Consistent / Granular Event Visibility
Fictitious Event -> Blue/Green Deployments, Prod Smoke Testing, Etc.

Generic Event Mediator



Required Features & Functionality

Design Criteria

- 1) AVRO Event Serialization
- 2) JSON Validation of Event Body*
- 3) Centralized Event Registry
- 4) Distributed Tracing of Events*
- 5) Sensitive Data Redaction*
- 6) Turn / Key Self-Service
- 7) Cloud Agnostic*
- 8) Permanent Event Storage
- 9) Flexible Ingestion Intake*
- 10) Pre-built Monitoring / Dashboards
- 11) Synthetic Events*

Leverage 3rd Party Frameworks or Build Custom?

Assessment:

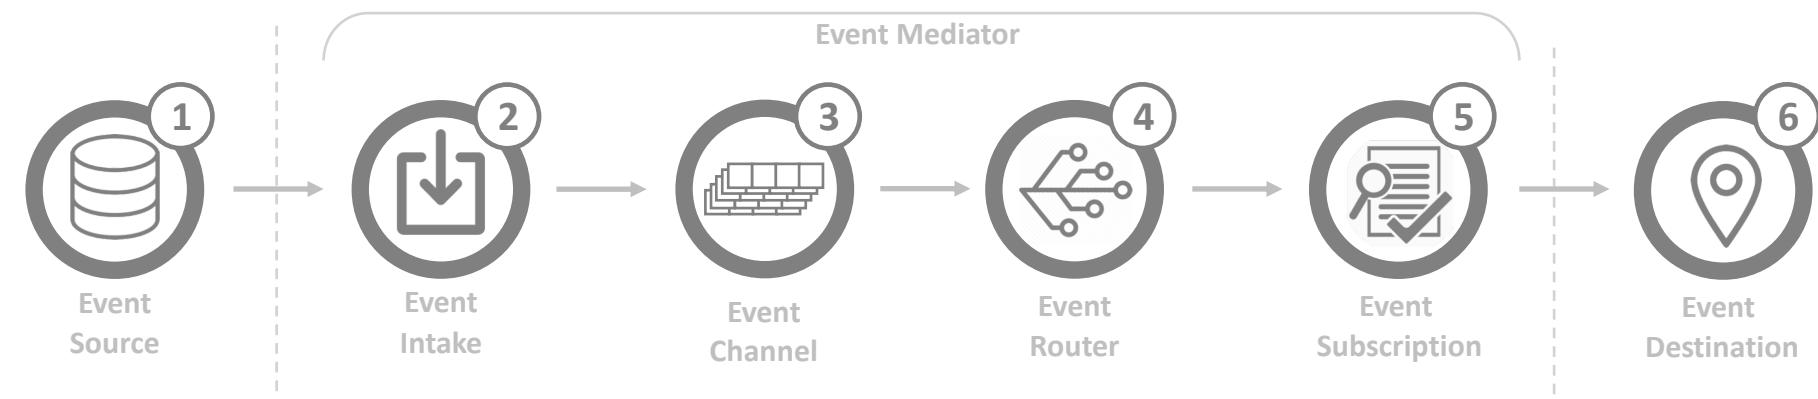
Majority of frameworks focused engineering effort on how to get data into framework as easily as possible

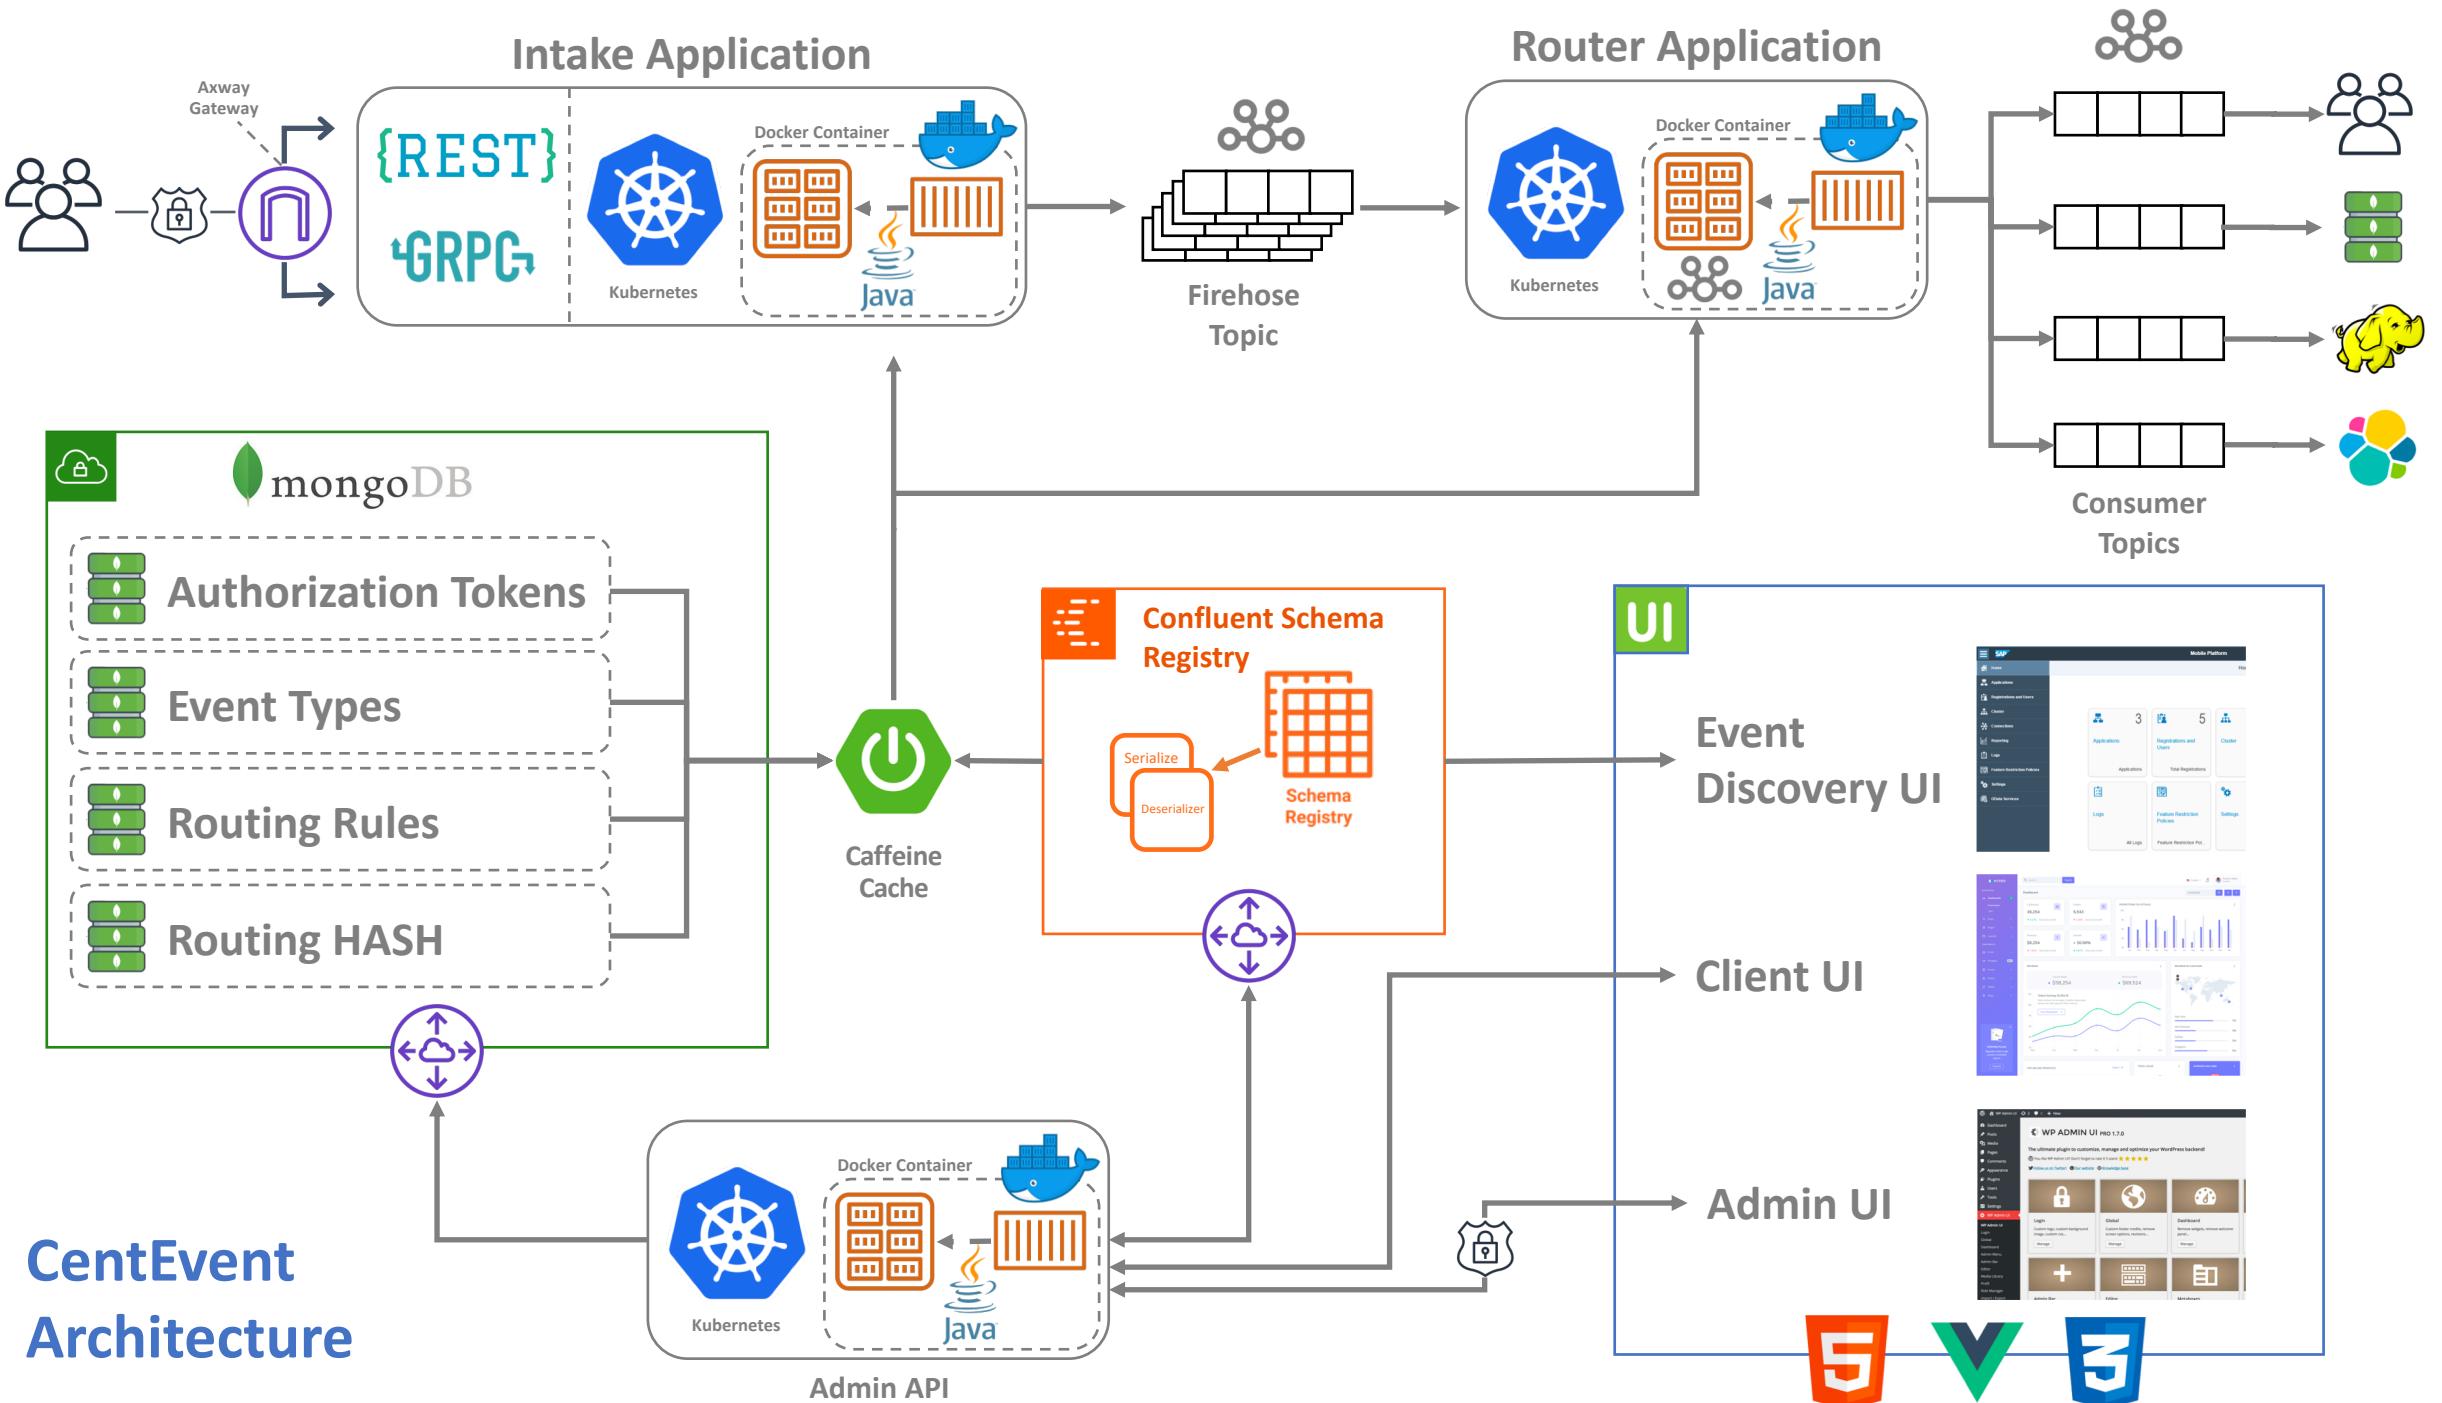
- Higher Data Ingest = Increased Billables (SaaS)
- Too many gaps with current features*

Decision:

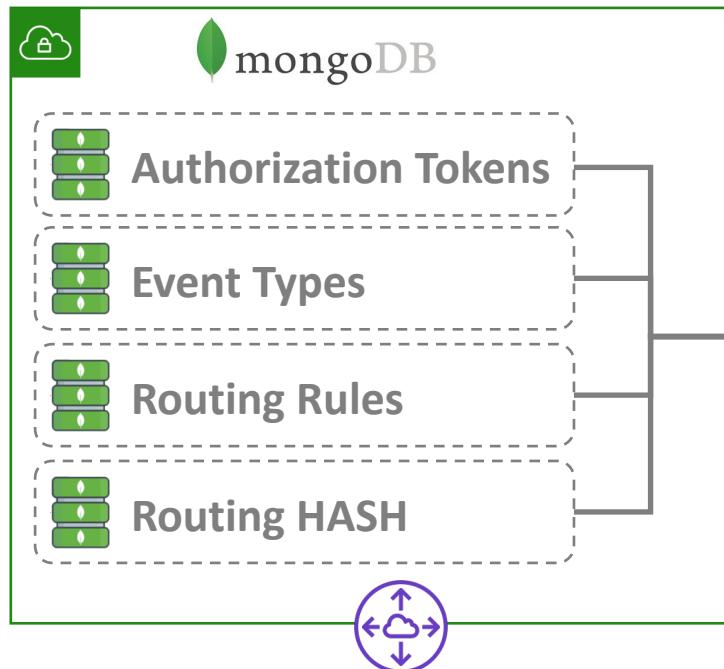
Build Centralized Eventing Framework for Enterprise use across all Centene Domains

Generic Event Mediator

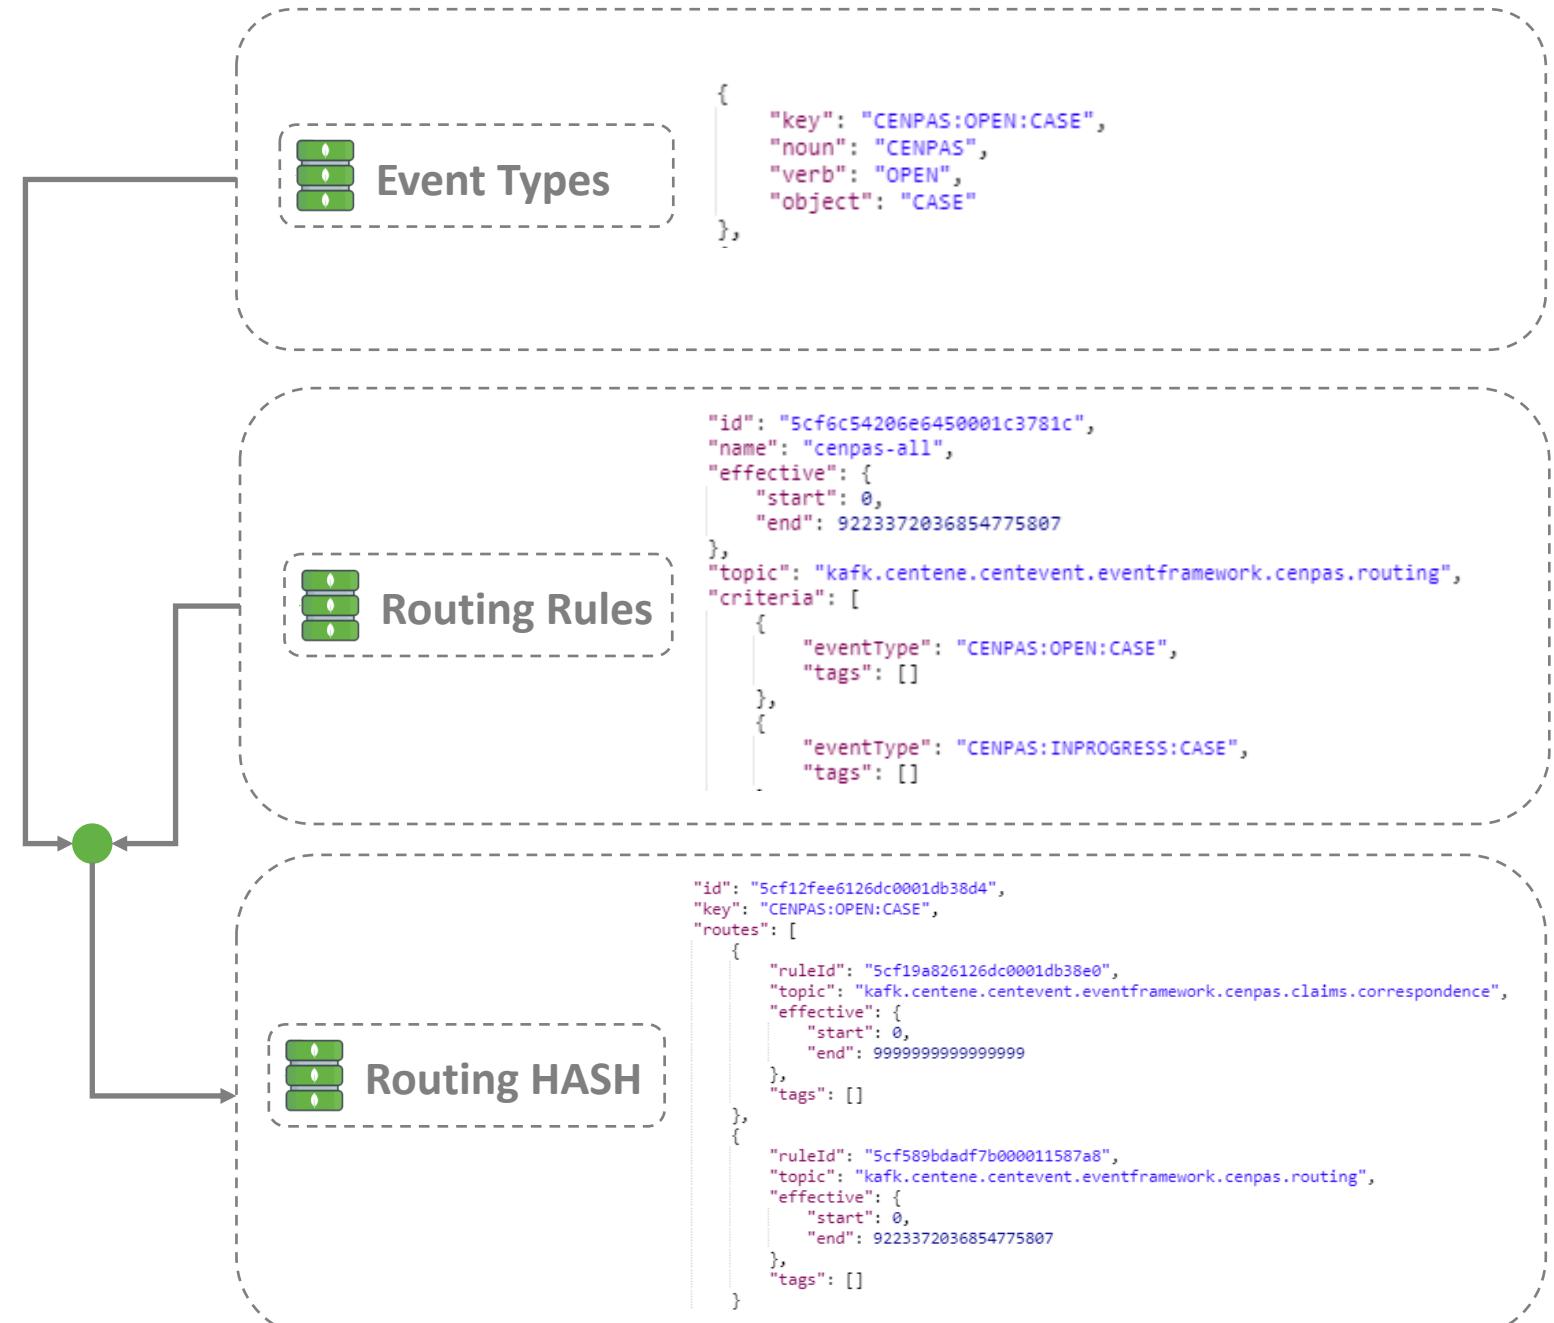
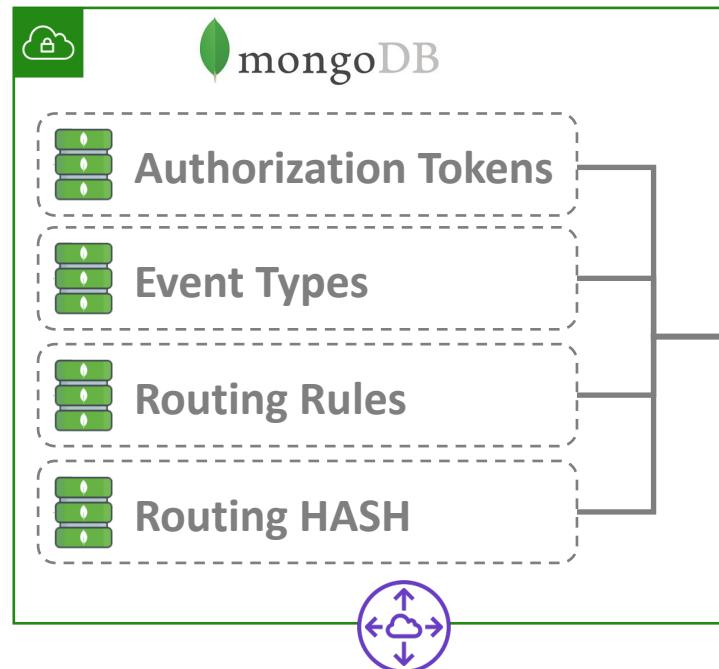




Configuration / Admin Collections



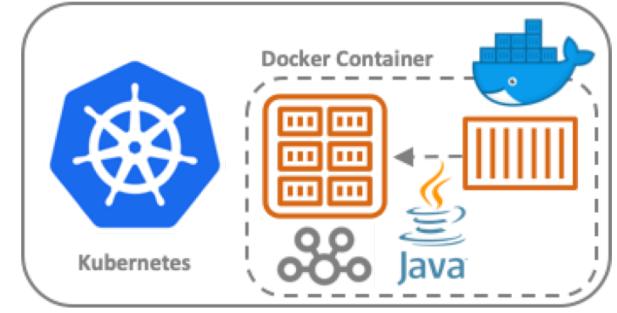
Configuration / Admin Collections



K-Streams Application Code

- 1) Fetch the Event Key --> Key = N:V:O
- 2) Join Hashed Routing Rules: maps N:V:O to Consumers that have Subscribed
- 3) Discard Events if no Consumer has Subscribed
- 4) Generate Duplicate Event message; one for each subscribed Consumers
- 5) Validate Consumers effective date is within Event date
- 6) Validate the Event contains tags the Consumer specifies
- 7) Validate the Consumers PHI permissions are appropriate
- 8) Place the Event destination topic onto the Kafka message header (metadata)
- 9) Send the event to the destination topic specified in the message header

Router Application

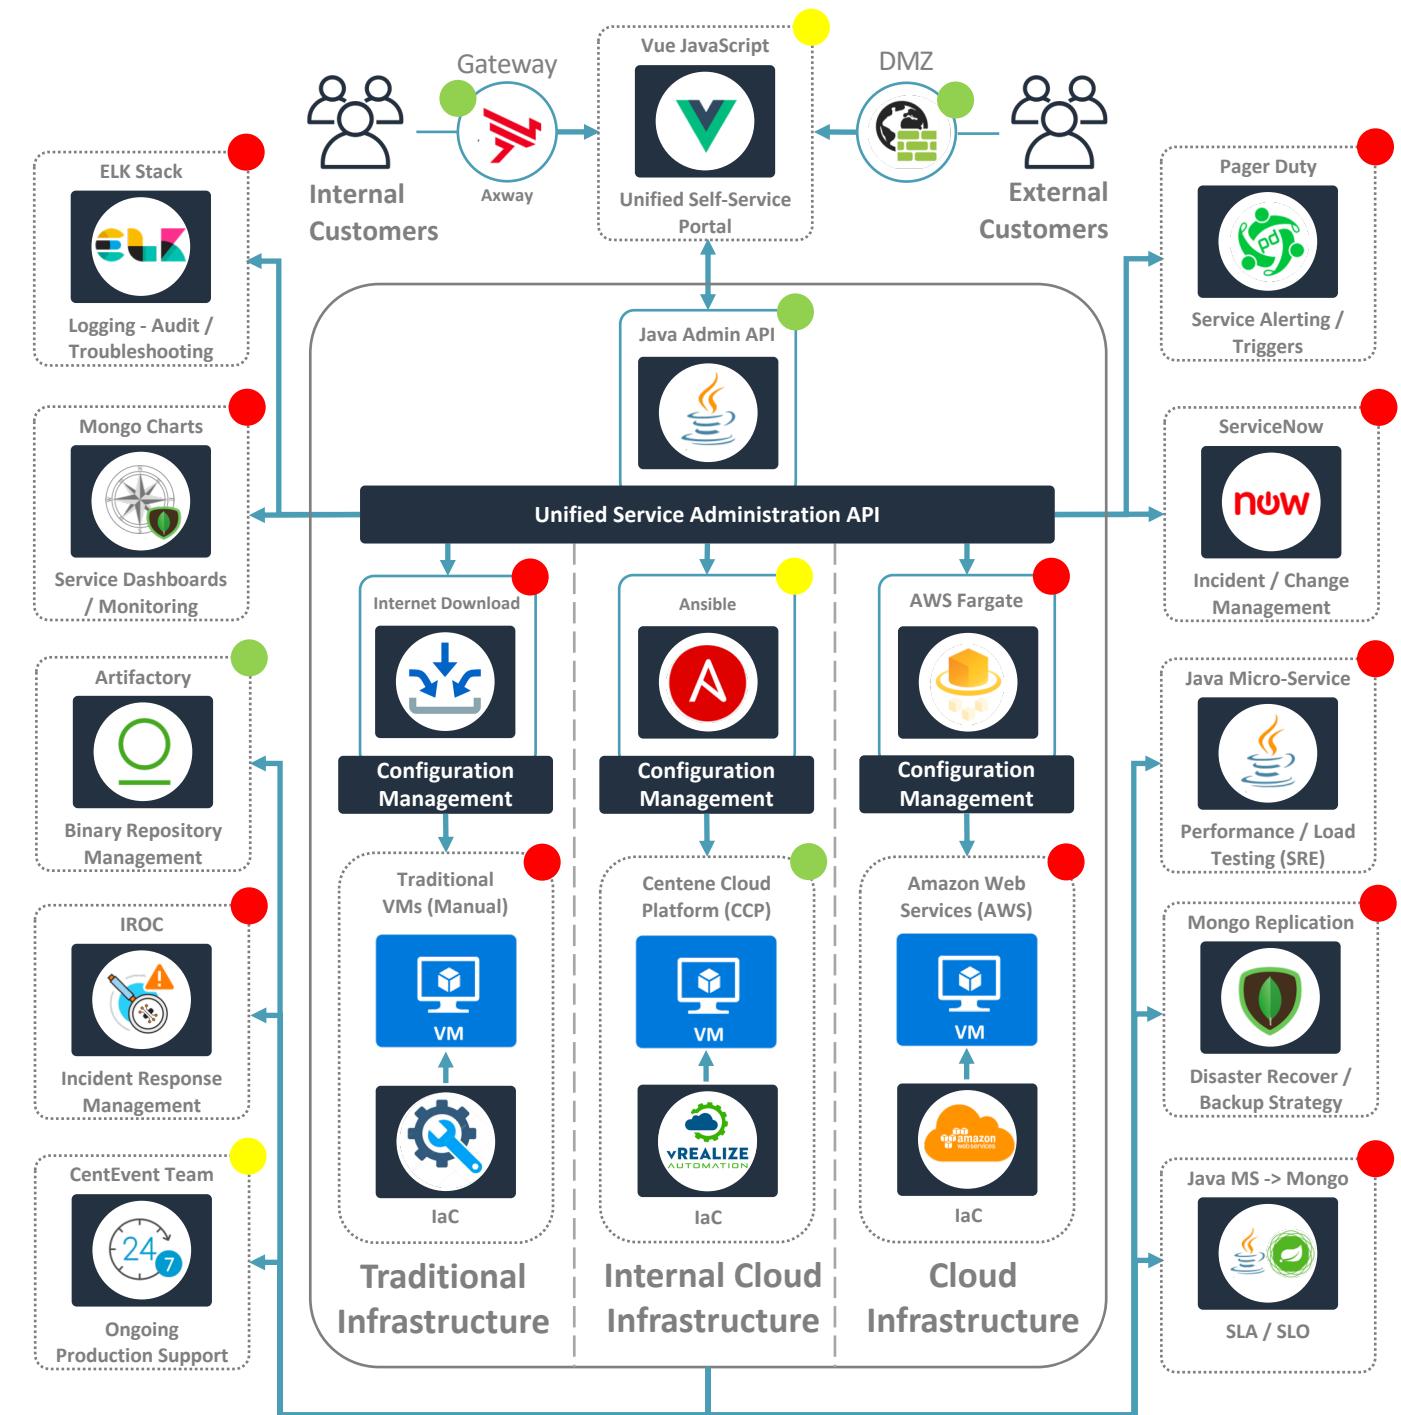


```
public void routeEvent(KStream<String, Event> stream) {  
    stream.selectKey((id, event) -> event.getMetadata().getType()) KStream<EventType, Event>  
        .map((key, value) -> new KeyValue<>(value, eventTypeRoutingService.findEventTypeRouting(EventTypeUtil.toString(key)).getRoutes())) KStream<Event, List<EventRouteDocument>>  
        .peek((k, v) -> log.debug("key: {}, routes: {}", k, v)) KStream<Event, List<EventRouteDocument>>  
        .filter((key, value) -> !value.isEmpty()) KStream<Event, List<EventRouteDocument>>  
        .flatMapValues((readOnlyKey, value) -> value) KStream<Event, EventRouteDocument>  
        .filter(this::effectiveDateRangeFilter) KStream<Event, EventRouteDocument>  
        .filter(this::tagFilter) KStream<Event, EventRouteDocument>  
        .filter(this::phiFilter) KStream<Event, EventRouteDocument>  
        .map(this::updateRoutedTopic) KStream<capture of ? extends EventKey, capture of ? extends Event>  
        .peek((k, v) -> log.debug("routing complete, sending to {}", v.getMetadata().getRoutedTopic())) KStream<capture of ? extends EventKey, capture of ? extends Event>  
        .to((key, value, recordContext) -> value.getMetadata().getRoutedTopic());  
}
```

Eventing as a Service - CentEvent

- VM Provisioning / Management
- Infrastructure as Code (IaC)
- Configuration Management
- Service Administration API
- Service Dashboards / Monitoring
- Incident / Change Management
- Service Alerting / Triggers
- Service Logging (Audit / Troubleshooting)
- Self-Service Portal
- Automated Performance / Load Tests (SRE)
- Disaster Recovery / Backup Strategy
- SLA / SLO (Defined and Captured)
- Binary Repository Management
- Incident Response Management (IROC)
- Ongoing Production Support

● Complete
● In-progress
● Incomplete



CentEvent Roadmap

Operations

Product Maturity

Demand

Q4 - 2019

Q1 - 2020

Q2 - 2020

Oct

Nov

Dec

Jan

Feb

Mar

Apr

May

June

SLA-SLO (Mongo Charts)

Event RTR – Connector + Pipeline

IROC – Run Books + APIs

Regression Testing - Producer

Phase II : Consumer

ELK Dashboards - Phase I

ELK Dashboard - Phase II

Consumer Notifications / Alerts

AWS POC - Phase I

AWS MVP - Phase II

AWS Deployment - Phase III

Lambda Function POC – Phase I

Lambda Deployment - Phase II

Self-Service Portal MVP – Phase I

Phase II : Deployment

Event Sourcing – Hydration Pipeline

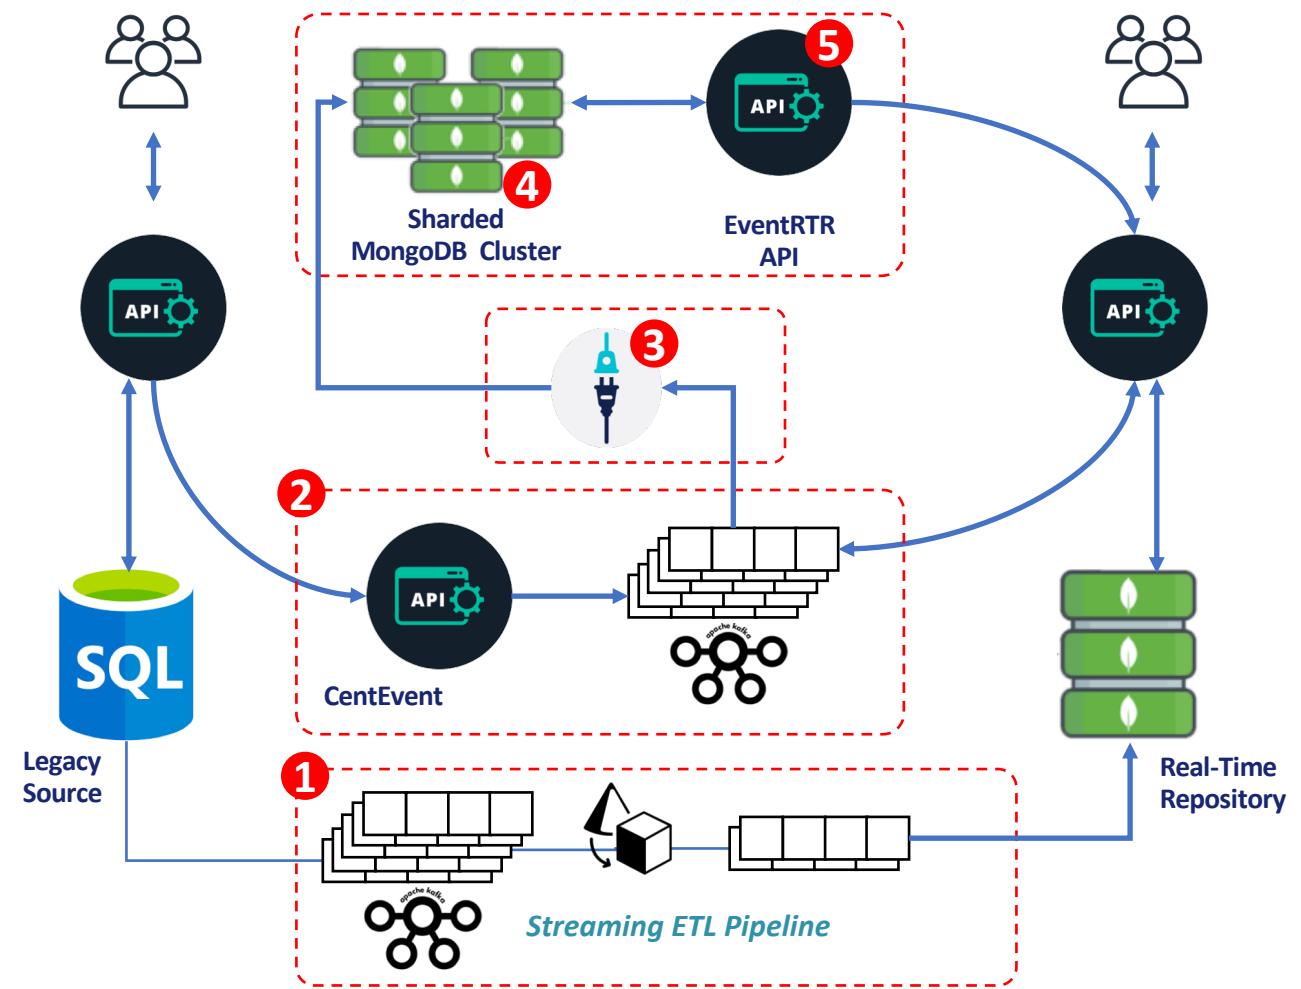
Universal Centene Integration Pattern

Key Components:

- 1) CDC Streaming ETL Pipelines
- 2) CentEvent (Event Mediator)
- 3) MongoDB Connector (modified)
- 4) MongoDB Cluster
- 5) Event RTR API

Framework Features:

 Synthetic Events	 Event Routing
 PHI Filtering	 ELK Monitoring
 Event Tagging	 Self-Service Portal
 Distributed Tracing	 Event Discovery UI



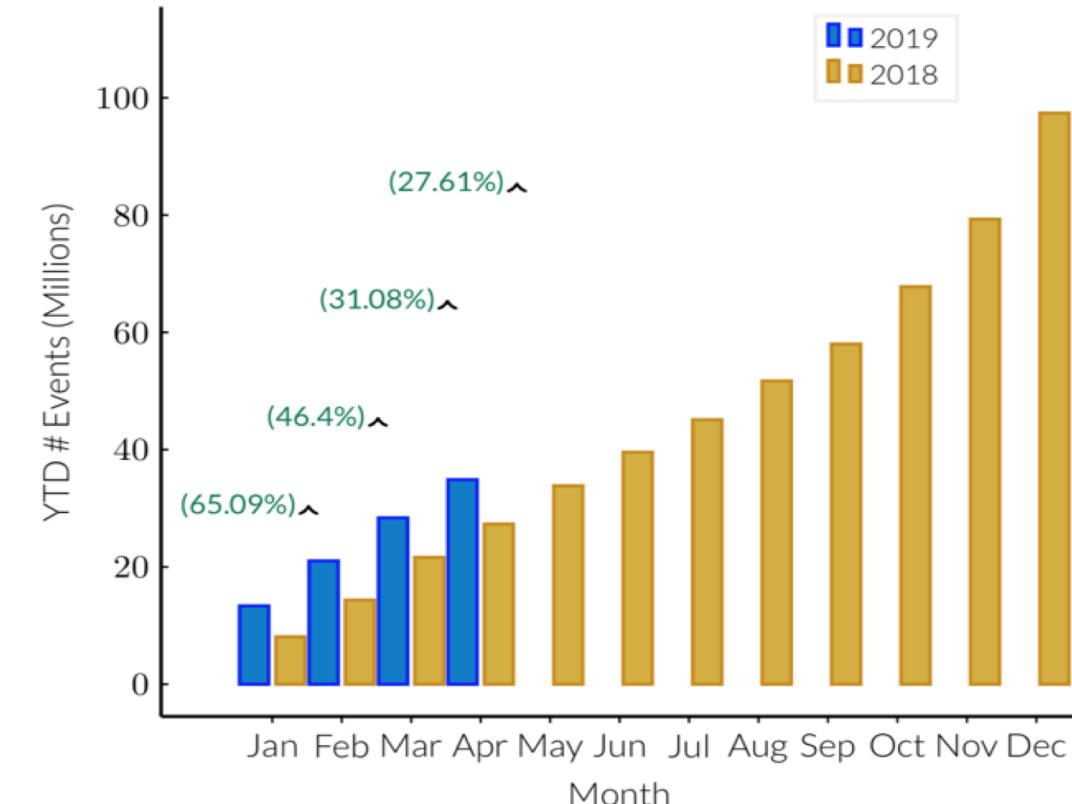
Application Refactor Use Case

Member Insight Platform Migration

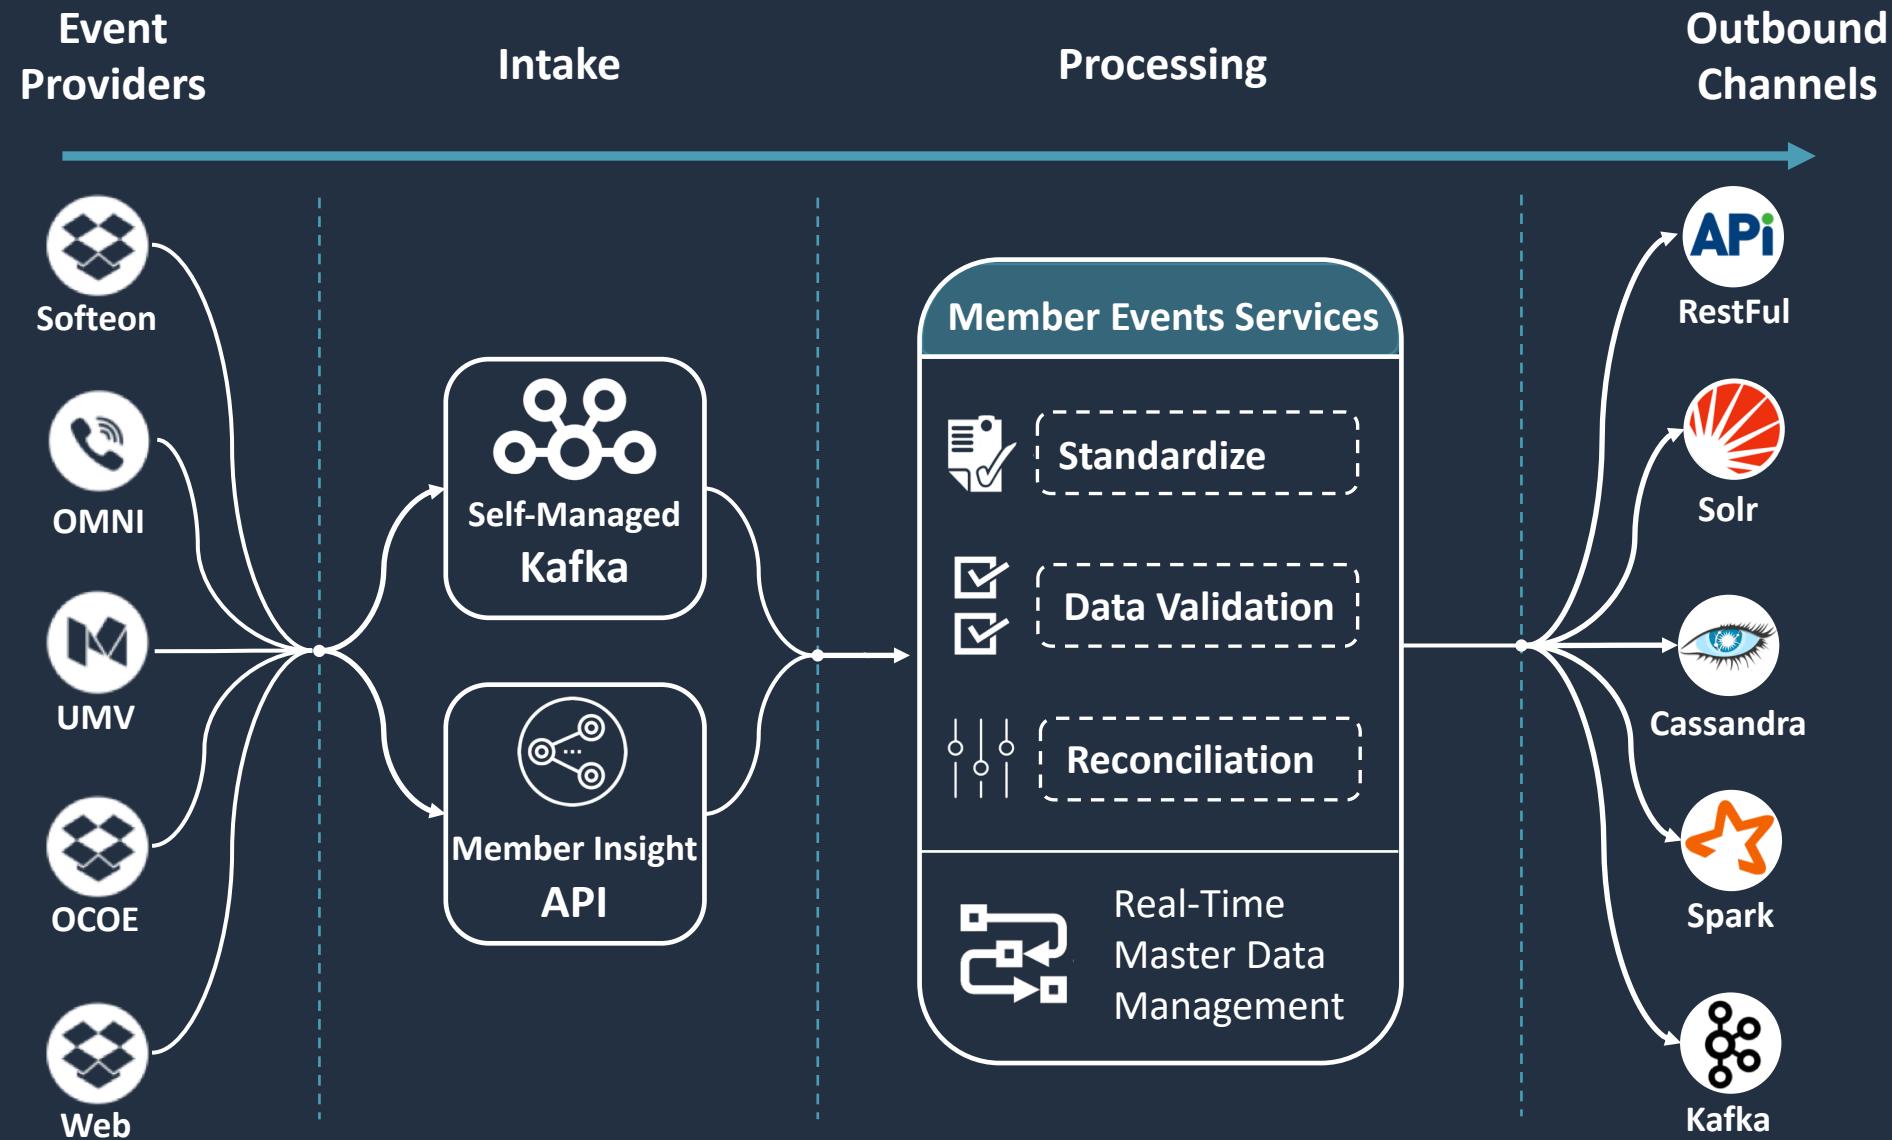
- Refactor from self-managed Kafka to KaaS
- Refactor from Member Insight API to CentEvent



Member Event Volume continues to grow significantly - increasing the importance of a stable and scalable underlying infrastructure of services



Member Insight Process Flow: Before



Member Insight Process Flow: After

