

@rmoff

# From zero to Hero with Kafka Connect

a.k.a.

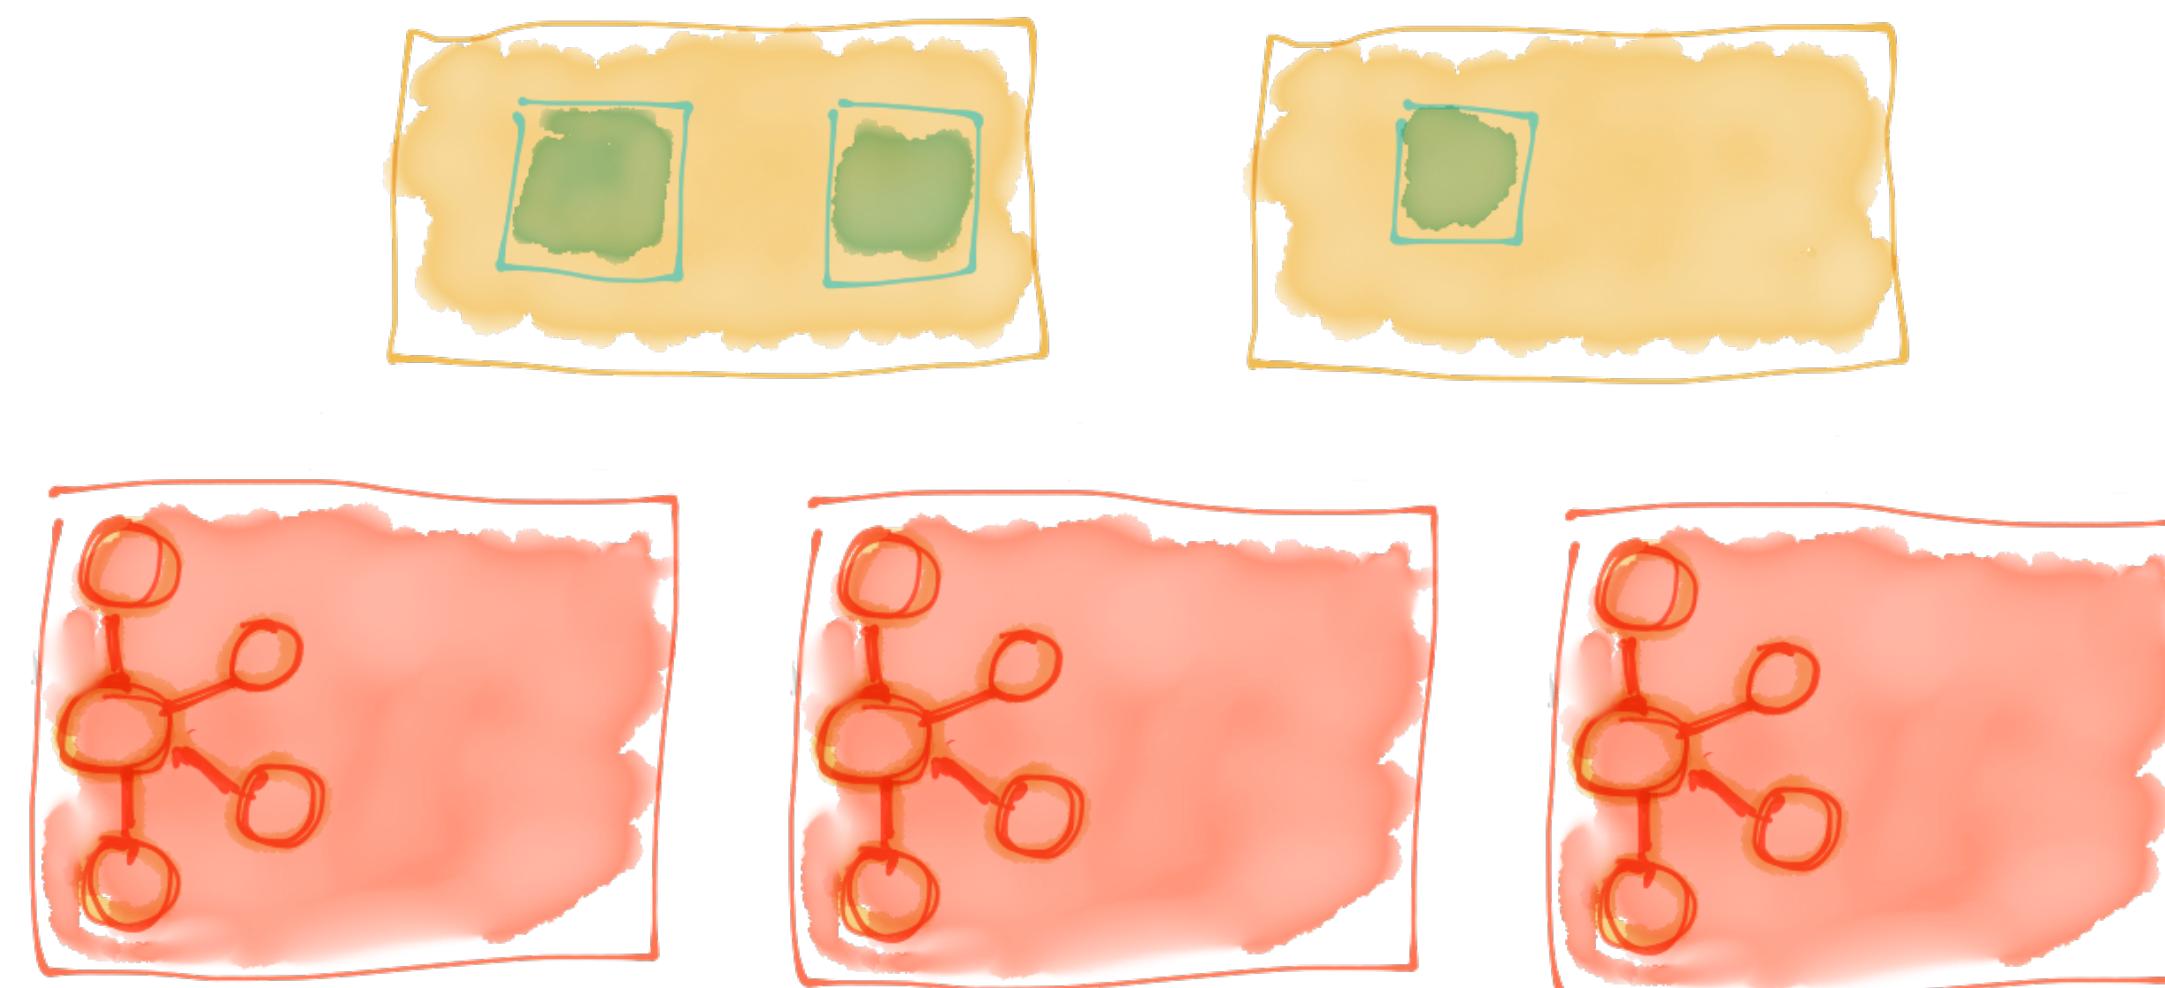
A practical guide to becoming l33t with Kafka Connect

# What is Kafka Connect?

# Streaming Integration with Kafka Connect



## Sources

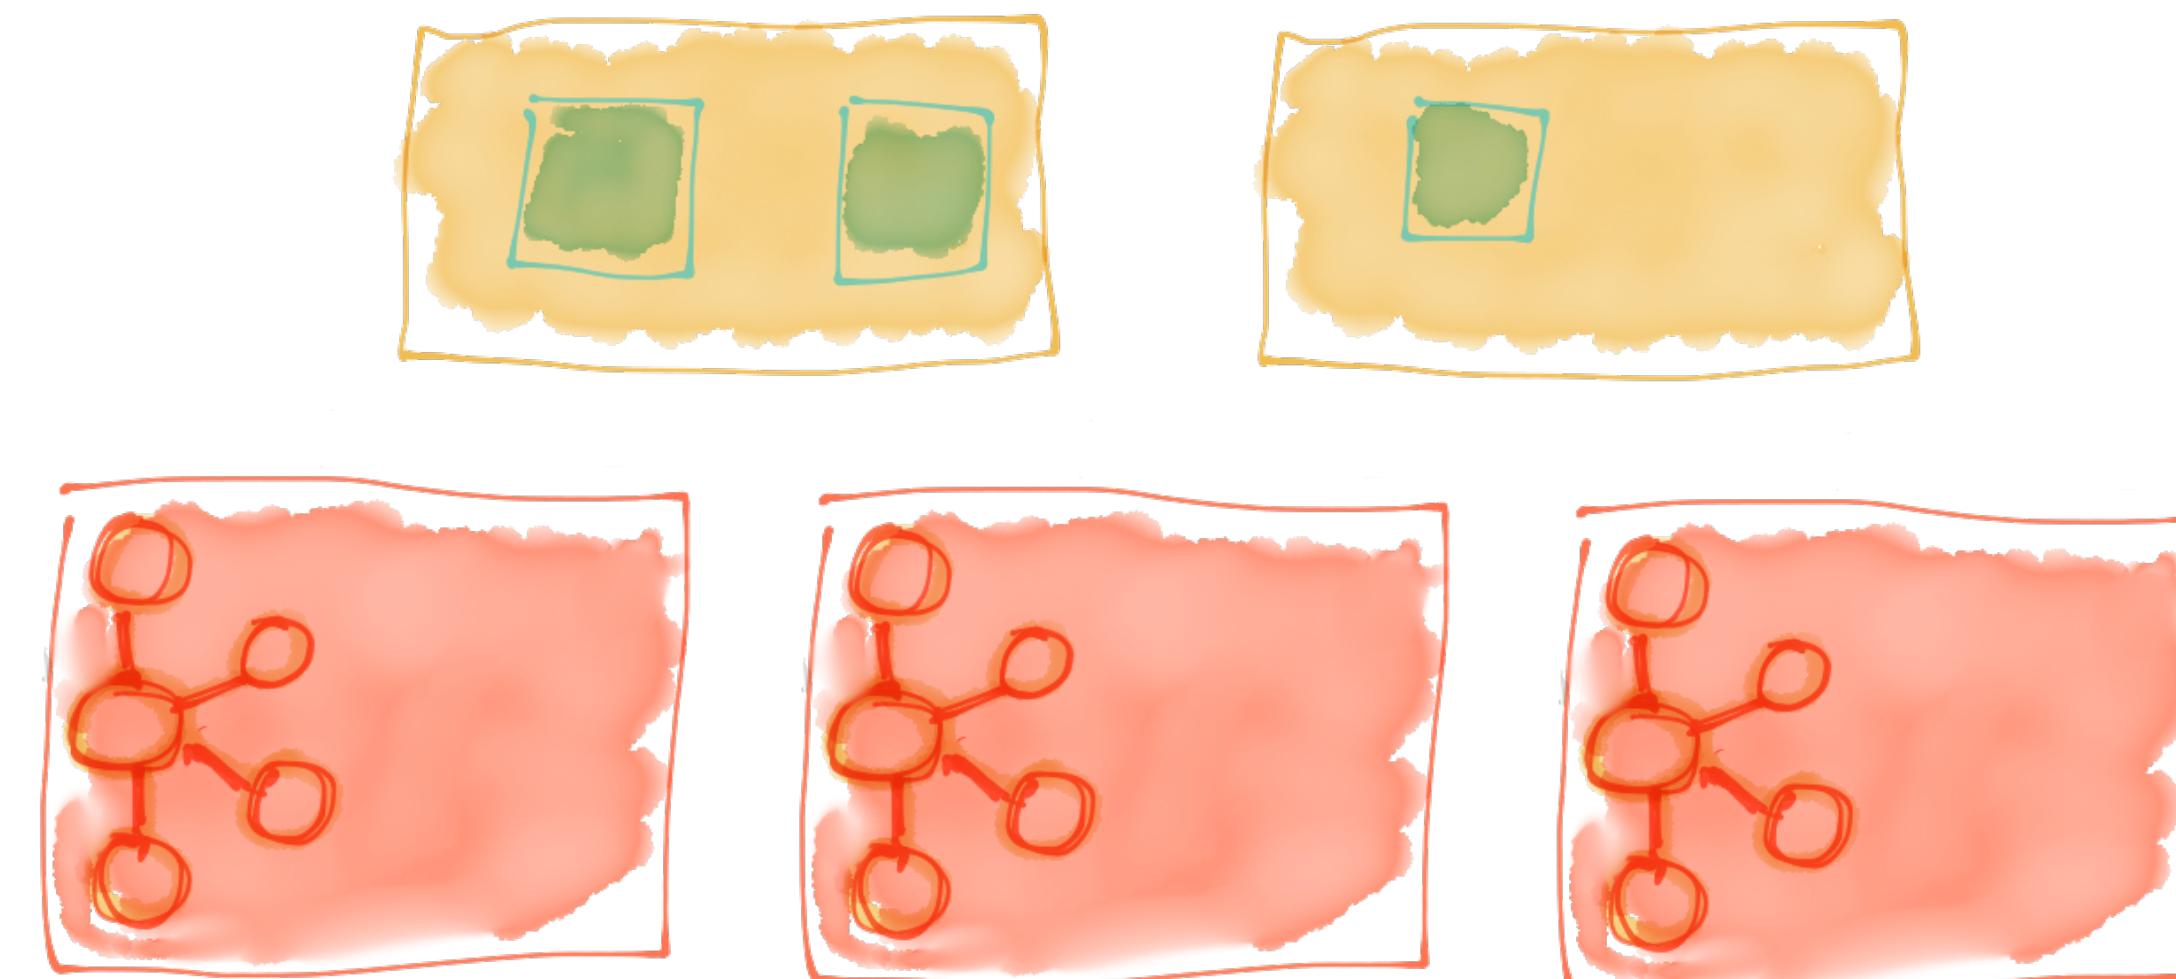


## Kafka Connect

## Kafka Brokers

# Streaming Integration with Kafka Connect

Sinks



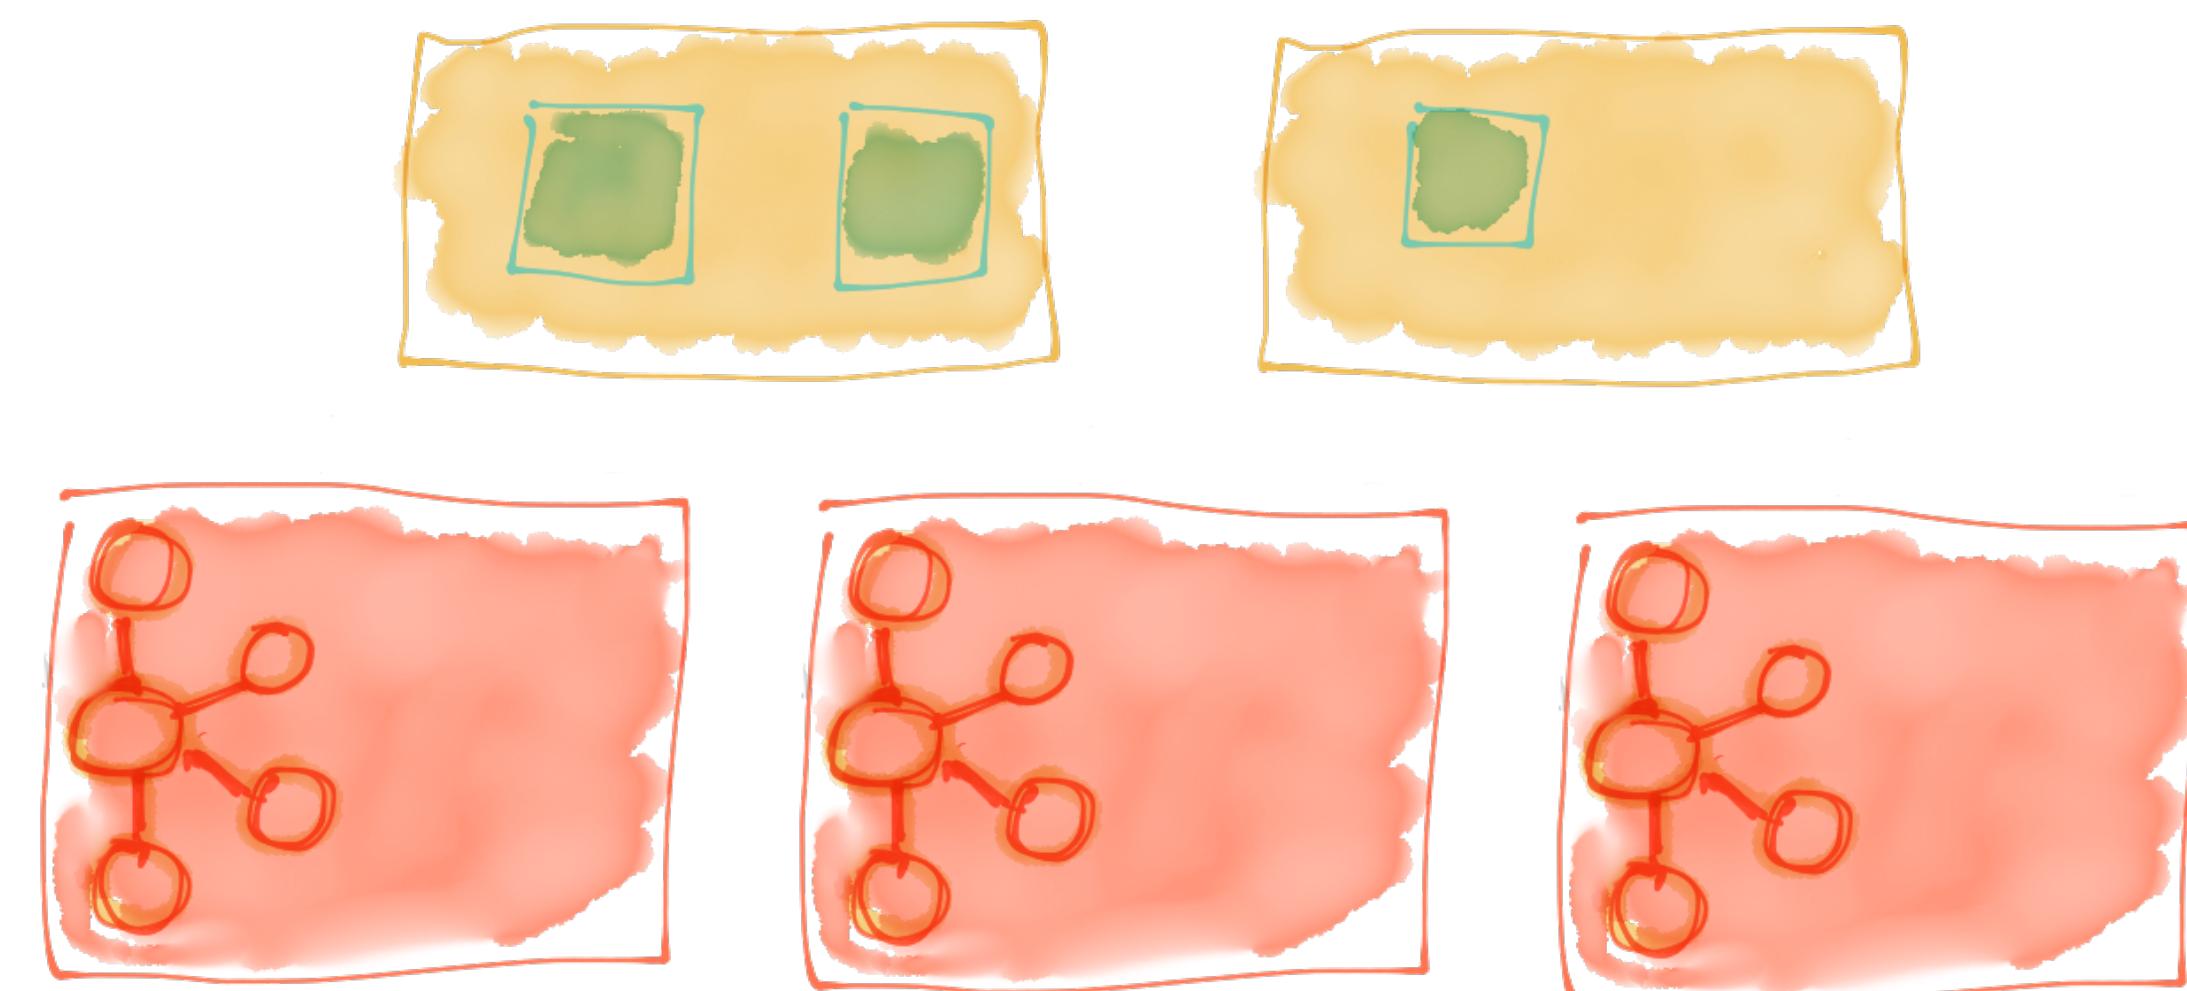
Kafka Connect

Kafka Brokers

# Streaming Integration with Kafka Connect



Kafka Connect



Kafka Brokers

# Look Ma, No Code!

{

```
"connector.class":
```

```
    "io.confluent.connect.jdbc.JdbcSourceConnector",
```

```
"connection.url":
```

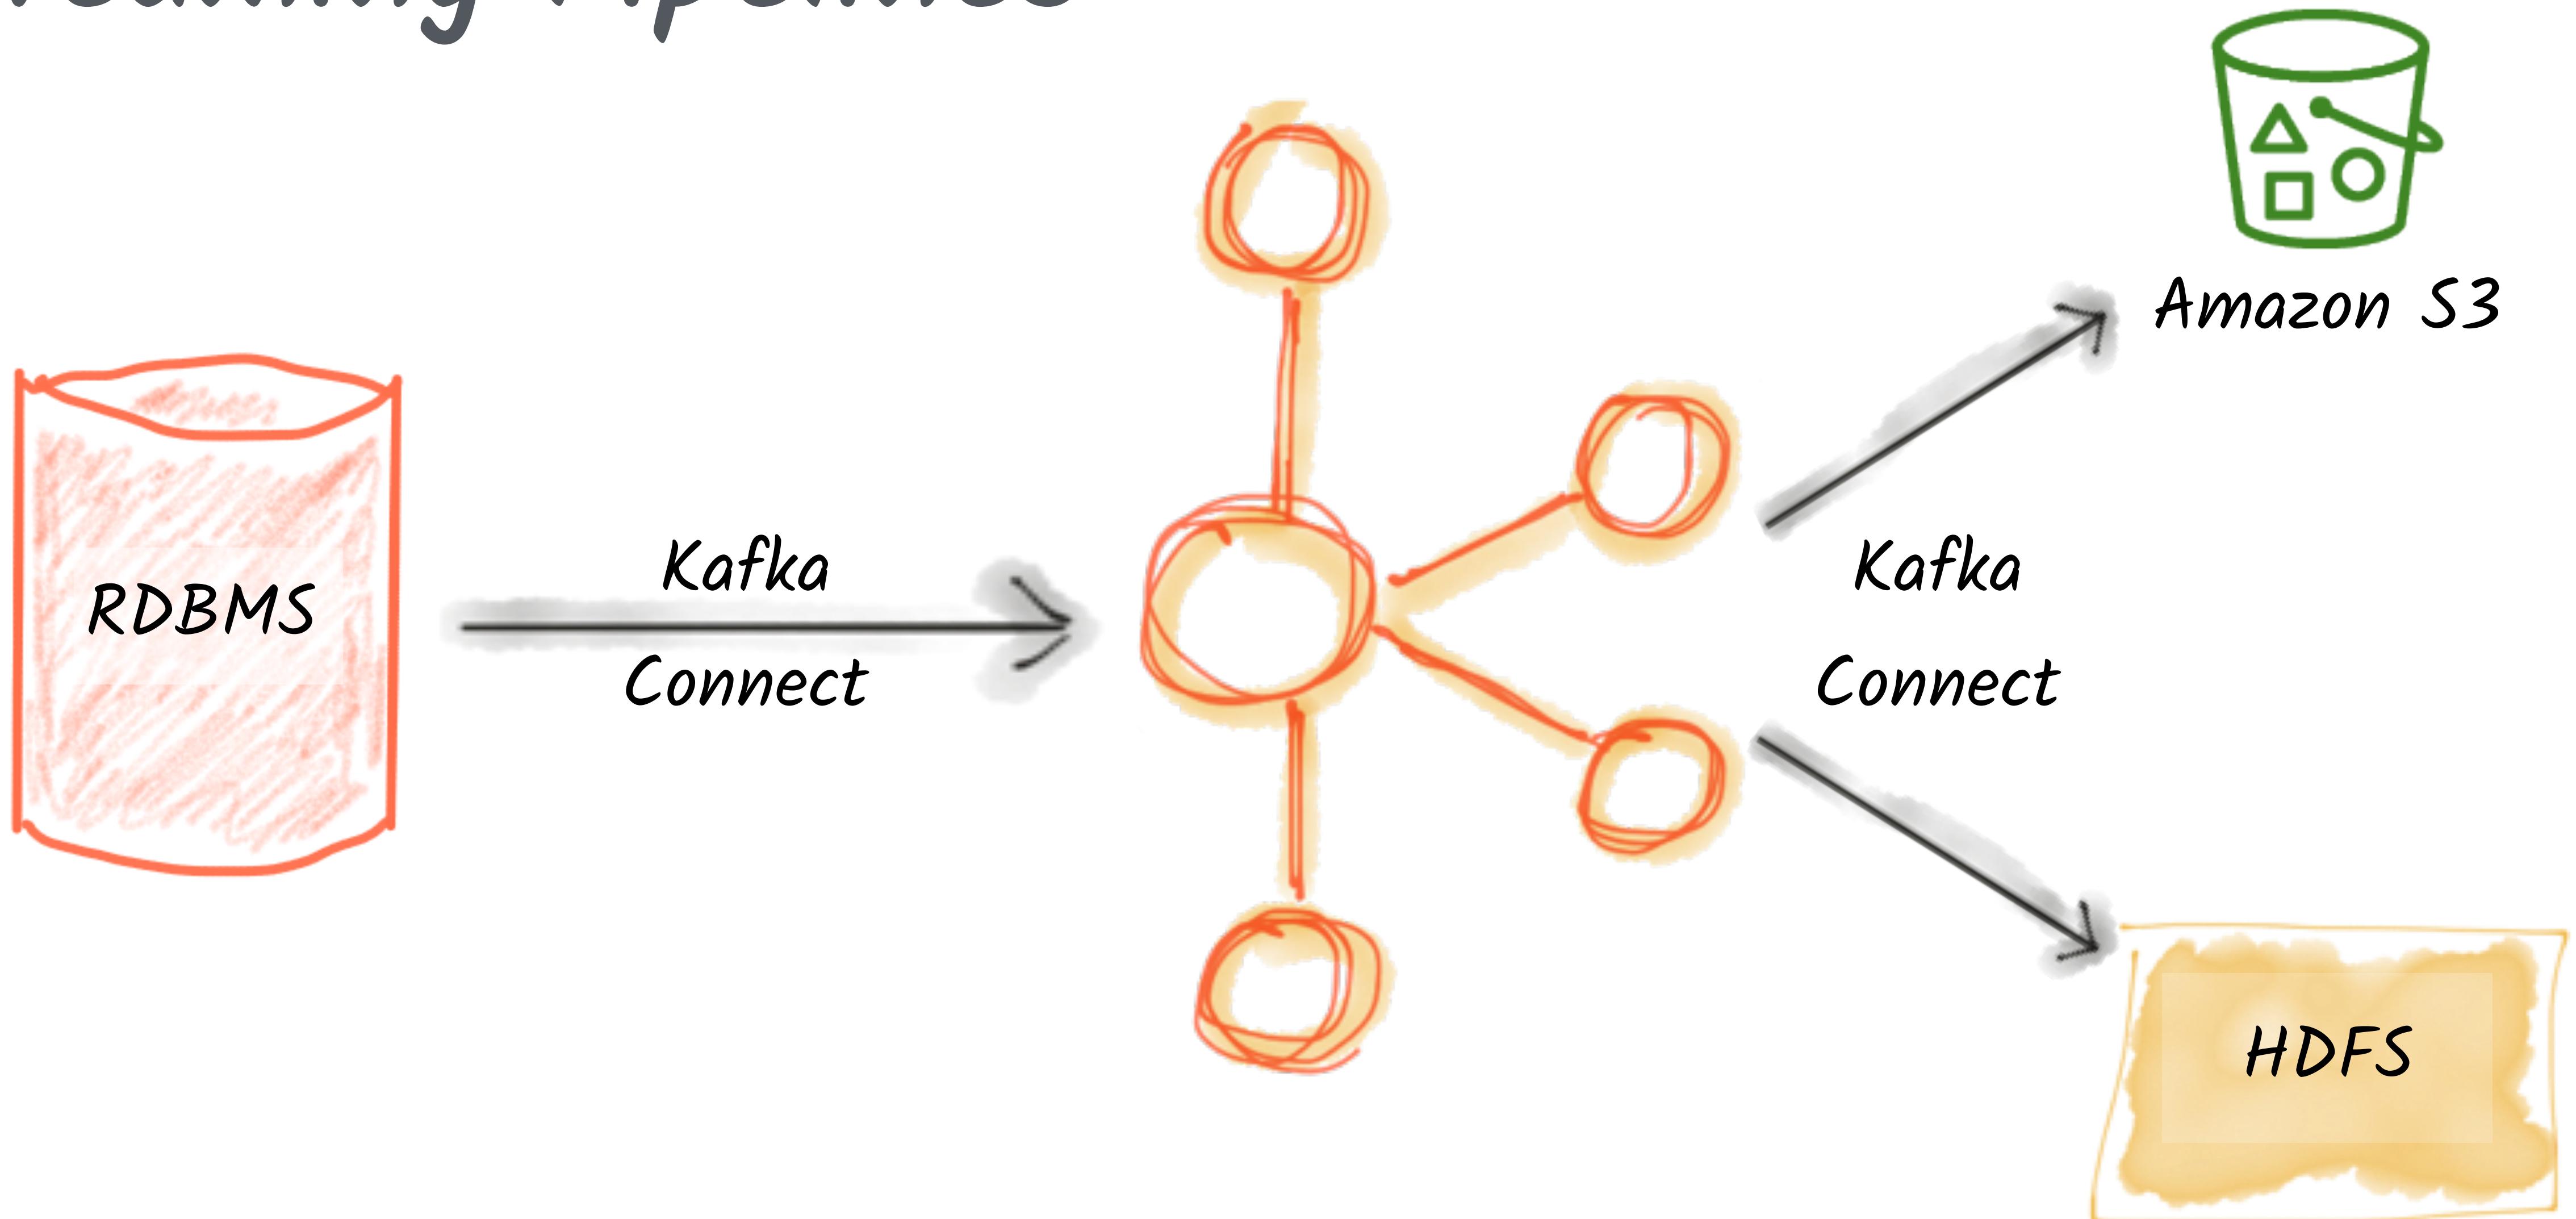
```
    "jdbc:mysql://asgard:3306/demo",
```

```
"table.whitelist":
```

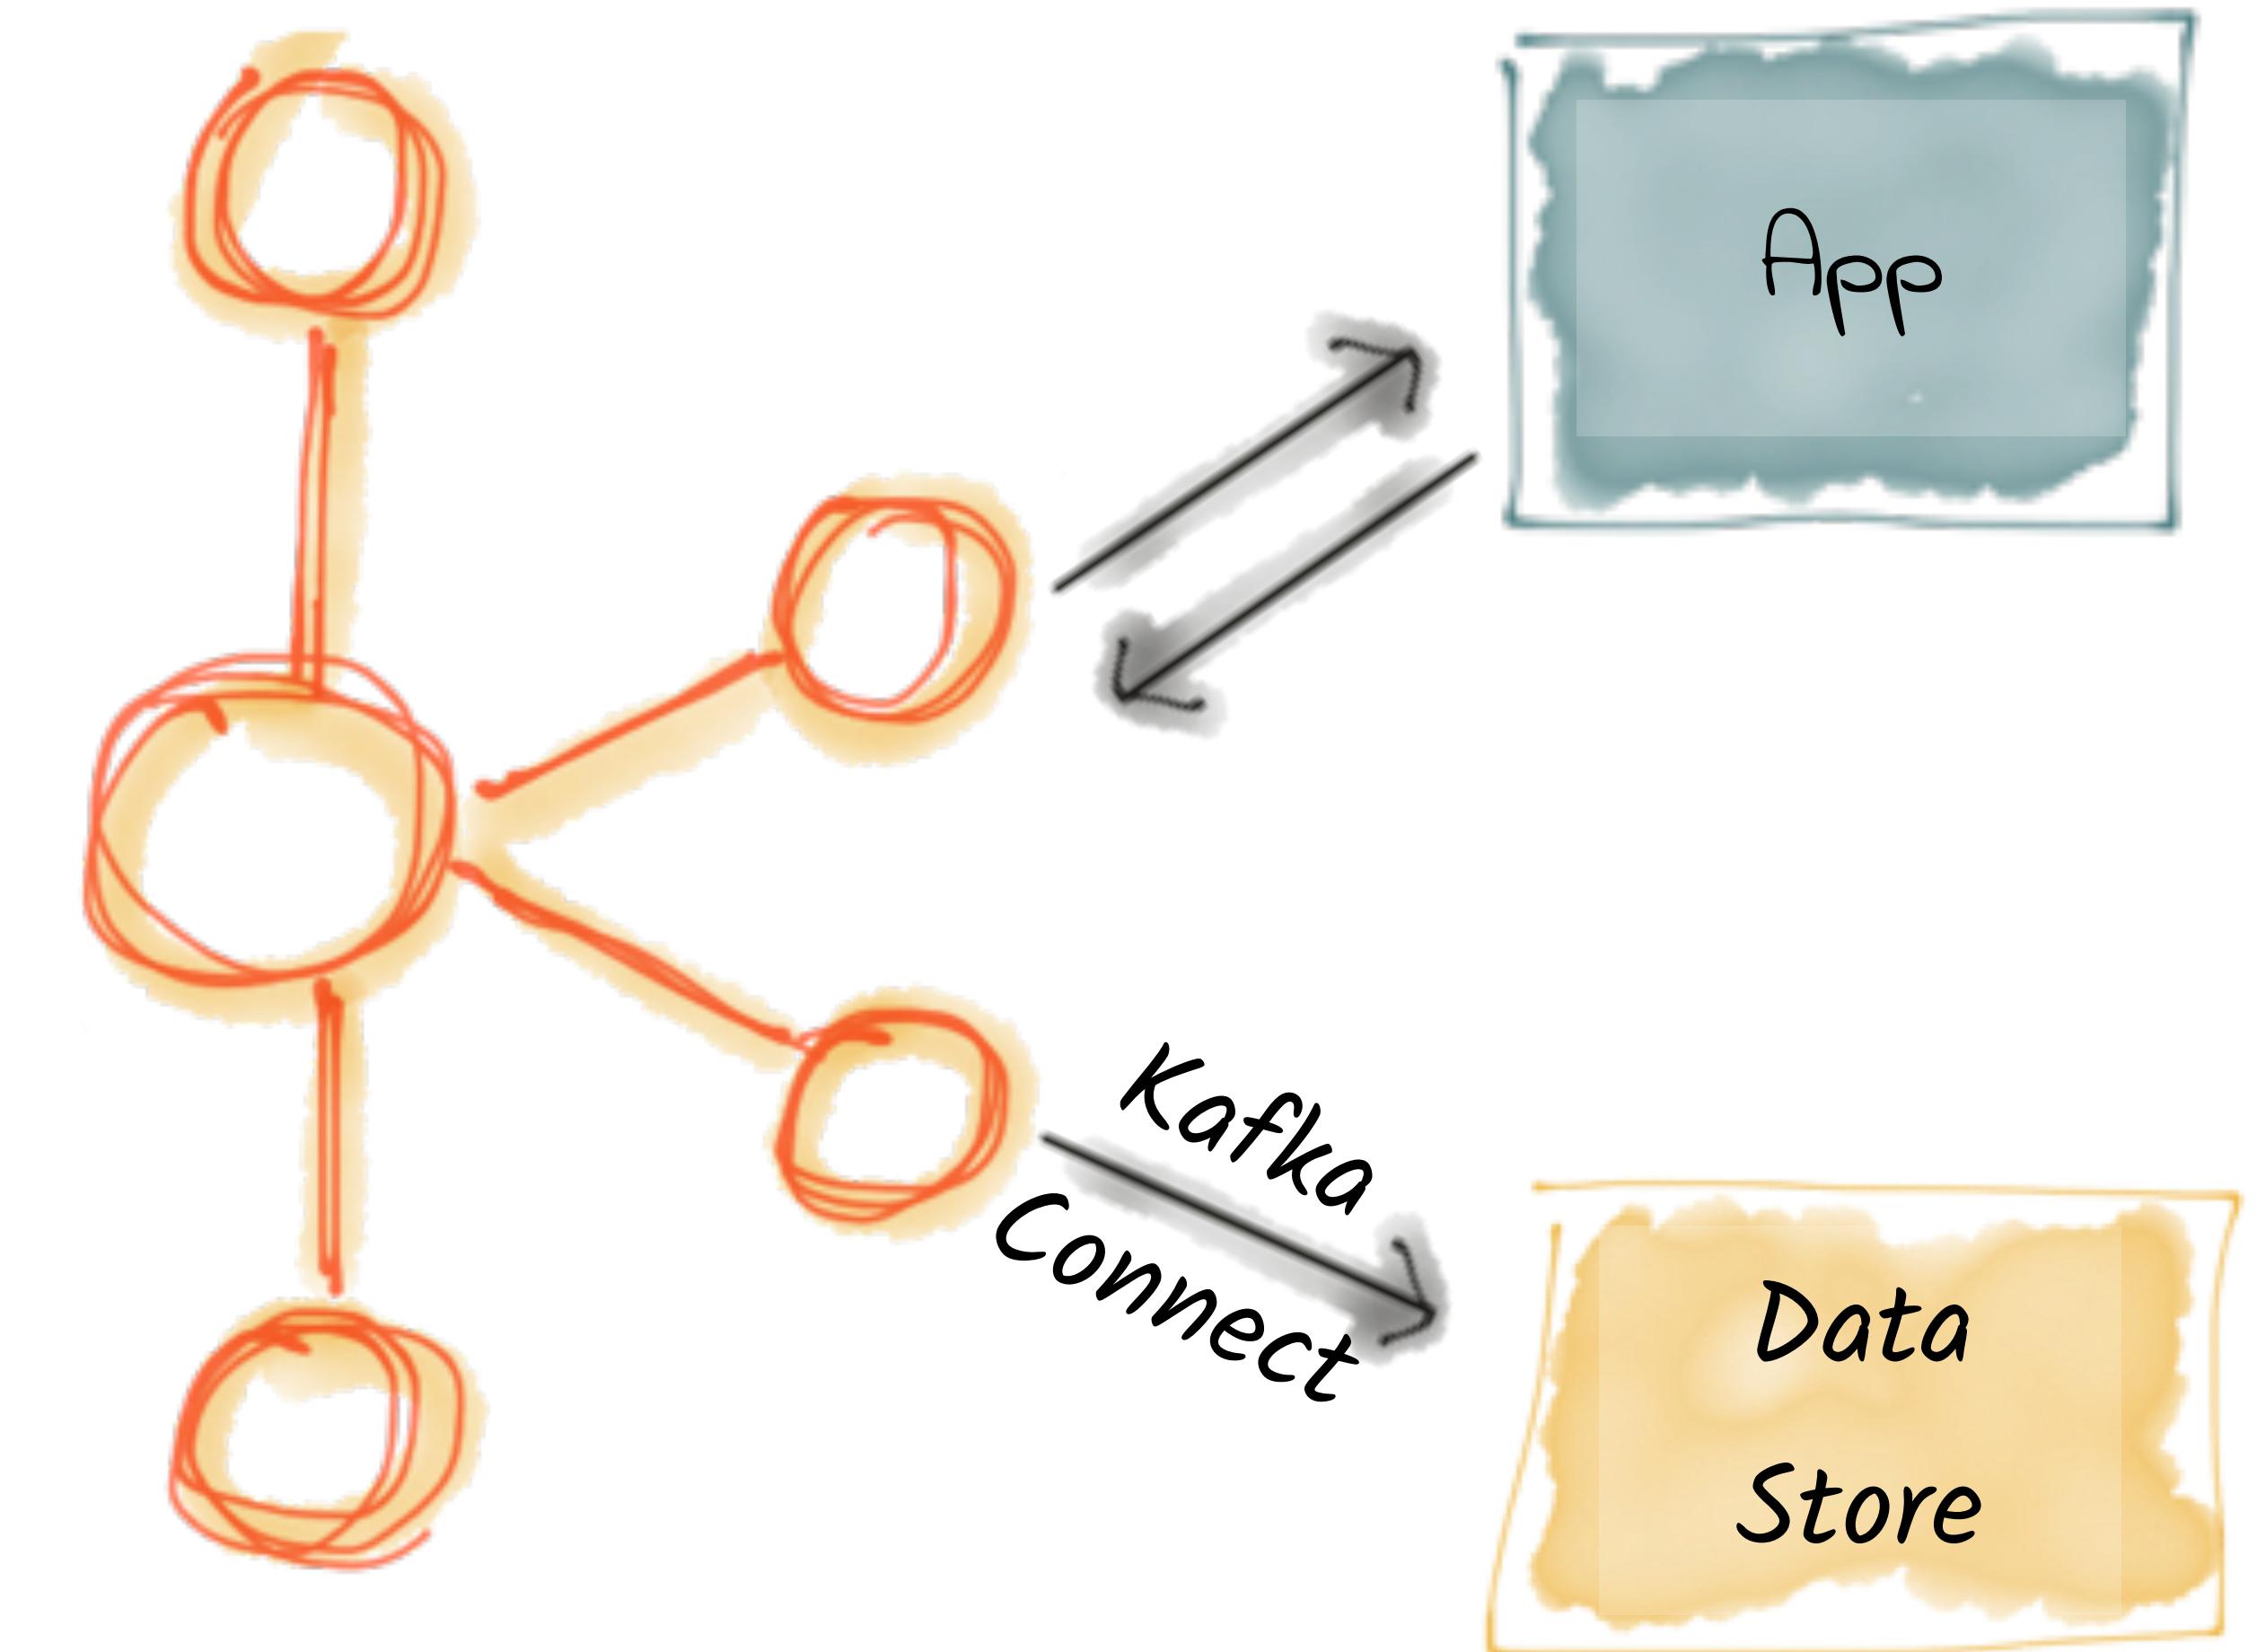
```
    "sales,orders,customers"
```

}

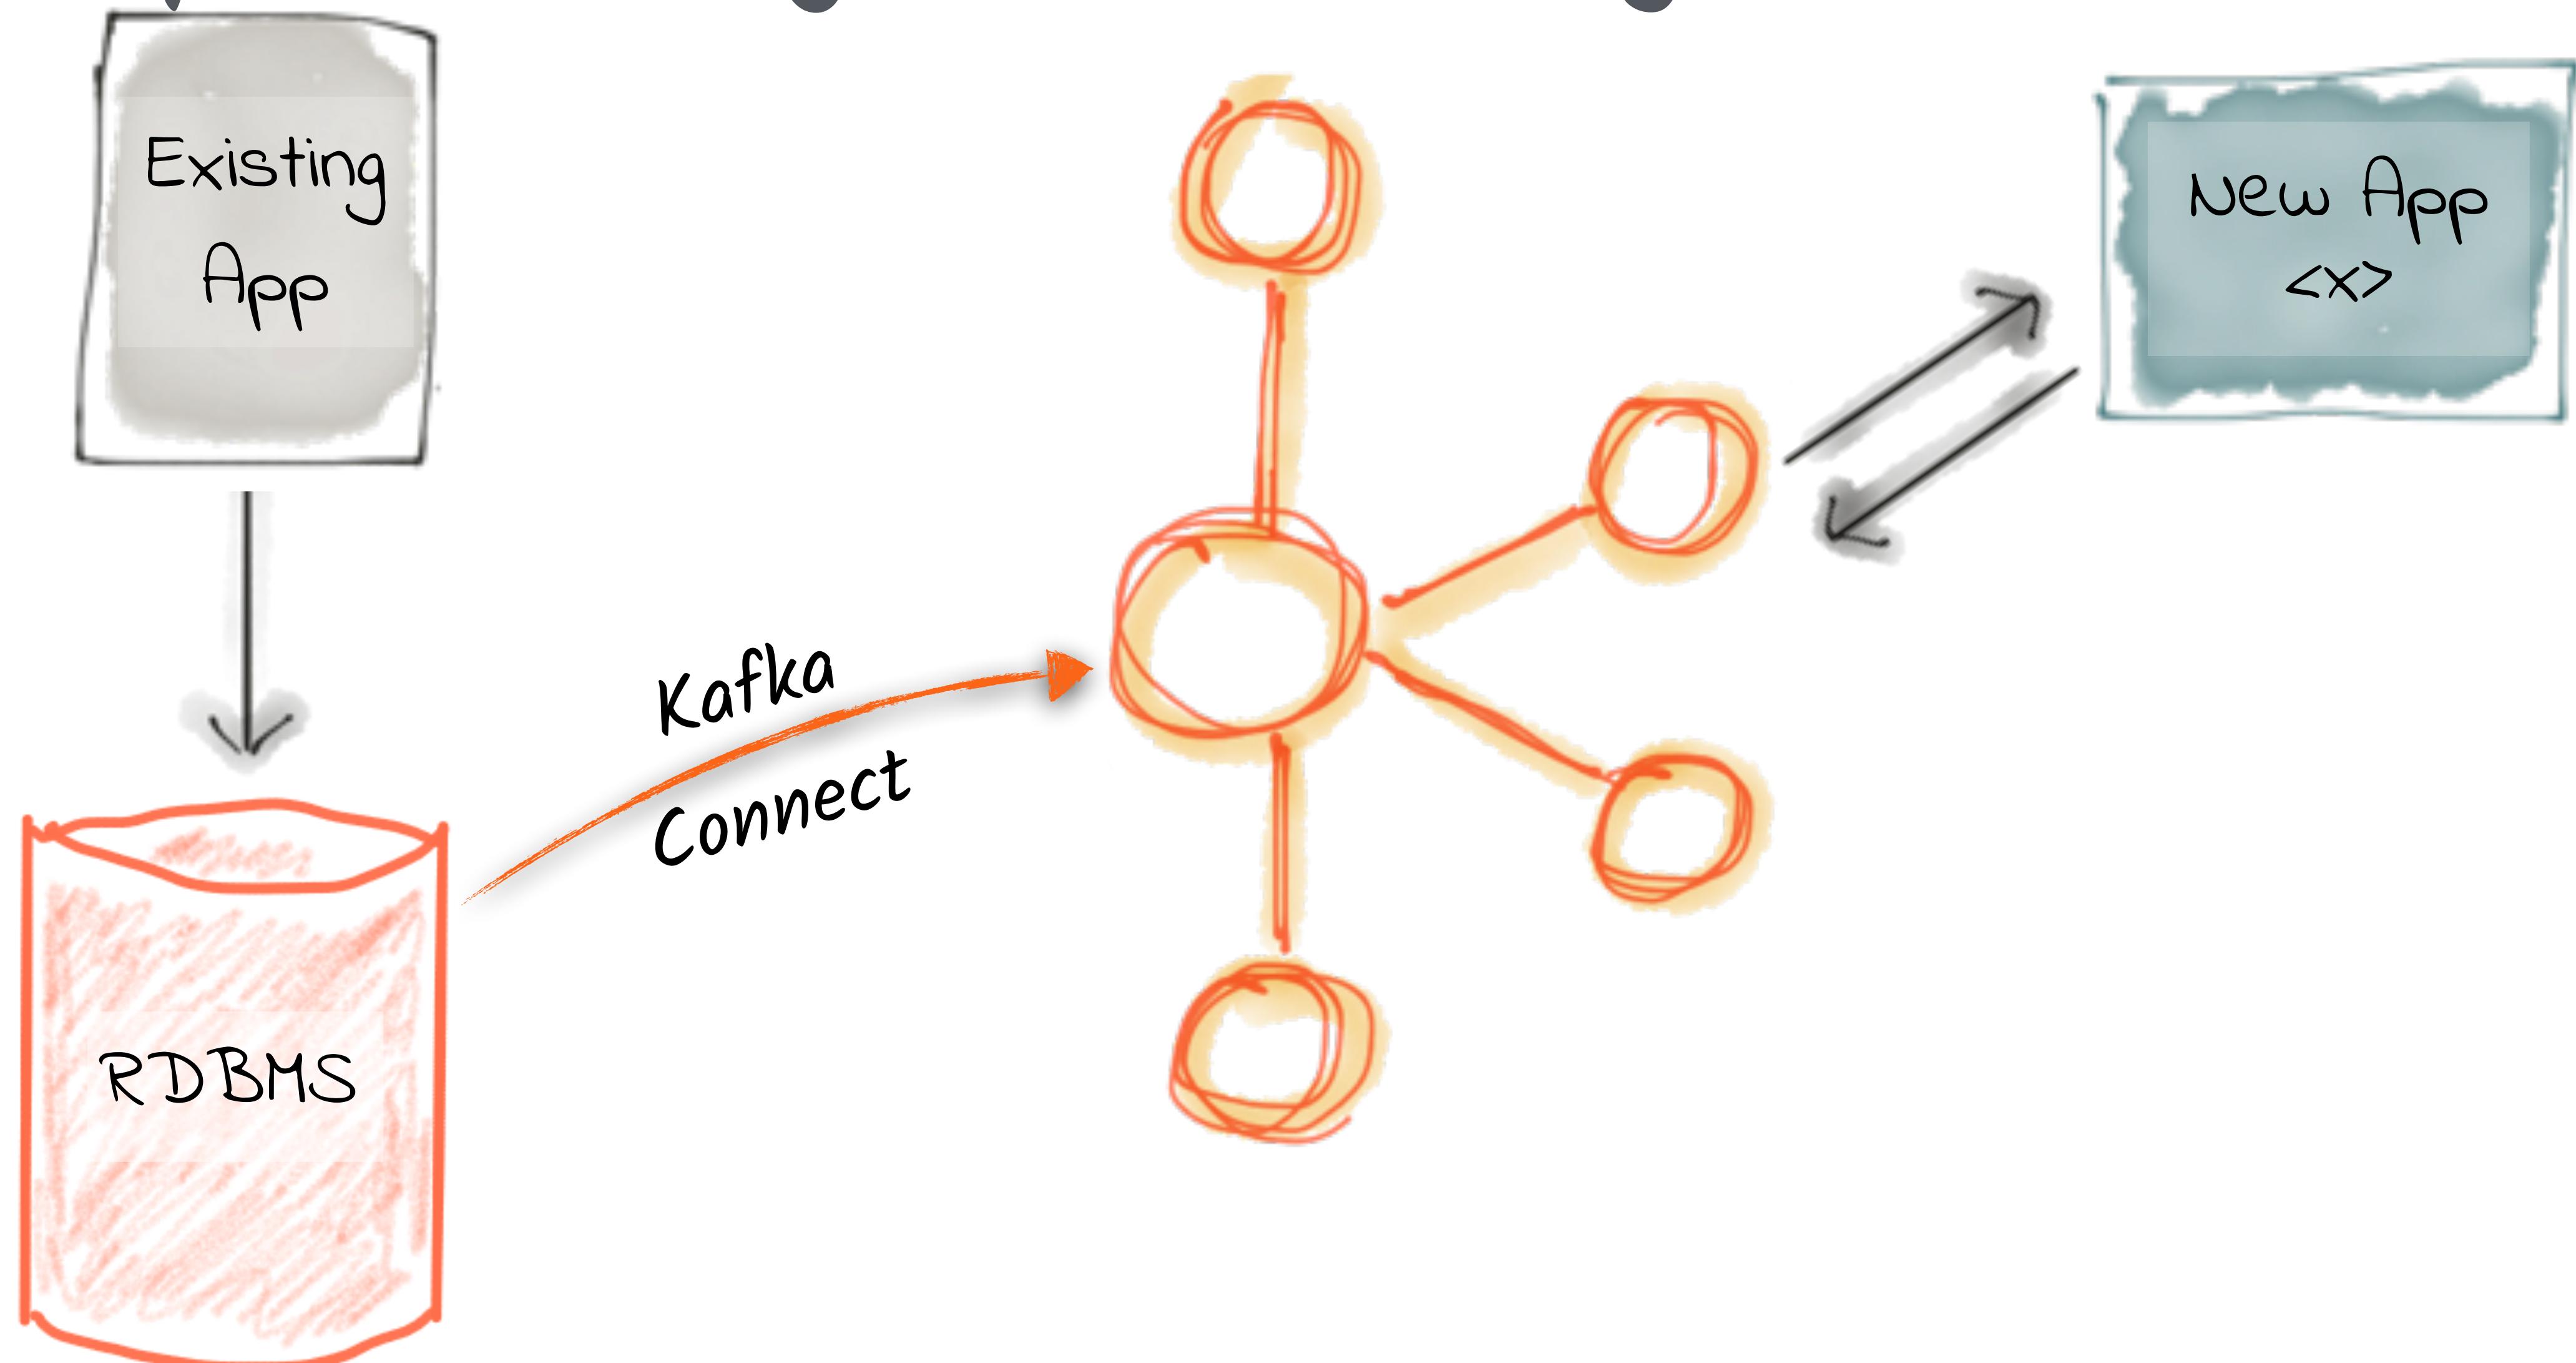
# Streaming Pipelines



# Writing to data stores from Kafka



# Evolve processing from old systems to new



# Demo

<http://rmoff.dev/kafka-connect-code>

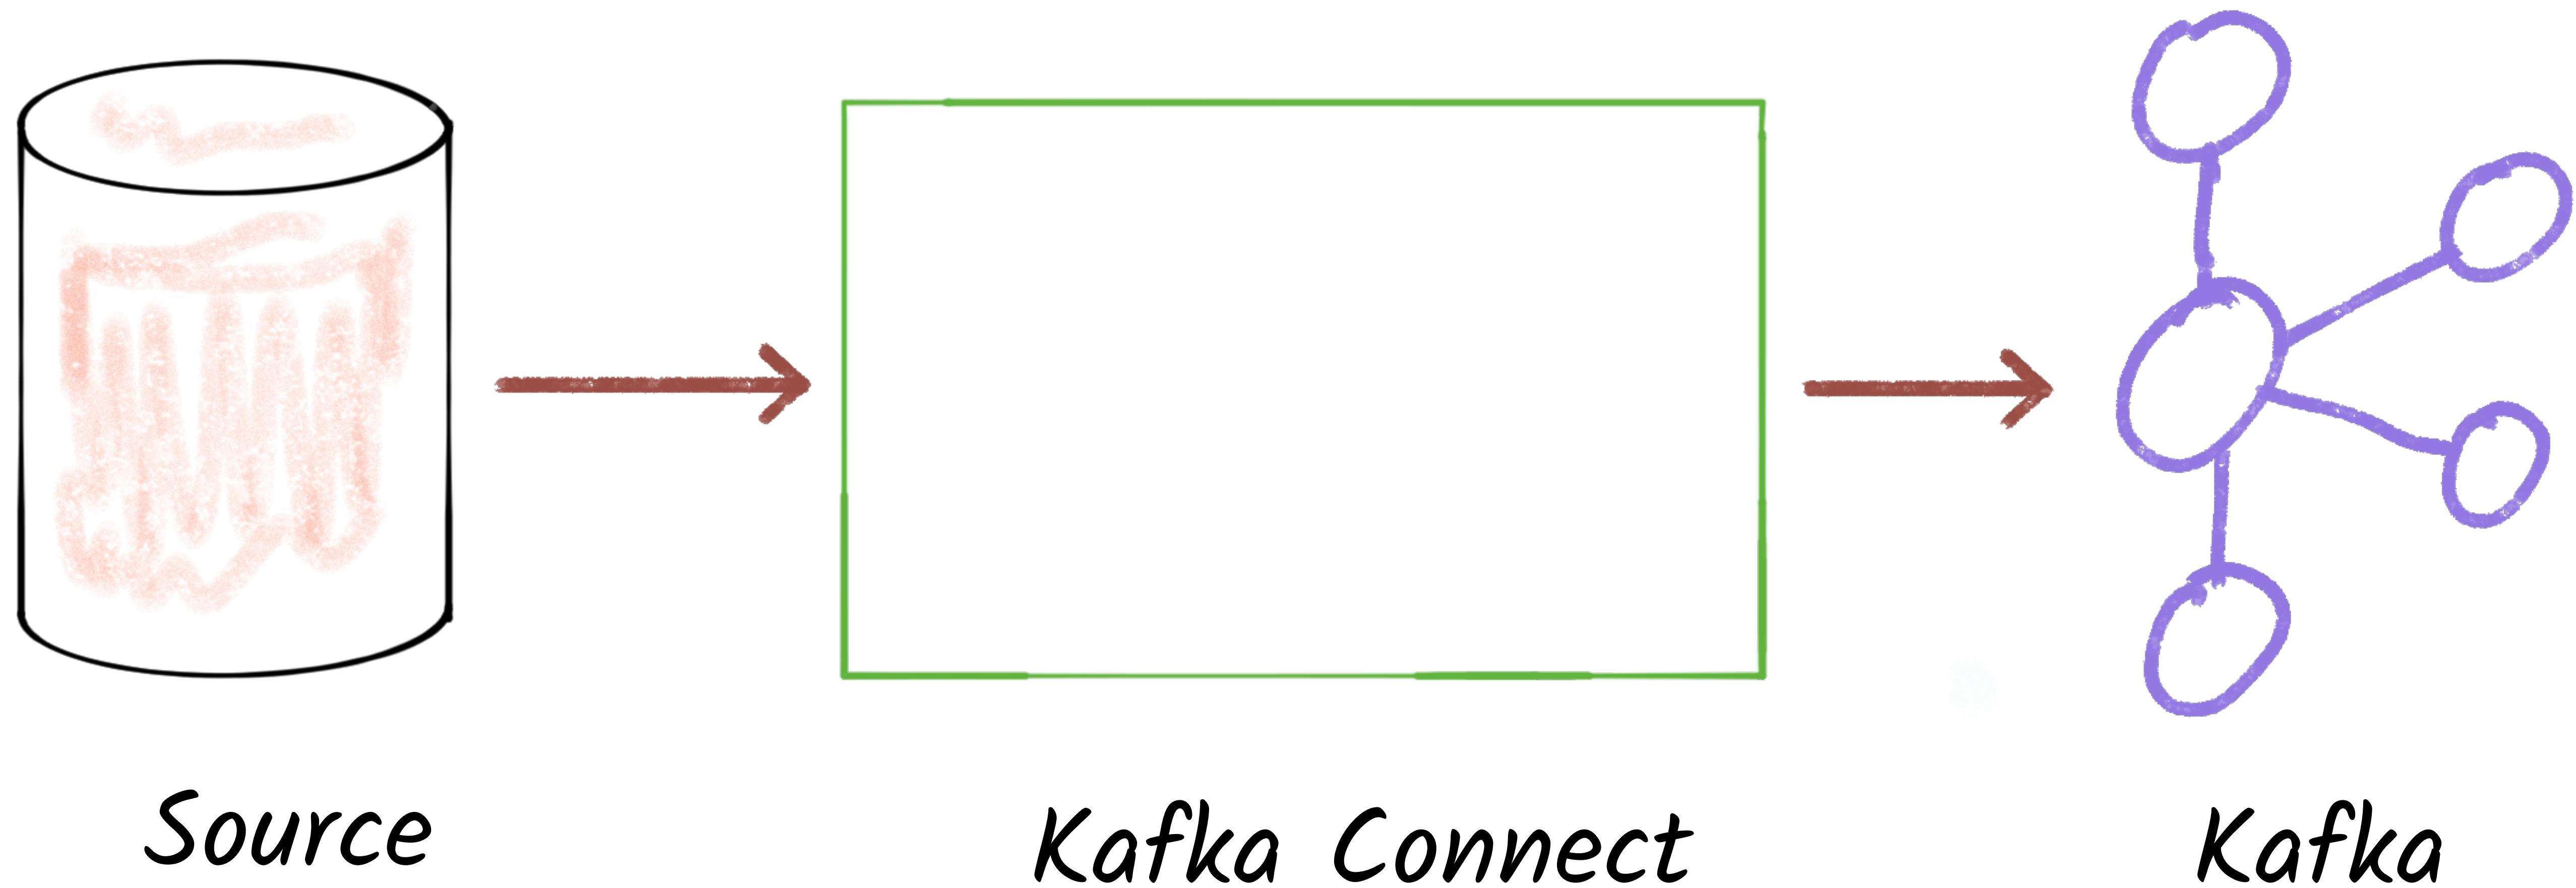
# Configuring

## Kafka

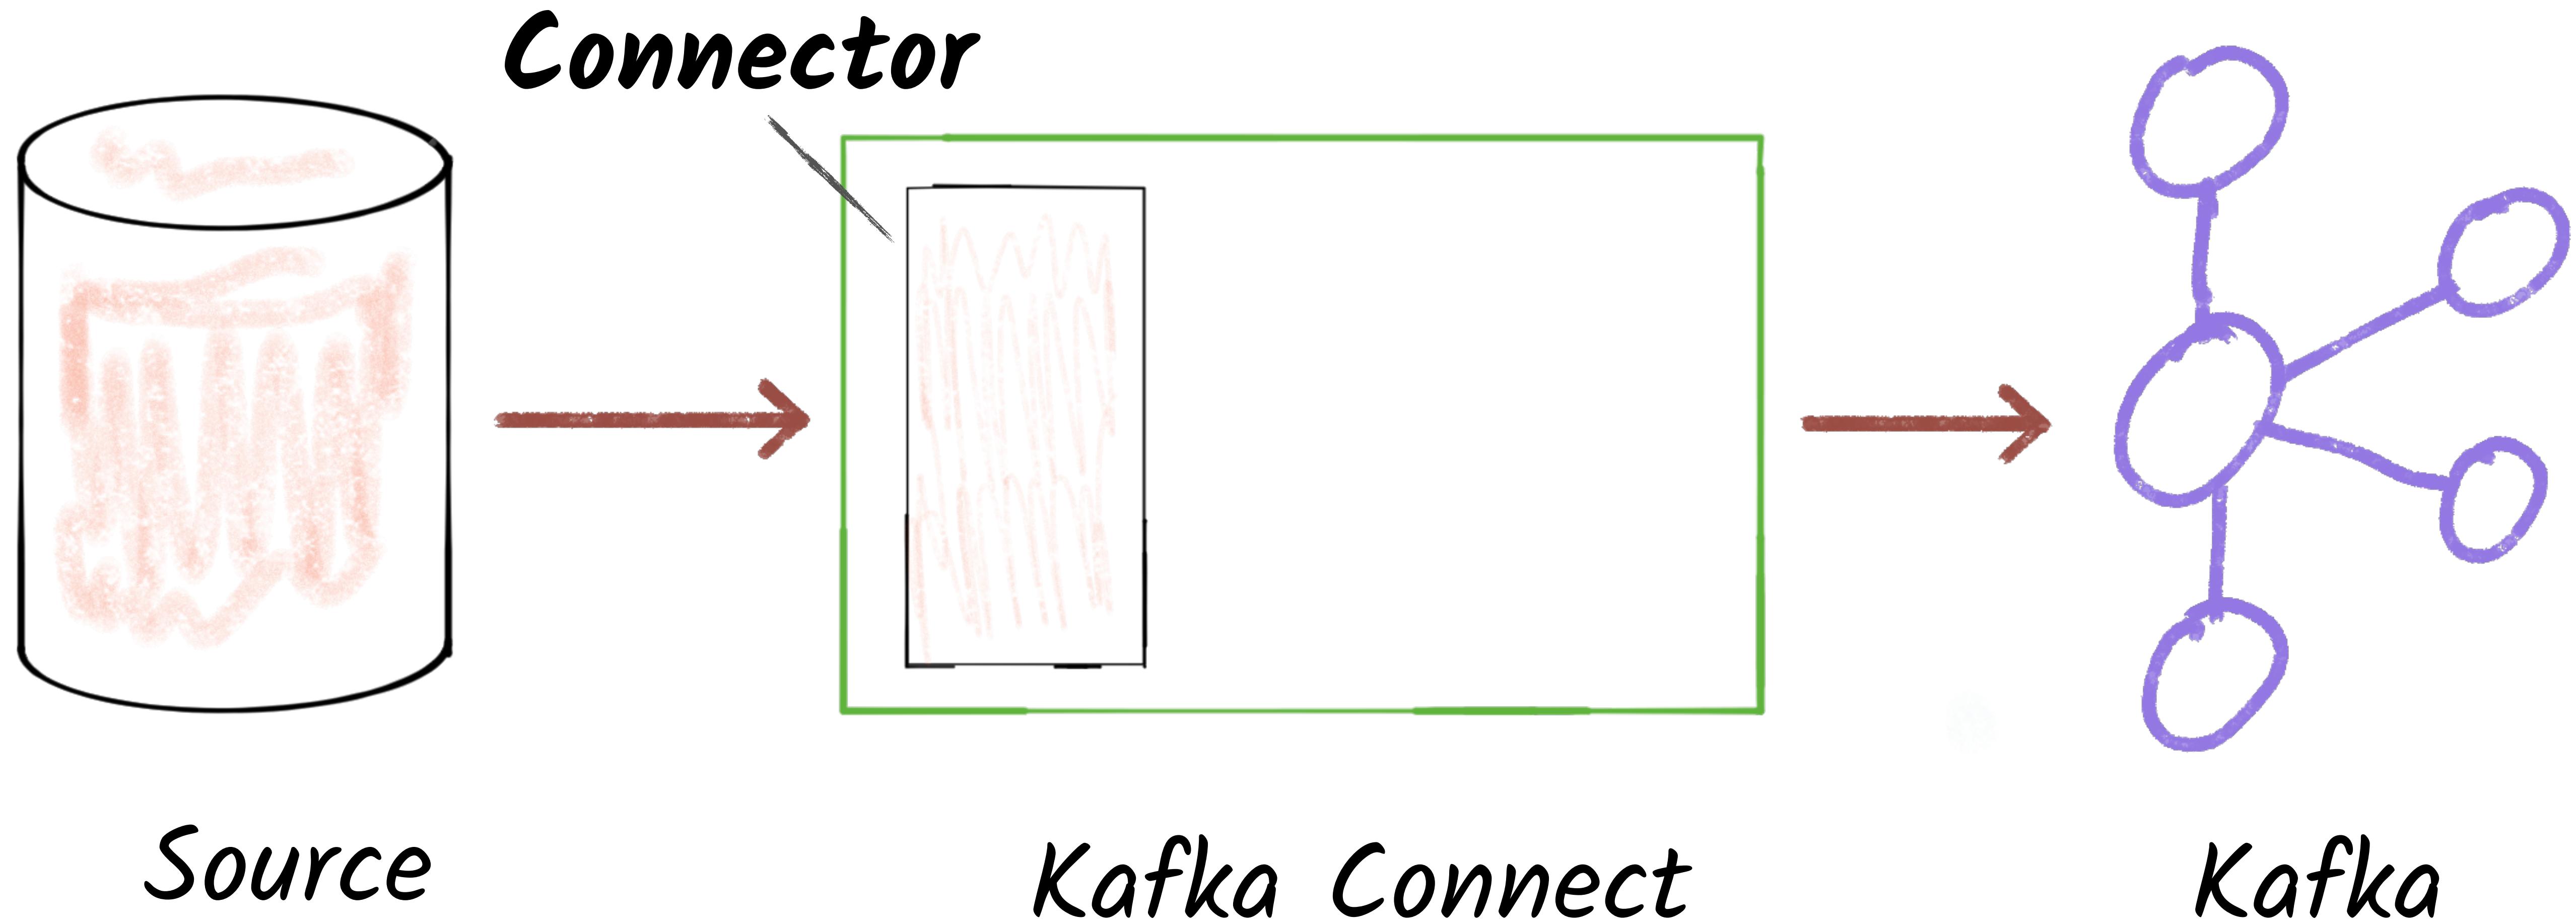
## Connect

Inside the API - connectors, transforms, converters

# Kafka Connect basics



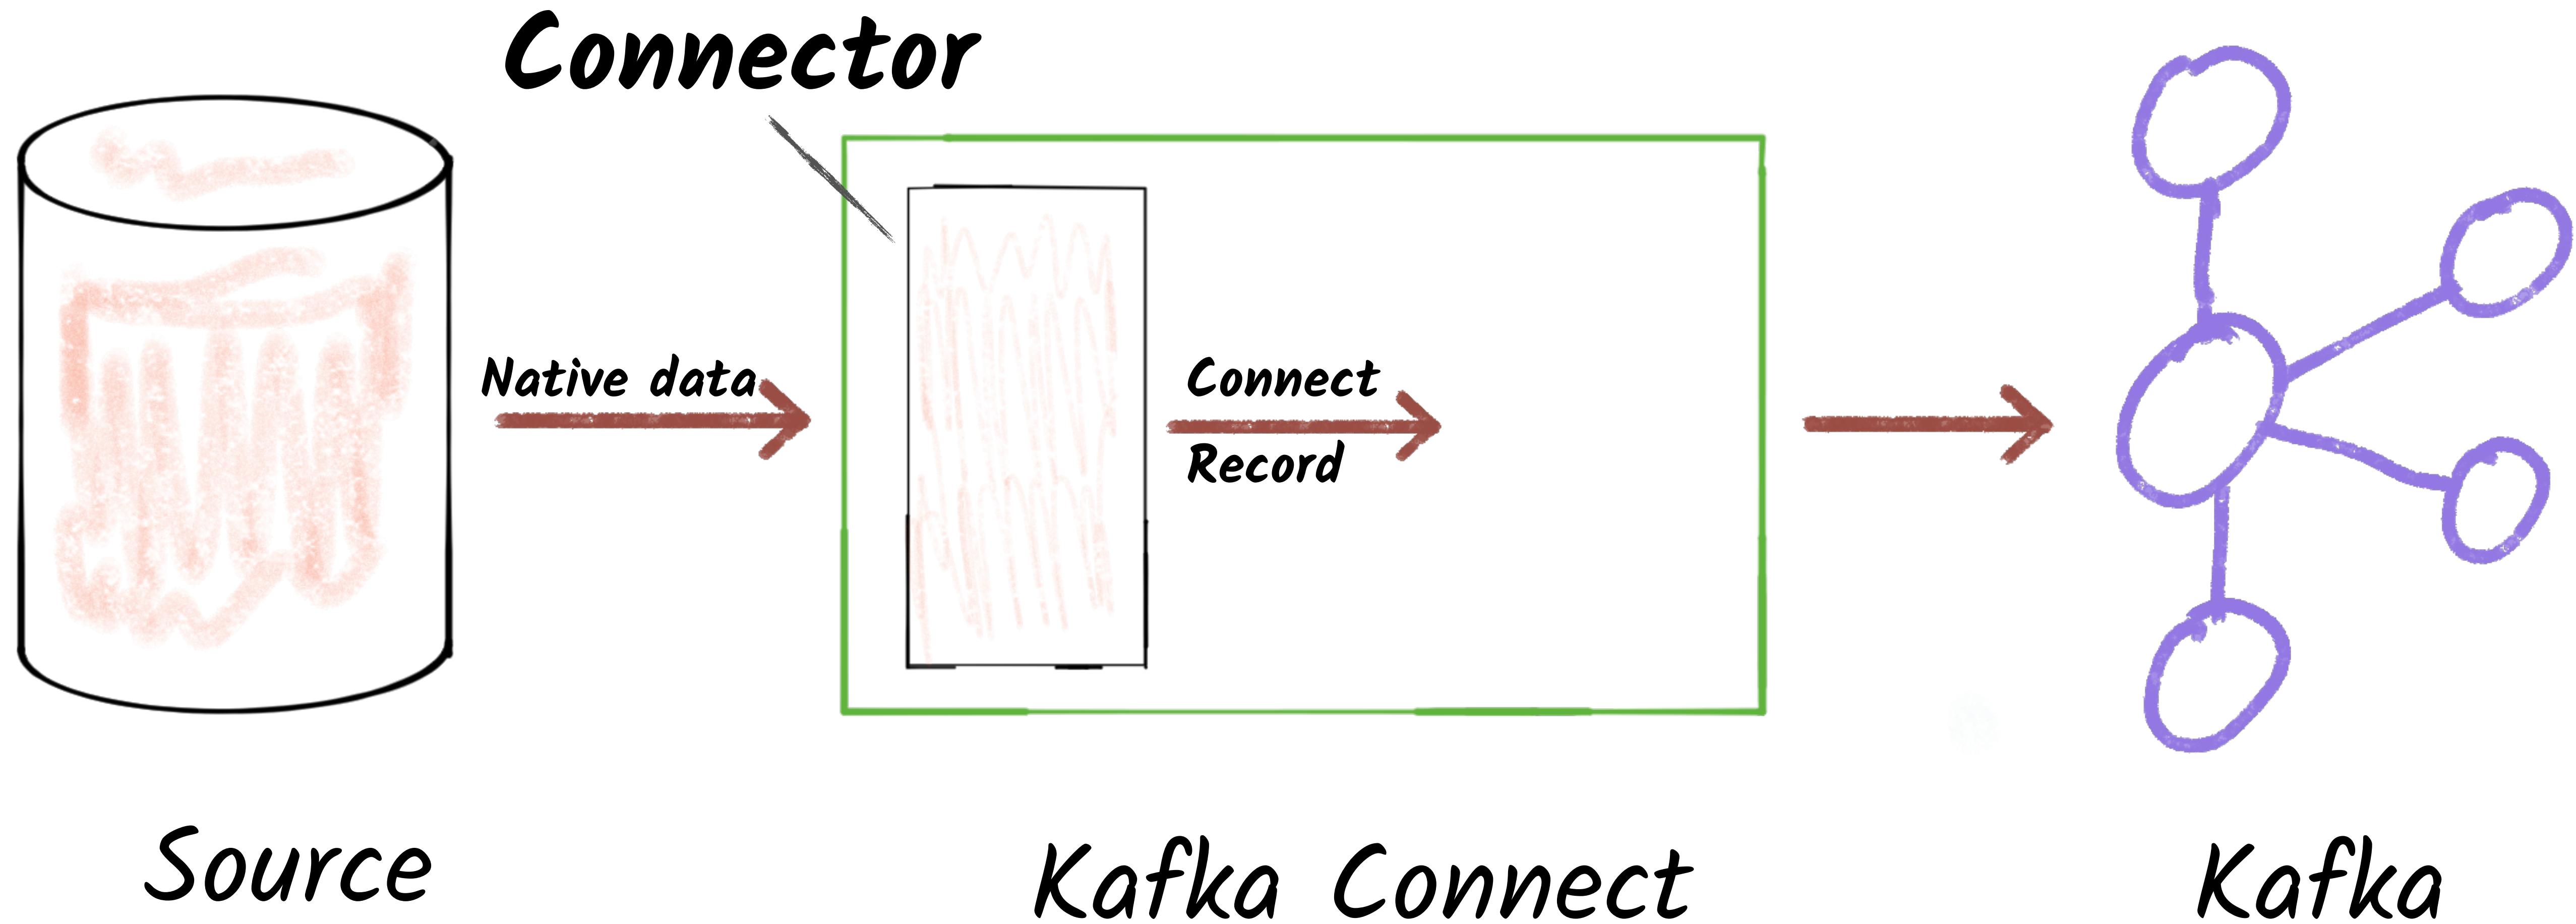
# Connectors



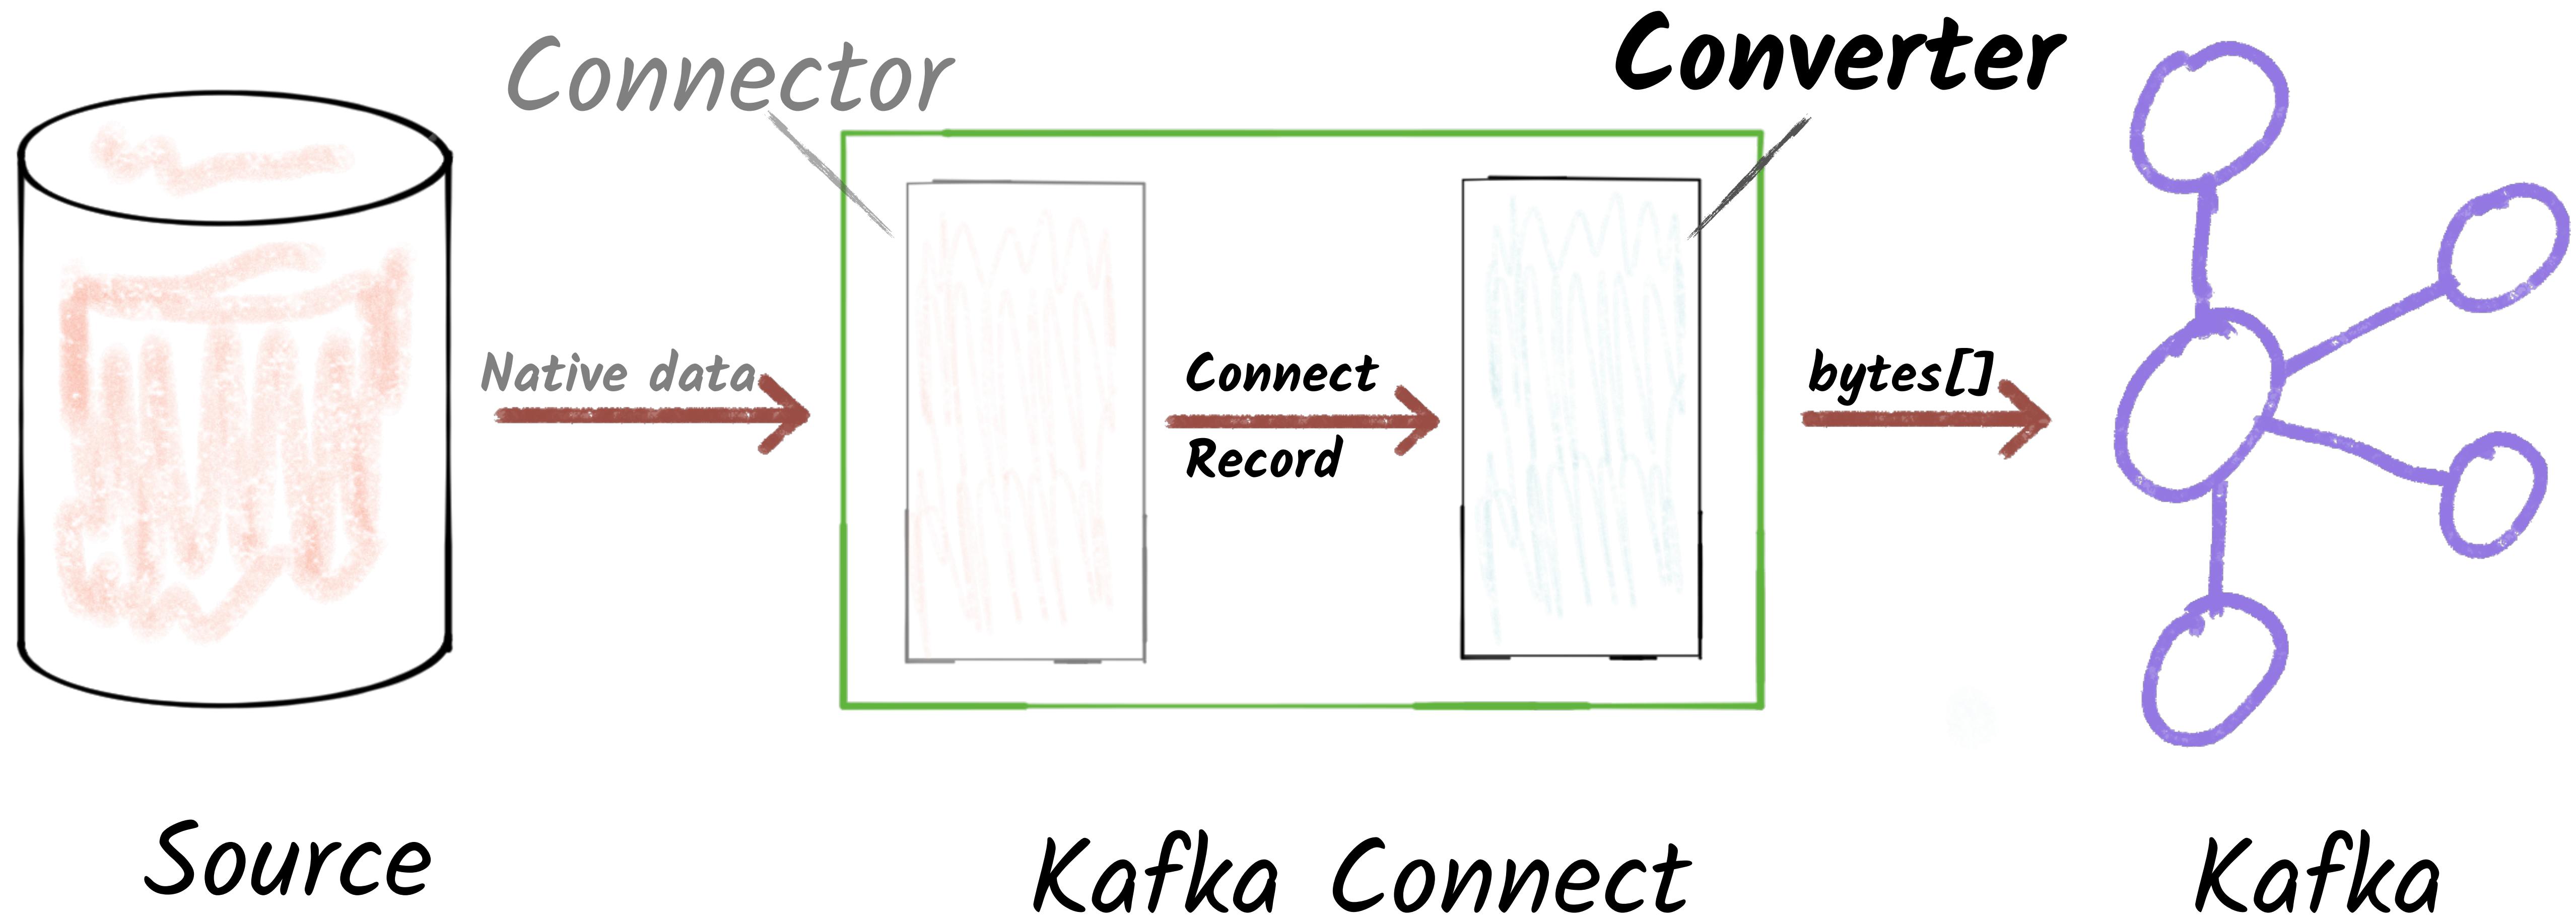
# Connectors

```
"config": {  
    [...]  
    "connector.class": "io.confluent.connect.jdbc.JdbcSinkConnector"  
    "connection.url": "jdbc:postgresql://postgres:5432/",  
    "topics": "asgard.demo.orders",  
}
```

# Connectors



# Converters



# Serialisation & Schemas

## Avro

-> Confluent  
Schema Registry

## Protobuf

## JSON

## CSV



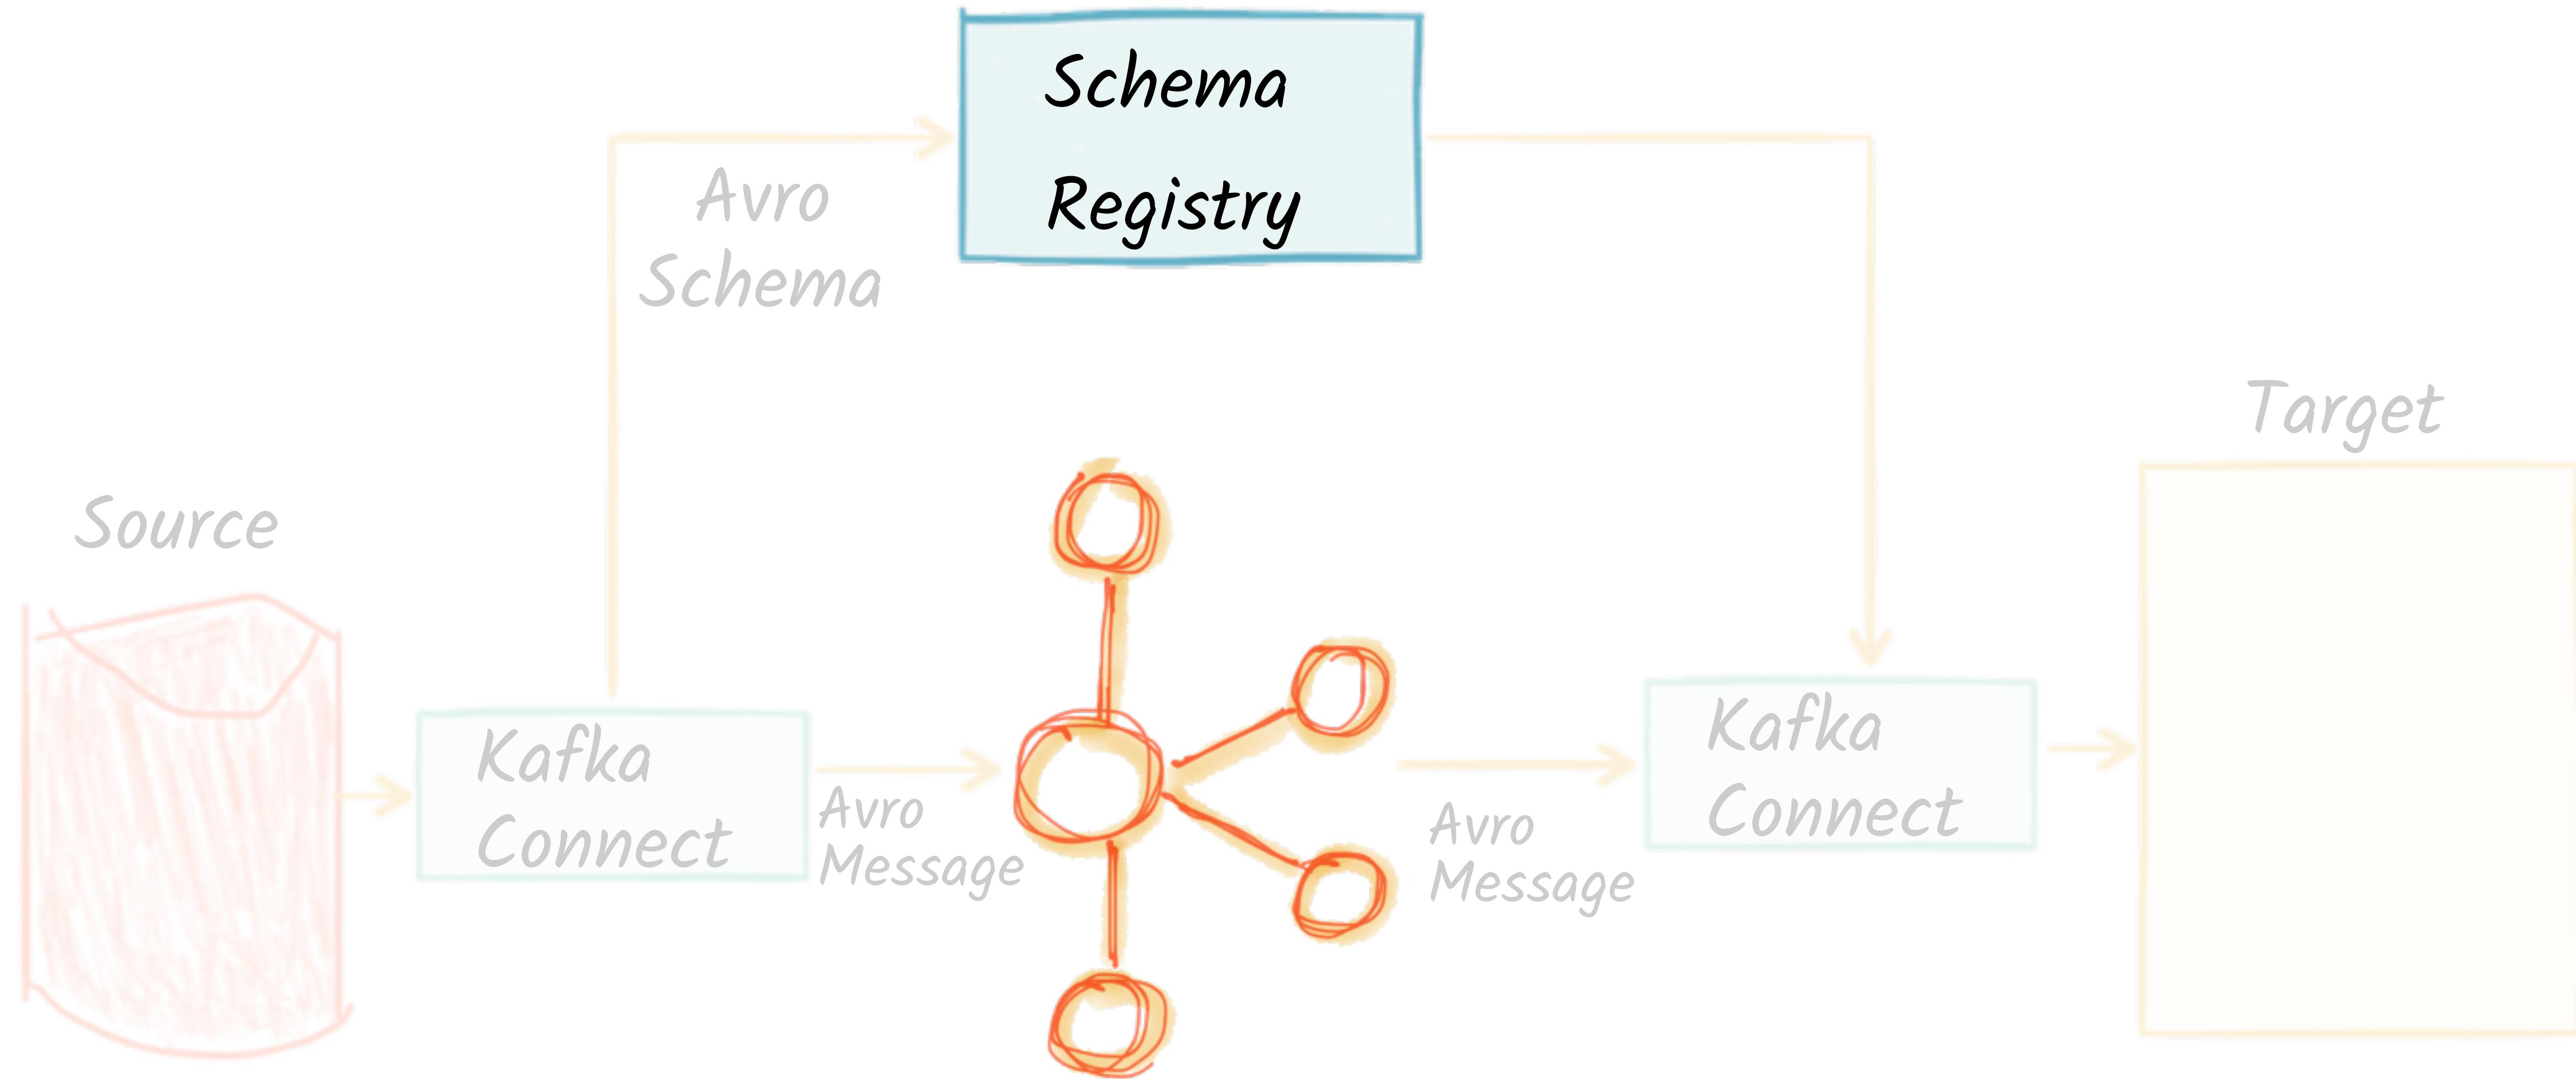
Gwen (Chen) Shapira  
@gwenshap

If your dev process doesn't validate schema compatibility somewhere between your IDE and production - you are screwed and don't know it.

5:50 AM - 5 Apr 2017

[https://qconnewyork.com/system/files/presentation-slides/qcon\\_17\\_-\\_schemas\\_and\\_apis.pdf](https://qconnewyork.com/system/files/presentation-slides/qcon_17_-_schemas_and_apis.pdf)

# The Confluent Schema Registry



# Converters

```
key.converter=io.confluent.connect.avro.AvroConverter  
key.converter.schema.registry.url=http://localhost:8081  
value.converter=io.confluent.connect.avro.AvroConverter  
value.converter.schema.registry.url=http://localhost:8081
```

*Set as a global default per-worker; optionally can be overridden per-connector*

# What about internal converters?

`value.converter=org.apache.kafka.connect.json.JsonConverter`

`internal.value.converter=org.apache.kafka.connect.json.JsonConverter`

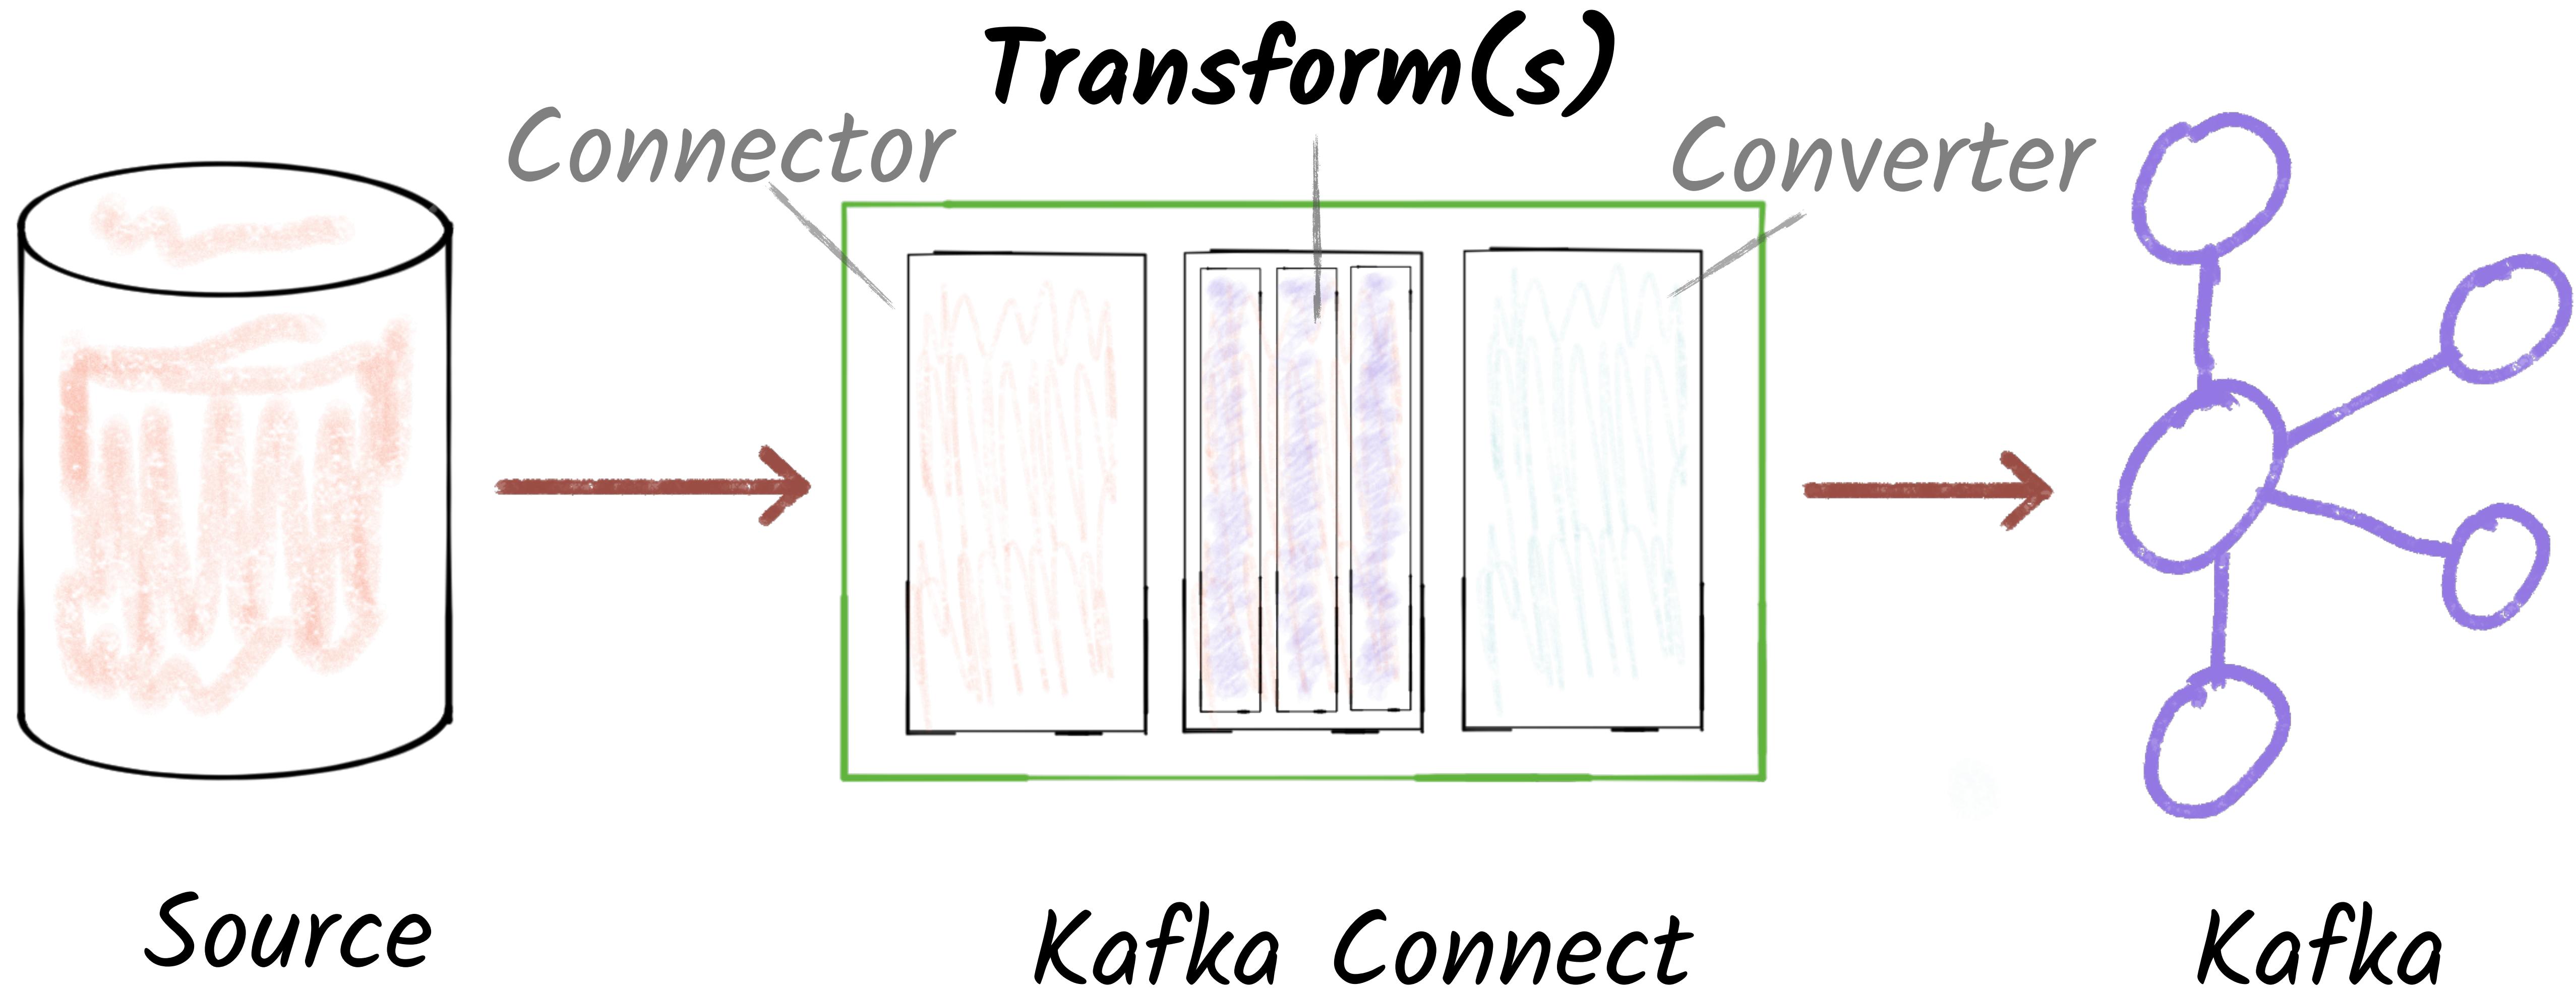
`key.internal.value.converter=org.apache.kafka.connect.json.JsonConverter`

`value.internal.value.converter=org.apache.kafka.connect.json.JsonConverter`

`key.internal.value.converter.bork.bork.bork=org.apache.kafka.connect.json.JsonConverter`

`key.internal.value.please.just.work.converter=org.apache.kafka.connect.json.JsonConverter`

# Single Message Transforms



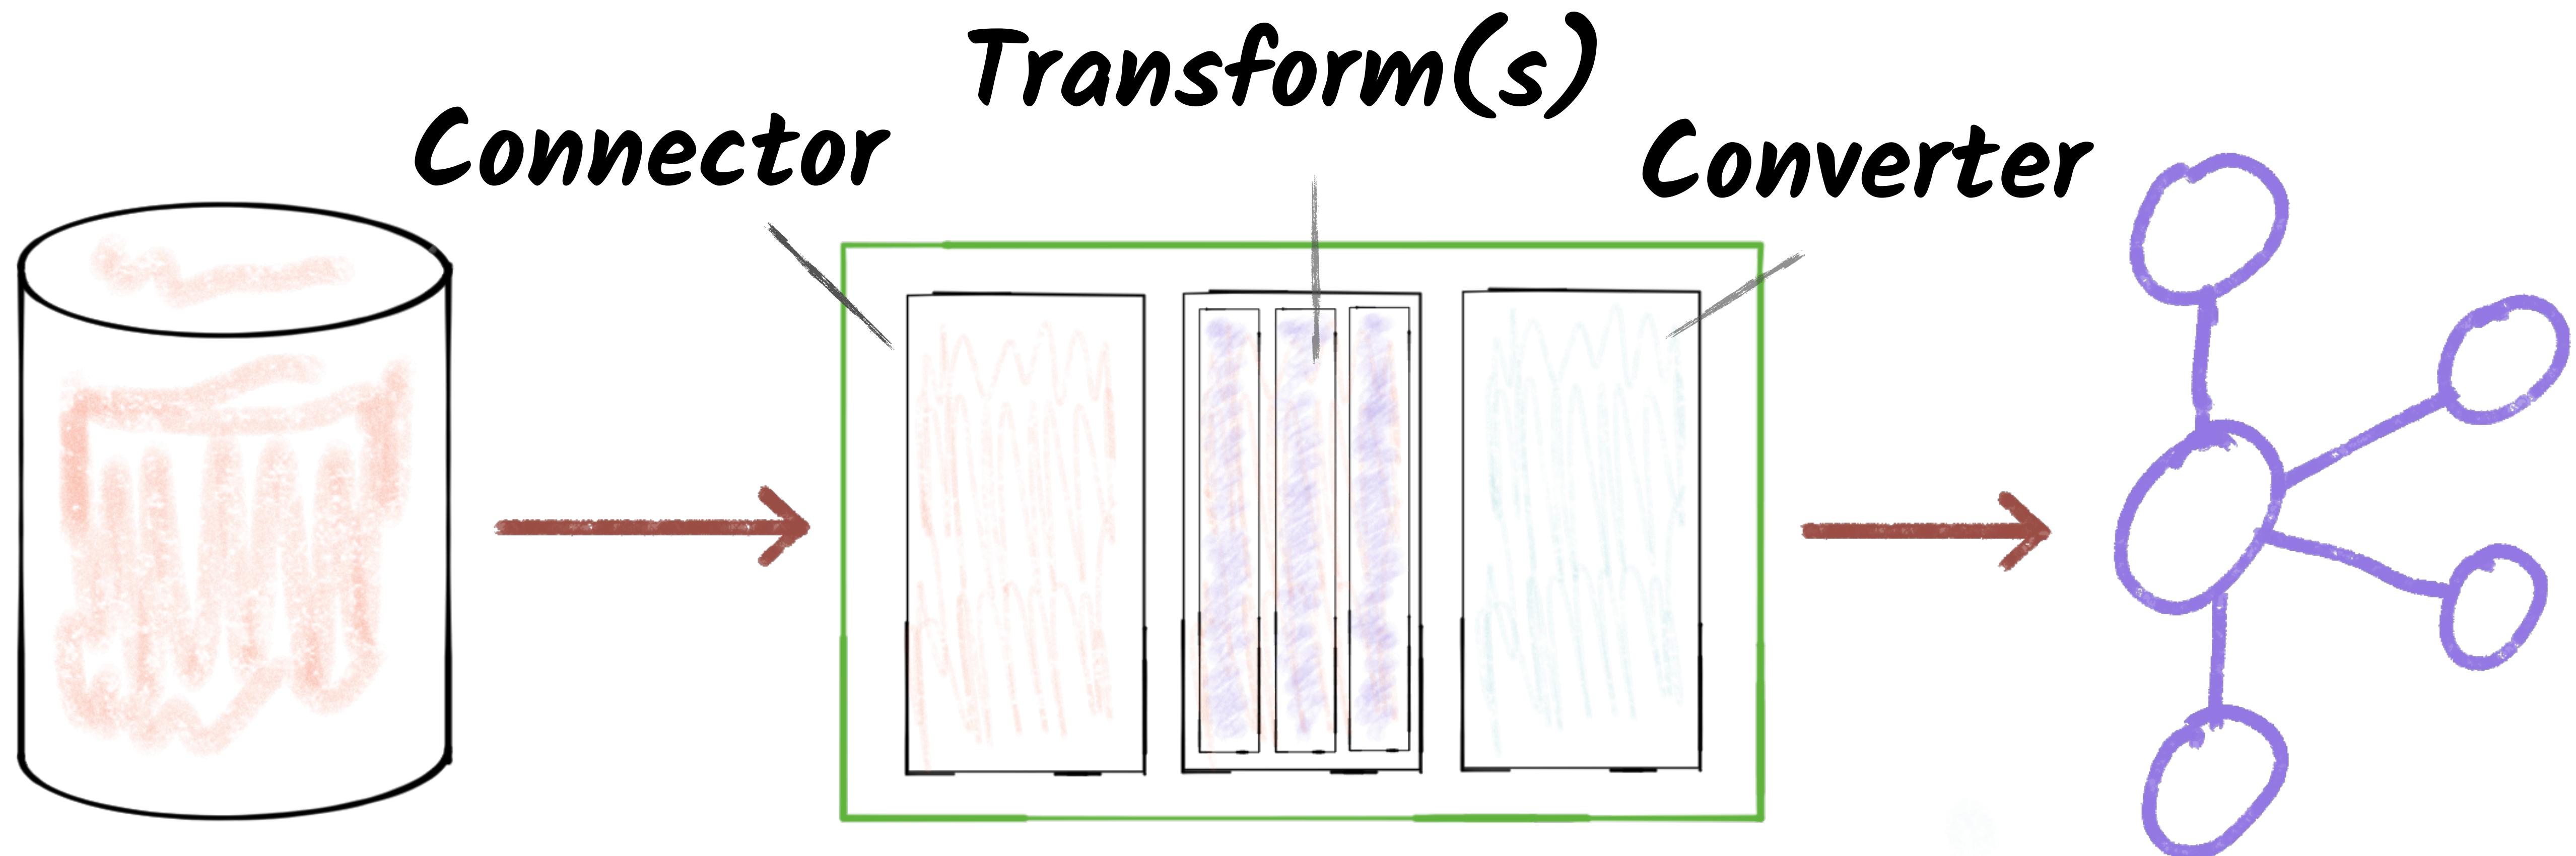
# Single Message Transforms

```
"config": {  
    [...]  
    "transforms": "addDateToTopic, labelFooBar",  
    "transforms.addDateToTopic.type": "org.apache.kafka.connect.transforms.TimestampRouter",  
    "transforms.addDateToTopic.topic.format": "${topic}-${timestamp}",  
    "transforms.addDateToTopic.timestamp.format": "YYYYMM",  
    "transforms.labelFooBar.type": "org.apache.kafka.connect.transforms.ReplaceField$Value",  
    "transforms.labelFooBar.renames": "delivery_address:shipping_address",  
}
```

*Do these transforms*

*Transforms config*    *Config per transform*

# Extensible



# Confluent Hub

CONFLUENT HUB

## Discover and share Connectors and more

Search Connectors 

All Verified Sources Sinks Community

Confluent Supported



**Debezium MongoDB CDC Connector**  
Debezium Community

[Read More](#)

Confluent Supported



**Debezium MySQL CDC Connector**  
Debezium Community

[Read More](#)

Confluent Supported



**Debezium PostgreSQL CDC Connector**  
Debezium Community

[Read More](#)

Confluent Supported



**Debezium SQL Server CDC Connector**  
Debezium Community

[Read More](#)

[hub.confluent.io](https://hub.confluent.io)

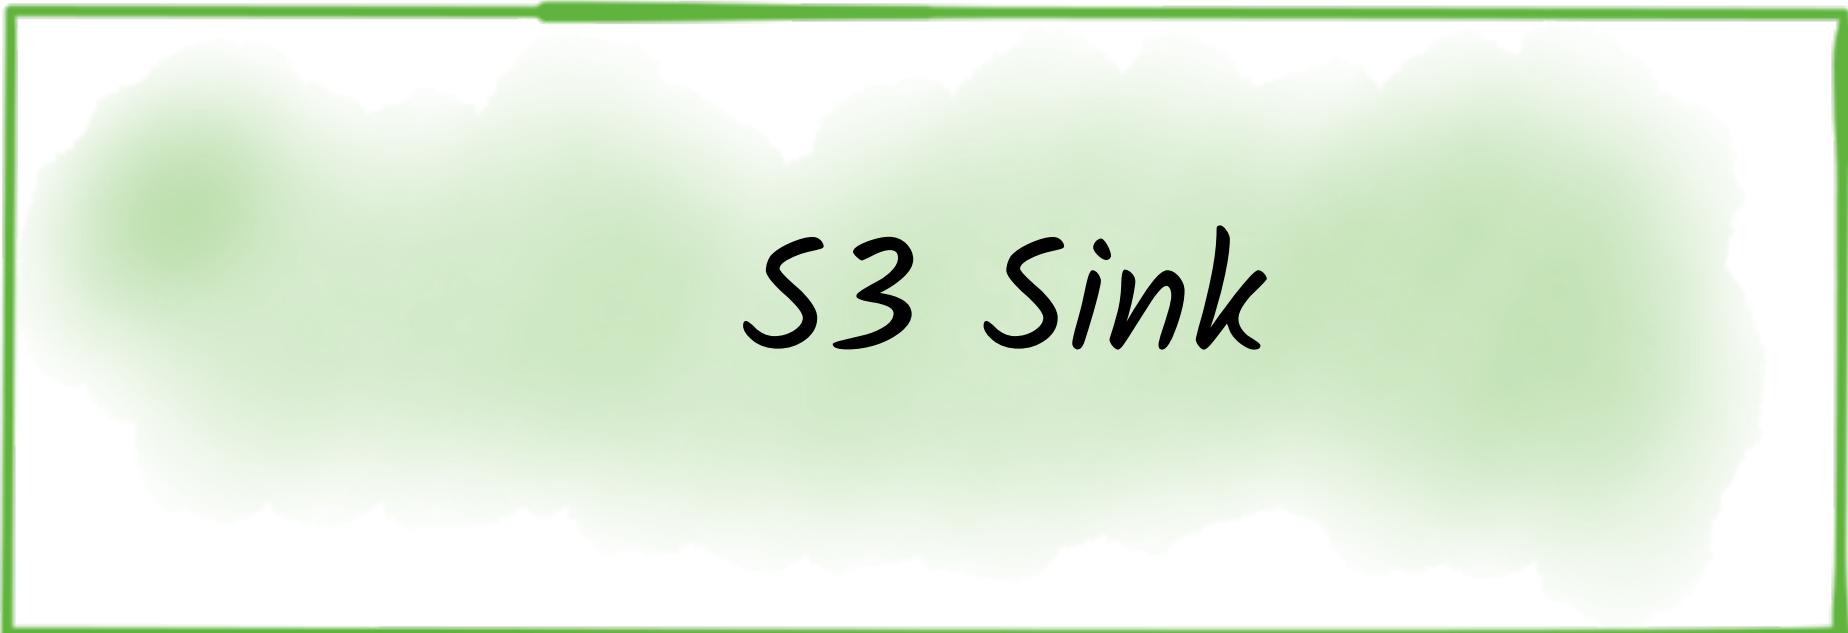
# Deploying

# Kafka

# Connect

Connectors, Tasks, and Workers

# Connectors and Tasks



# Connectors and Tasks

JDBC Source

S3 Sink

S3 Task #1

JDBC Task #1

# Connectors and Tasks

JDBC Source

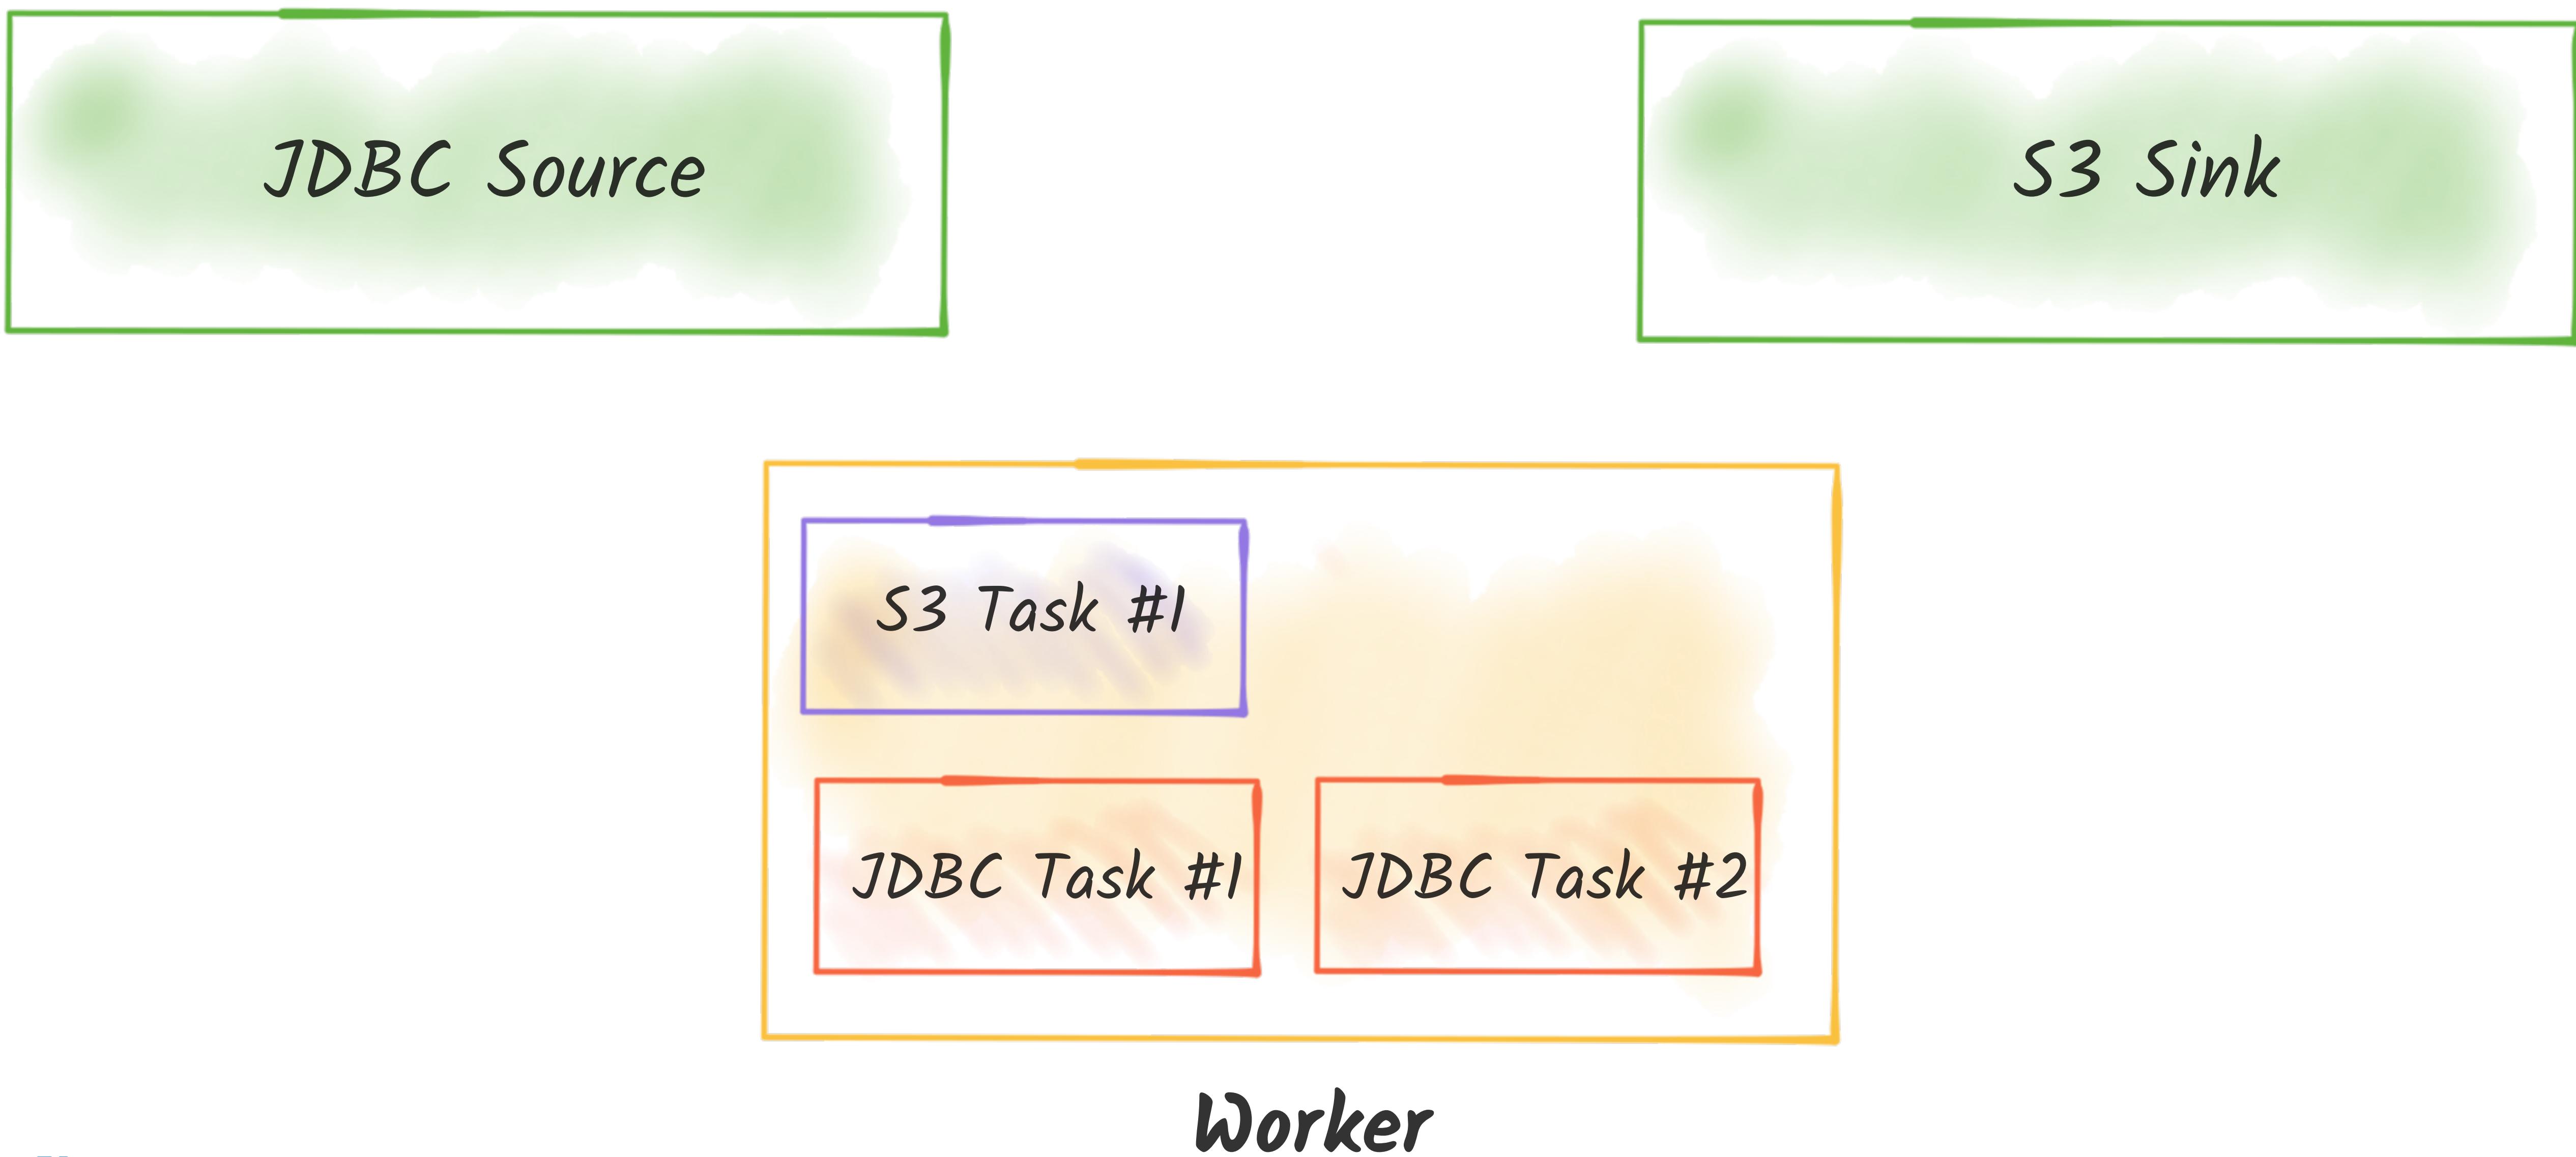
S3 Sink

S3 Task #1

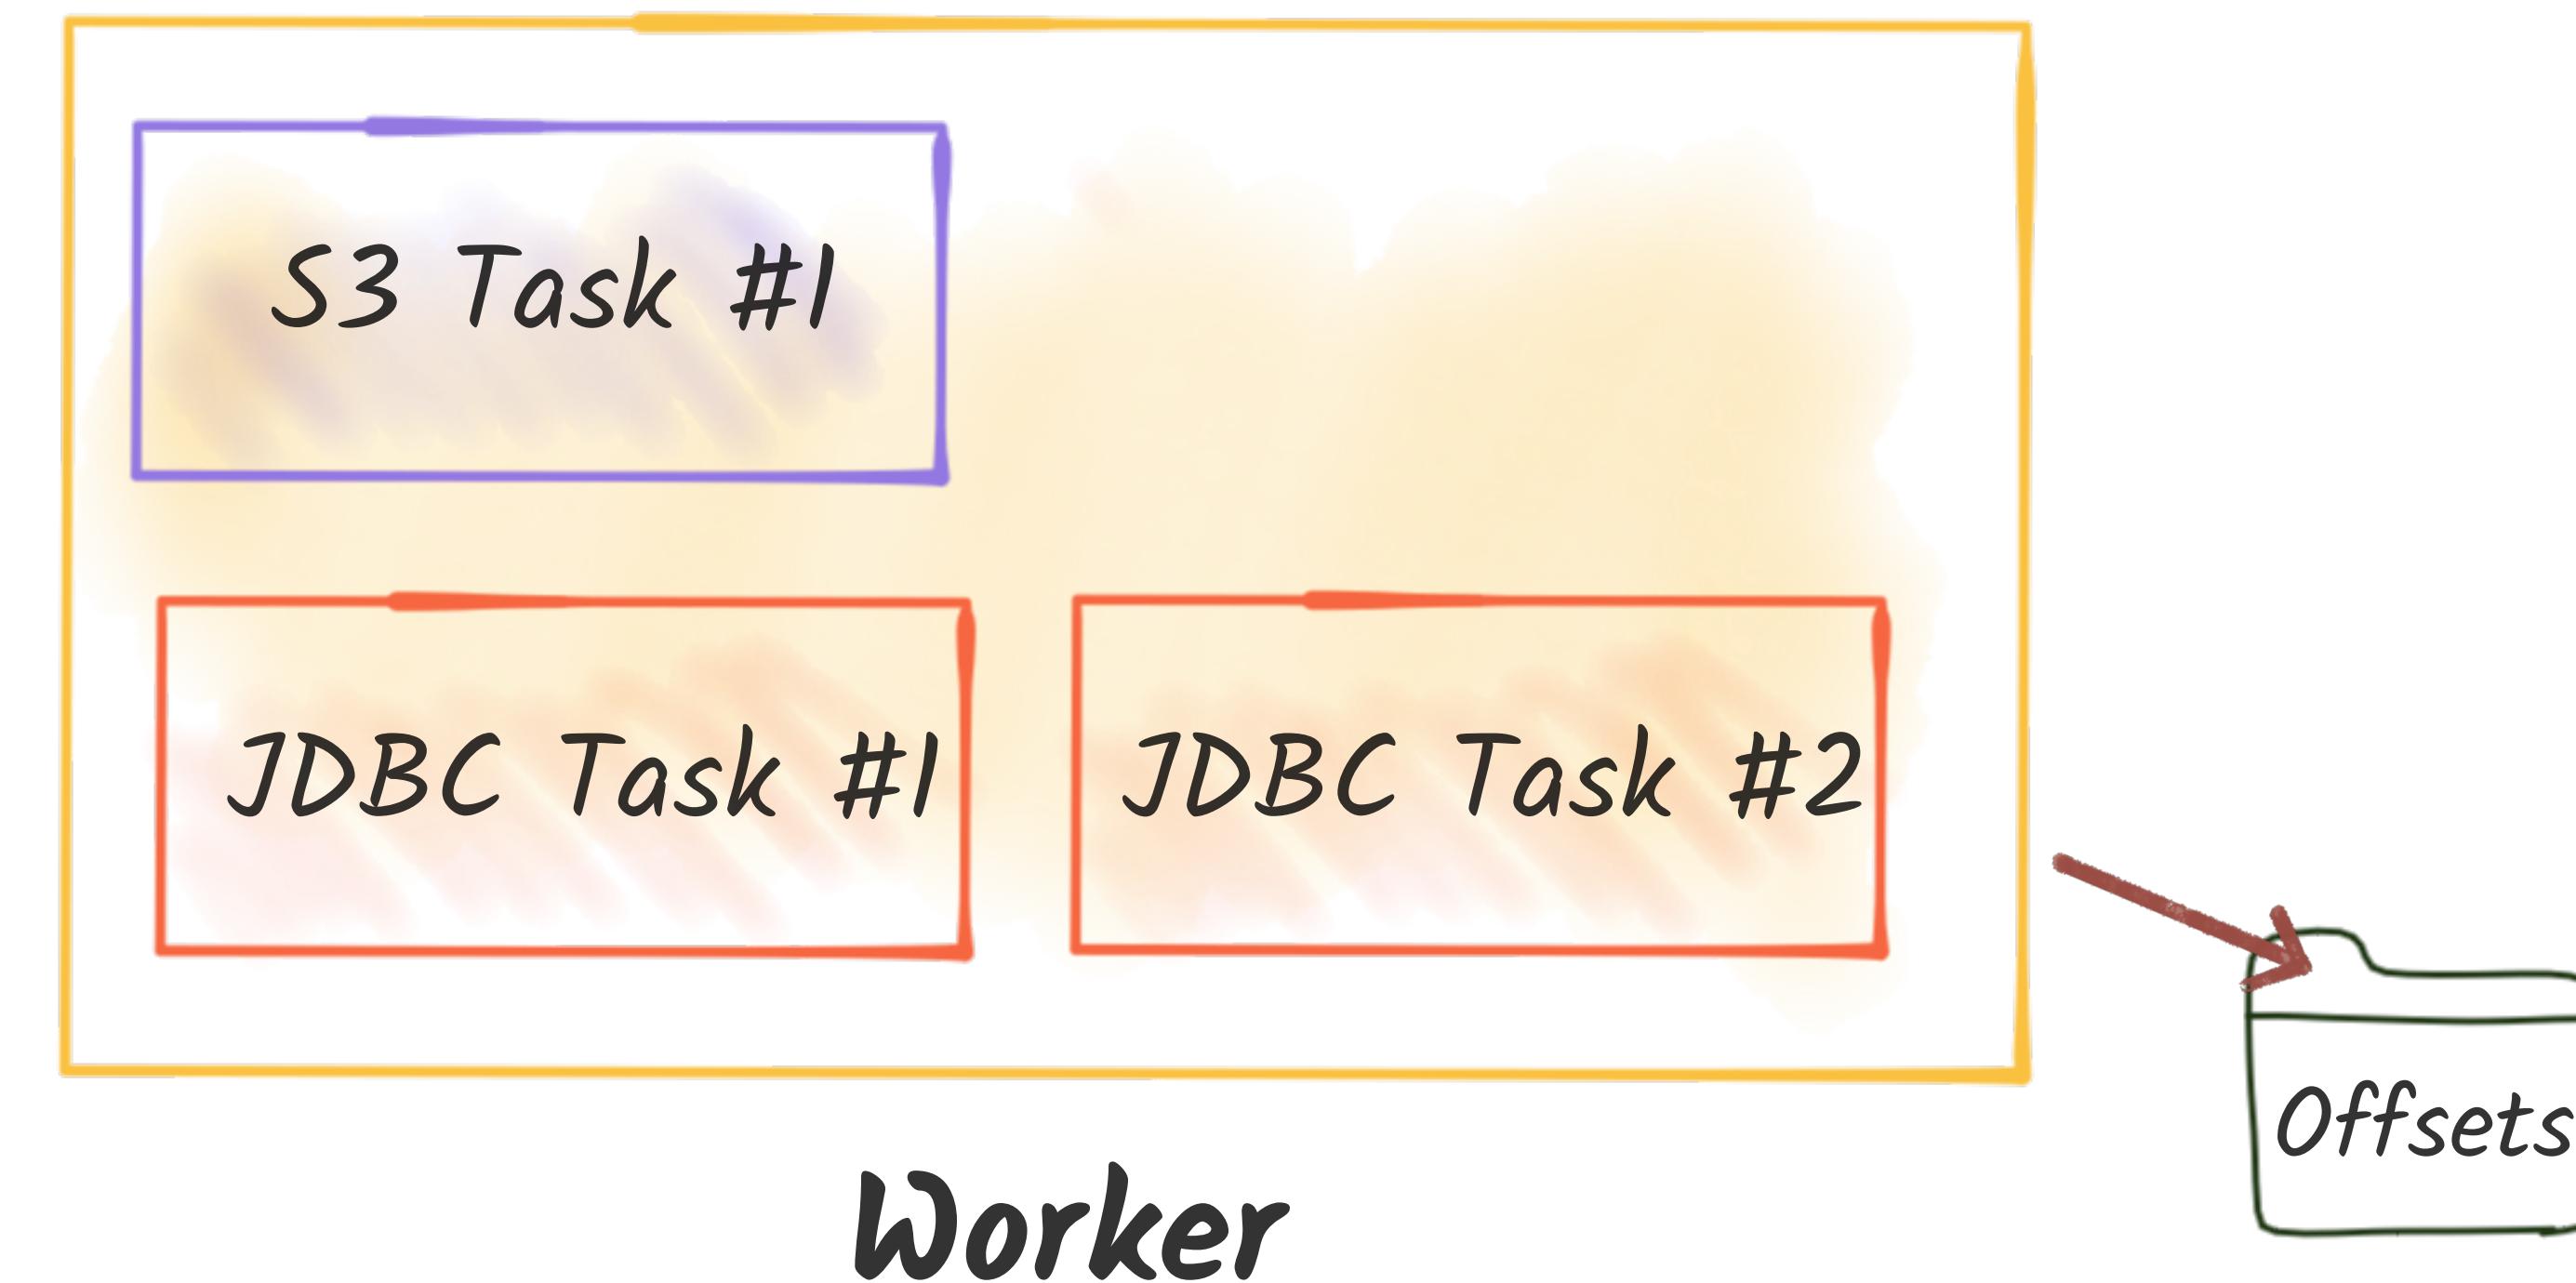
JDBC Task #1

JDBC Task #2

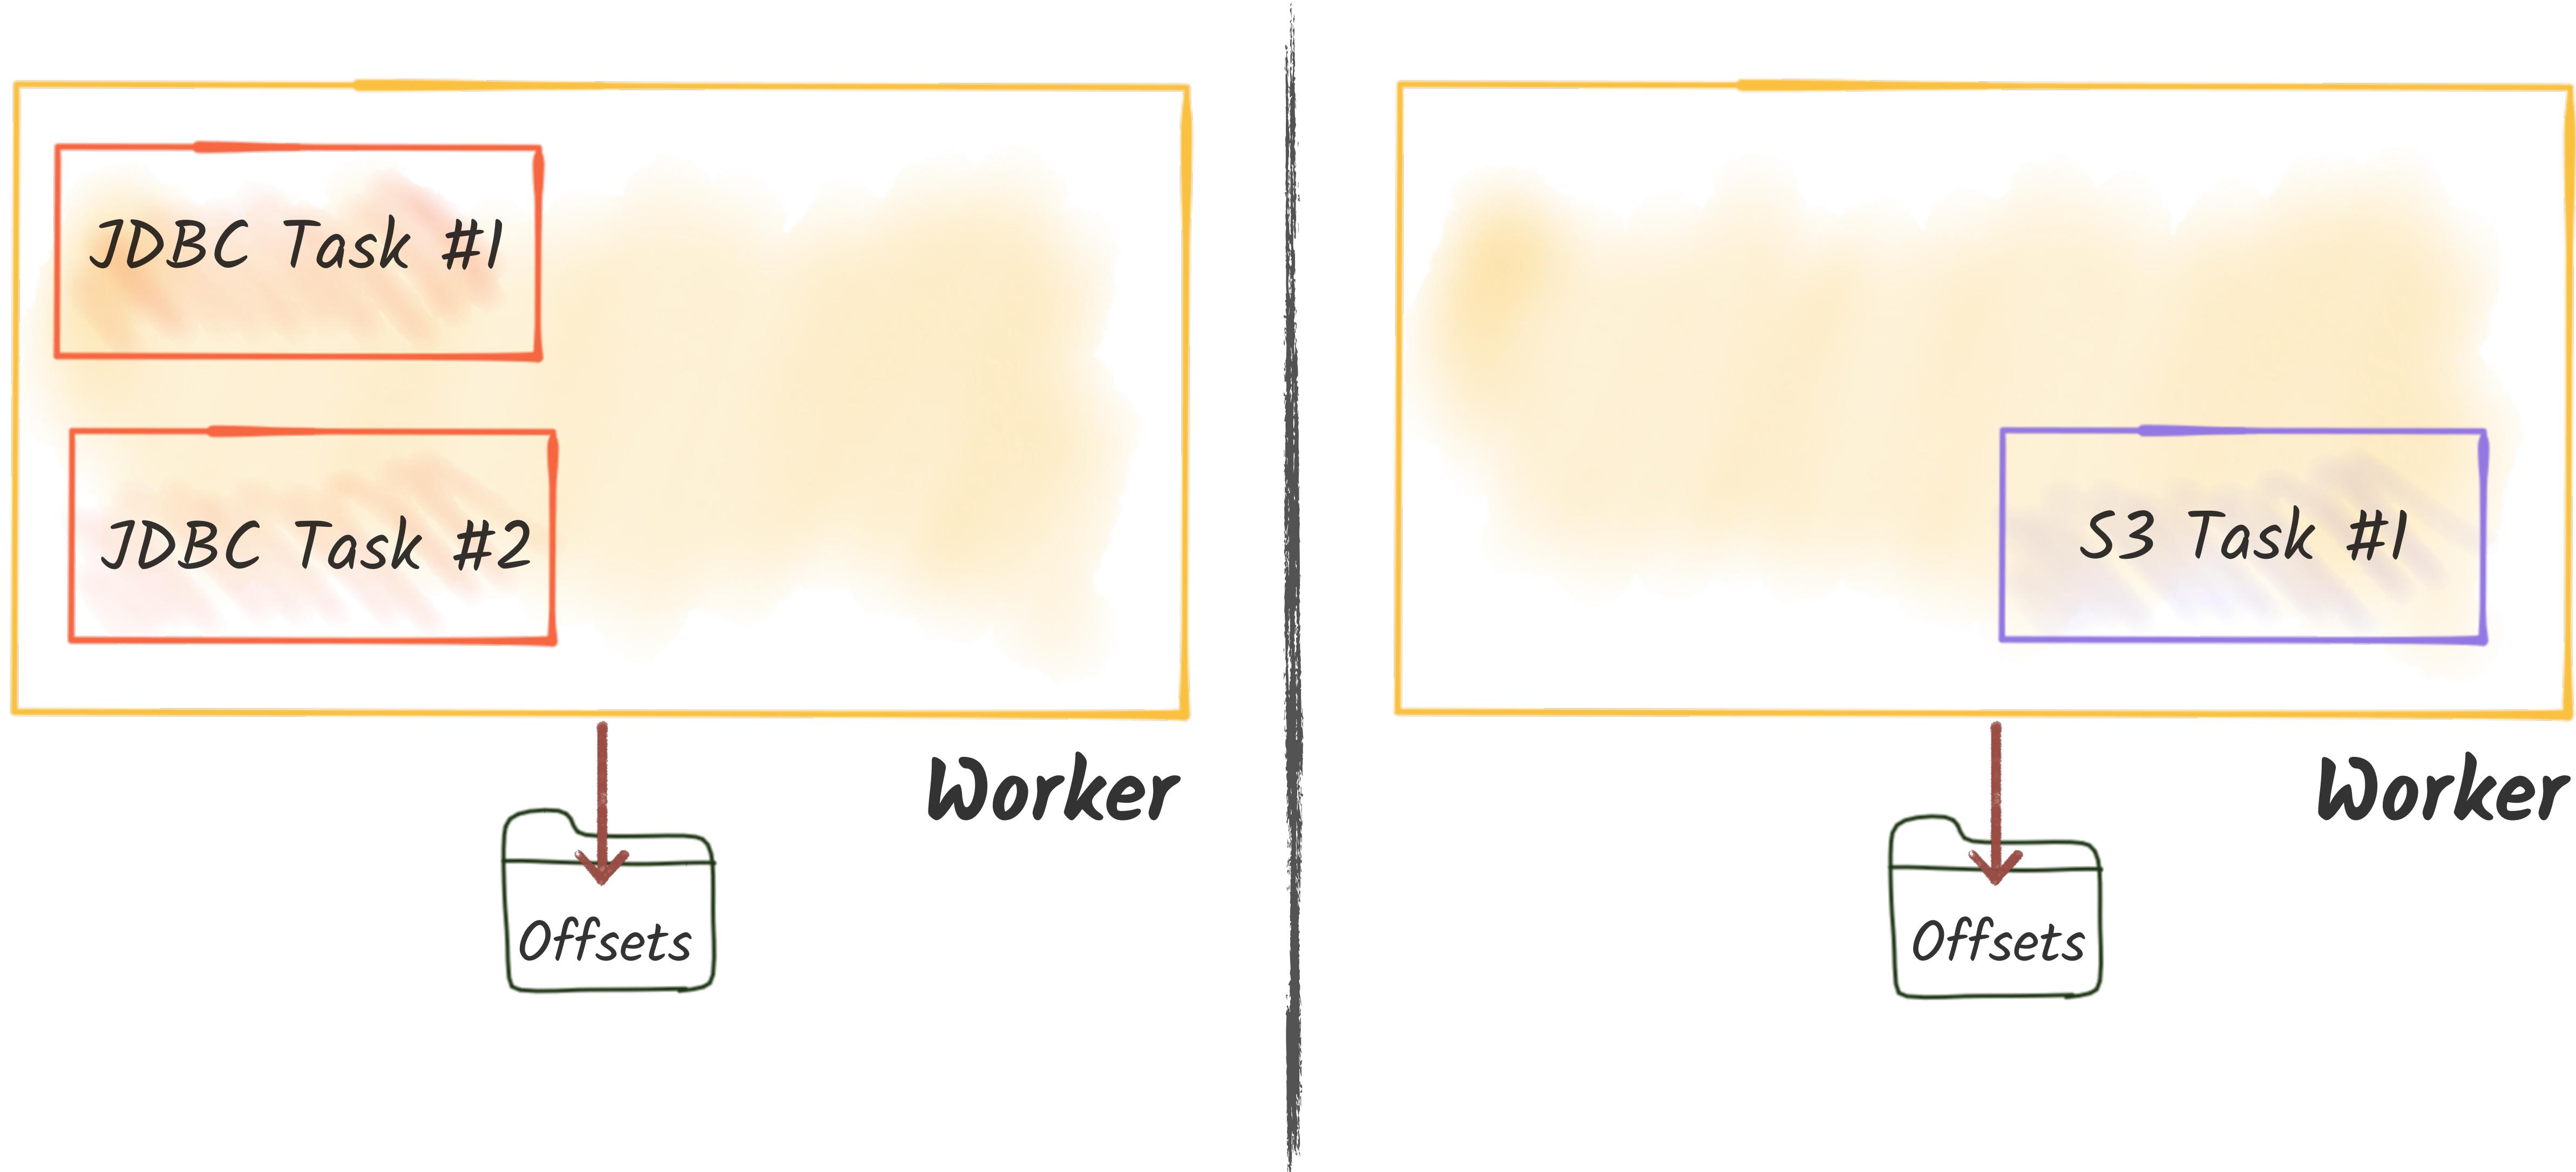
# Tasks and Workers



# Kafka Connect Standalone Worker

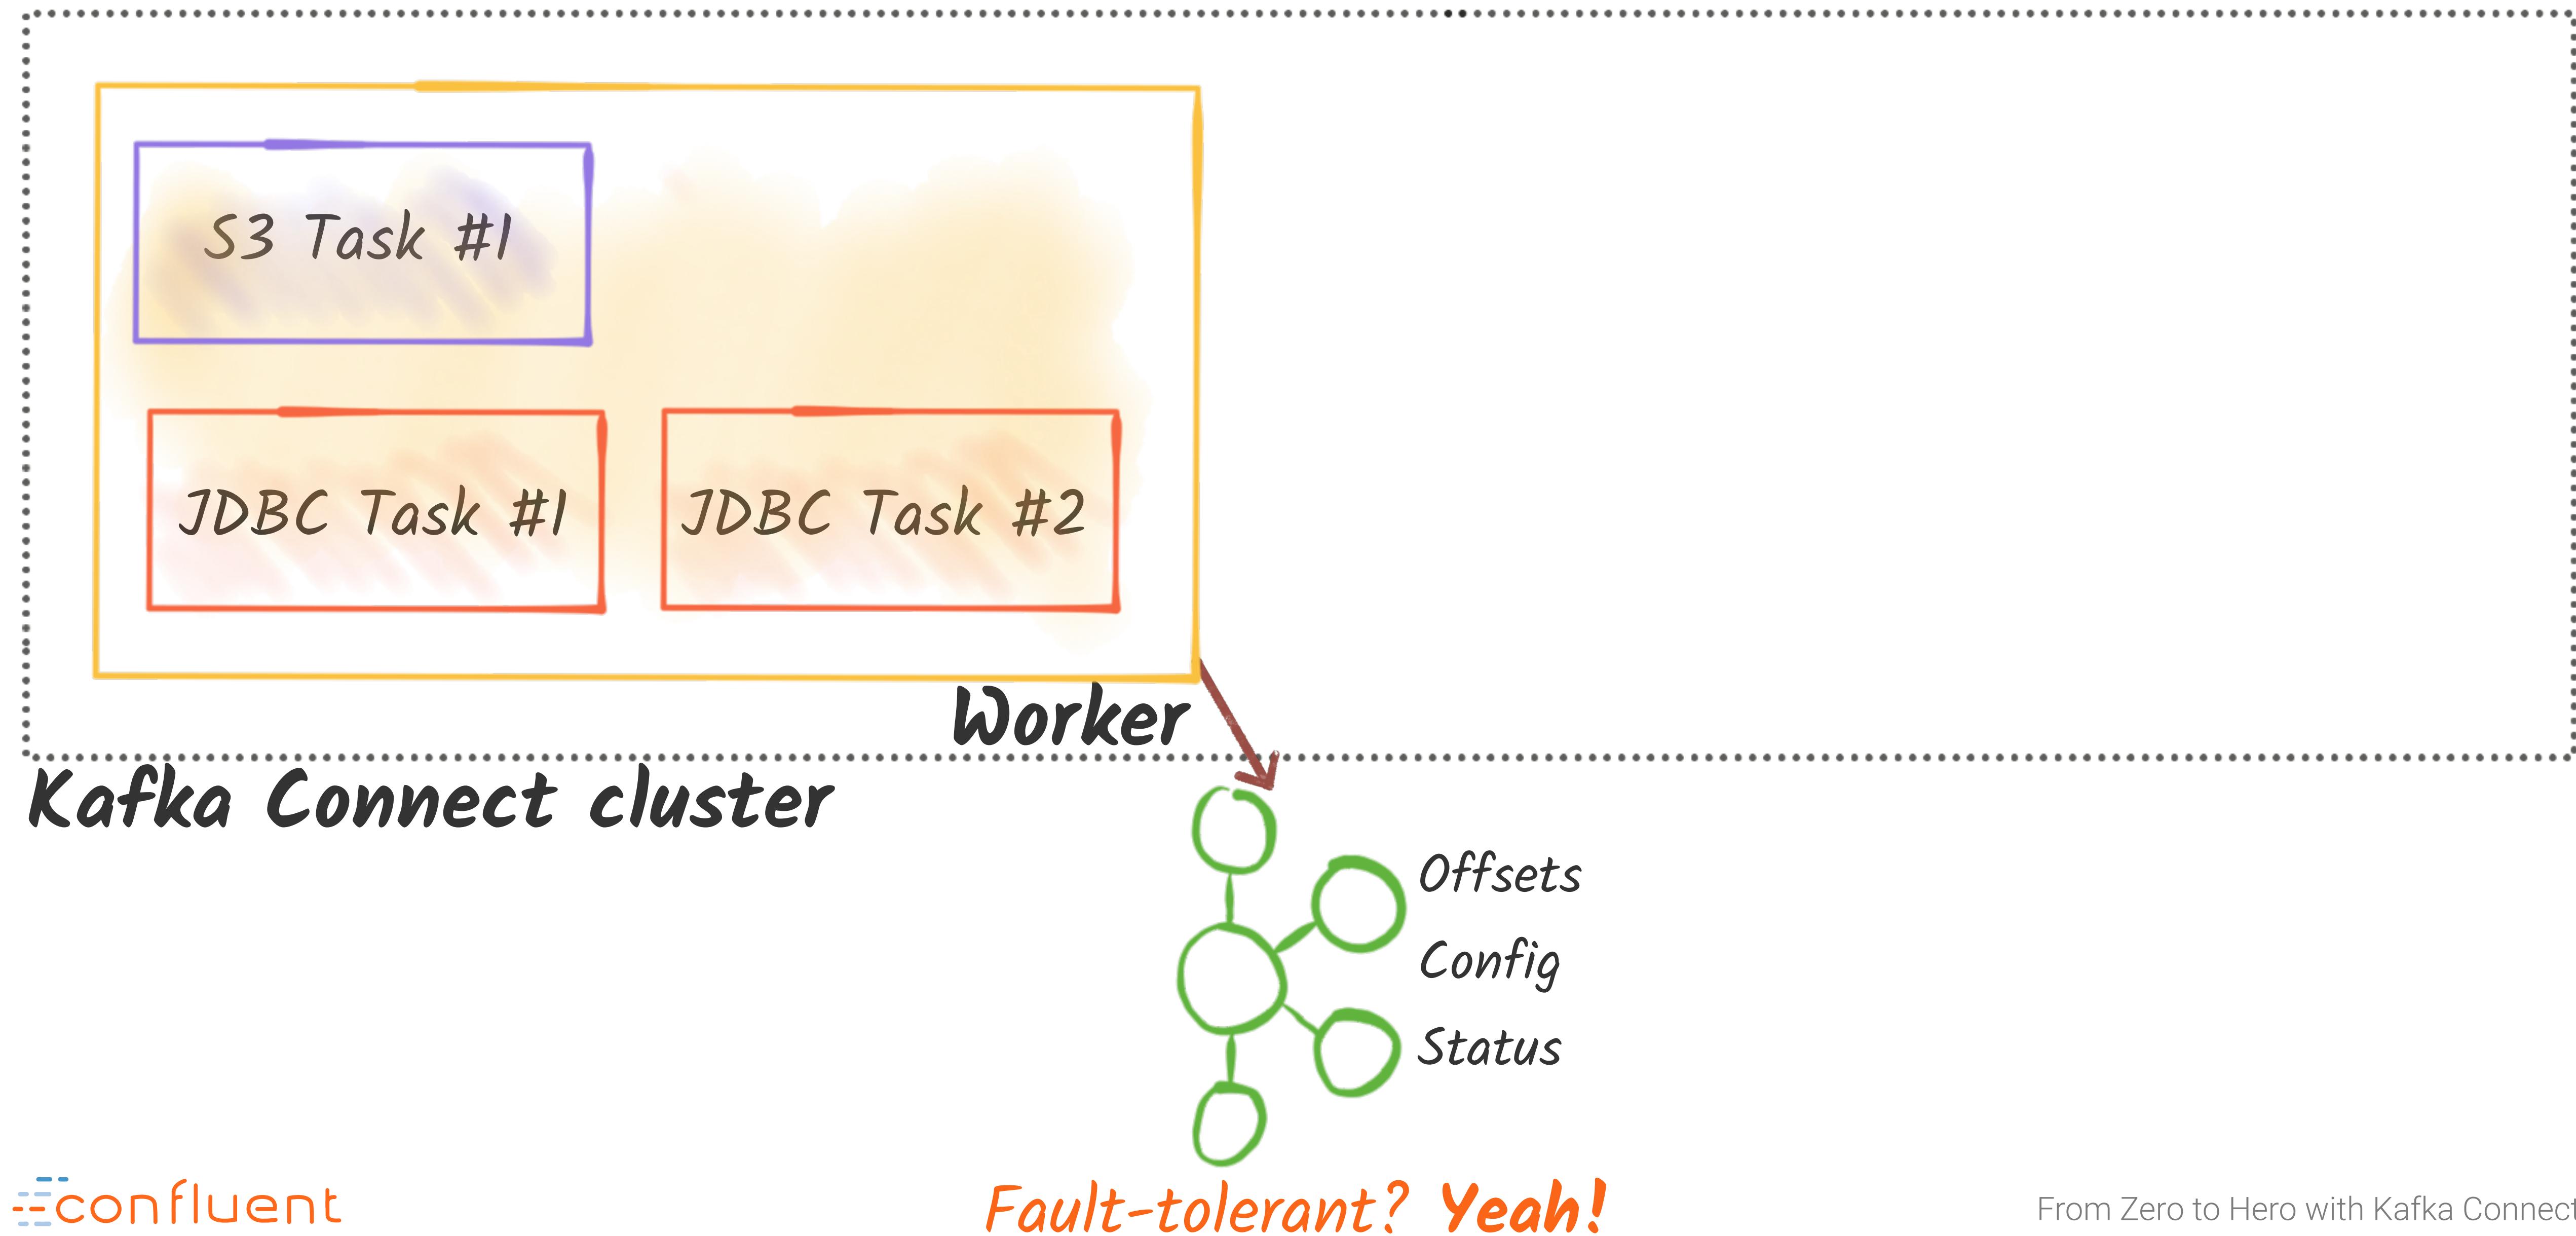


# "Scaling" the Standalone Worker

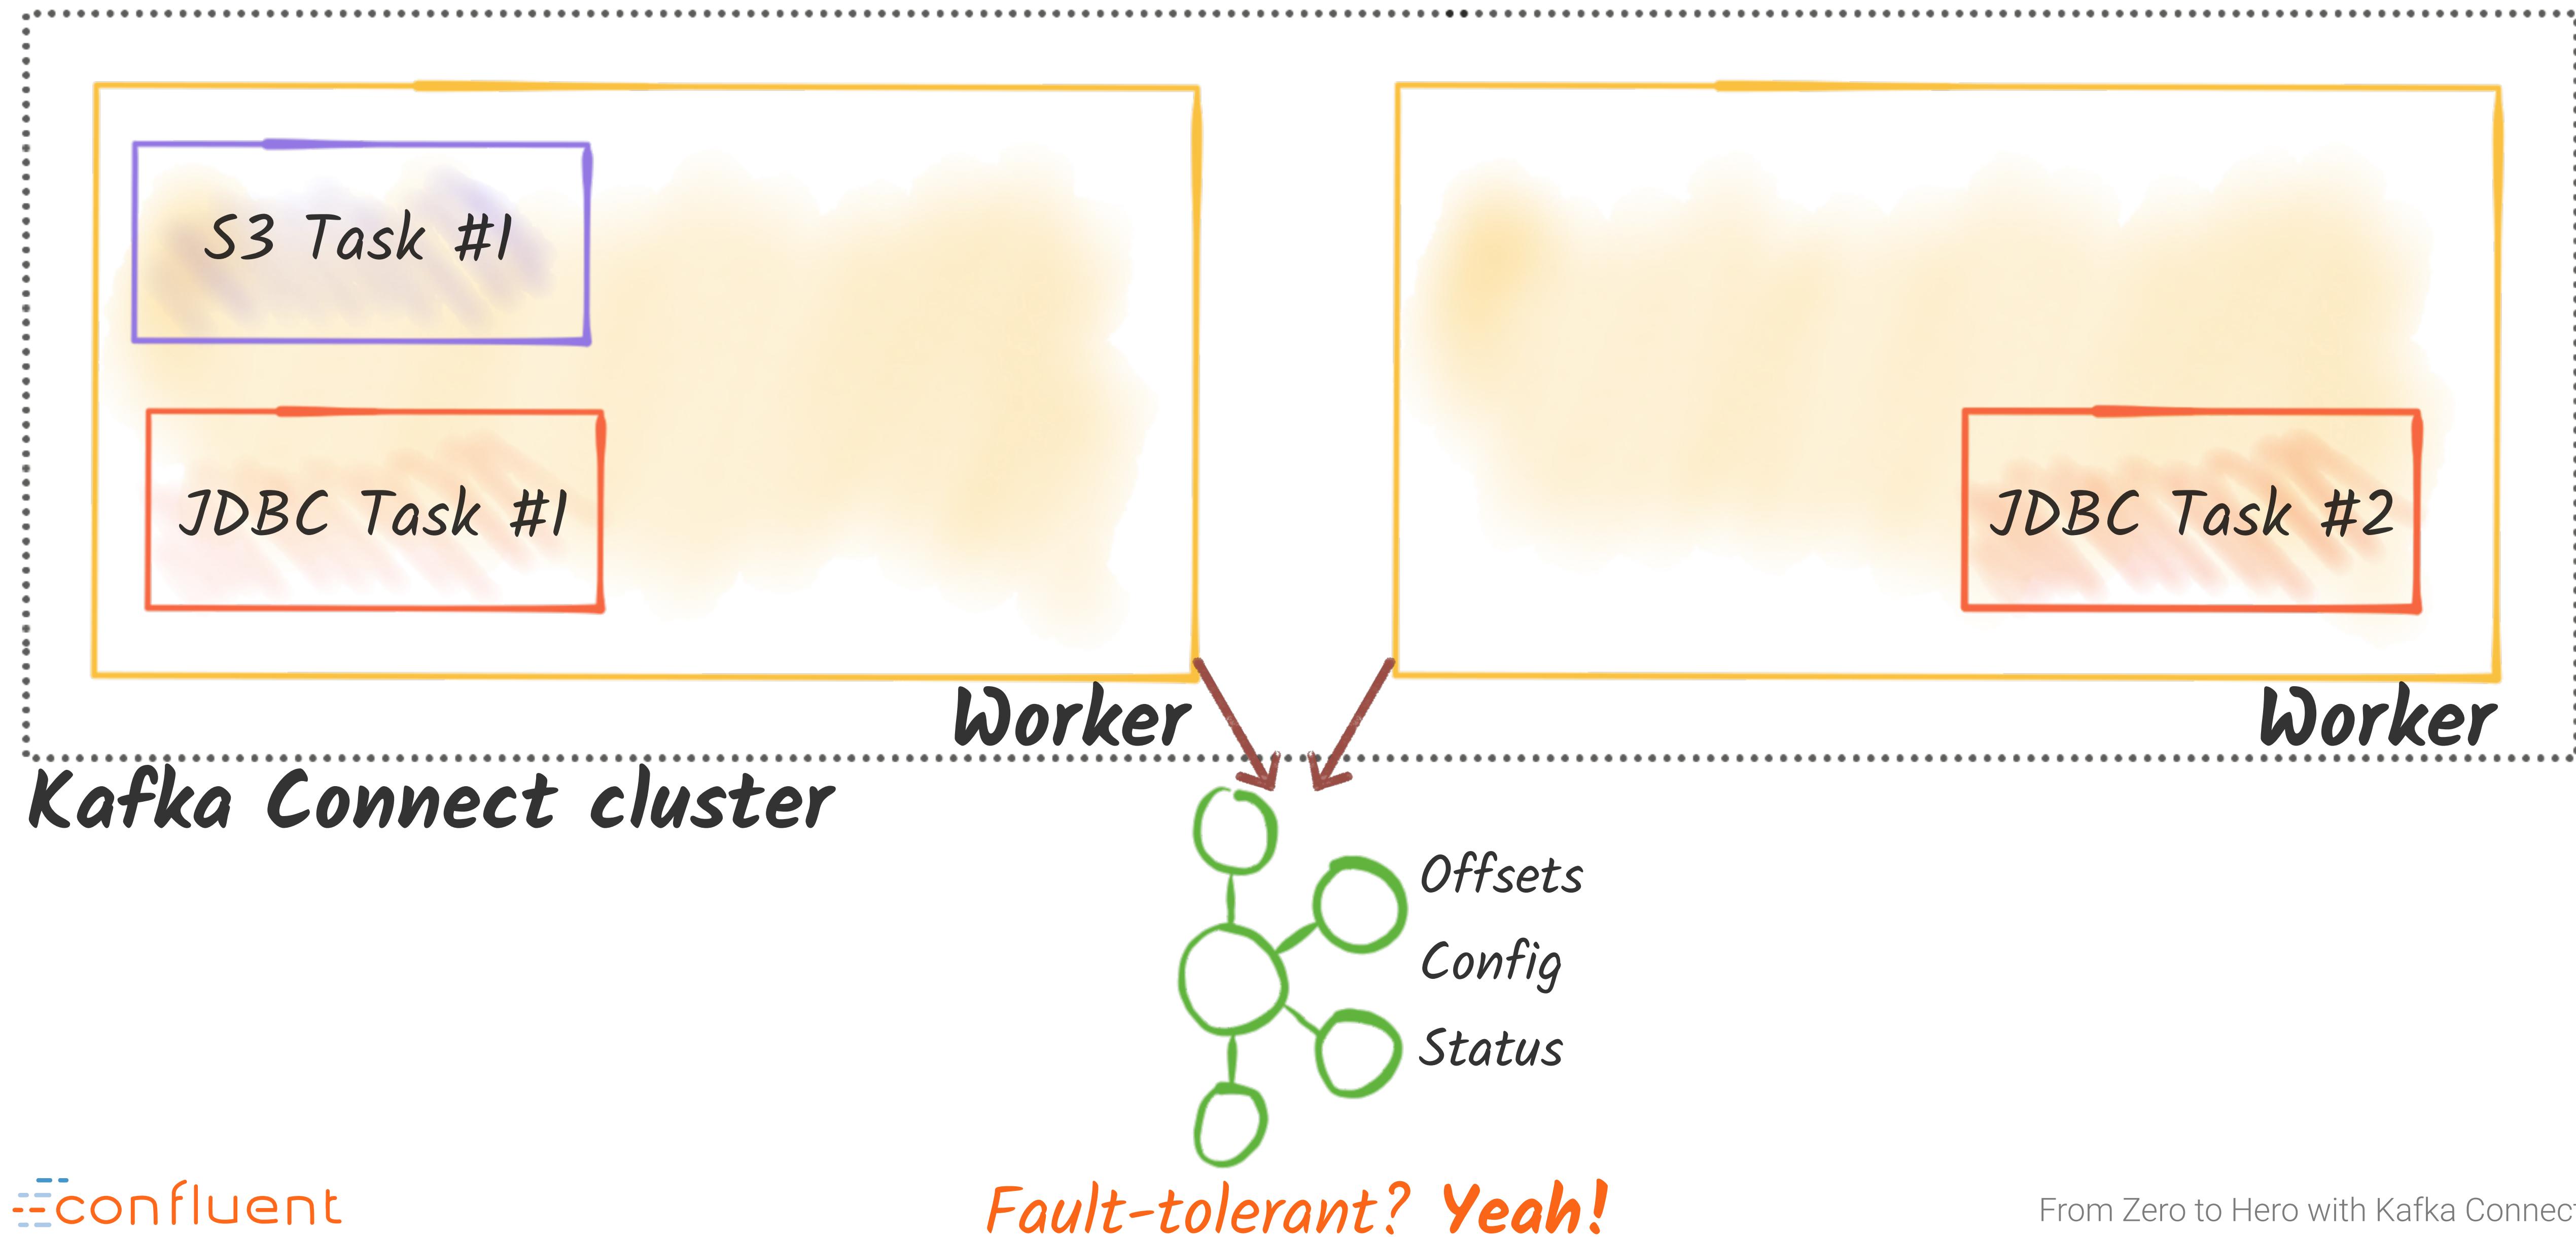


Fault-tolerant? **Nope.**

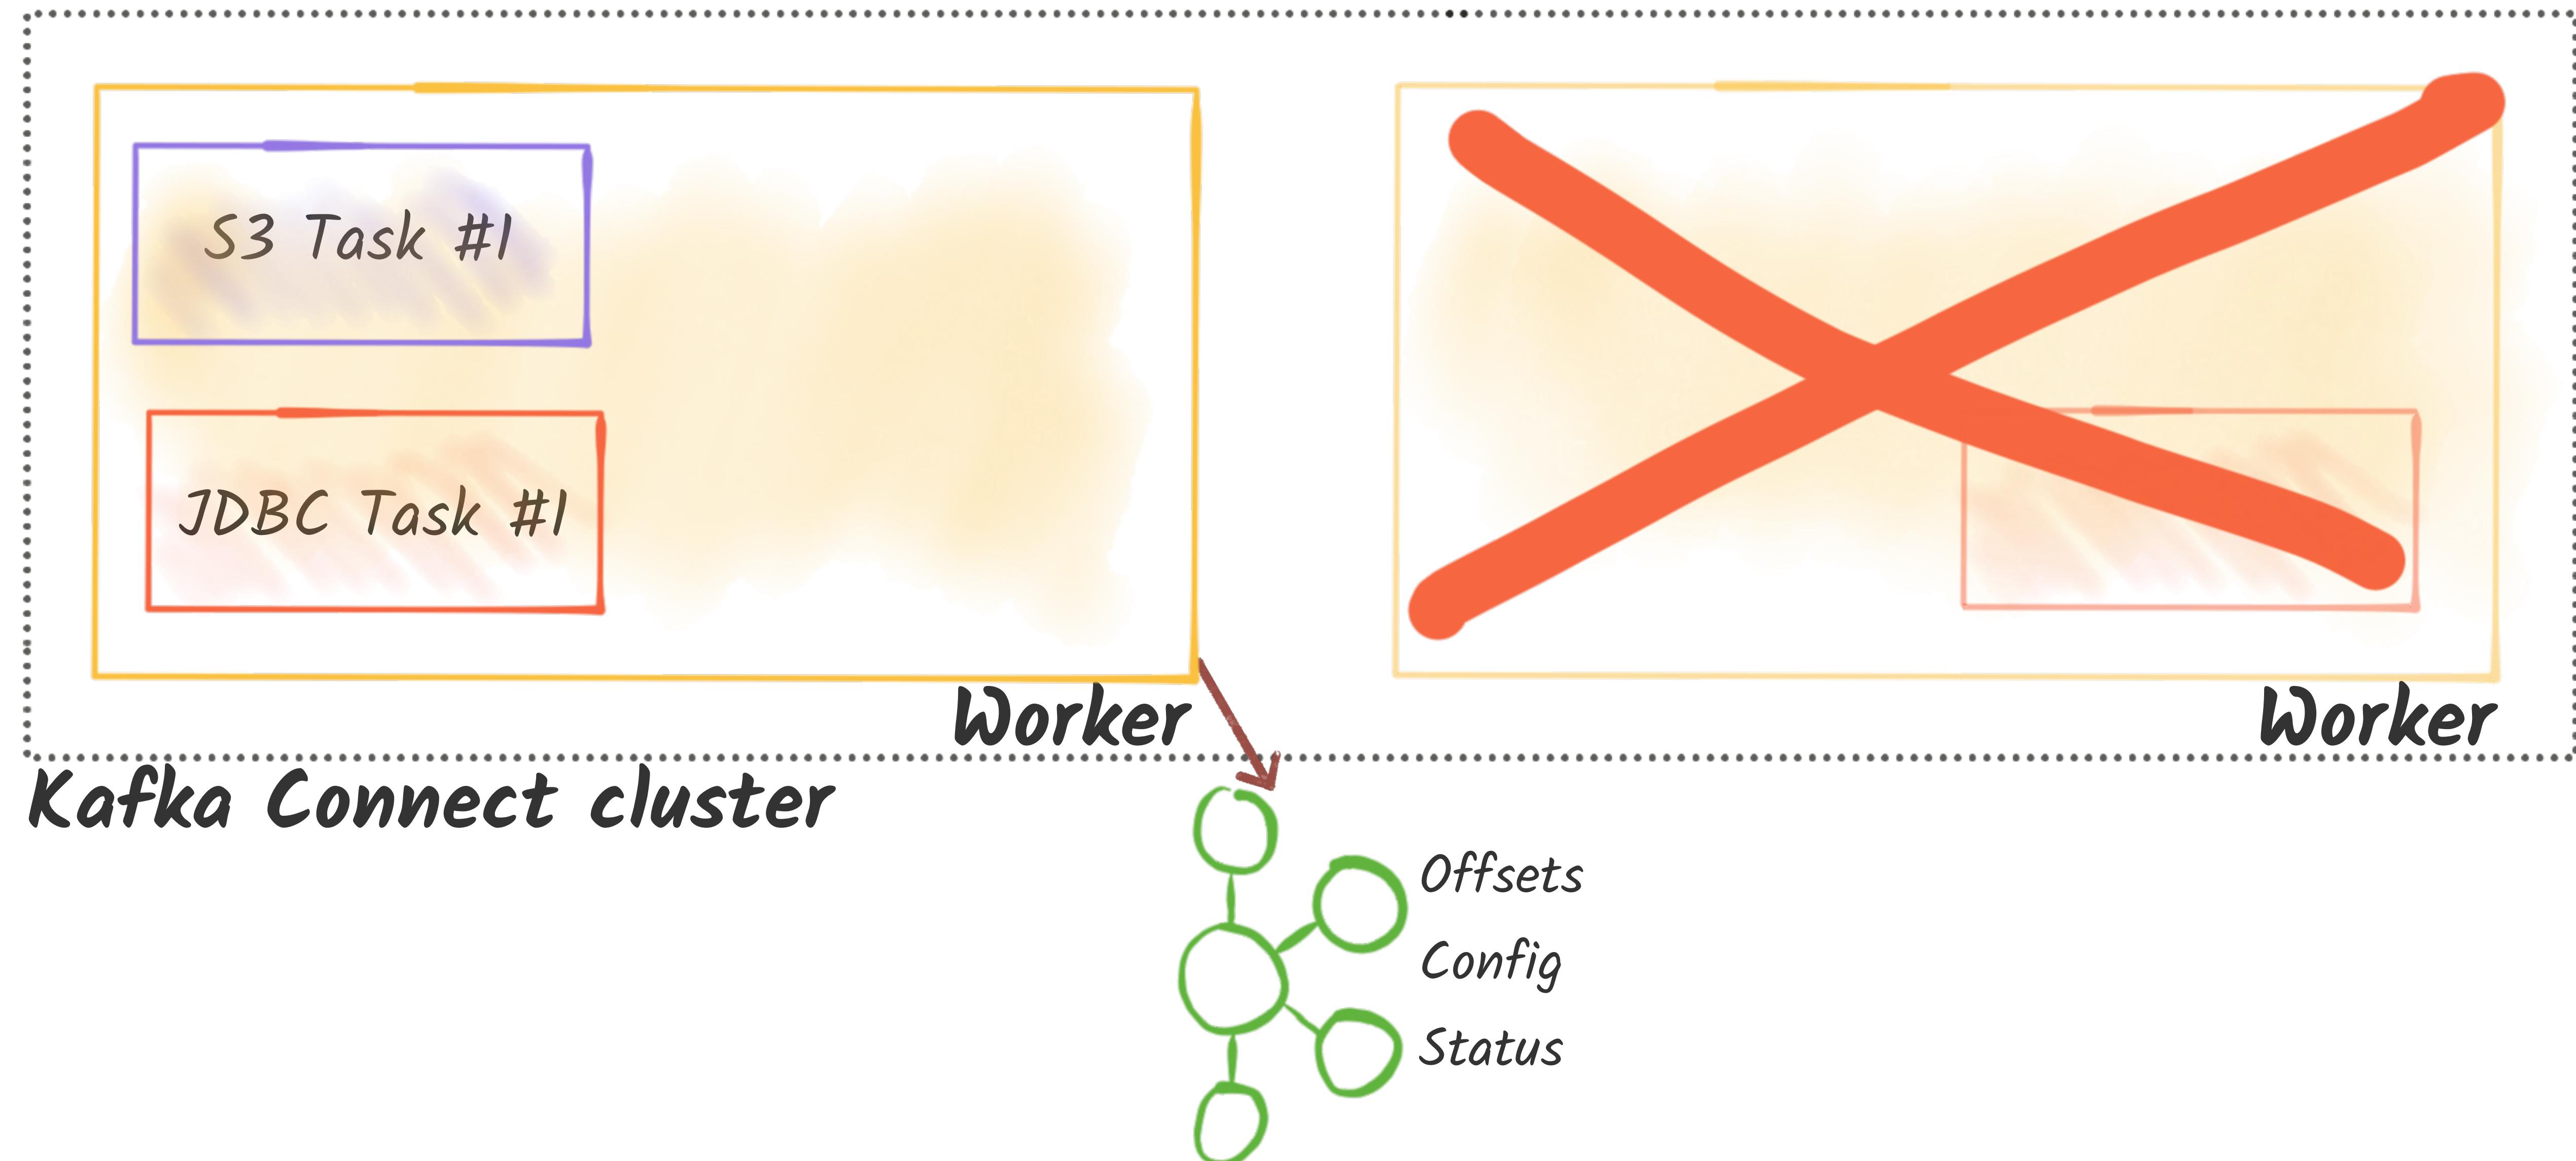
# Kafka Connect Distributed Worker



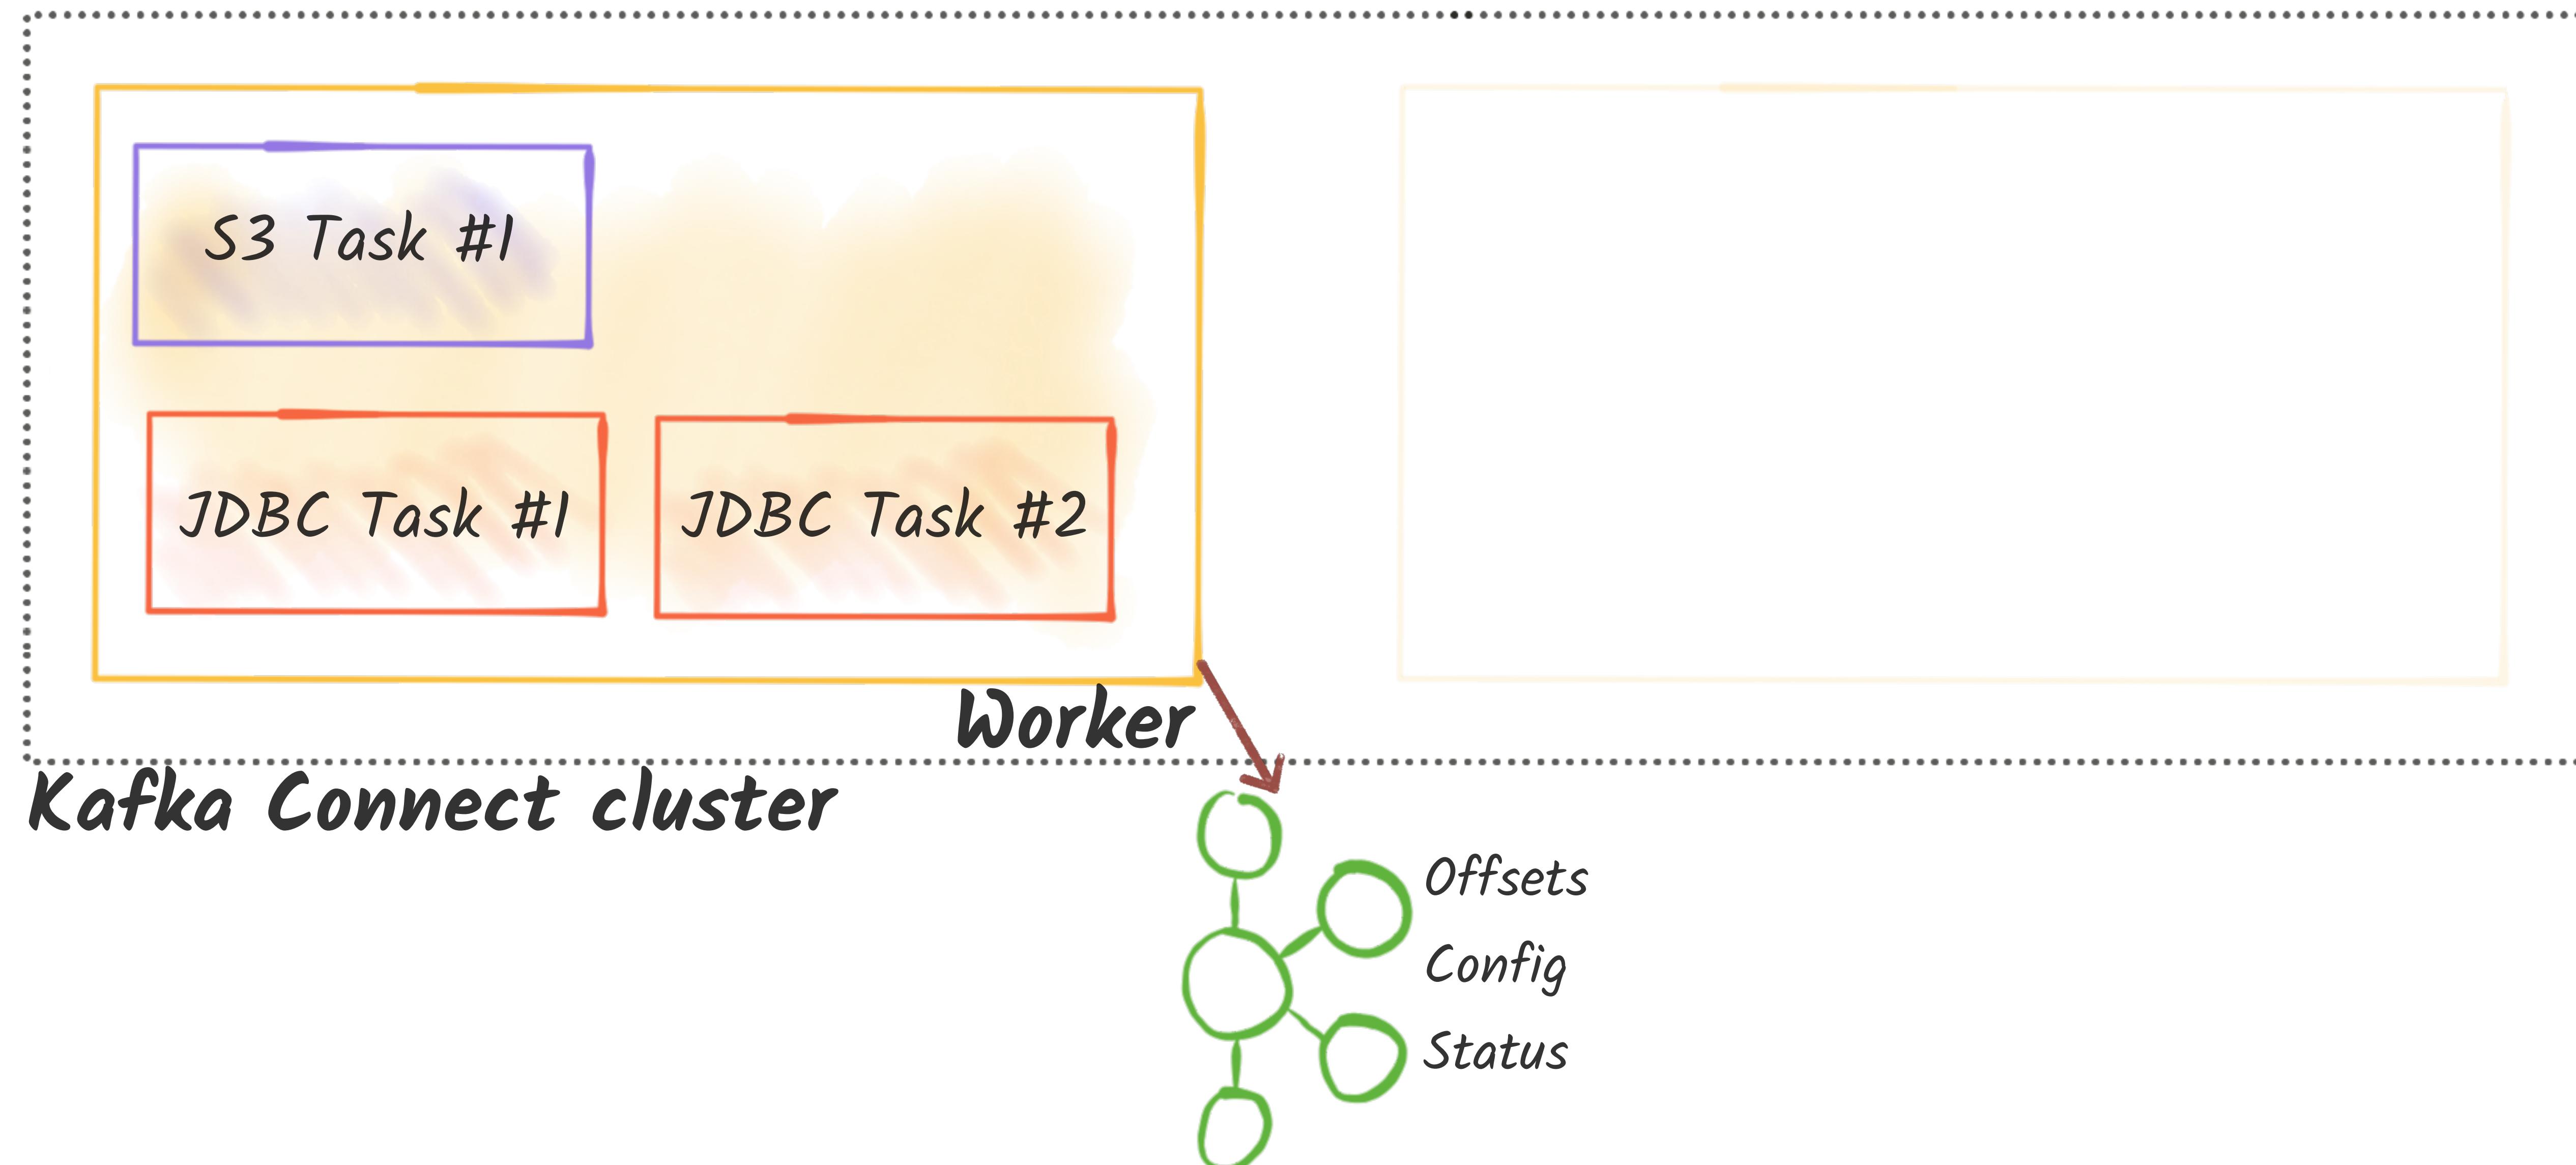
# Scaling the Distributed Worker



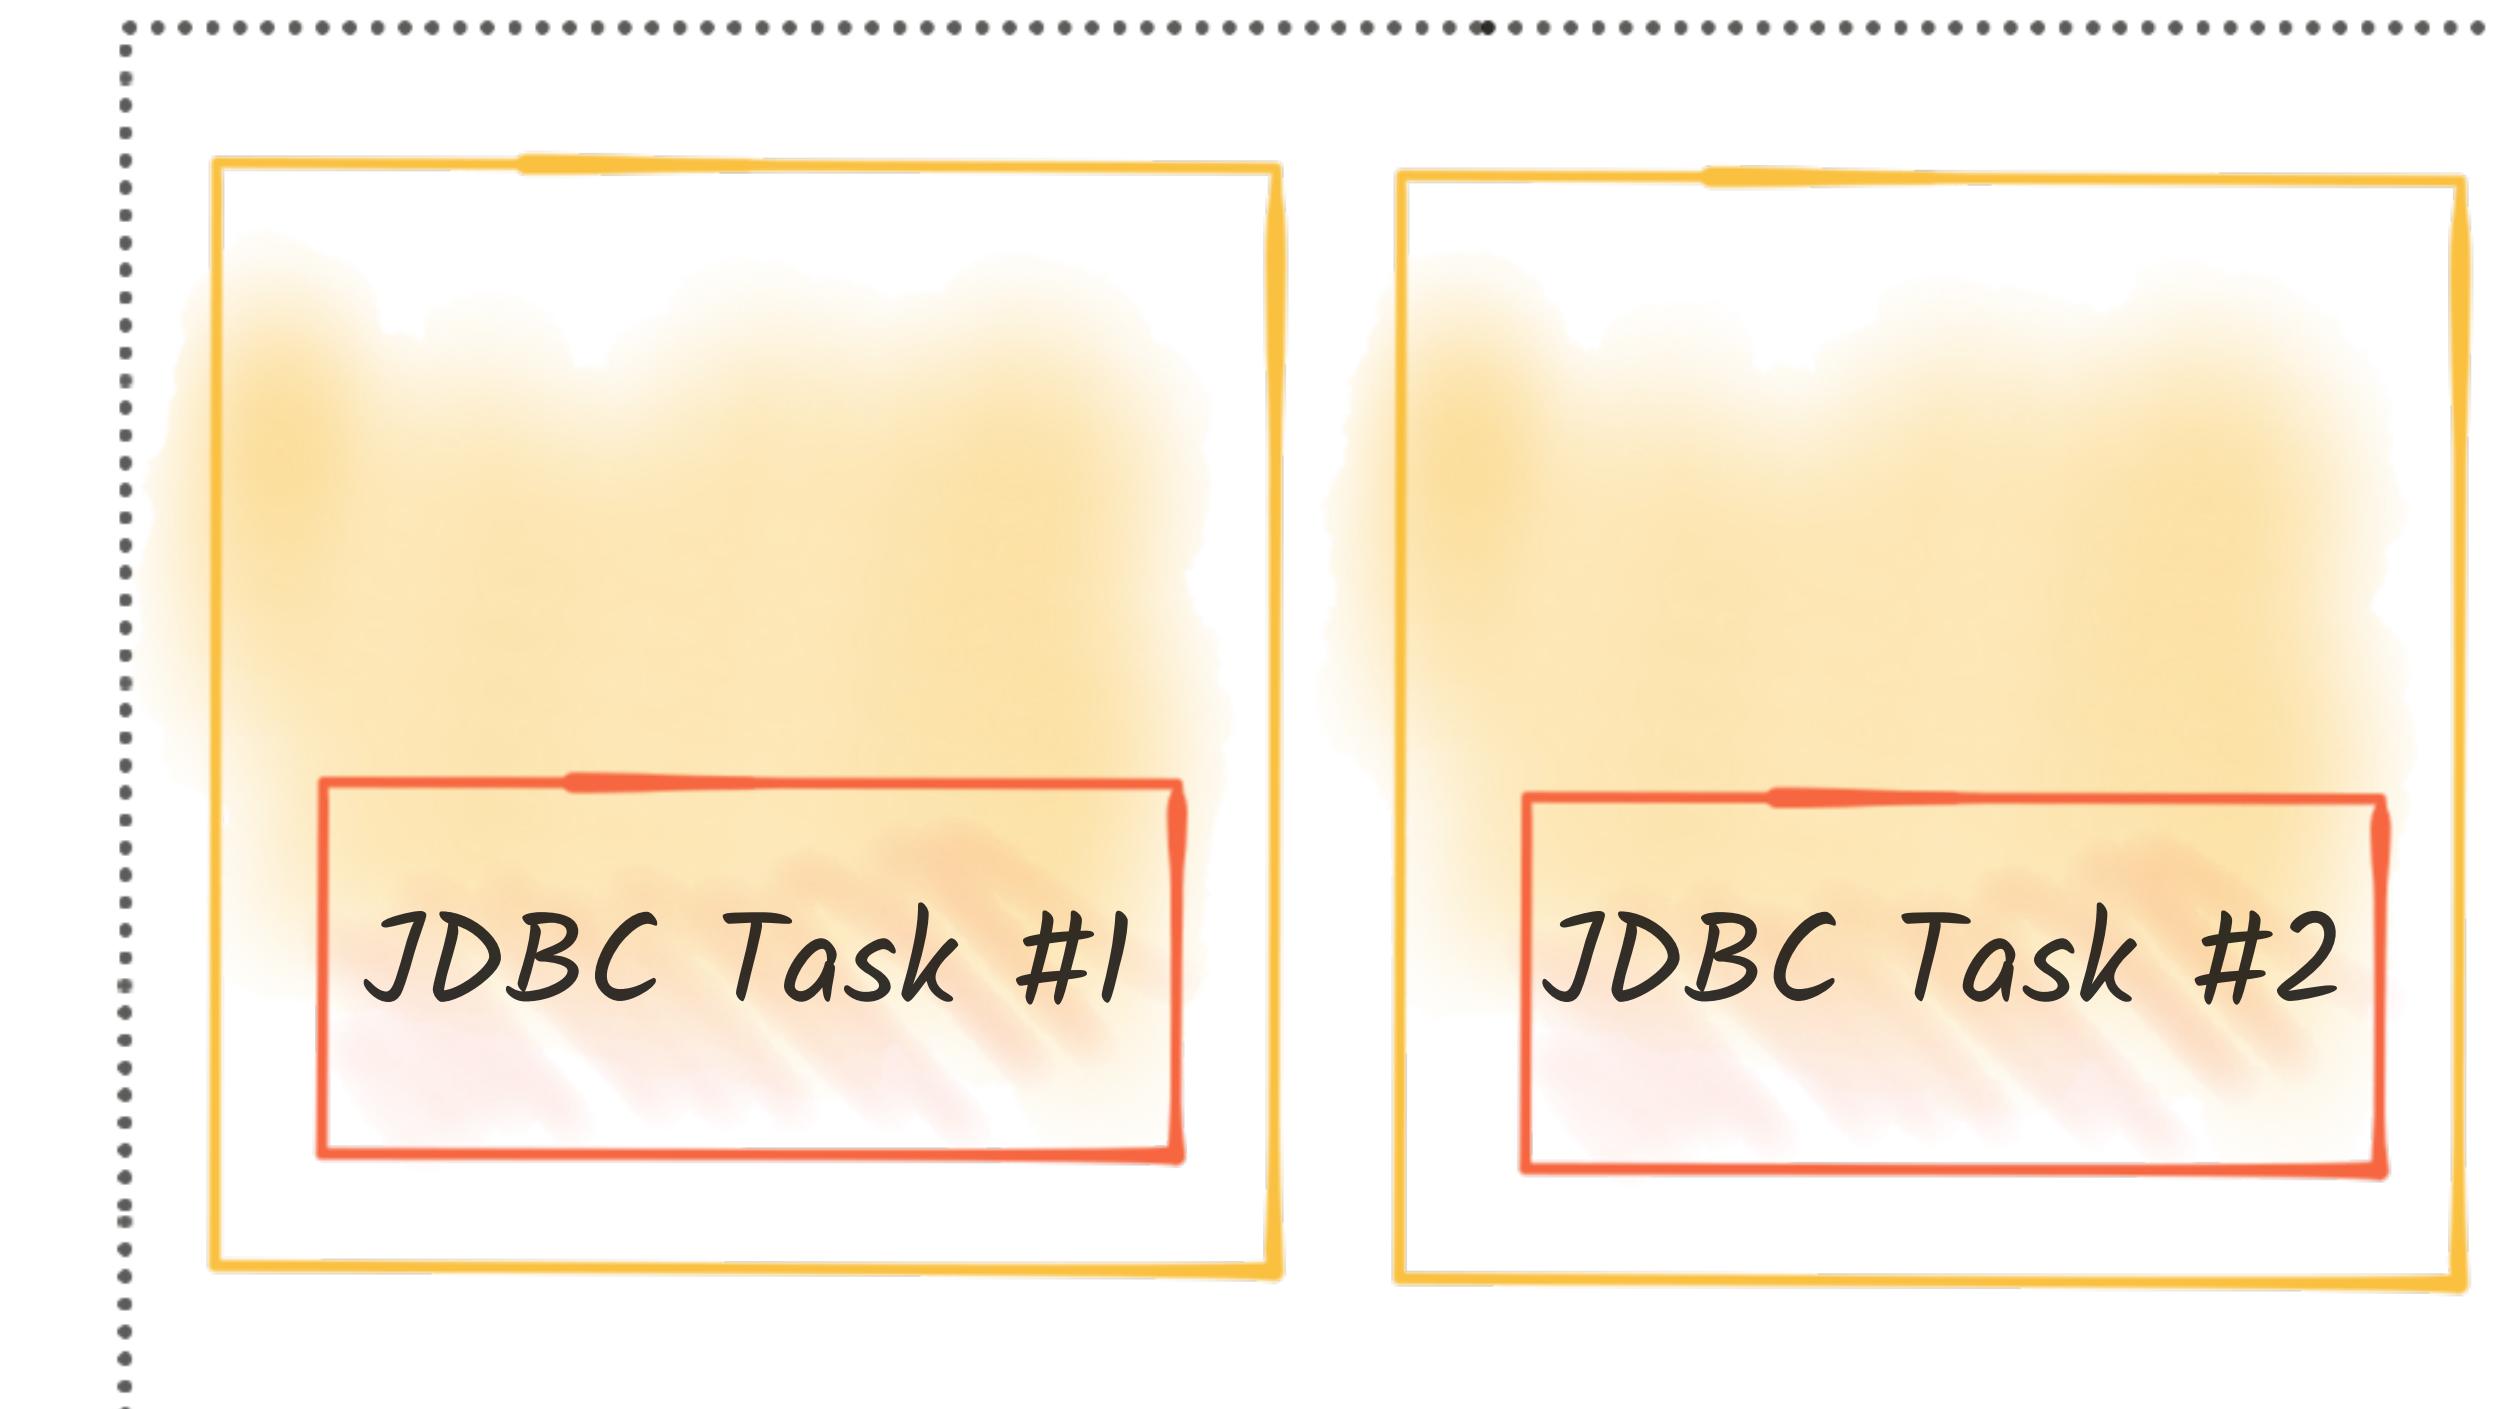
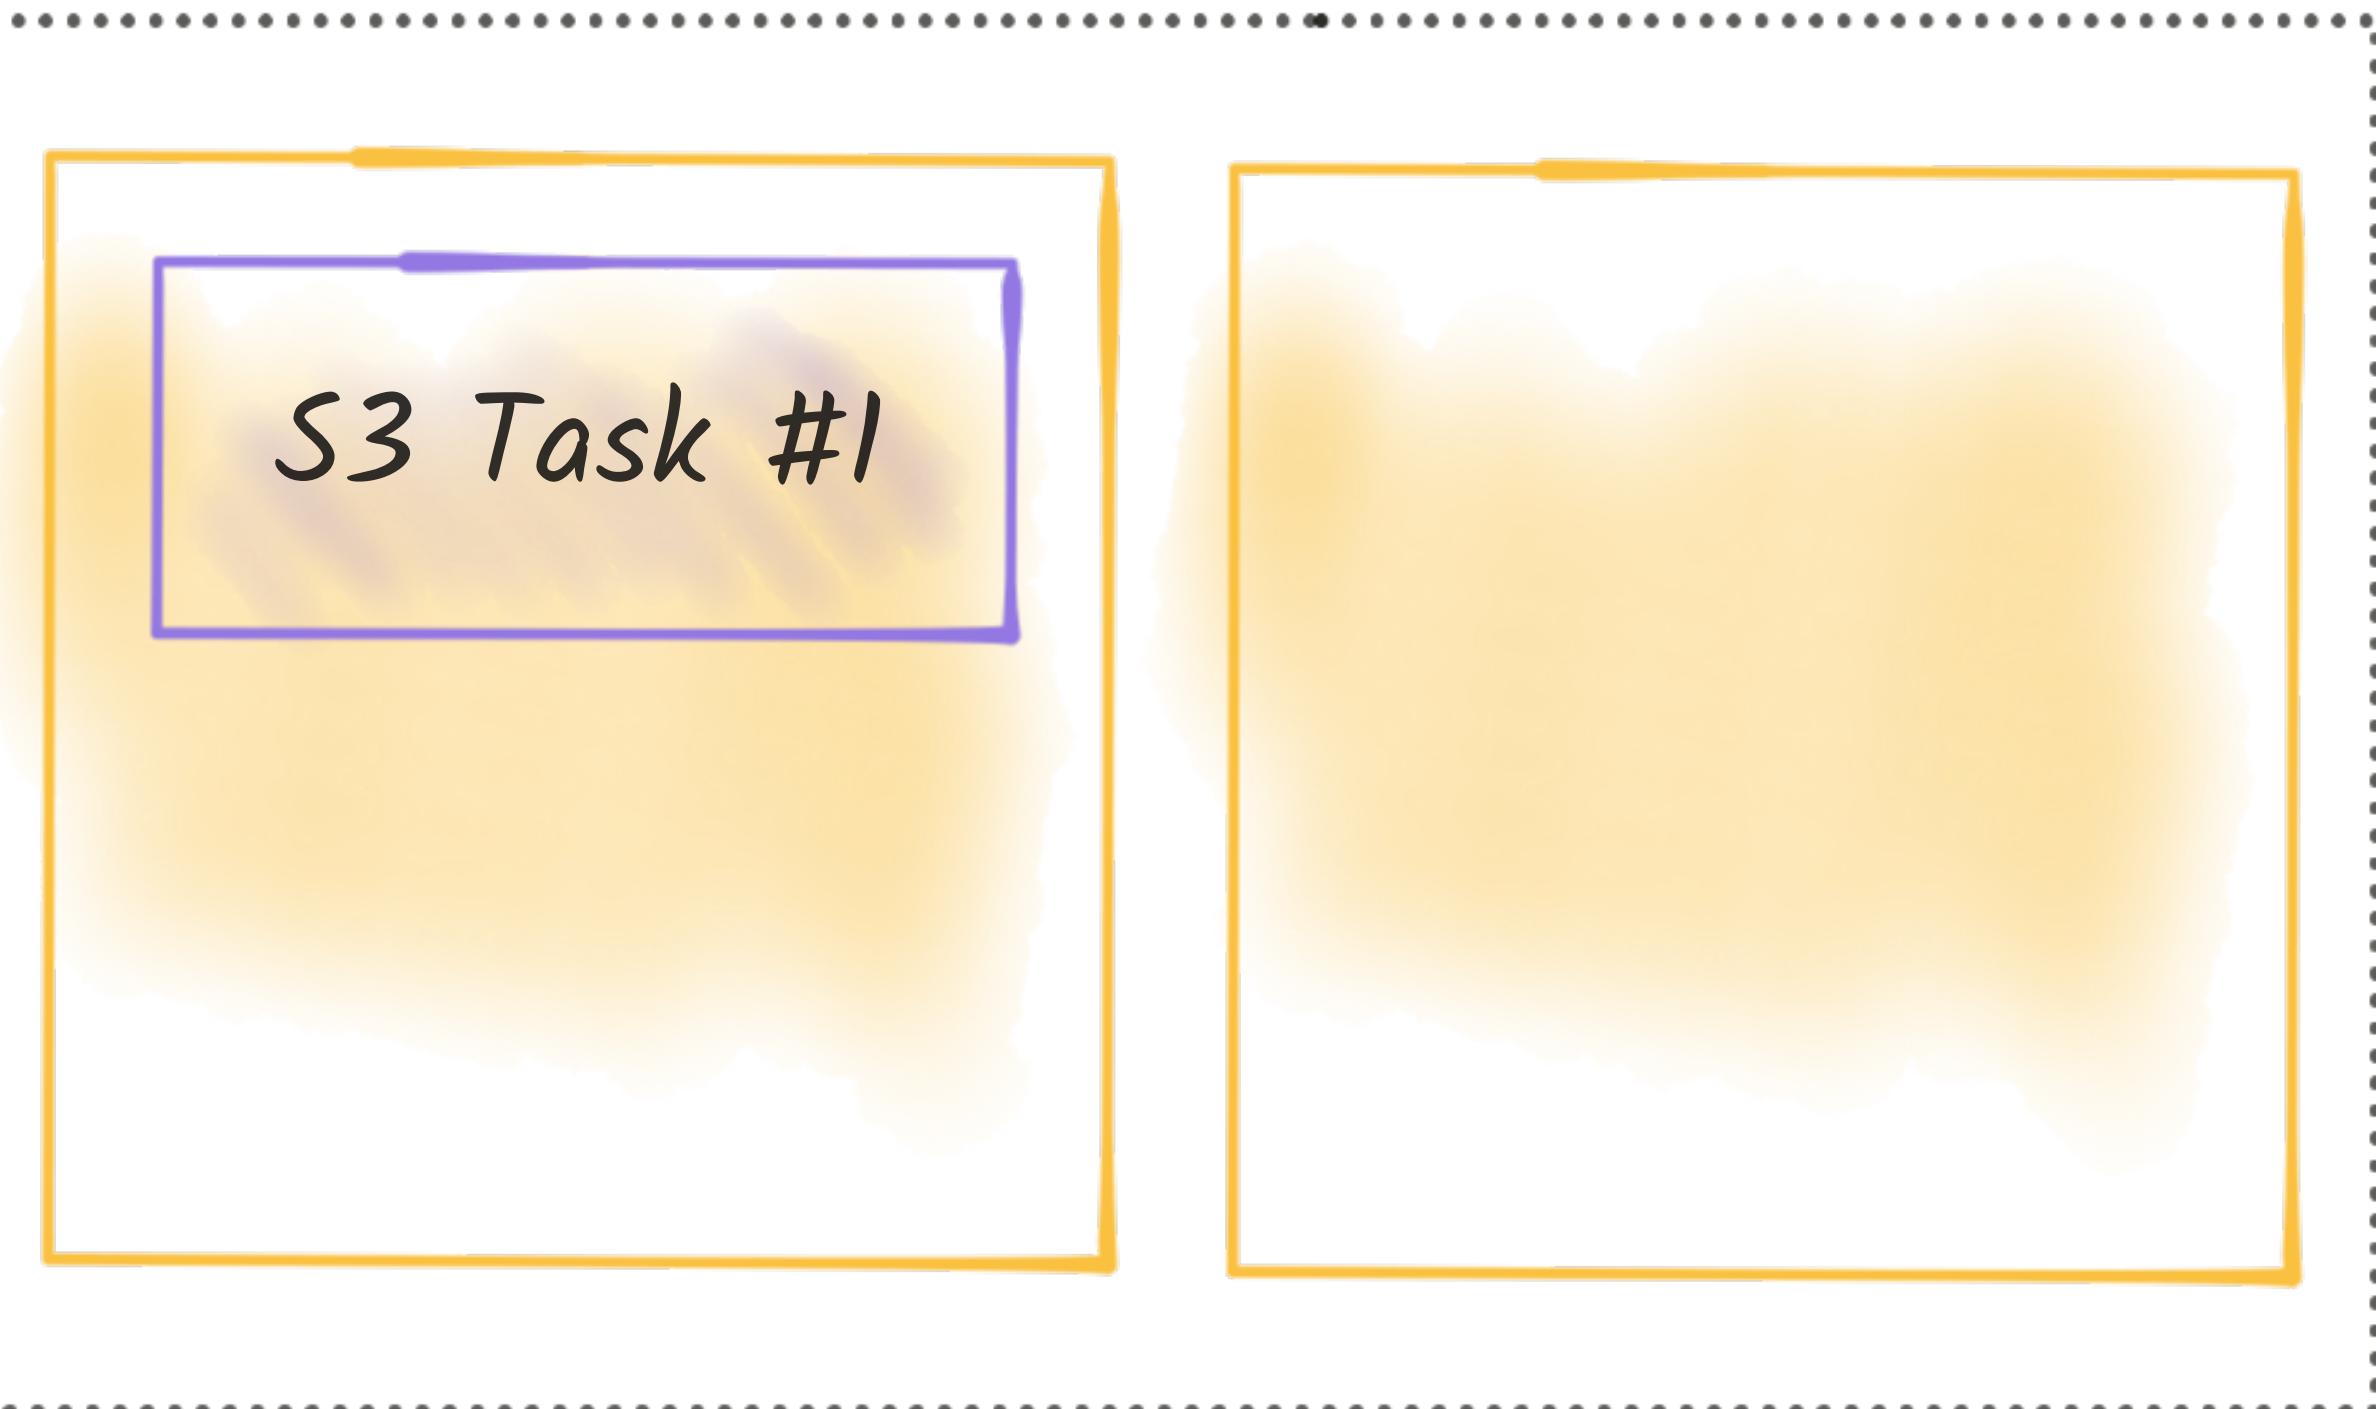
# Distributed Worker - fault tolerance



# Distributed Worker - fault tolerance



# Multiple Distributed Clusters



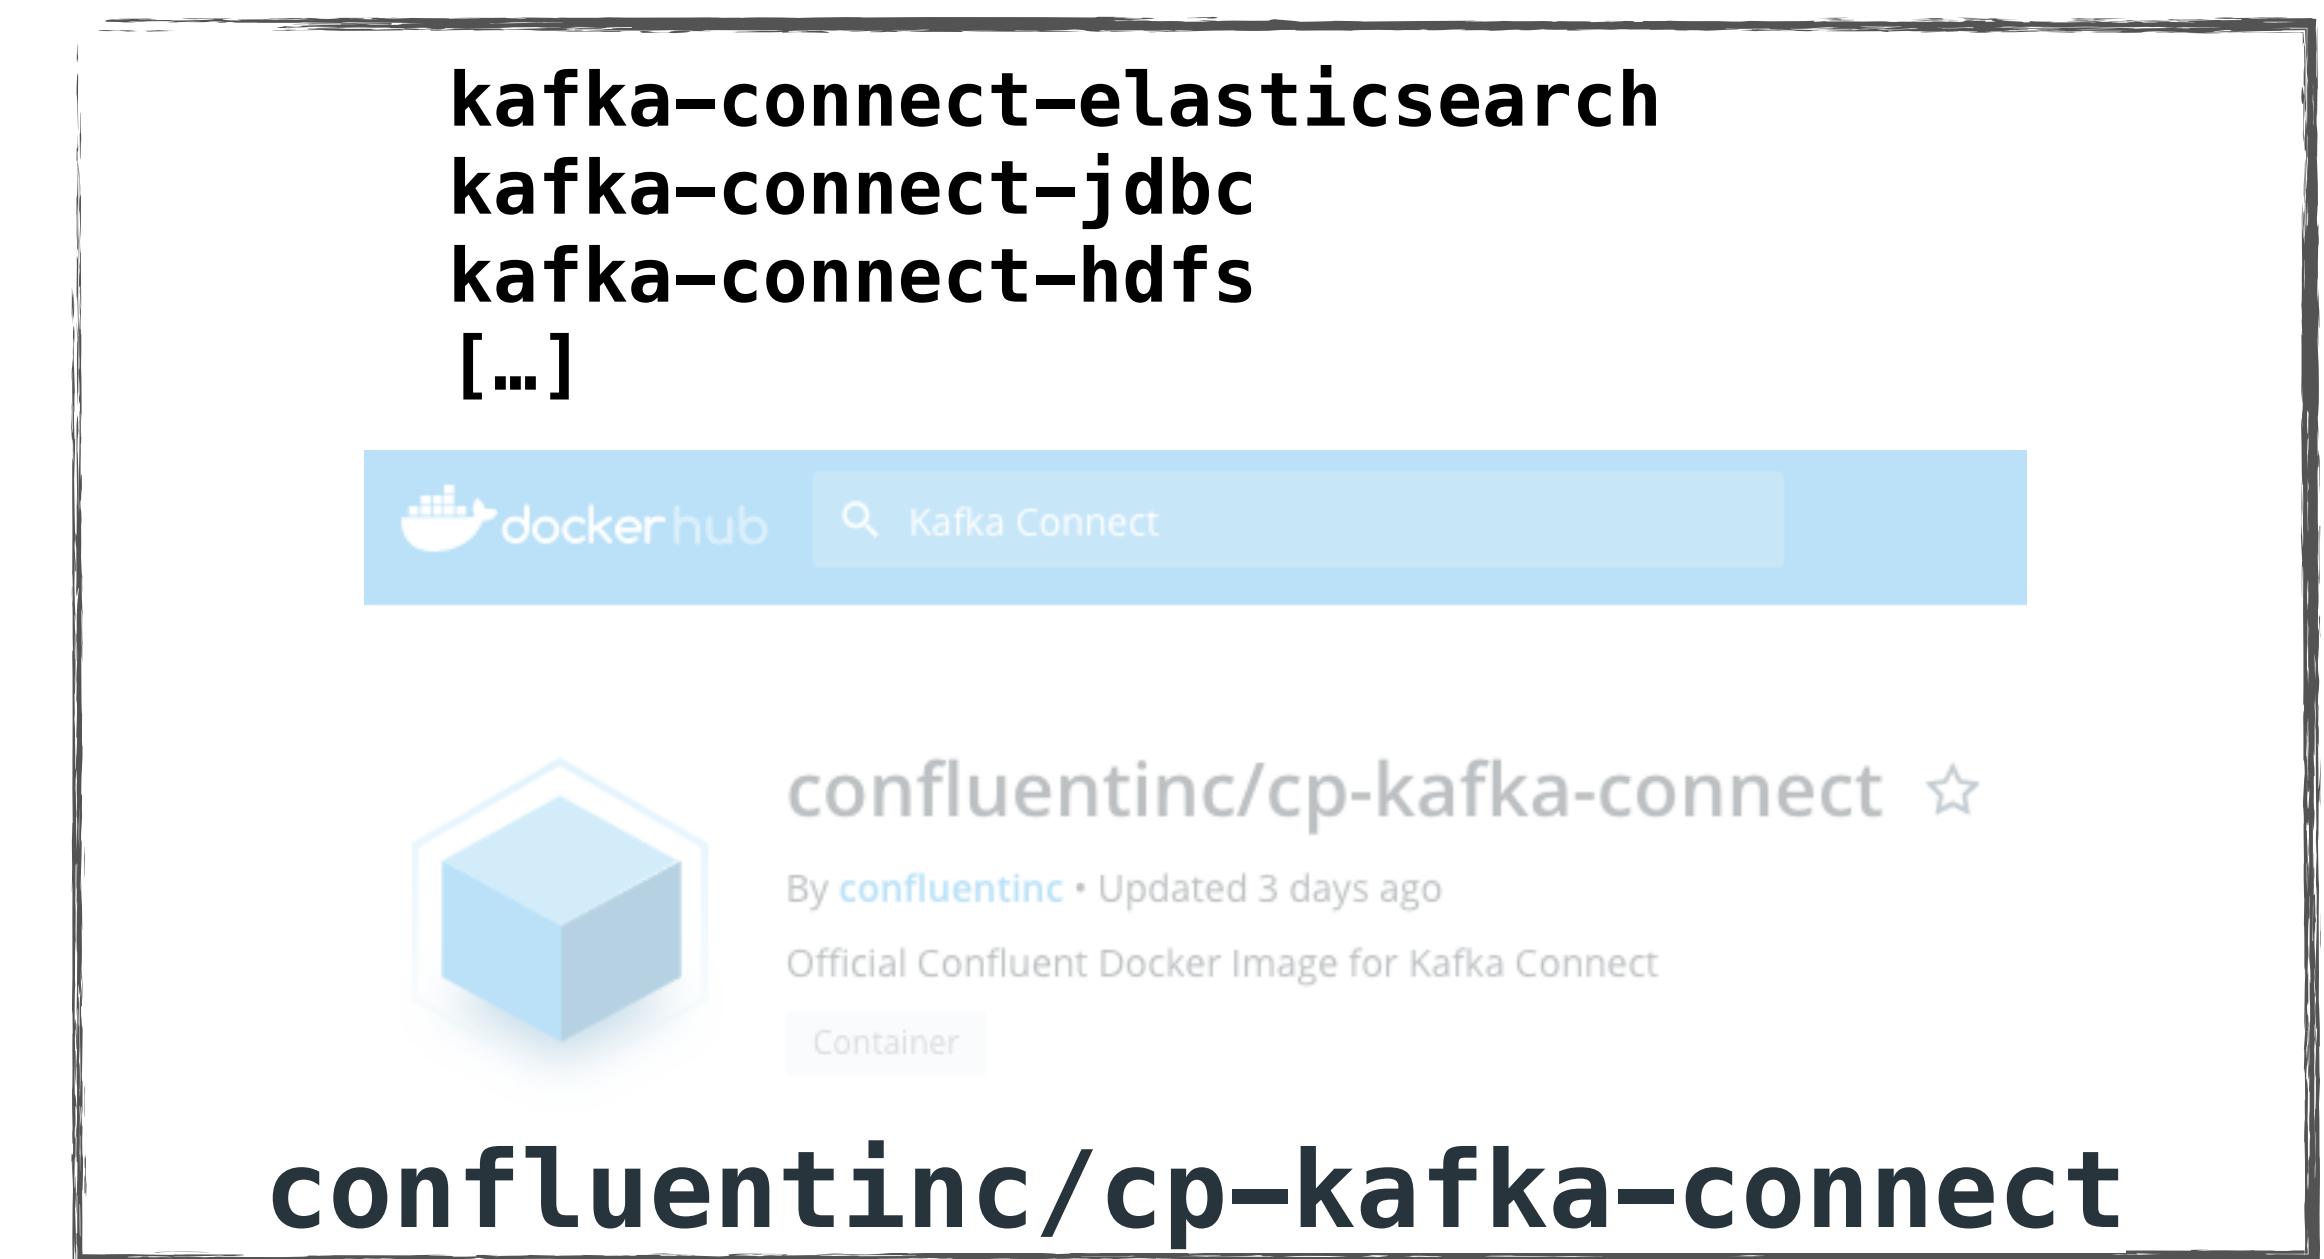
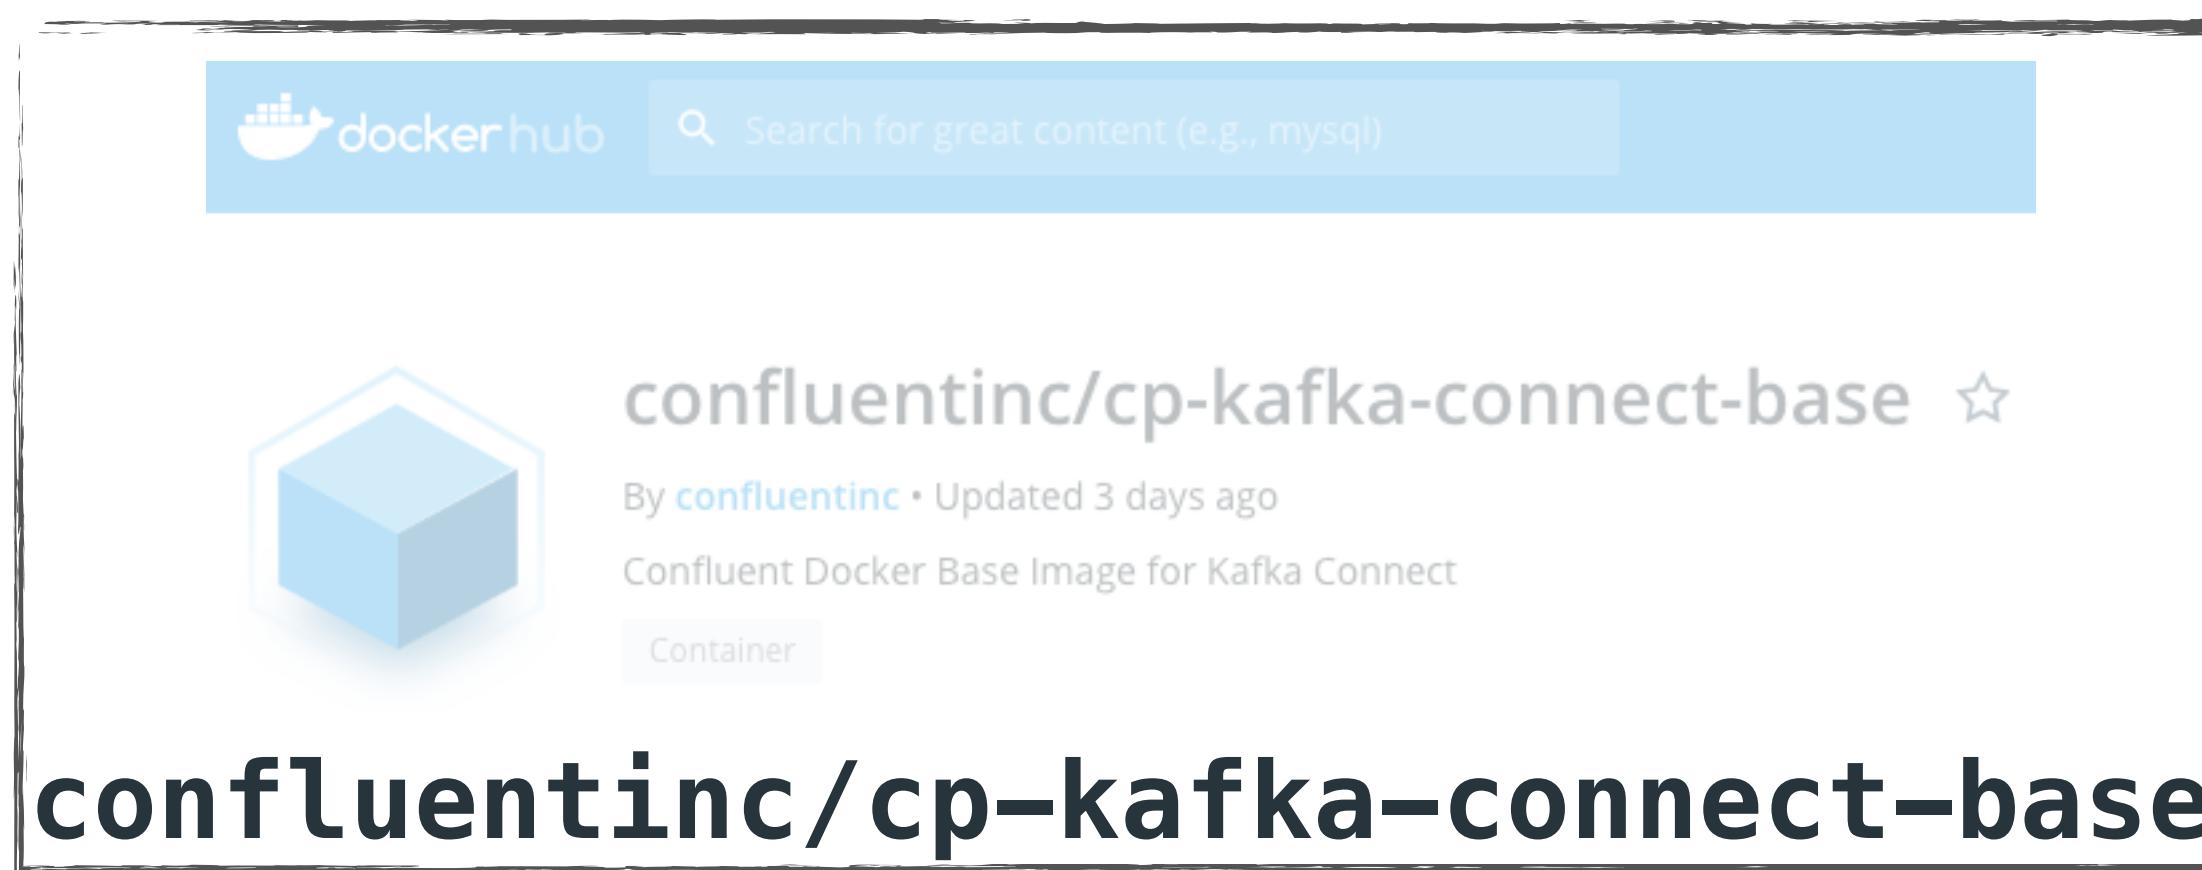
Kafka Connect cluster #1

Kafka Connect cluster #2



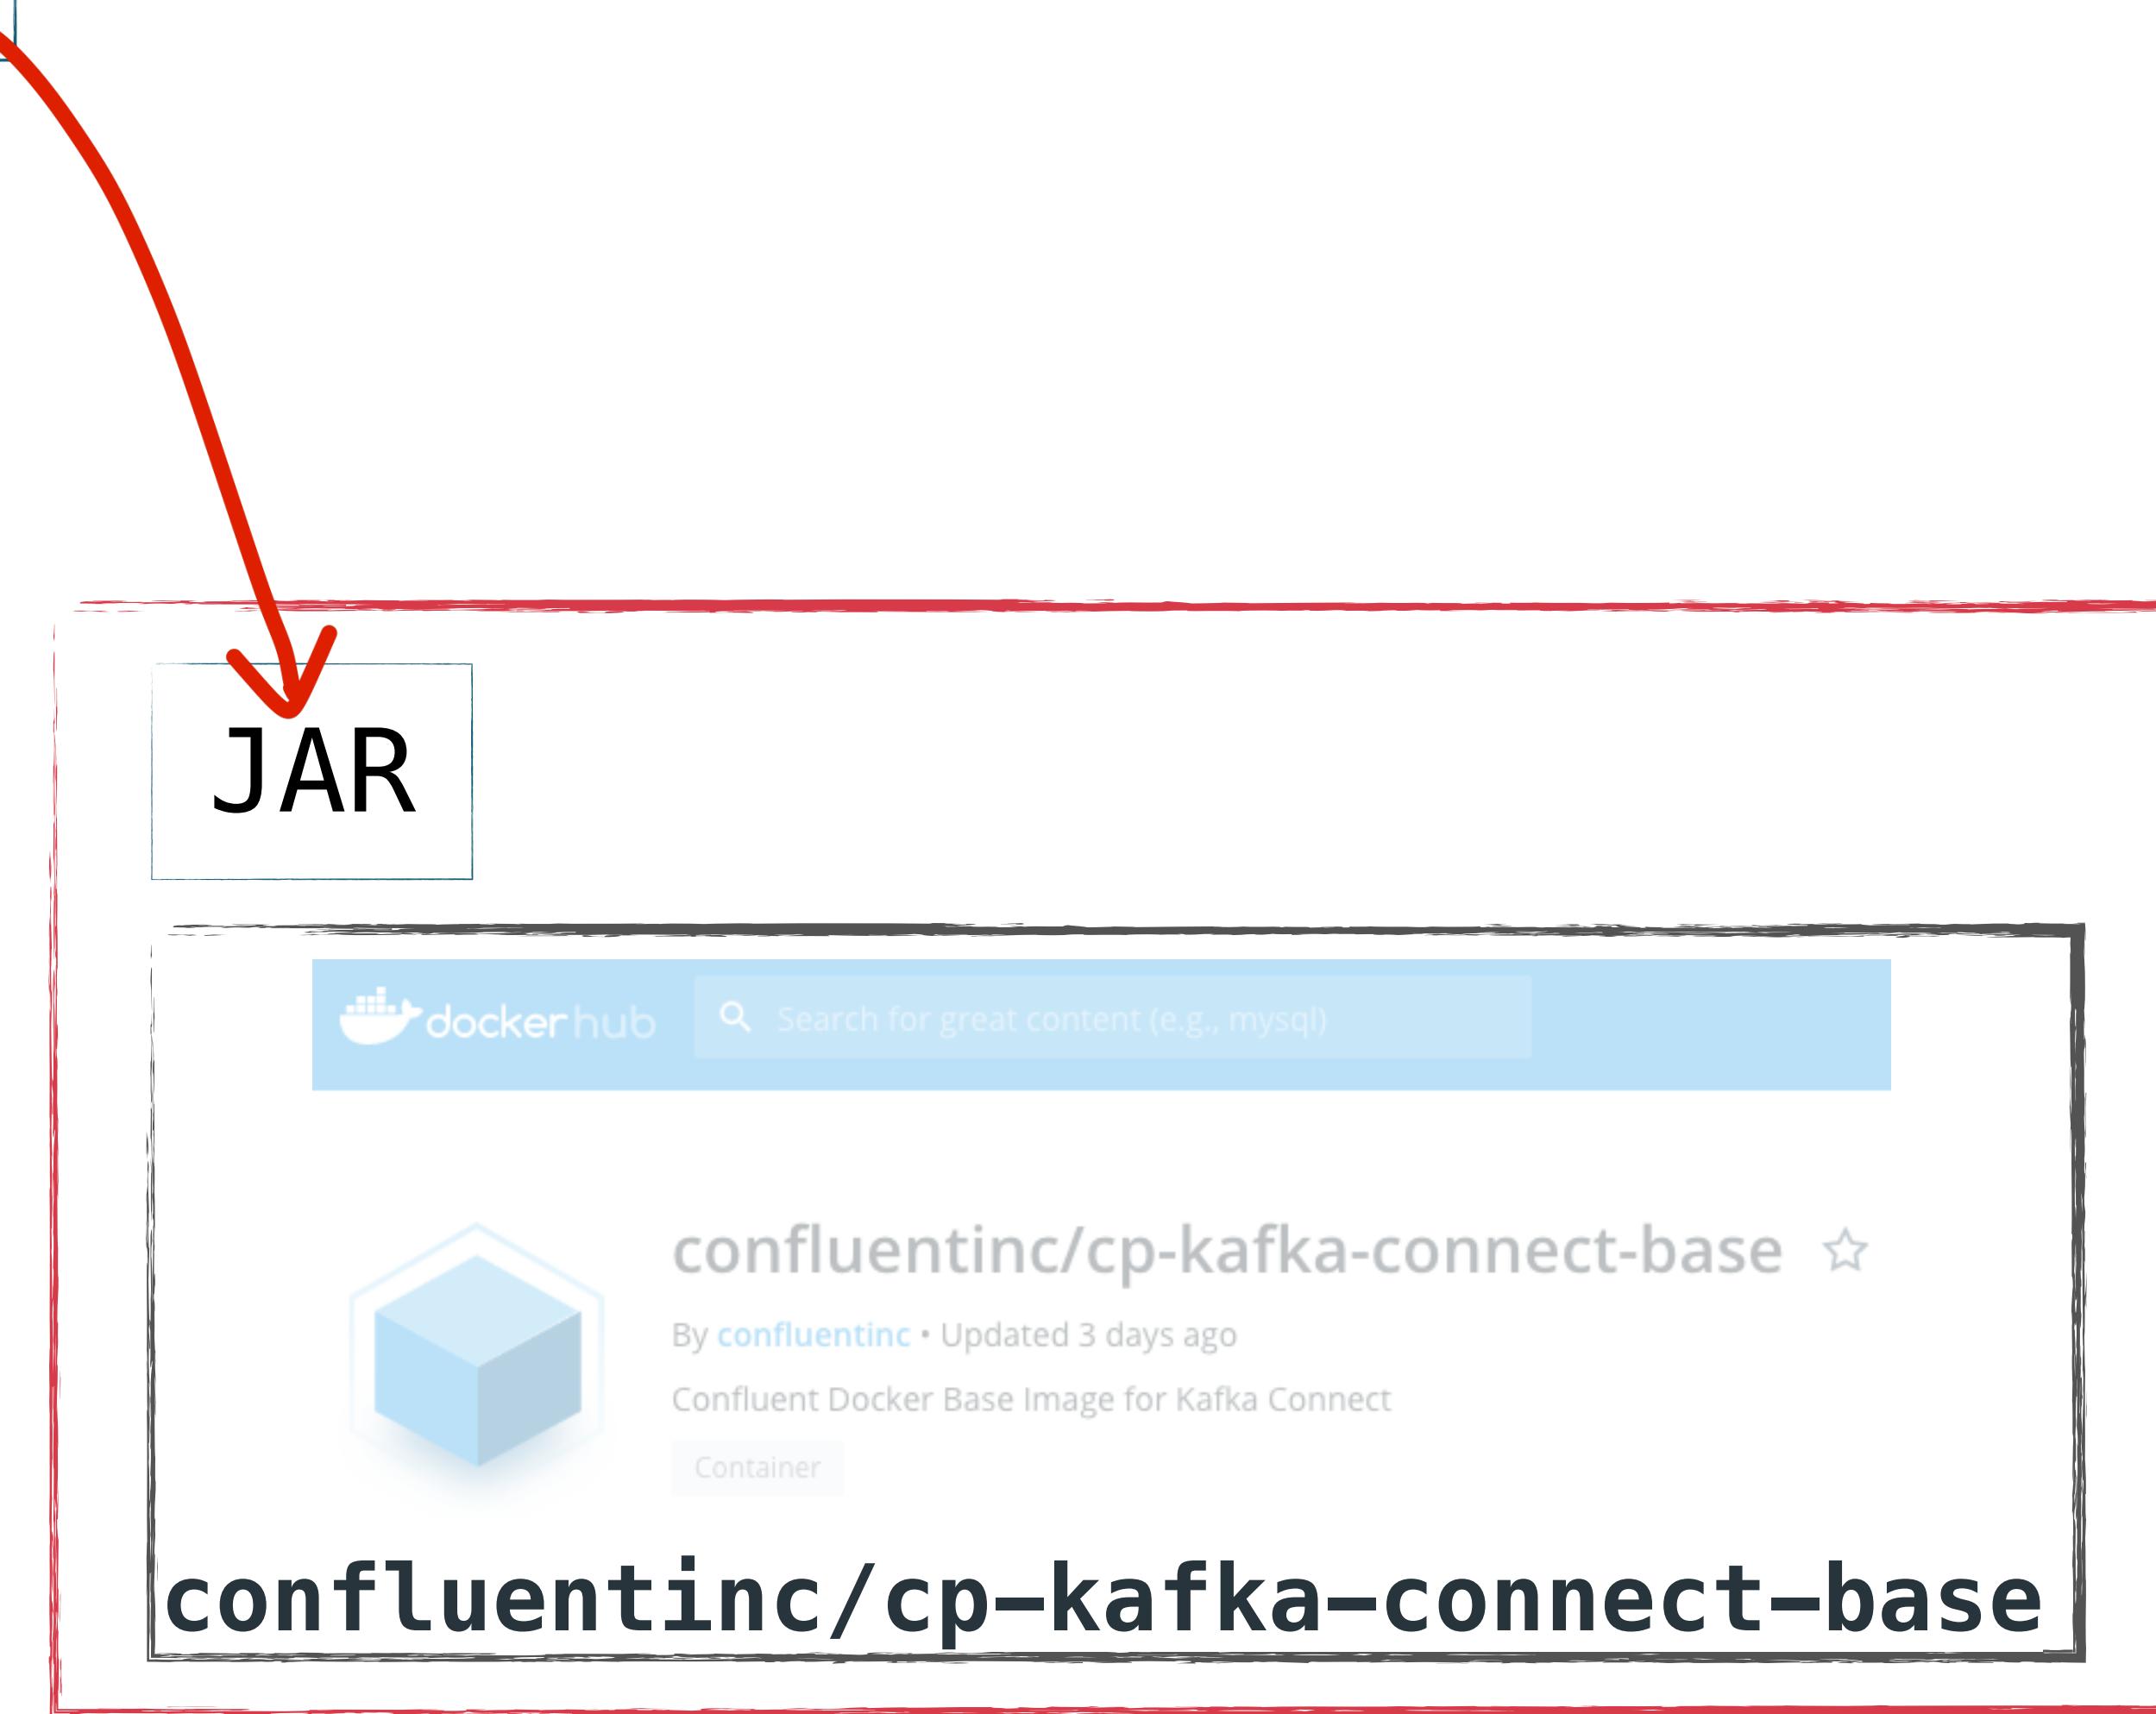
# Containers

# Kafka Connect images on Docker Hub



# Adding connectors to a container

Confluent Hub



# At runtime

kafka-connect:

```
image: confluentinc/cp-kafka-connect:5.2.1
```

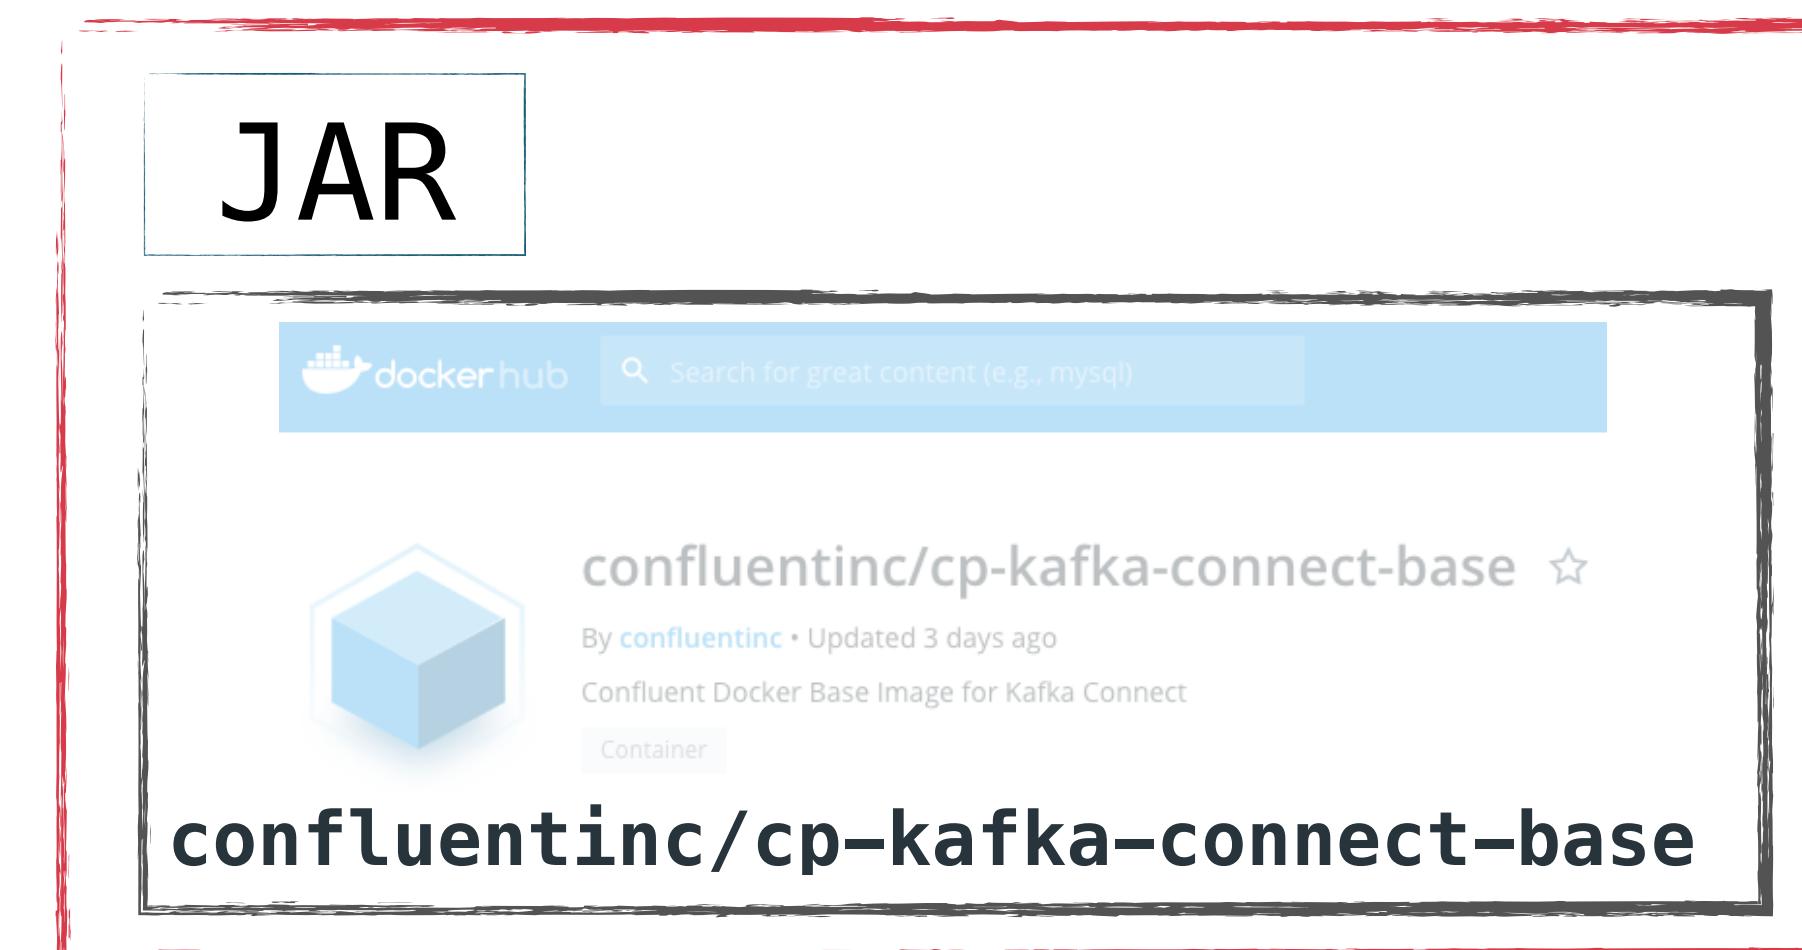
environment:

```
CONNECT_PLUGIN_PATH: '/usr/share/java,/usr/share/confluent-hub-components'
```

command:

- bash
- -c
- |

```
confluent-hub install --no-prompt neo4j/kafka-connect-neo4j:1.0.0  
/etc/confluent/docker/run
```



<http://rmoff.dev/ksln19-connect-docker>

From Zero to Hero with Kafka Connect

# Build a new image

```
FROM confluentinc/cp-kafka-connect:5.2.1
ENV CONNECT_PLUGIN_PATH="/usr/share/java,/usr/share/confluent-hub-components"
RUN confluent-hub install --no-prompt neo4j/kafka-connect-neo4j:1.0.0
```



# Automating connector creation

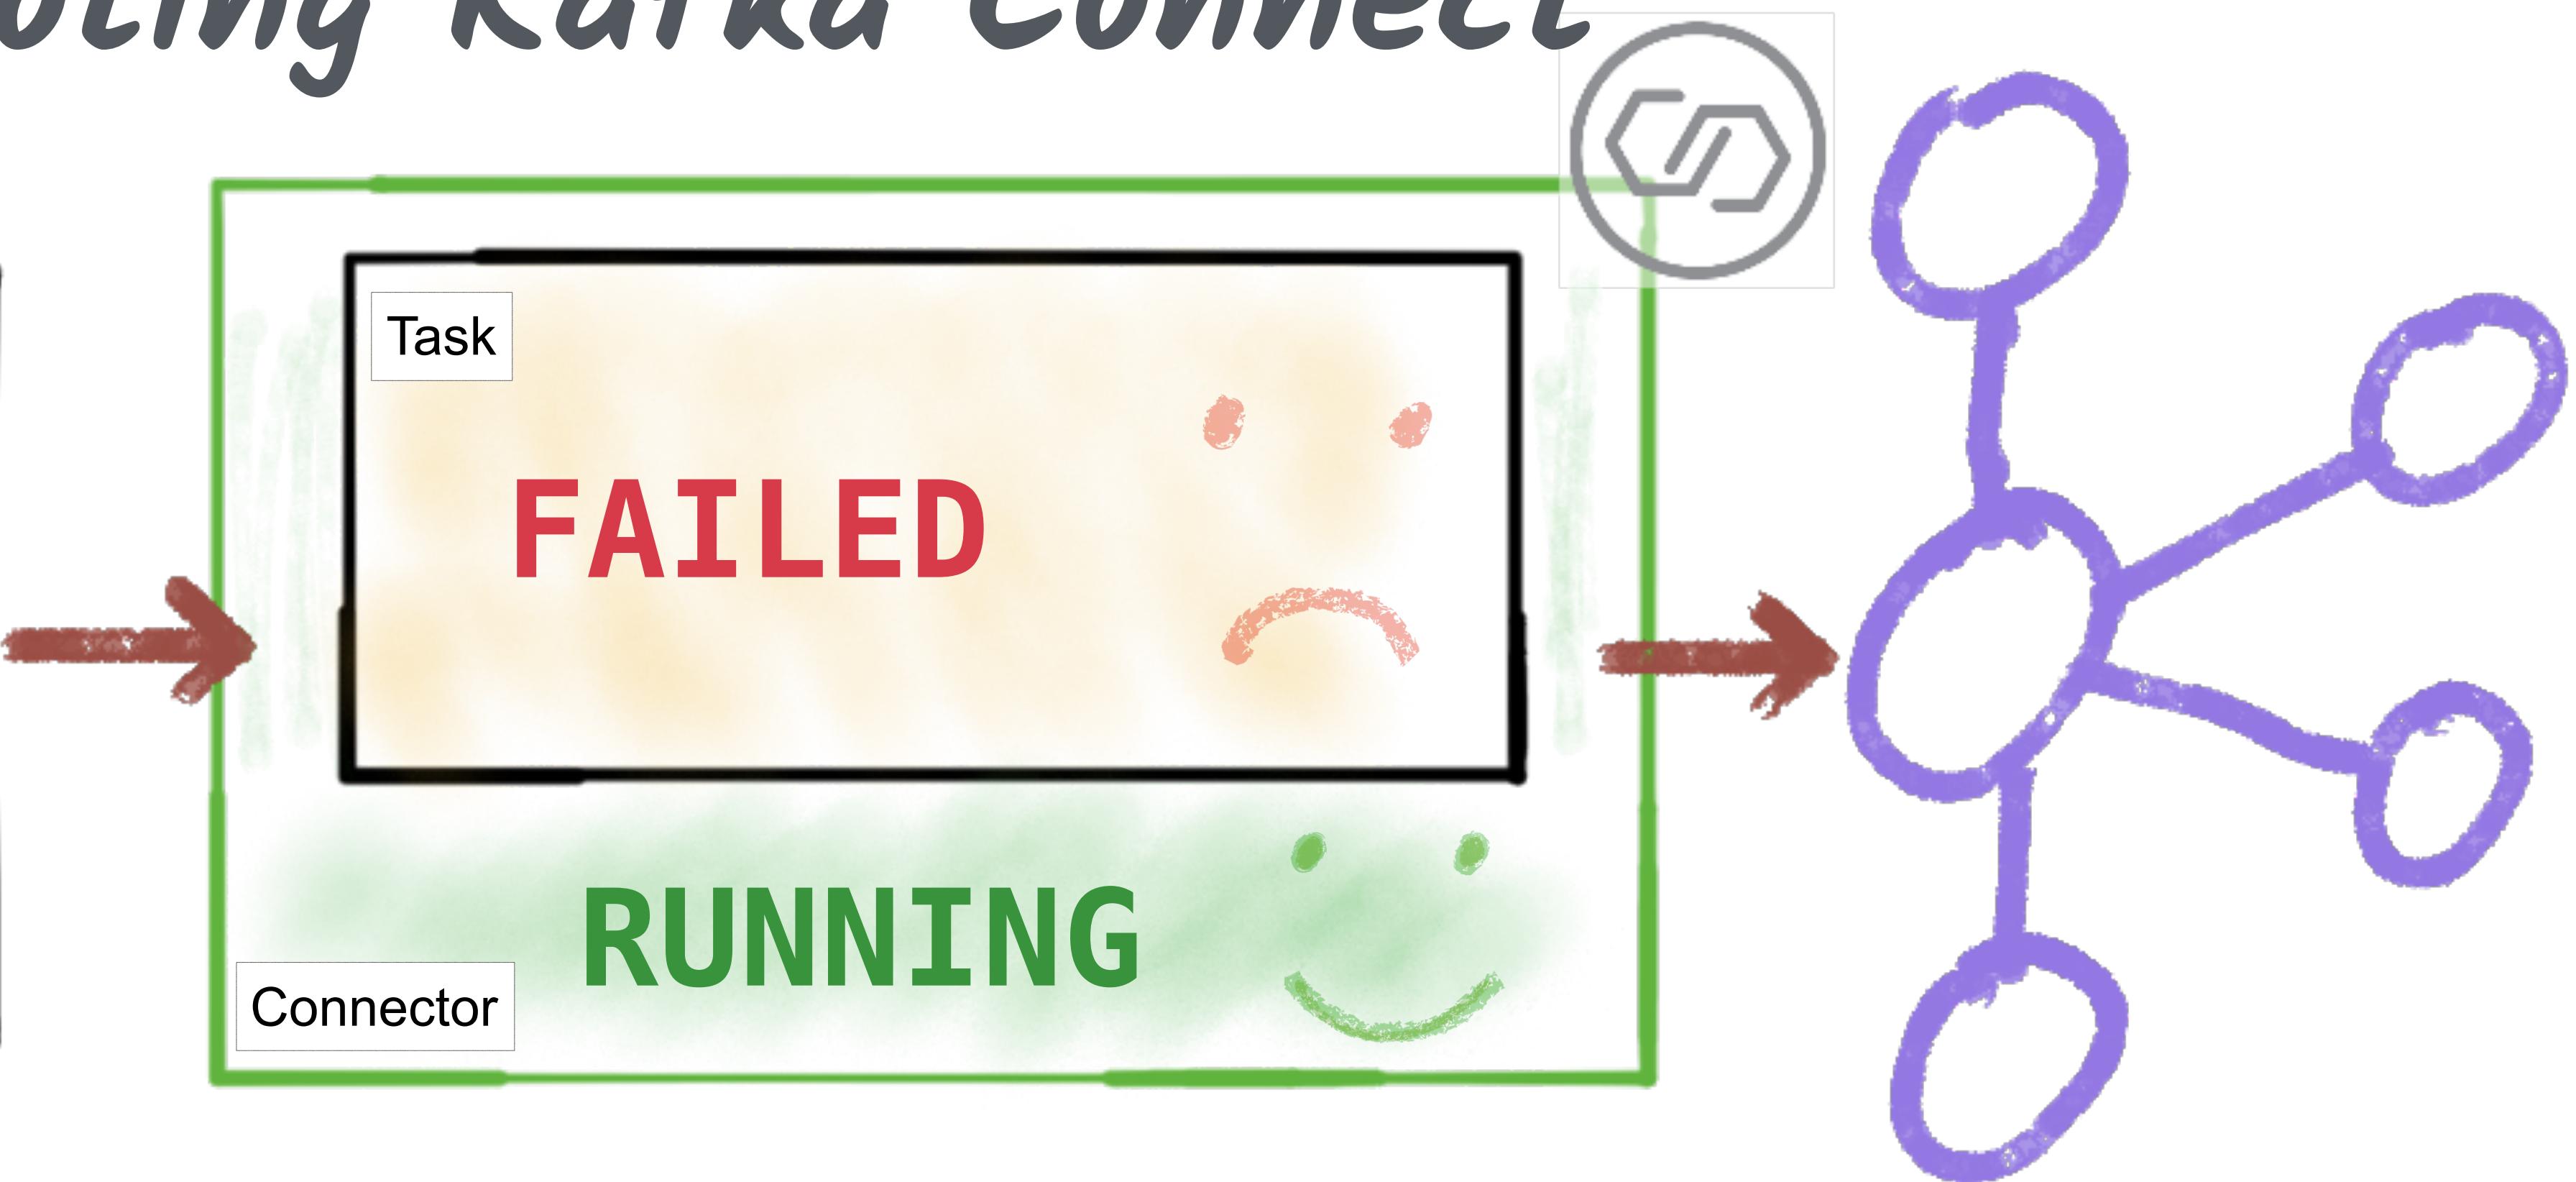
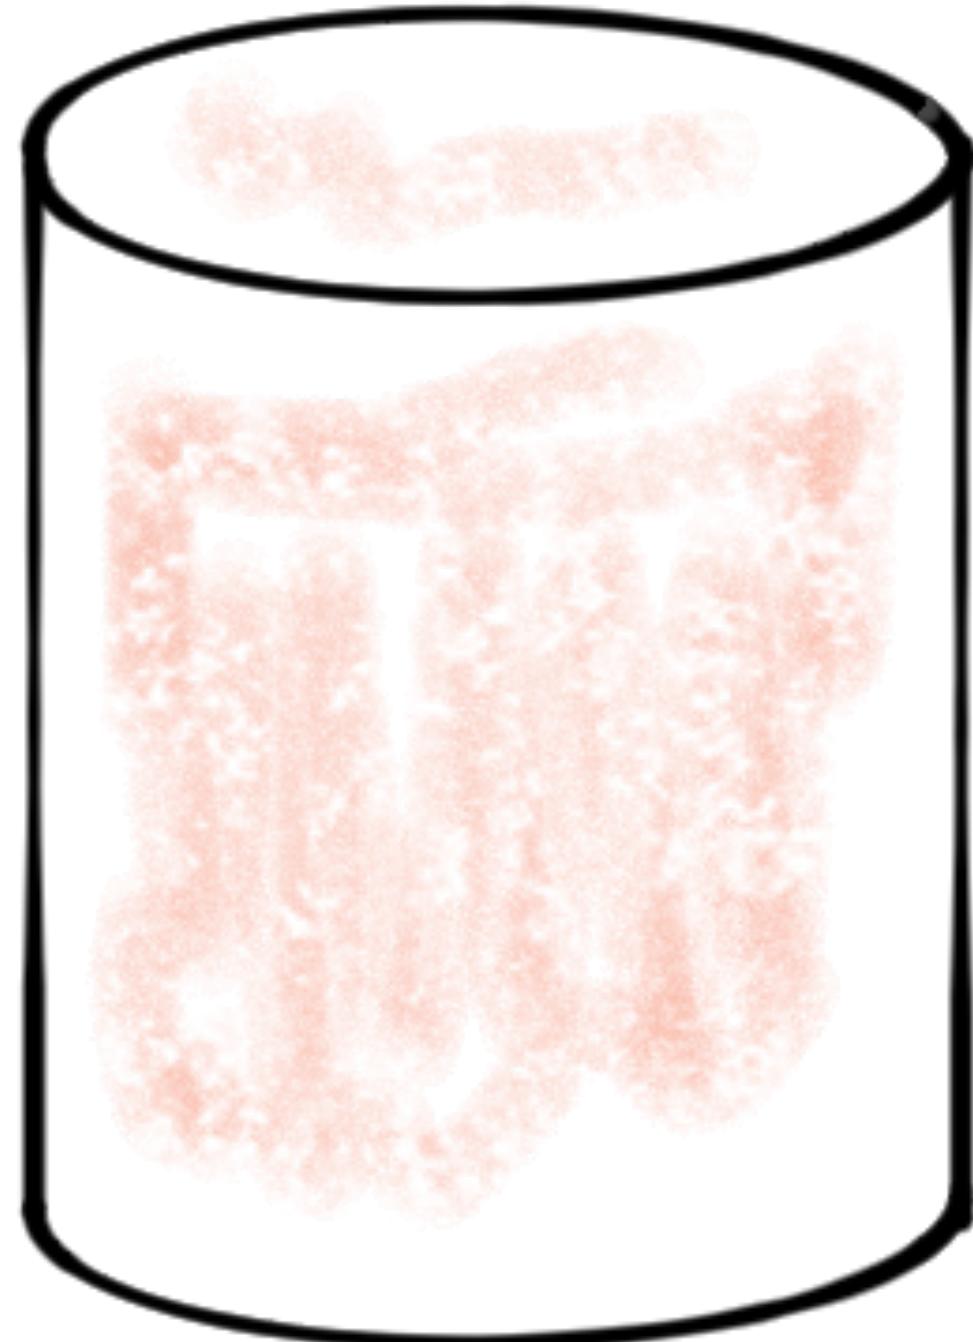
```
# # Download JDBC drivers
cd /usr/share/java/kafka-connect-jdbc/
curl https://cdn.mysql.com/Downloads/Connector-J/mysql-connector-java-8.0.13.tar.gz | tar xz
#
# Now launch Kafka Connect
/etc/confluent/docker/run &
#
# Wait for Kafka Connect listener
while [ $$($ curl -s -o /dev/null -w %{http_code} http://$$CONNECT_REST_ADVERTISED_HOST_NAME:$...
echo -e $$($date) " Kafka Connect listener HTTP state: " $$($ curl -s -o /dev/null -w %{http_...
sleep 5
done
#
# Create JDBC Source connector
curl -X POST http://localhost:8083/connectors -H "Content-Type: application/json" -d '{
  "name": "jdbc_source_mysql_00",
  "config": {
    "connector.class": "io.confluent.connect.jdbc.JdbcSourceConnector",
    "connection.url": "jdbc:mysql://mysql:3306/demo",
    "connection.user": "connect_user",
    "connection.password": "asgard",
    "topic.prefix": "mysql-00-",
    "table.whitelist" : "demo.customers",
  }
}
#
# Don't let the container die
sleep infinity
```



<http://rmoff.dev/ksln19-connect-docker>

# Troubleshooting Kafka Connect

# Troubleshooting Kafka Connect



```
$ curl -s "http://localhost:8083/connectors/source-debezium-orders/status" | \
jq '.connector.state'  
"RUNNING"  
$ curl -s "http://localhost:8083/connectors/source-debezium-orders/status" | \
jq '.tasks[0].state'  
"FAILED"
```



# Troubleshooting Kafka Connect

```
curl -s "http://localhost:8083/connectors/source-debezium-orders-00/status"  
| jq '.tasks[0].trace'
```

```
"org.apache.kafka.connect.errors.ConnectException\n\tat  
io.debezium.connector.mysql.AbstractReader.wrap(AbstractReader.java:230)\n\tat  
io.debezium.connector.mysql.AbstractReader.failed(AbstractReader.java:197)\n\tat  
io.debezium.connector.mysql.BinlogReader$ReaderThreadLifecycleListener.onCommunicationFailure(BinlogReader.java:  
1018)\n\tat com.github.shyiko.mysql.binlog.BinaryLogClient.listenForEventPackets(BinaryLogClient.java:950)\n\tat  
com.github.shyiko.mysql.binlog.BinaryLogClient.connect(BinaryLogClient.java:580)\n\tat  
com.github.shyiko.mysql.binlog.BinaryLogClient$7.run(BinaryLogClient.java:825)\n\tat java.lang.Thread.run(Thread.java:  
748)\nCaused by: java.io.EOFException\n\tat  
com.github.shyiko.mysql.io.ByteArrayInputStream.read(ByteArrayInputStream.java:190)\n\tat  
com.github.shyiko.mysql.io.ByteArrayInputStream.readInteger(ByteArrayInputStream.java:46)\n\tat  
com.github.shyiko.mysql.event.serialization.EventHeaderV4Deserializer.deserialize(EventHeaderV4Deserializer.java  
:35)\n\tat  
com.github.shyiko.mysql.event.serialization.EventHeaderV4Deserializer.deserialize(EventHeaderV4Deserializer.java  
:27)\n\tat com.github.shyiko.mysql.event.serialization.EventDeserializer.nextEvent(EventDeserializer.java:  
212)\n\tat io.debezium.connector.mysql.BinlogReader$1.nextEvent(BinlogReader.java:224)\n\tat  
com.github.shyiko.mysql.binlog.BinaryLogClient.listenForEventPackets(BinaryLogClient.java:922)\n\t... 3 more\n"
```

*The log is the source of truth*

```
$ confluent log connect
```

```
$ docker-compose logs kafka-connect
```

```
$ cat /var/log/kafka/connect.log
```

# Kafka Connect

**"Task is being killed and will  
not recover until manually restarted"**

*Symptom not Cause*

# Error Handling

and

# Dead Letter

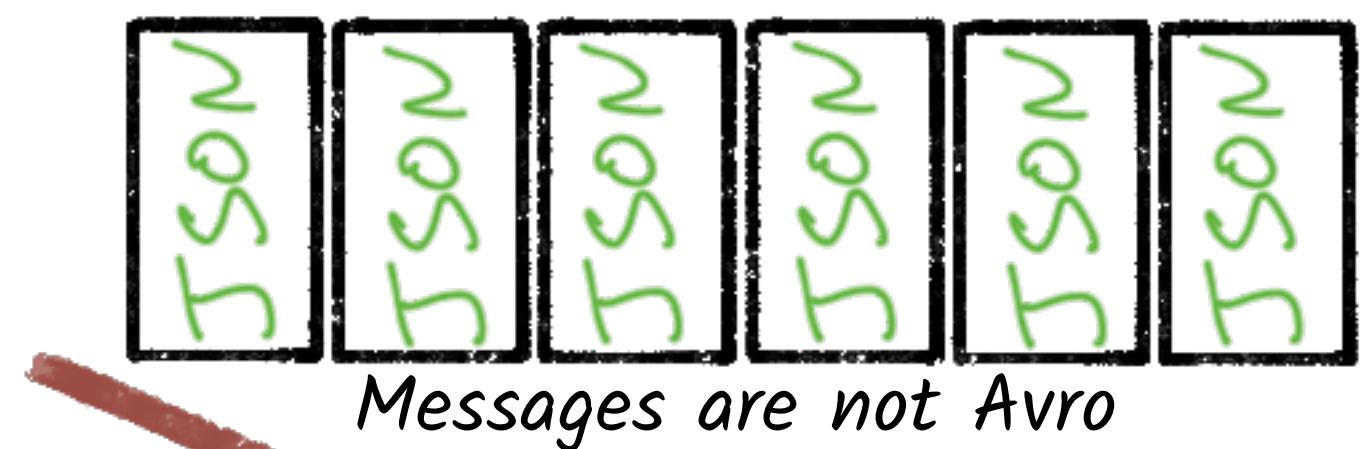
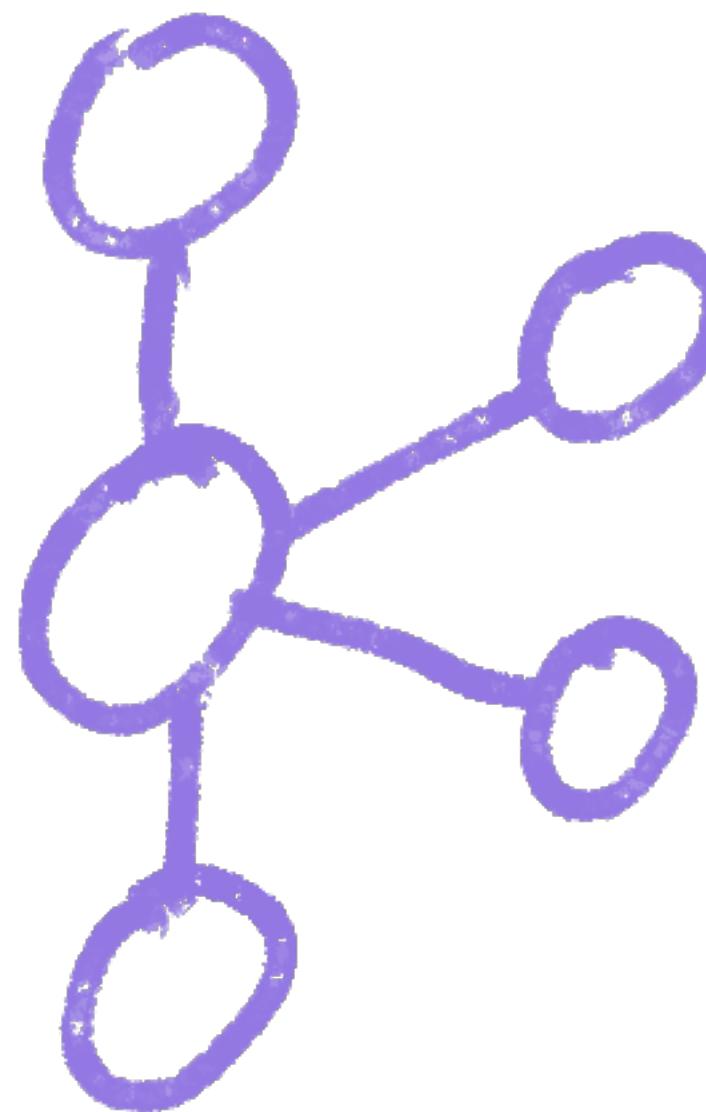
# Queues

*org.apache.kafka.common.errors.SerializationException:  
Unknown magic byte!*

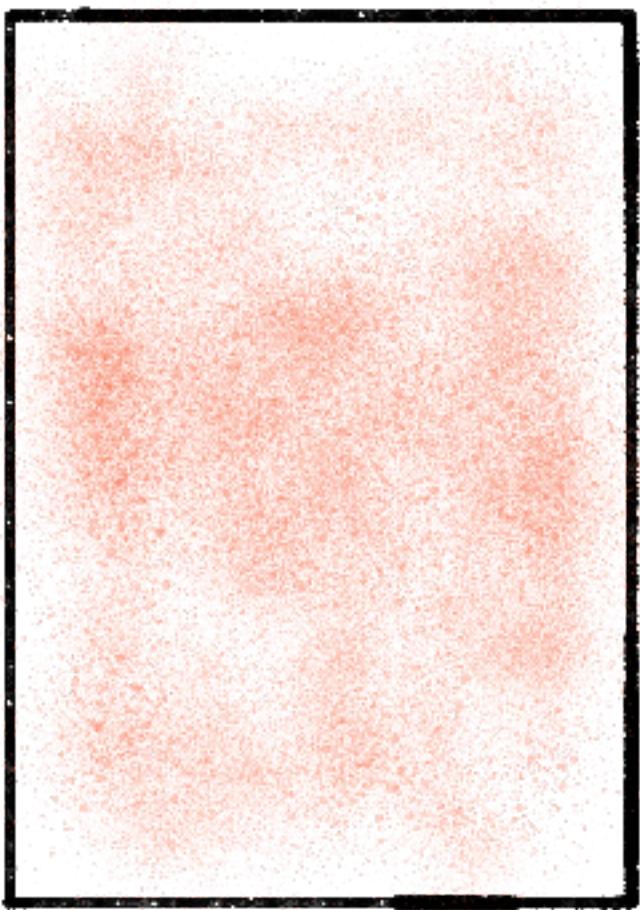
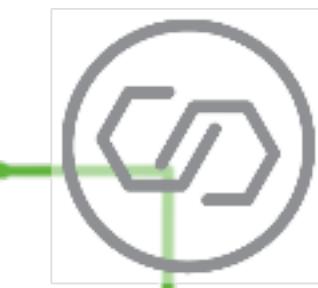
# Mismatched converters

`org.apache.kafka.common.errors.SerializationException:`

*Unknown magic byte!*



**"value.converter":  
"AvroConverter"**

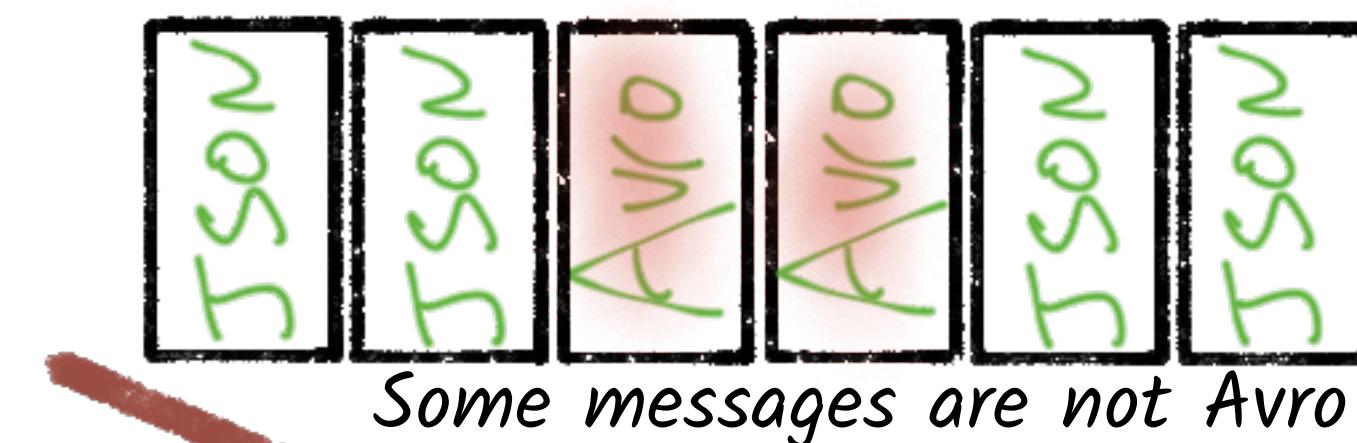
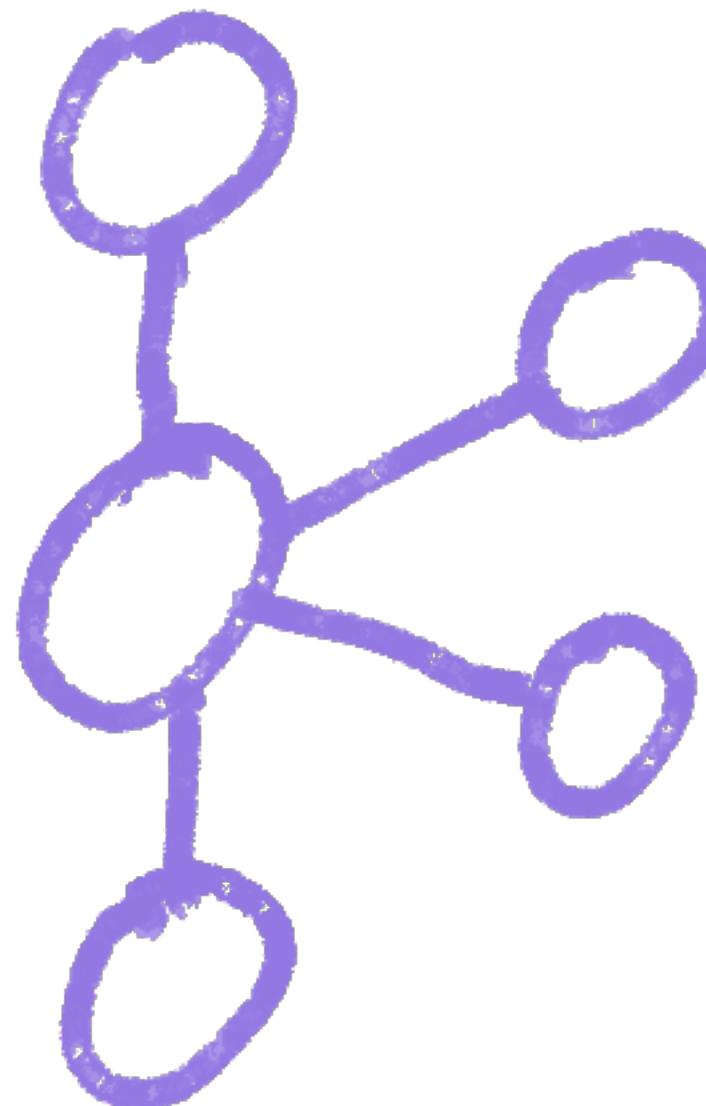


Use the correct Converter for the source data

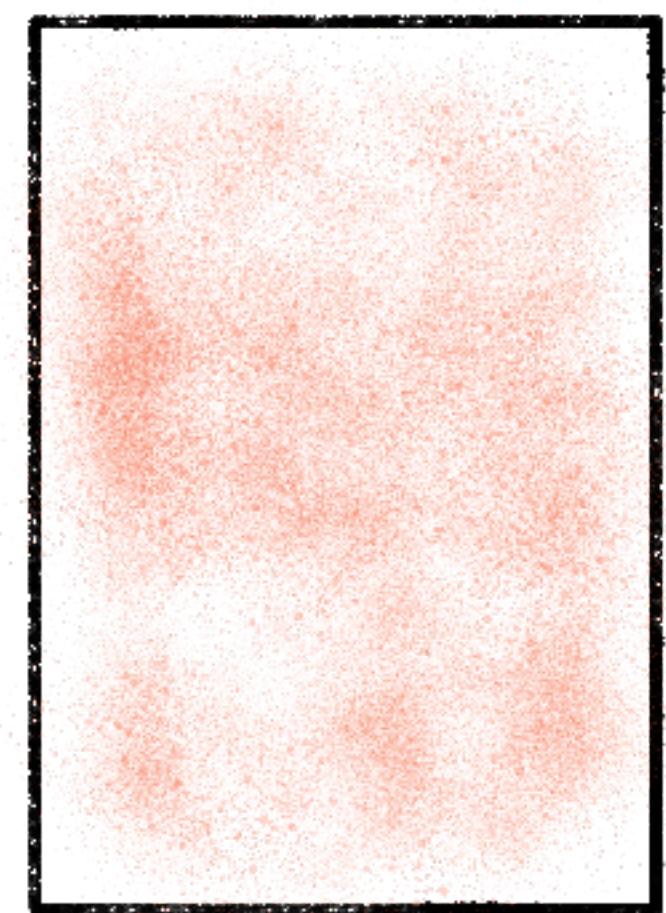
# Mixed serialisation methods

`org.apache.kafka.common.errors.SerializationException:`

*Unknown magic byte!*



`"value.converter":  
"AvroConverter"`



- ⓘ Use error handling to deal with bad messages

# Error Handling and DLQ

Handled

Convert

-> read/write from Kafka

-> [de]-serialisation

Transform

Not Handled

Start

-> Connections to a data store

Poll / Put

-> Read/Write from/to data store\*

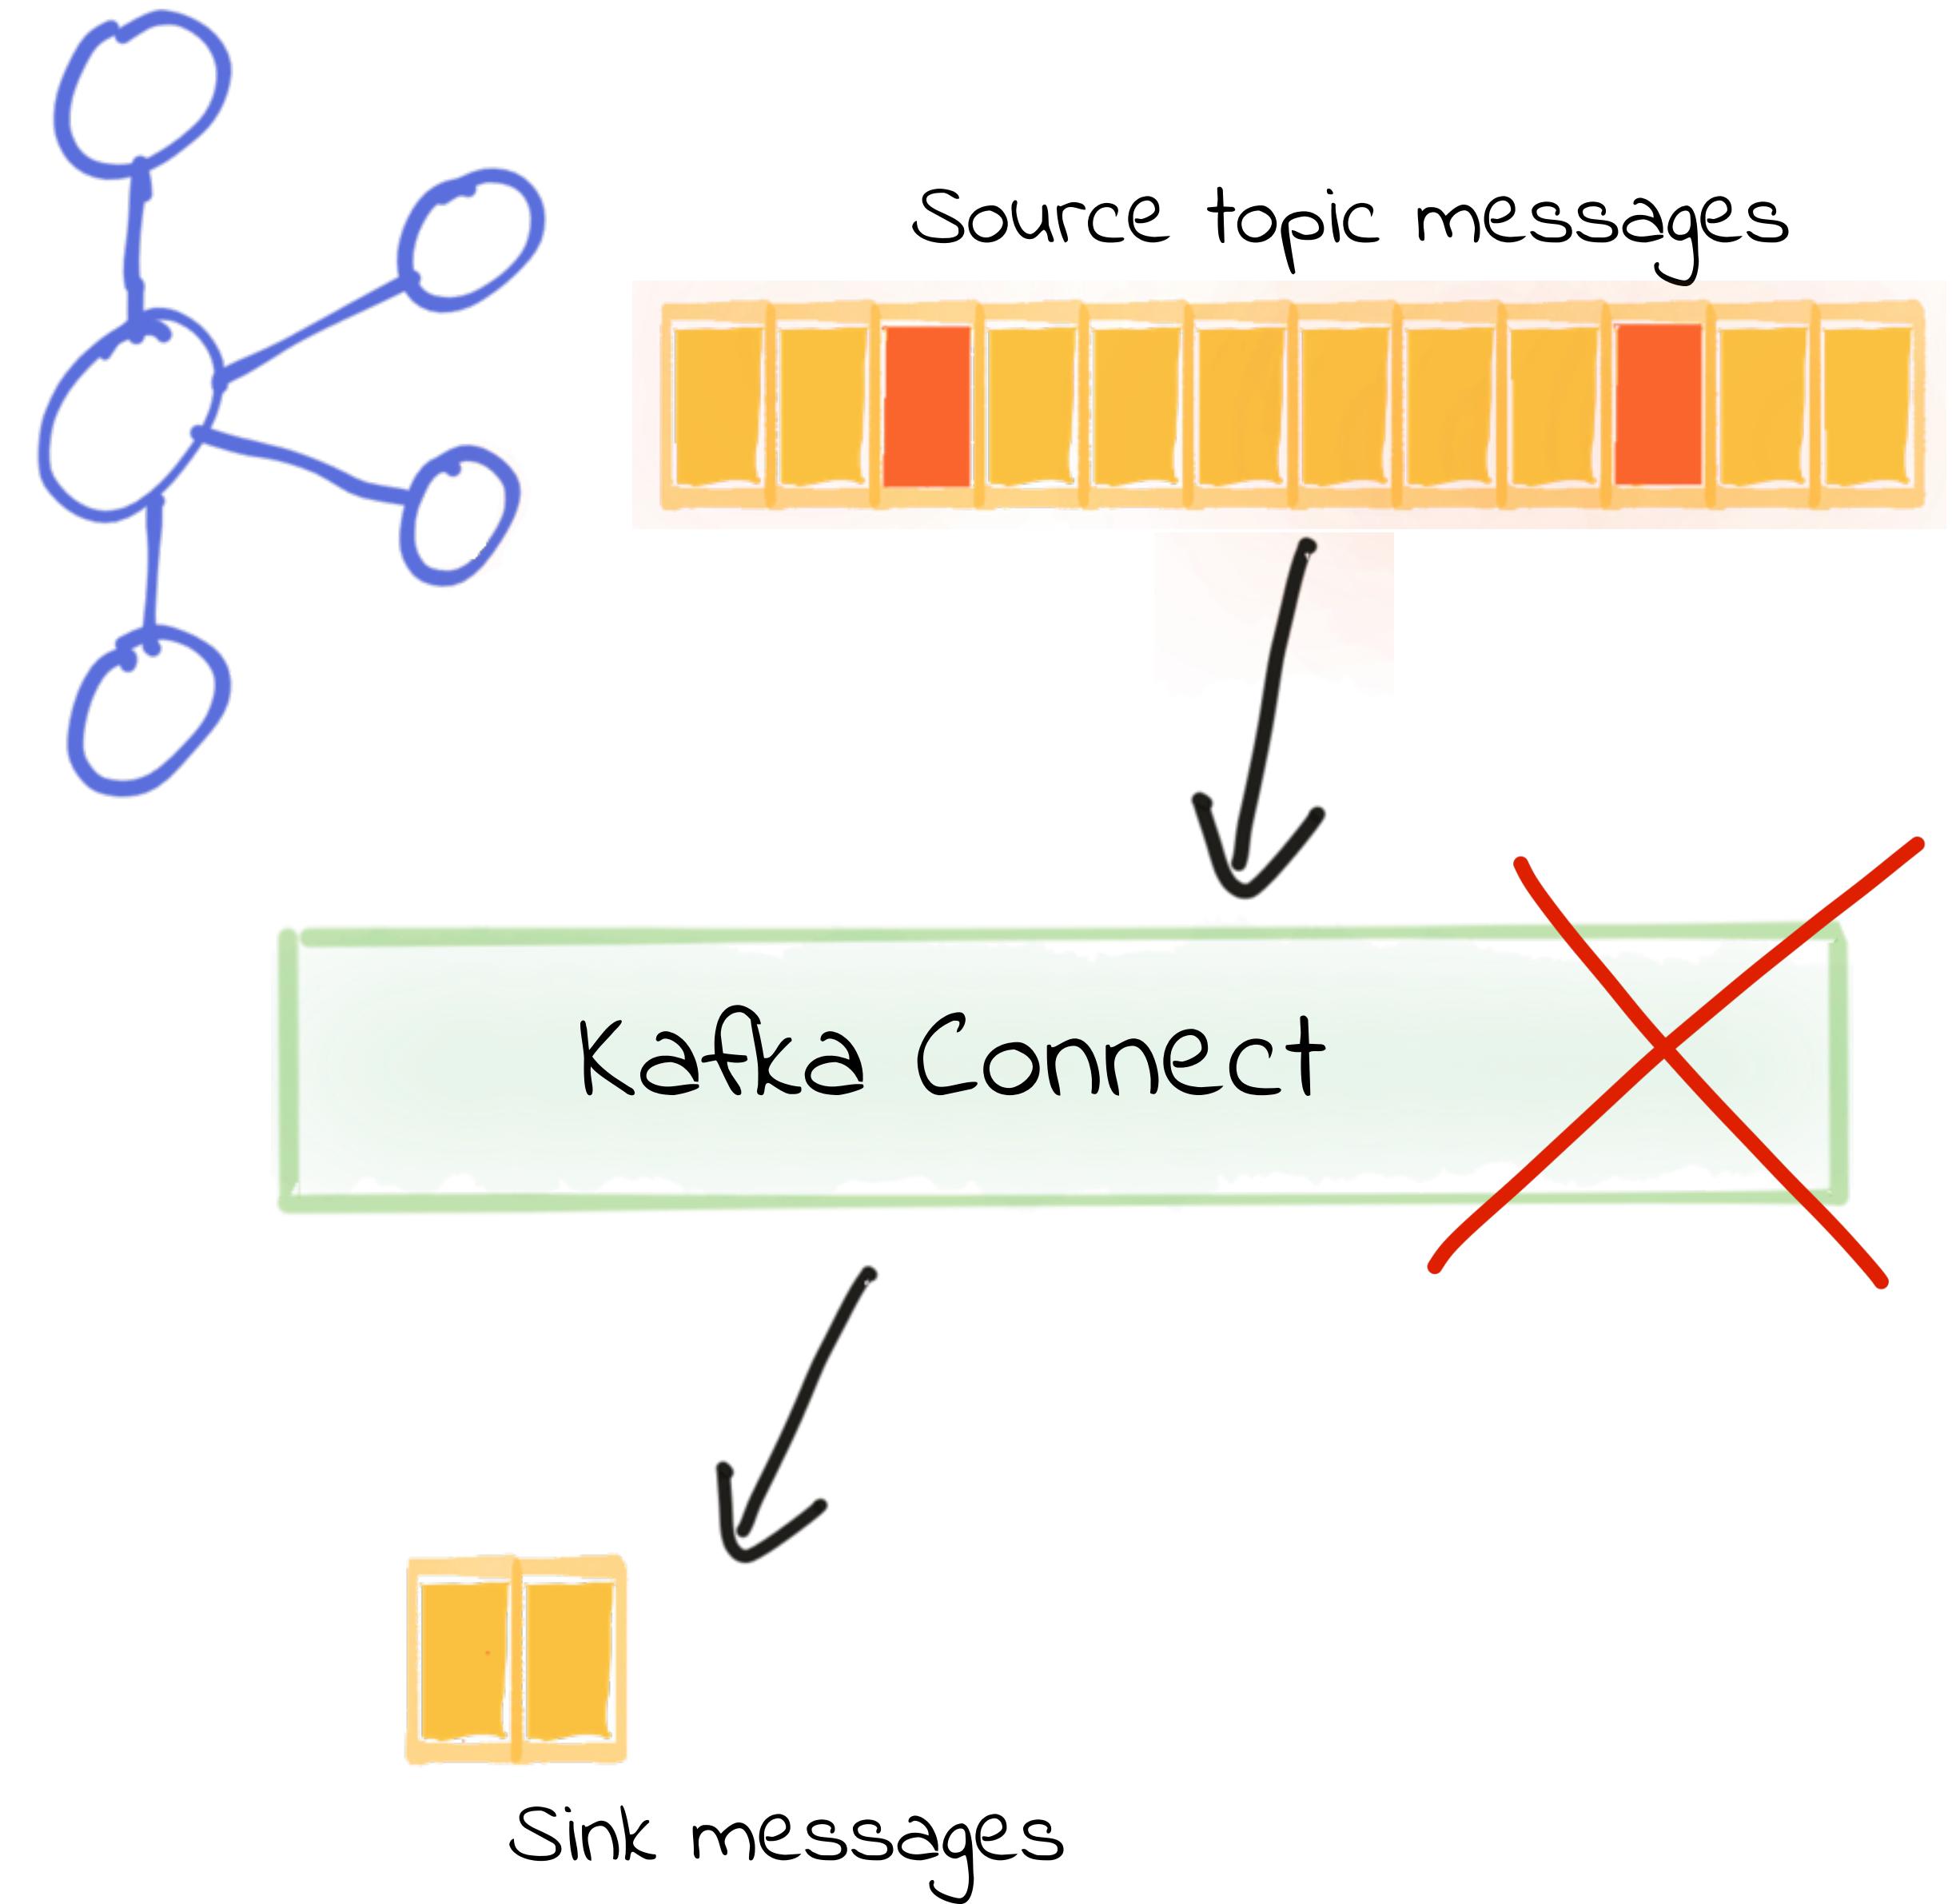
\* can be retried by Connect



<https://cnfl.io/connect-dlq>

 confluent

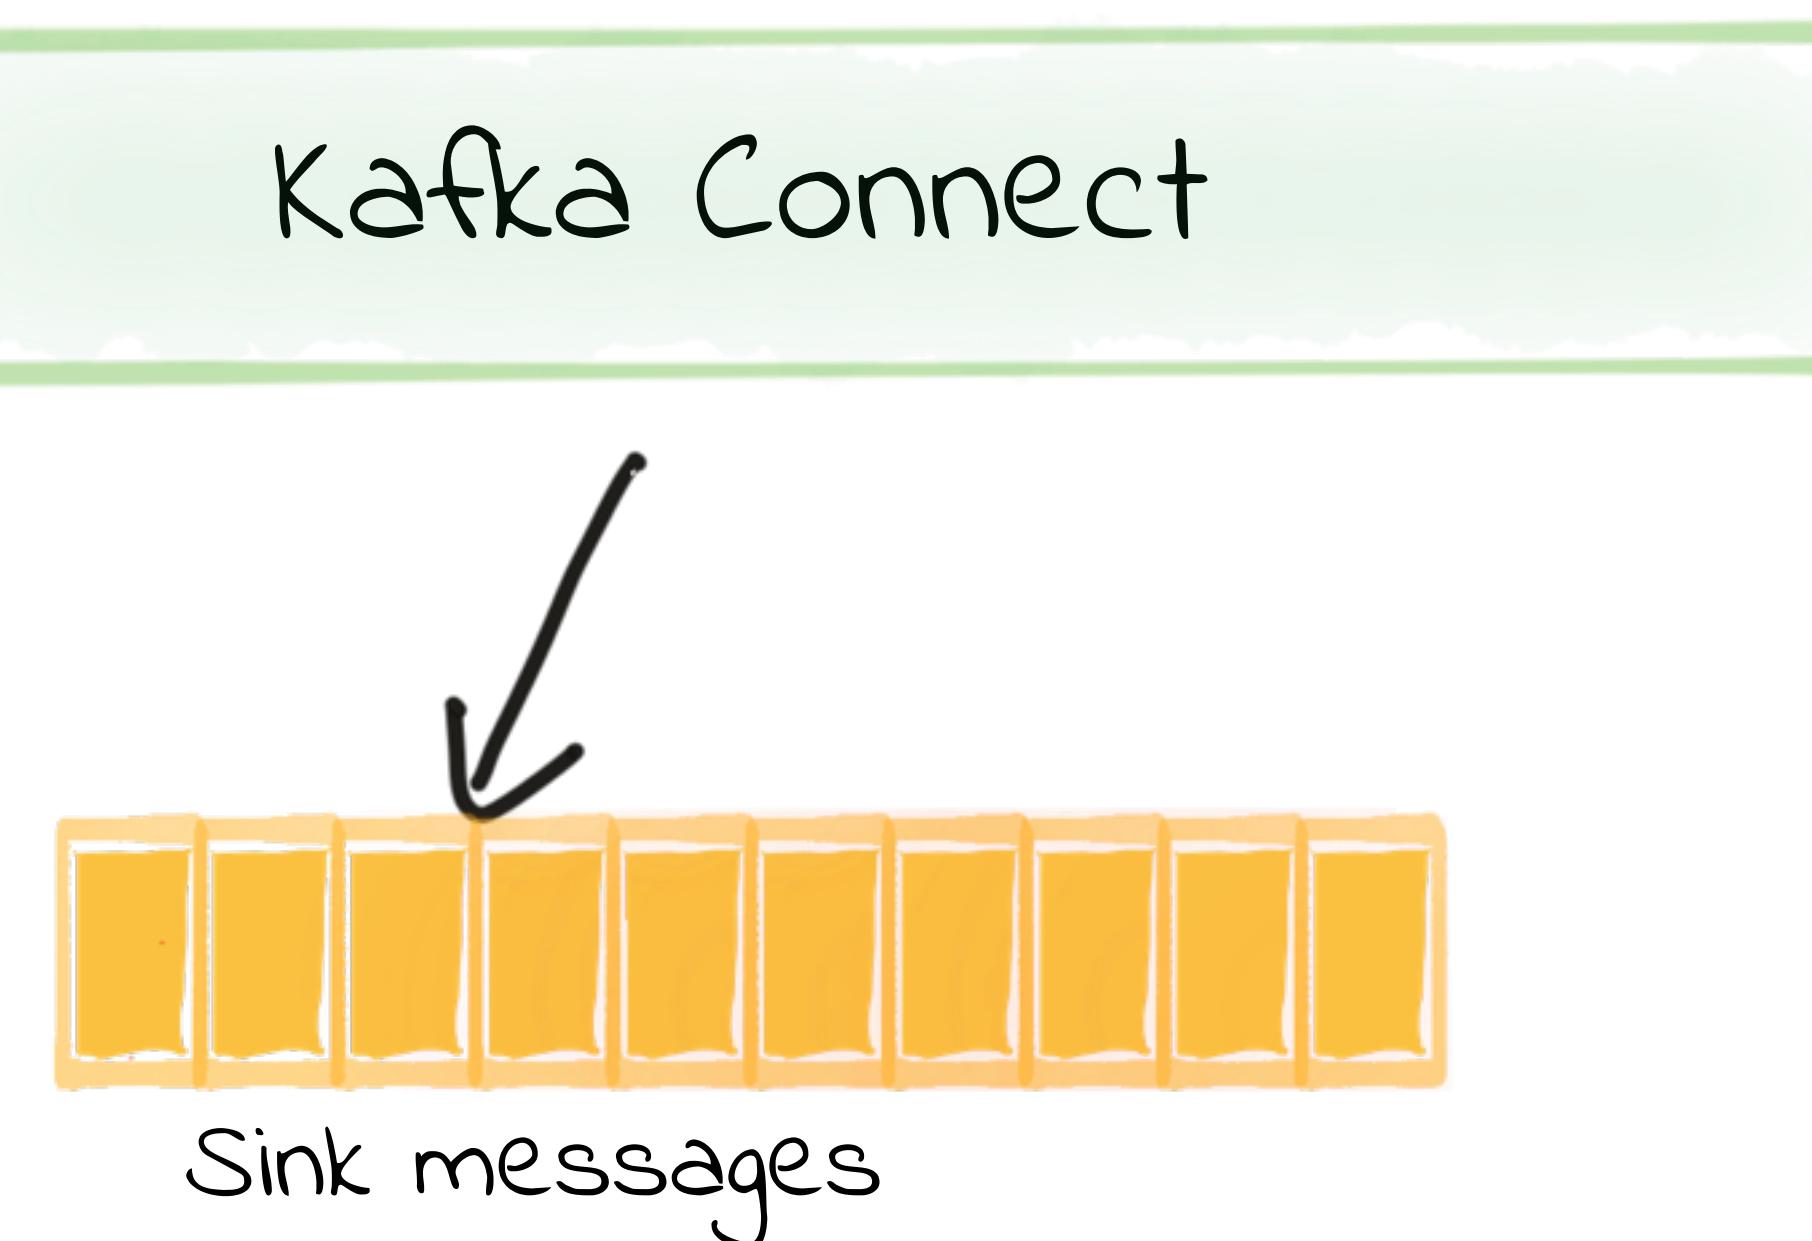
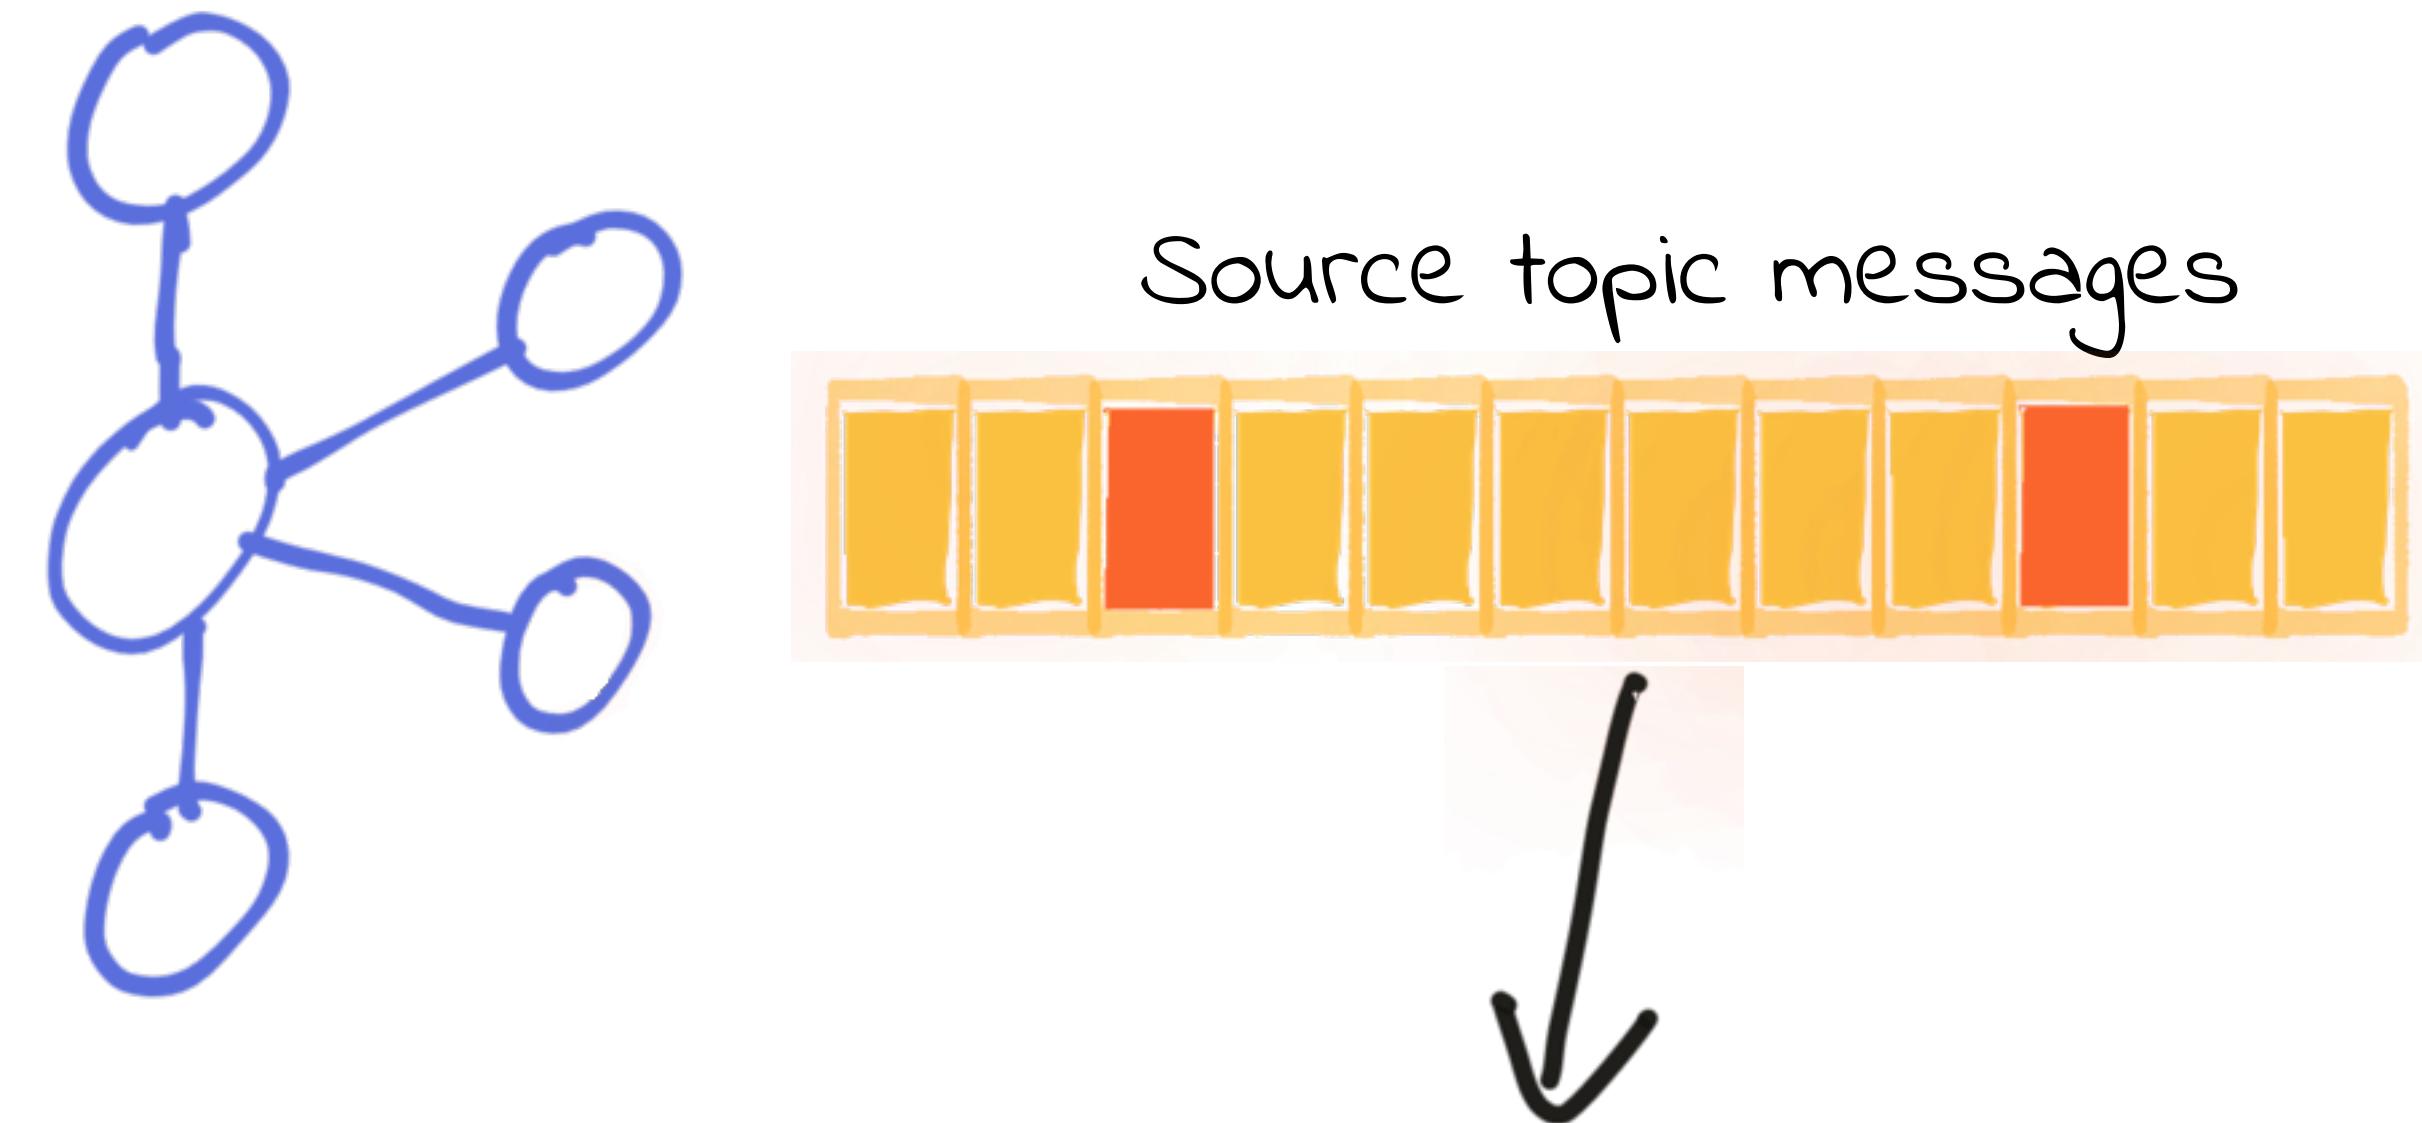
# Fail Fast



<https://cnfl.io/connect-dlq>

confluent

YOLO -\\_(ツ)\_/-



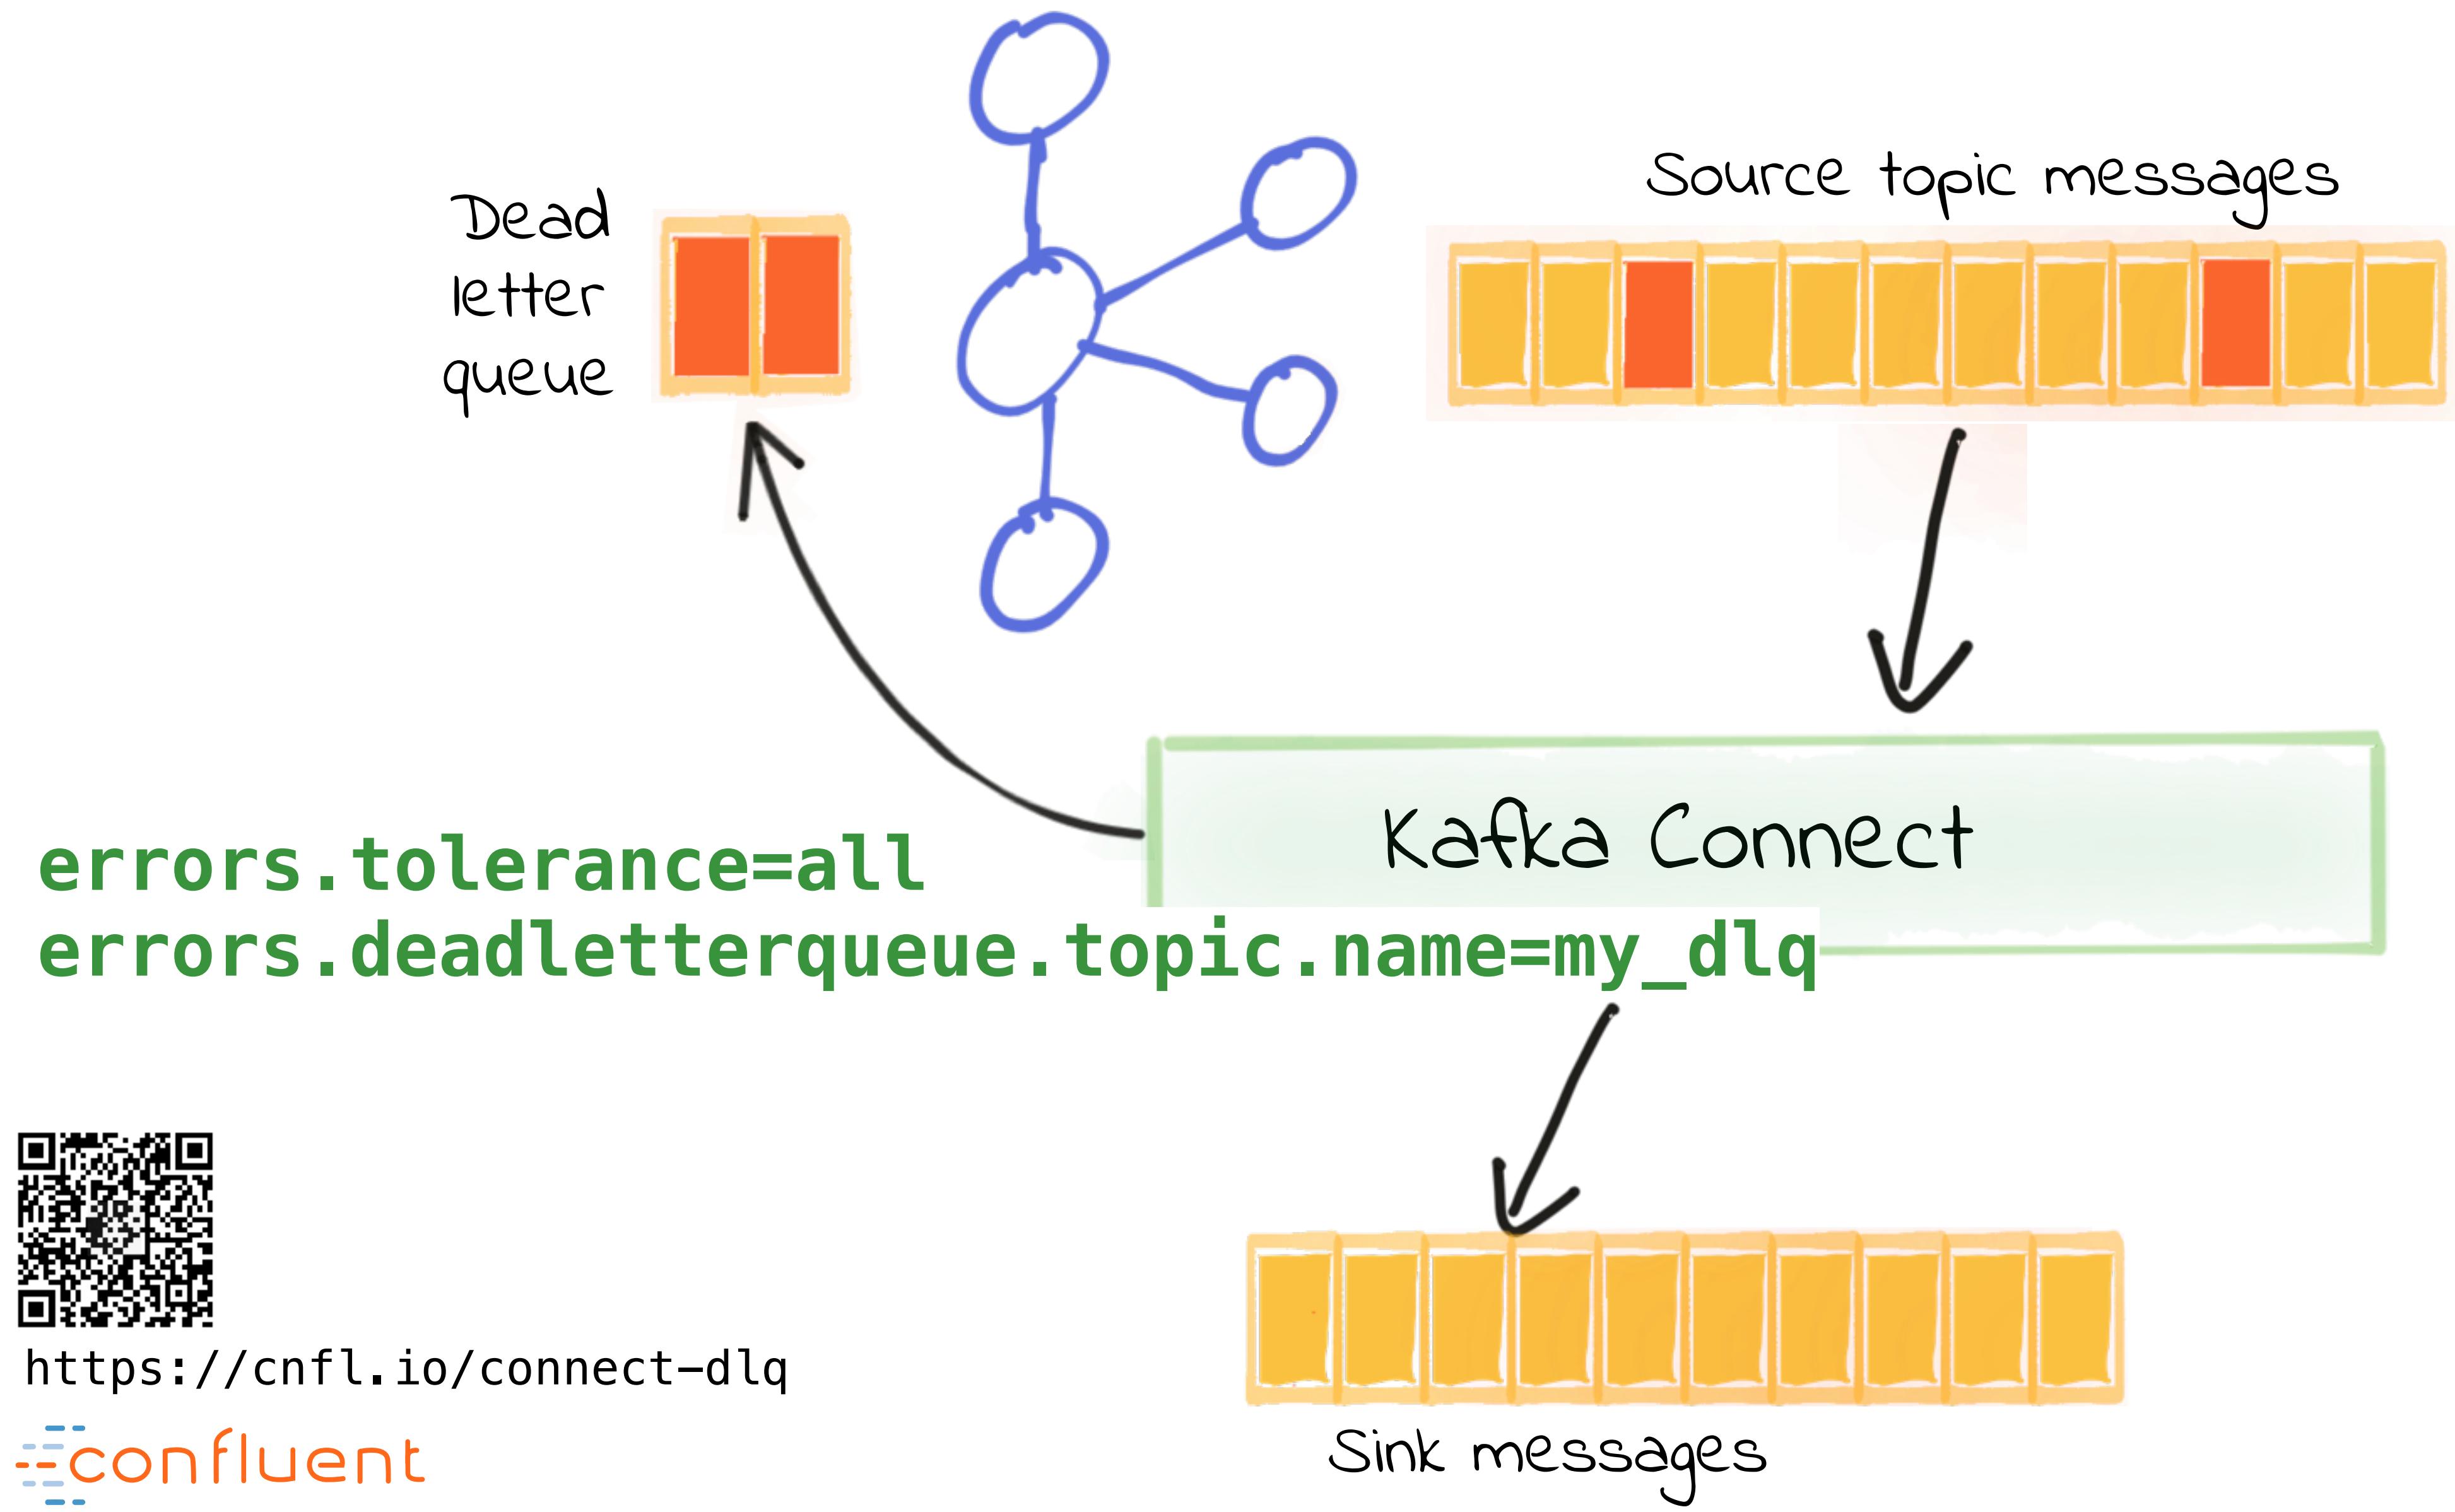
**errors.tolerance=all**



<https://cnfl.io/connect-dlq>

confluent

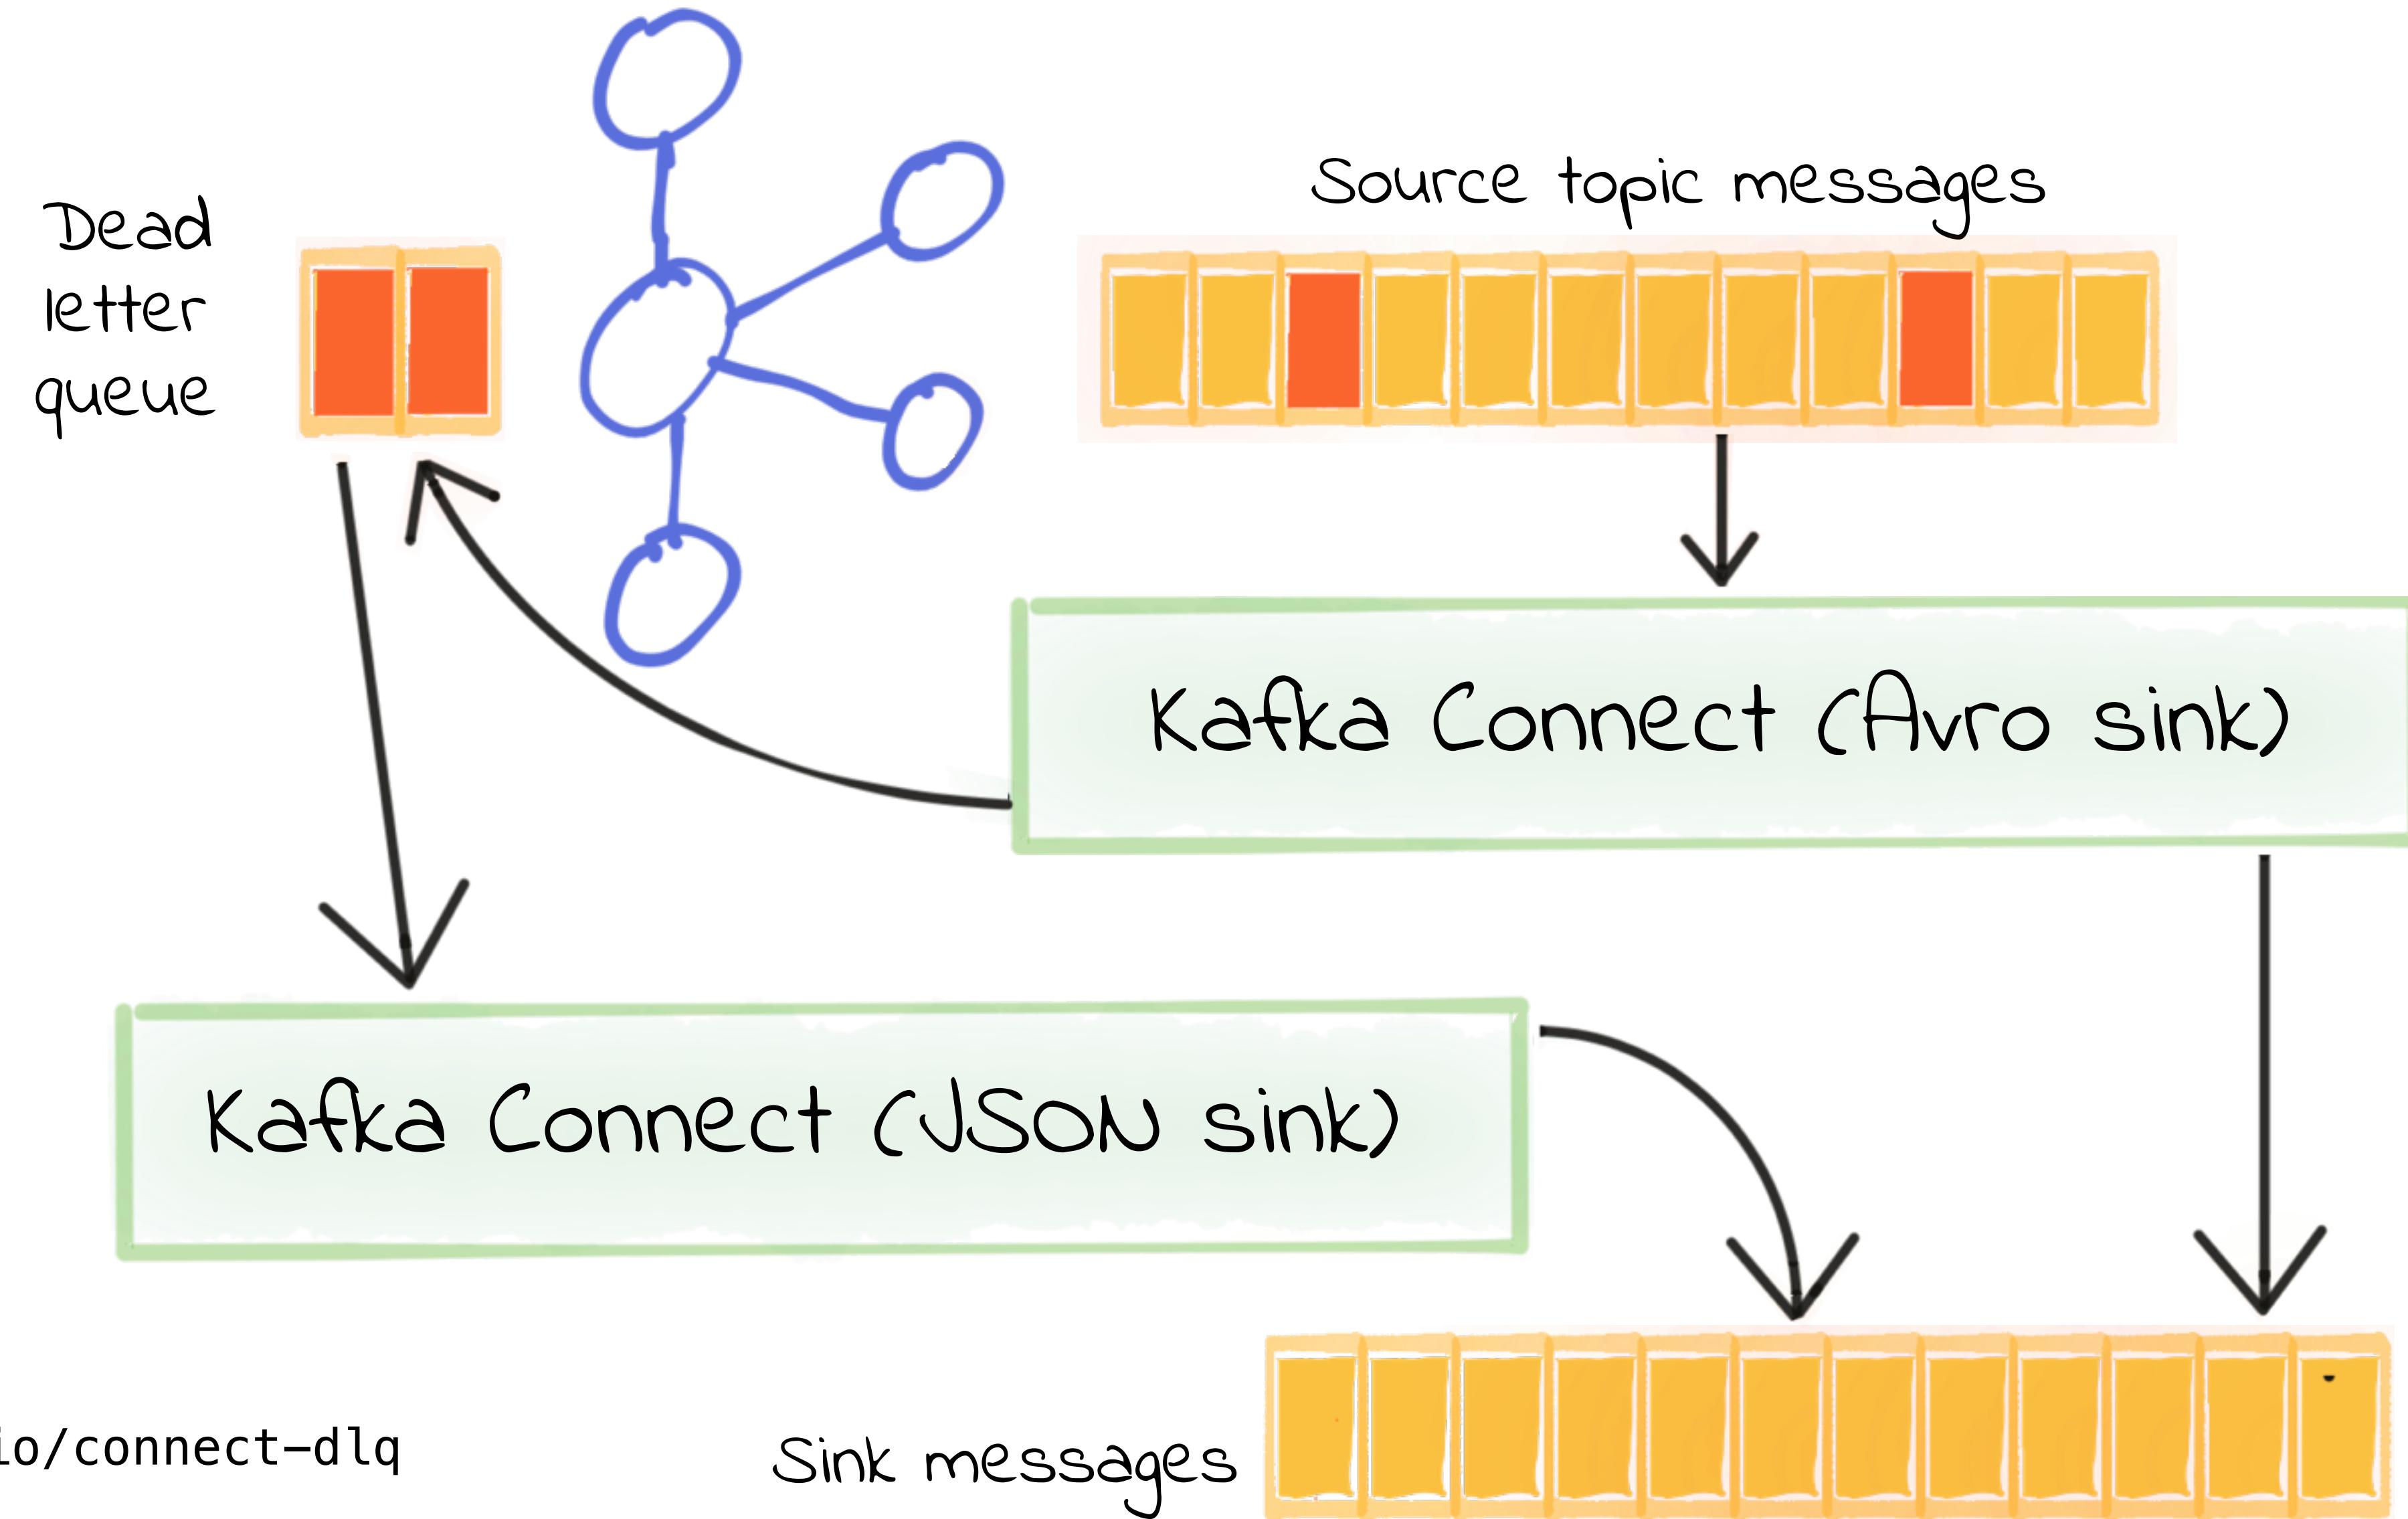
# Dead Letter Queue



<https://cnfl.io/connect-dlq>

confluent

# Re-processing the Dead Letter Queue



<https://cnfl.io/connect-dlq>

# Metrics and Monitoring

# REST API

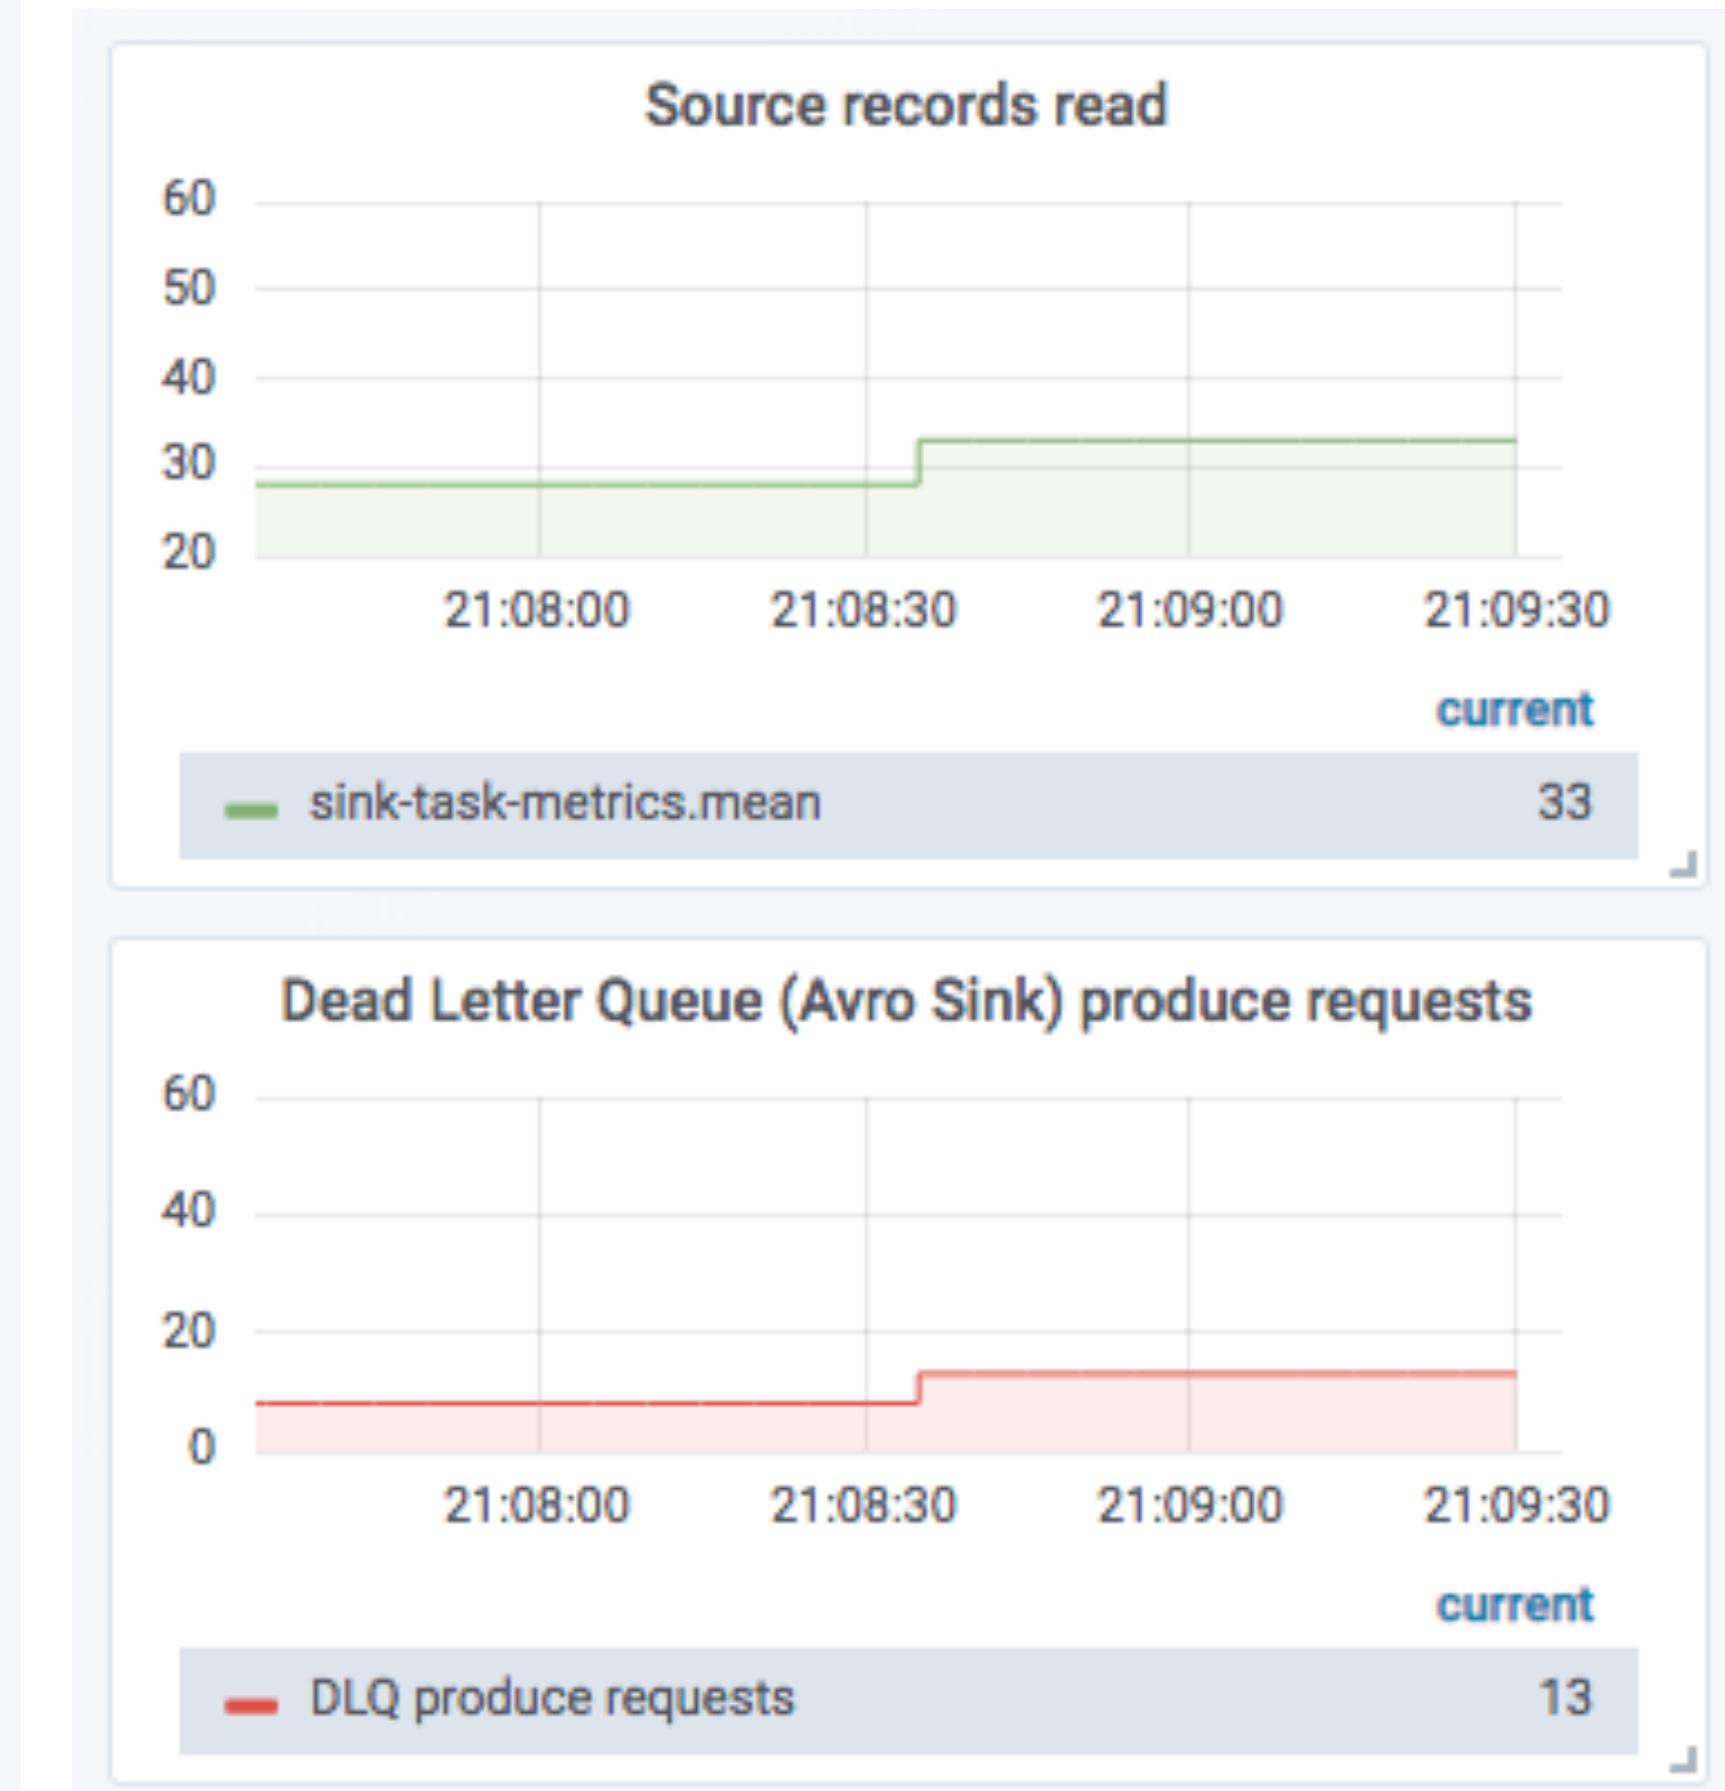
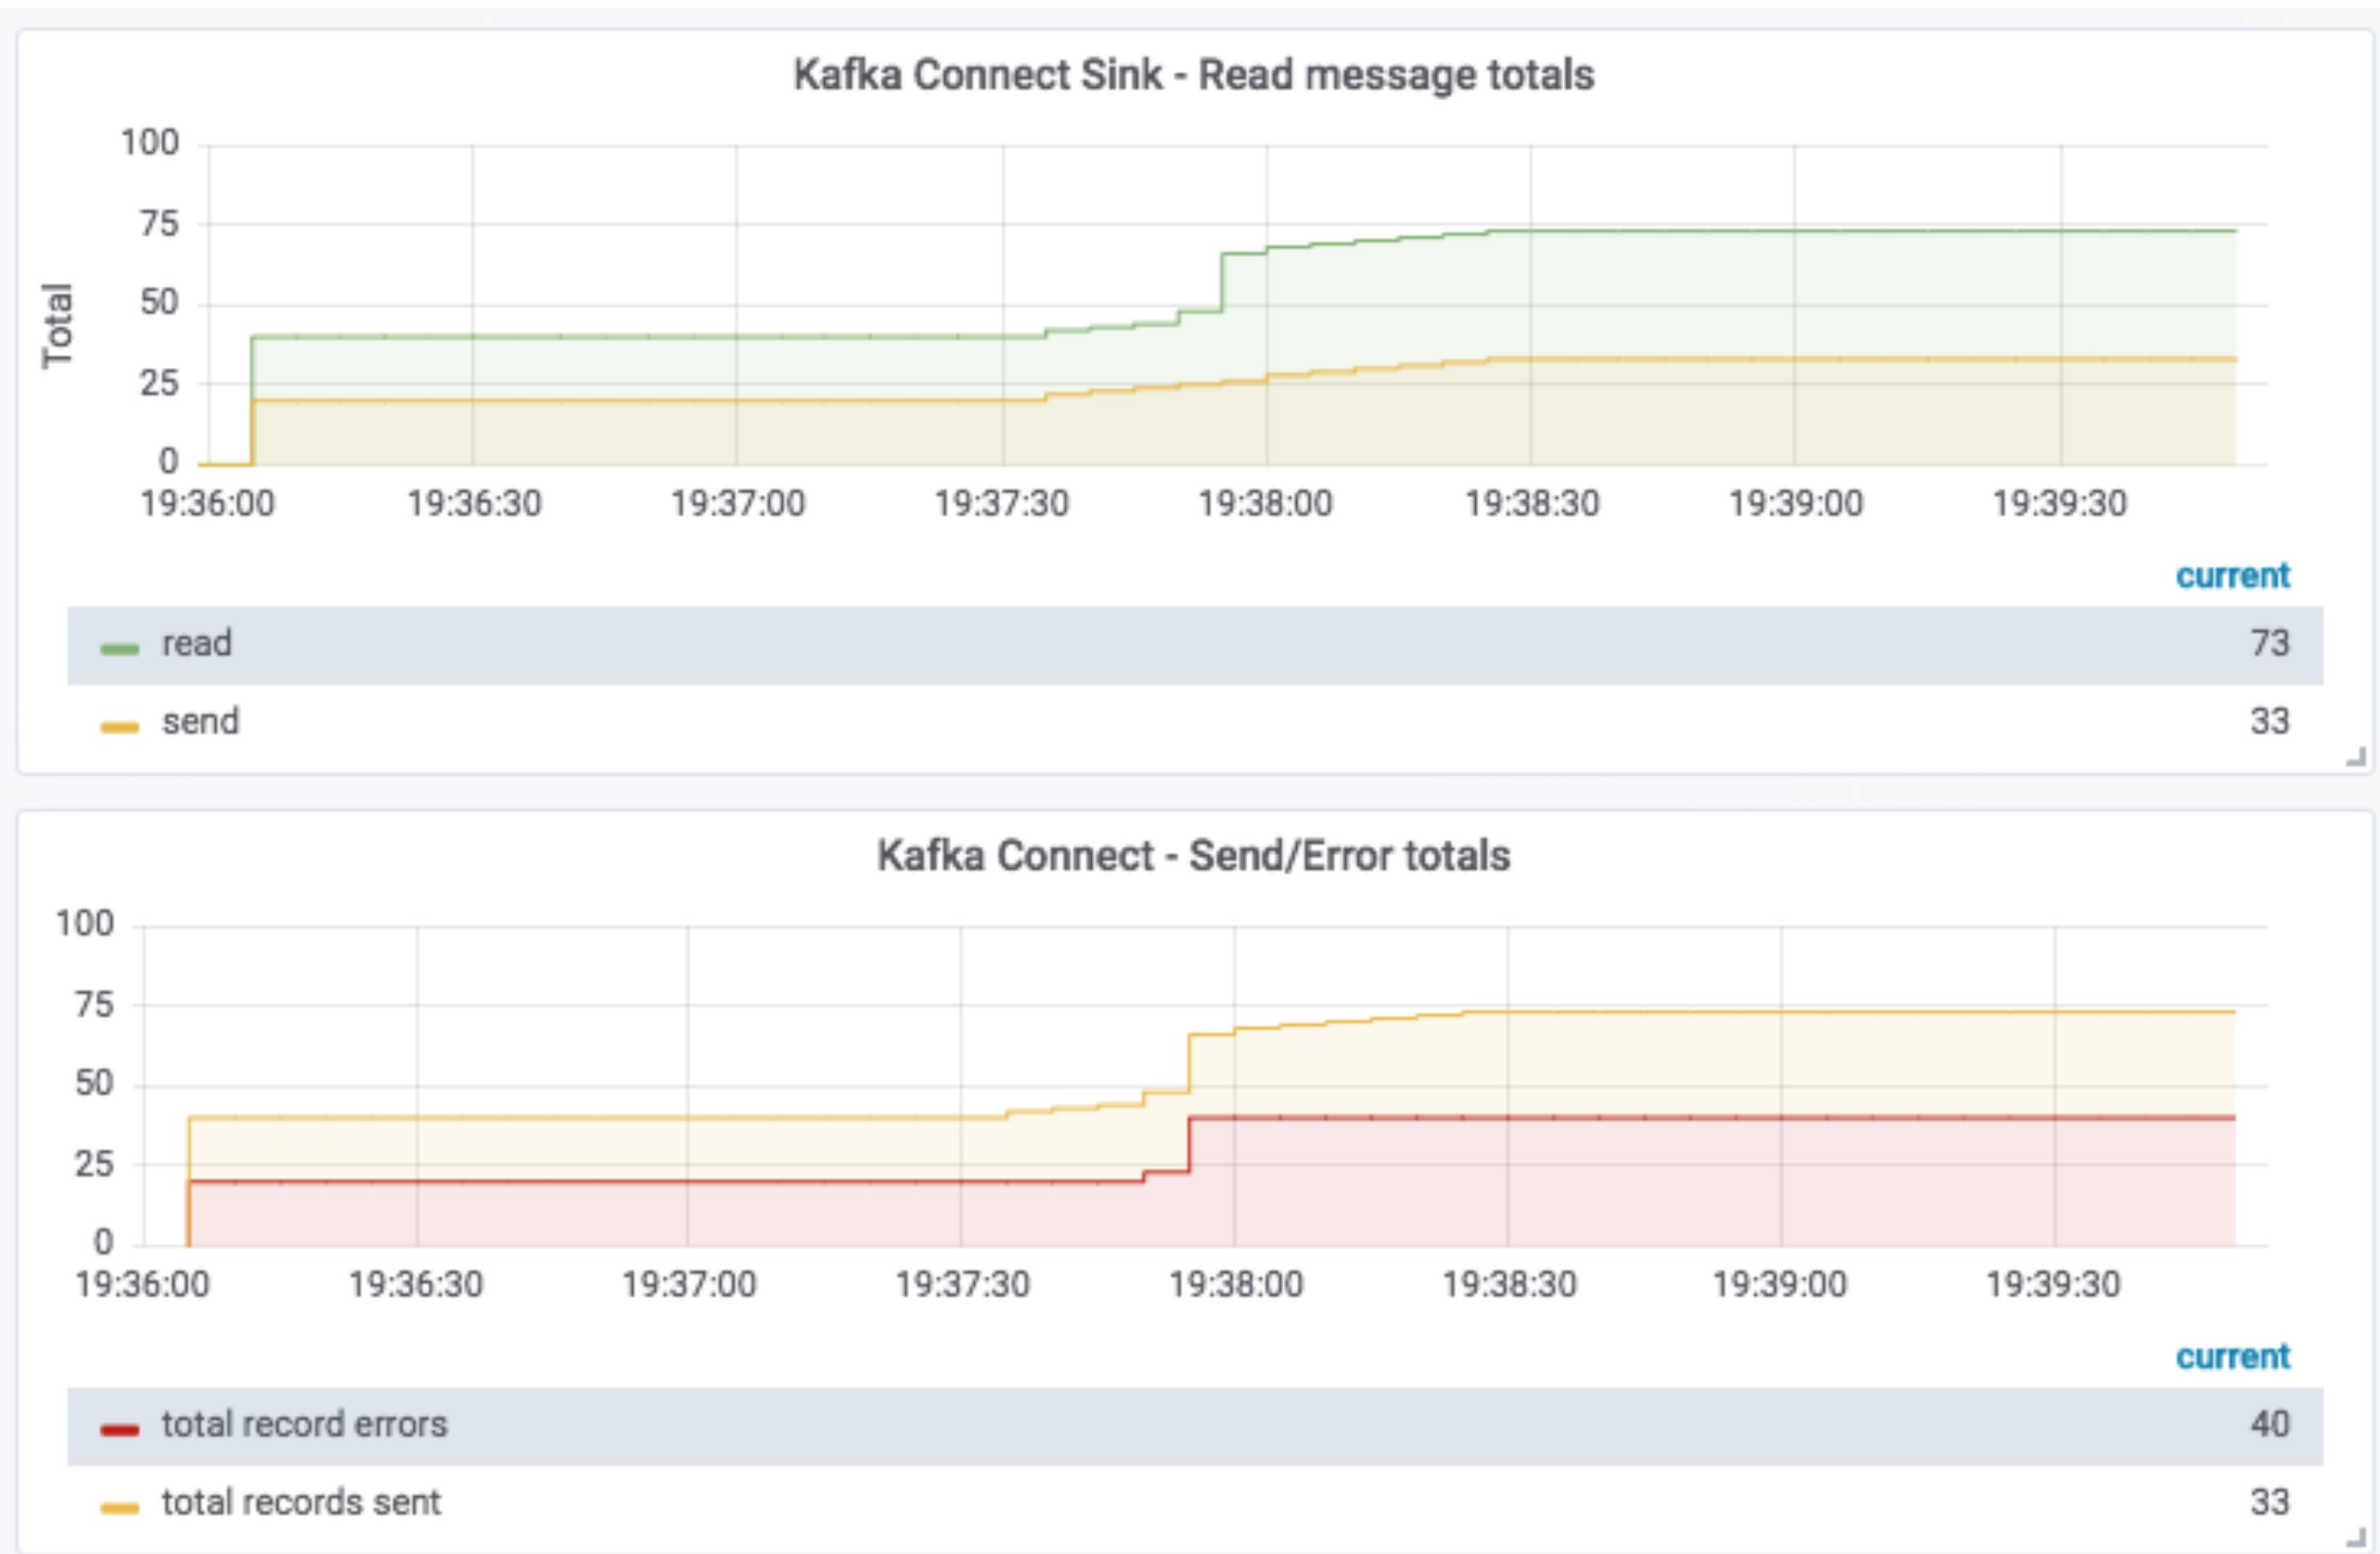
```
[~] Robin@asgard02-2.local
$ curl -s "http://localhost:8083/connectors" | \
  jq '.[]' | \
  xargs -I{connector_name} curl -s "http://localhost:8083/connectors/{connector_name}/status" | \
  jq -c -M '[.name,.connector.state,.tasks[].state]|join(":|:")' | \
  column -s : -t | \
  sed 's/\"//g' | \
  sort

sink-elastic-orders-00      | RUNNING    | RUNNING
sink-elastic-orders-01      | RUNNING    | RUNNING
source-debezium-orders-00   | RUNNING    | RUNNING
[~] Robin@asgard02-2.local
$
```

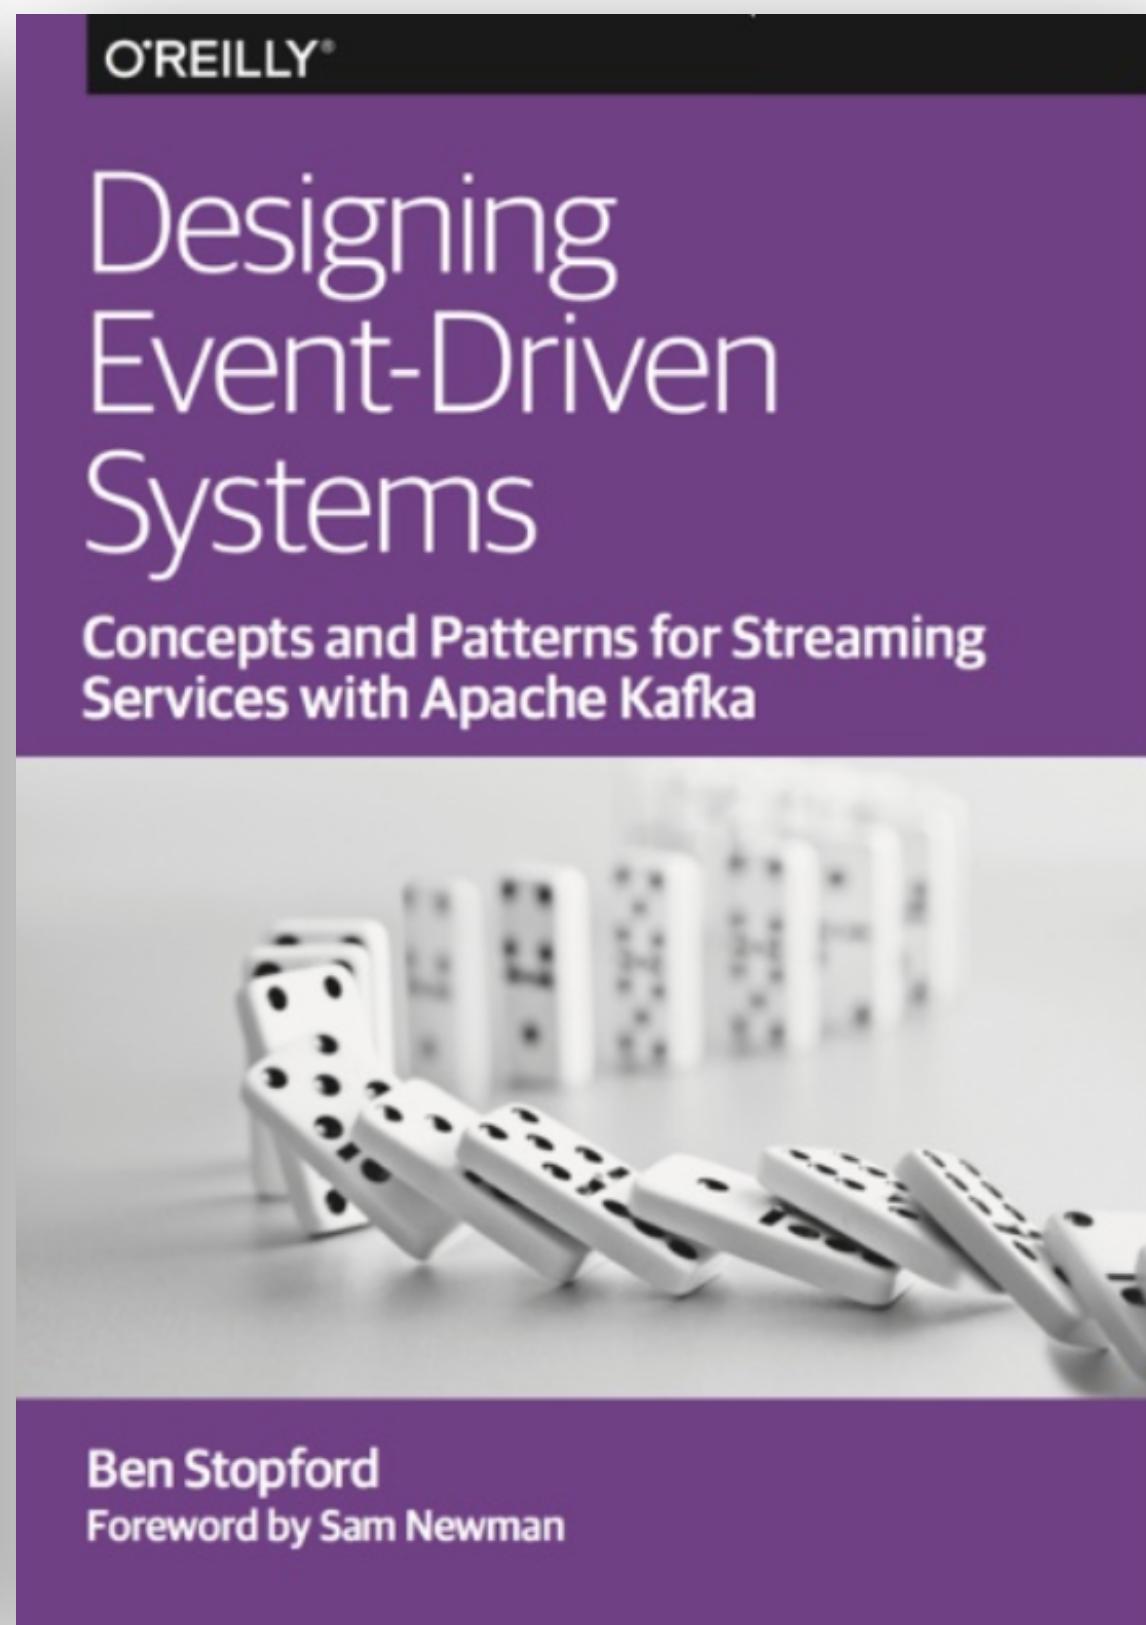
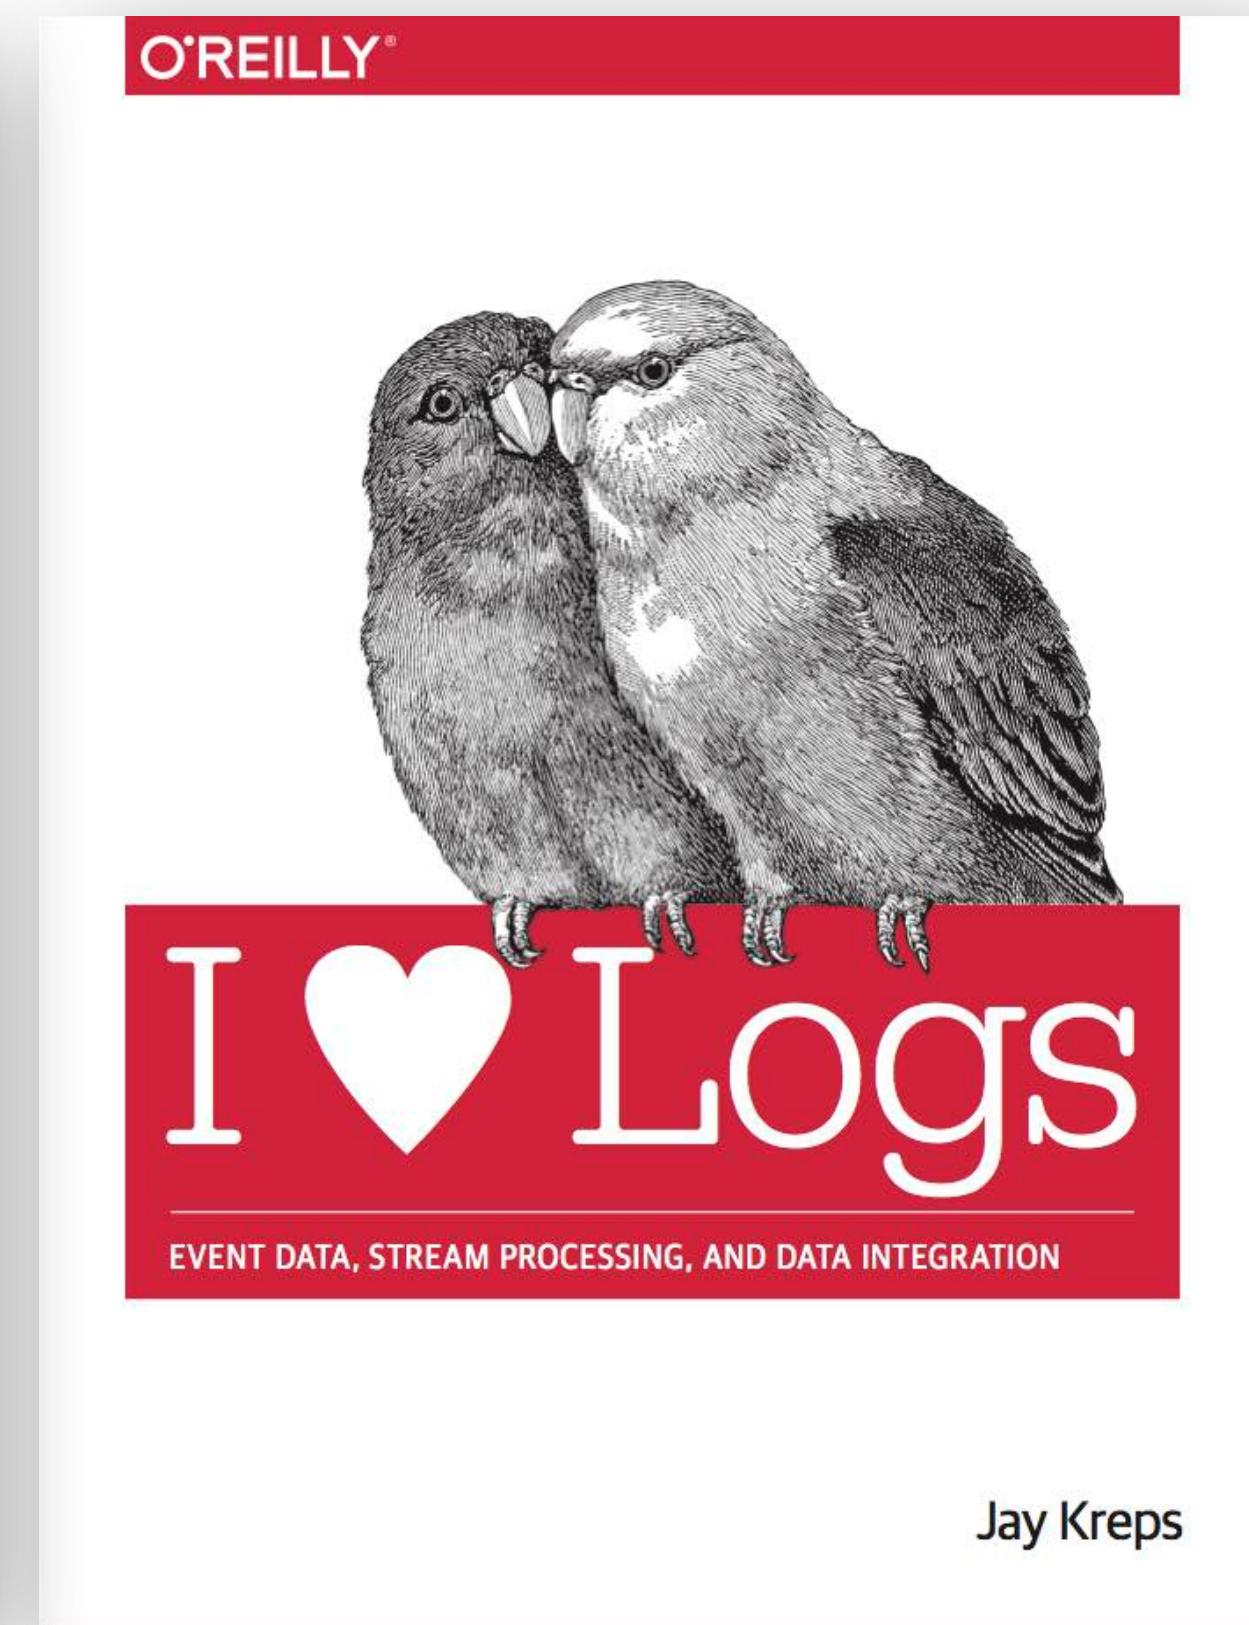
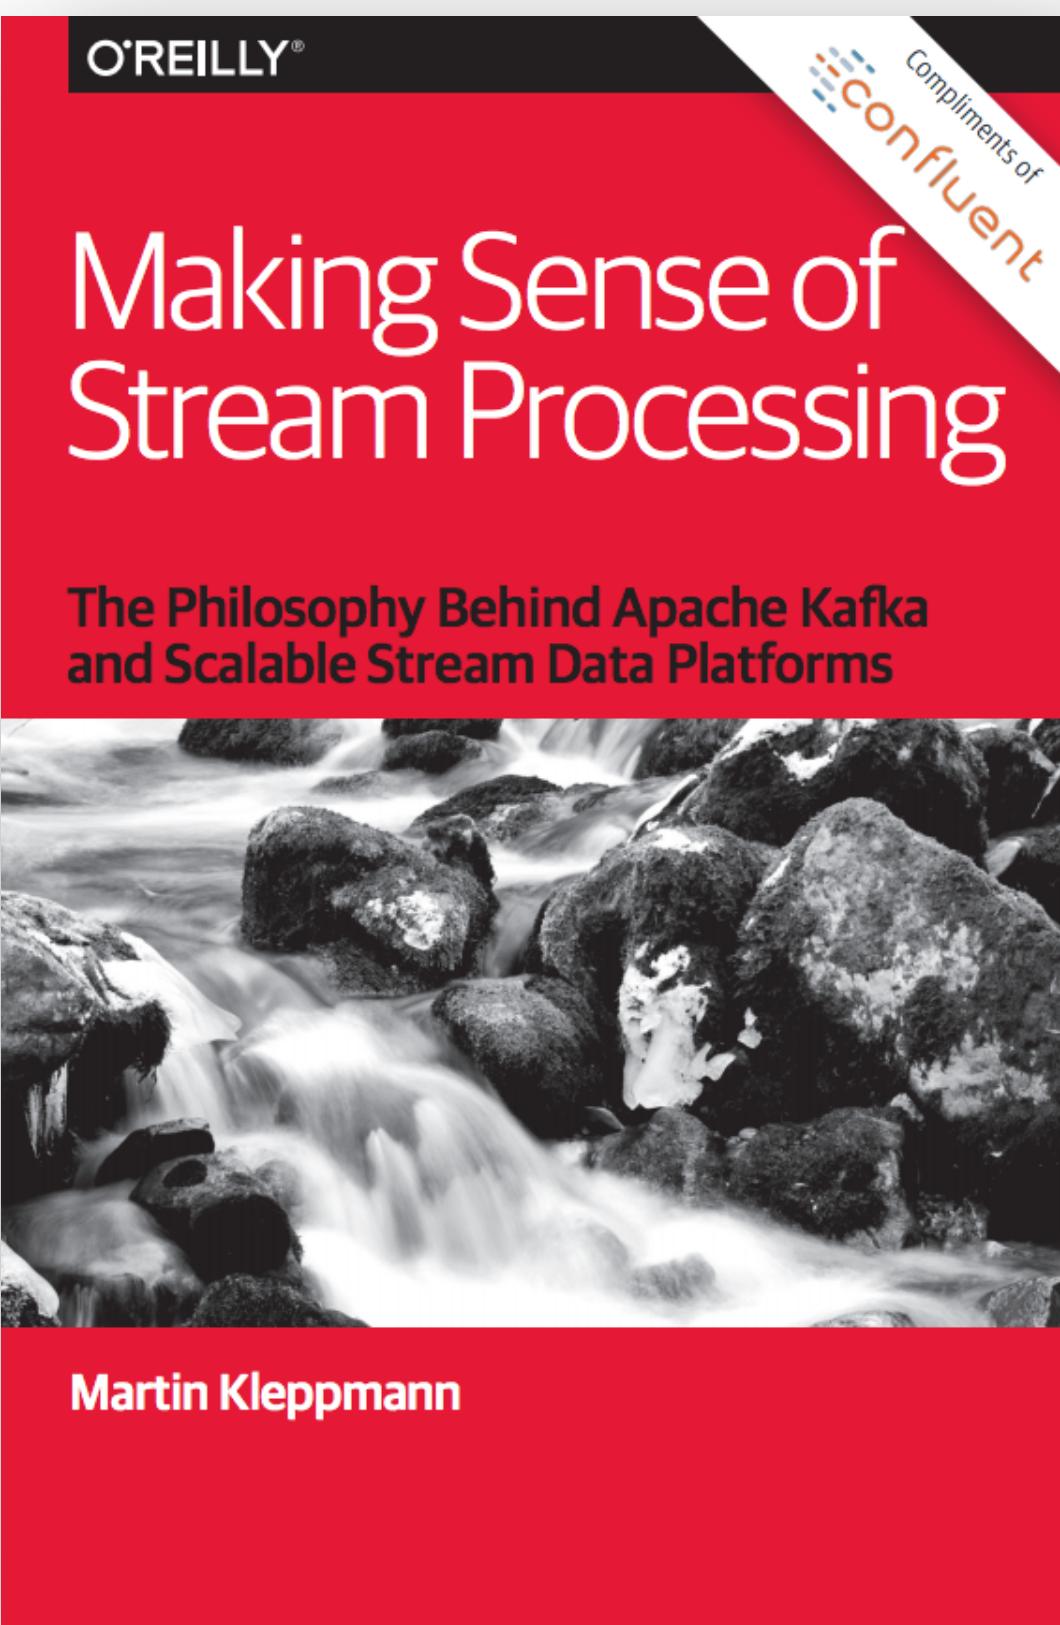
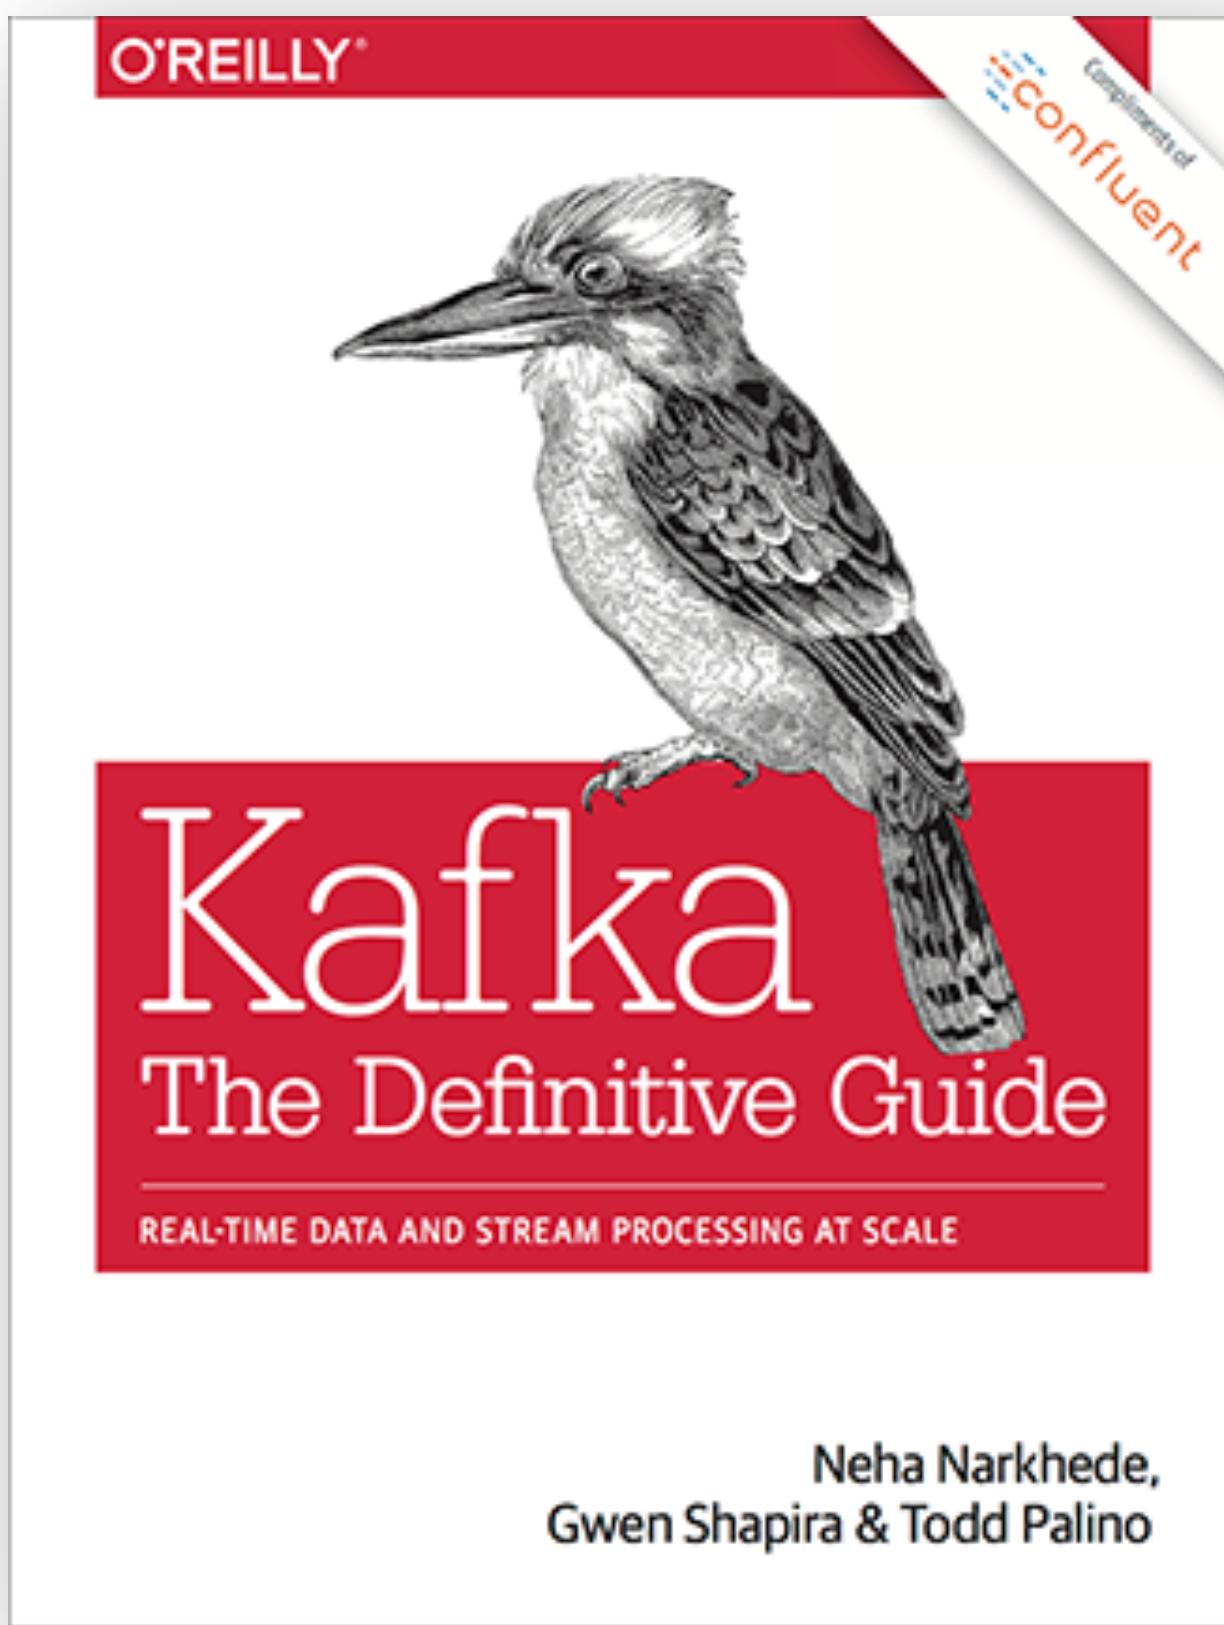


<http://go.rmoff.net/connector-status>

# JMX



# <http://cnfl.io/book-bundle>





# **Confluent Community**

## **Slack group**

**<http://cnfl.io/slack>**

@rmoff



<http://talks.rmoff.net/>