

O'REILLY®

Java Concurrency



About Me – Career

- New Relic, Principal Engineer
- jClarity, Co-founder
 - Sold to Microsoft
- Deutsche Bank
 - Chief Architect (Listed Derivatives)
- Morgan Stanley
 - Google IPO
- Sporting Bet
 - Chief Architect



About Me – Community

- Java Champion
- JavaOne Rock Star Speaker
- Java Community Process Executive Committee
- London Java Community
 - Organising Team
 - Co-founder, AdoptOpenJDK



Installing Java

- We will use AdoptOpenJDK - Java 8
<https://adoptopenjdk.net/>
- We'll use “OpenJDK 8 (LTS)” and the HotSpot JVM (or Java 11)
- Download this repo from GitHub:
- <https://github.com/kittylst/optimizing-java>
- The examples are in the repo
- My personal site (articles, links): <https://www.kittylst.com>



Synchronization

- Simple Hardware Model
- Exercise: Effect of Hardware Caching
- The Need for Synchronization
- Exercise: Counters
- Java's Original Thread & Synchronization API
- Synchronization Summary
- Q&A



Simple Hardware Model

- Mechanical Sympathy
- Interesting Graphs and their Consequences
- The Modern CPU

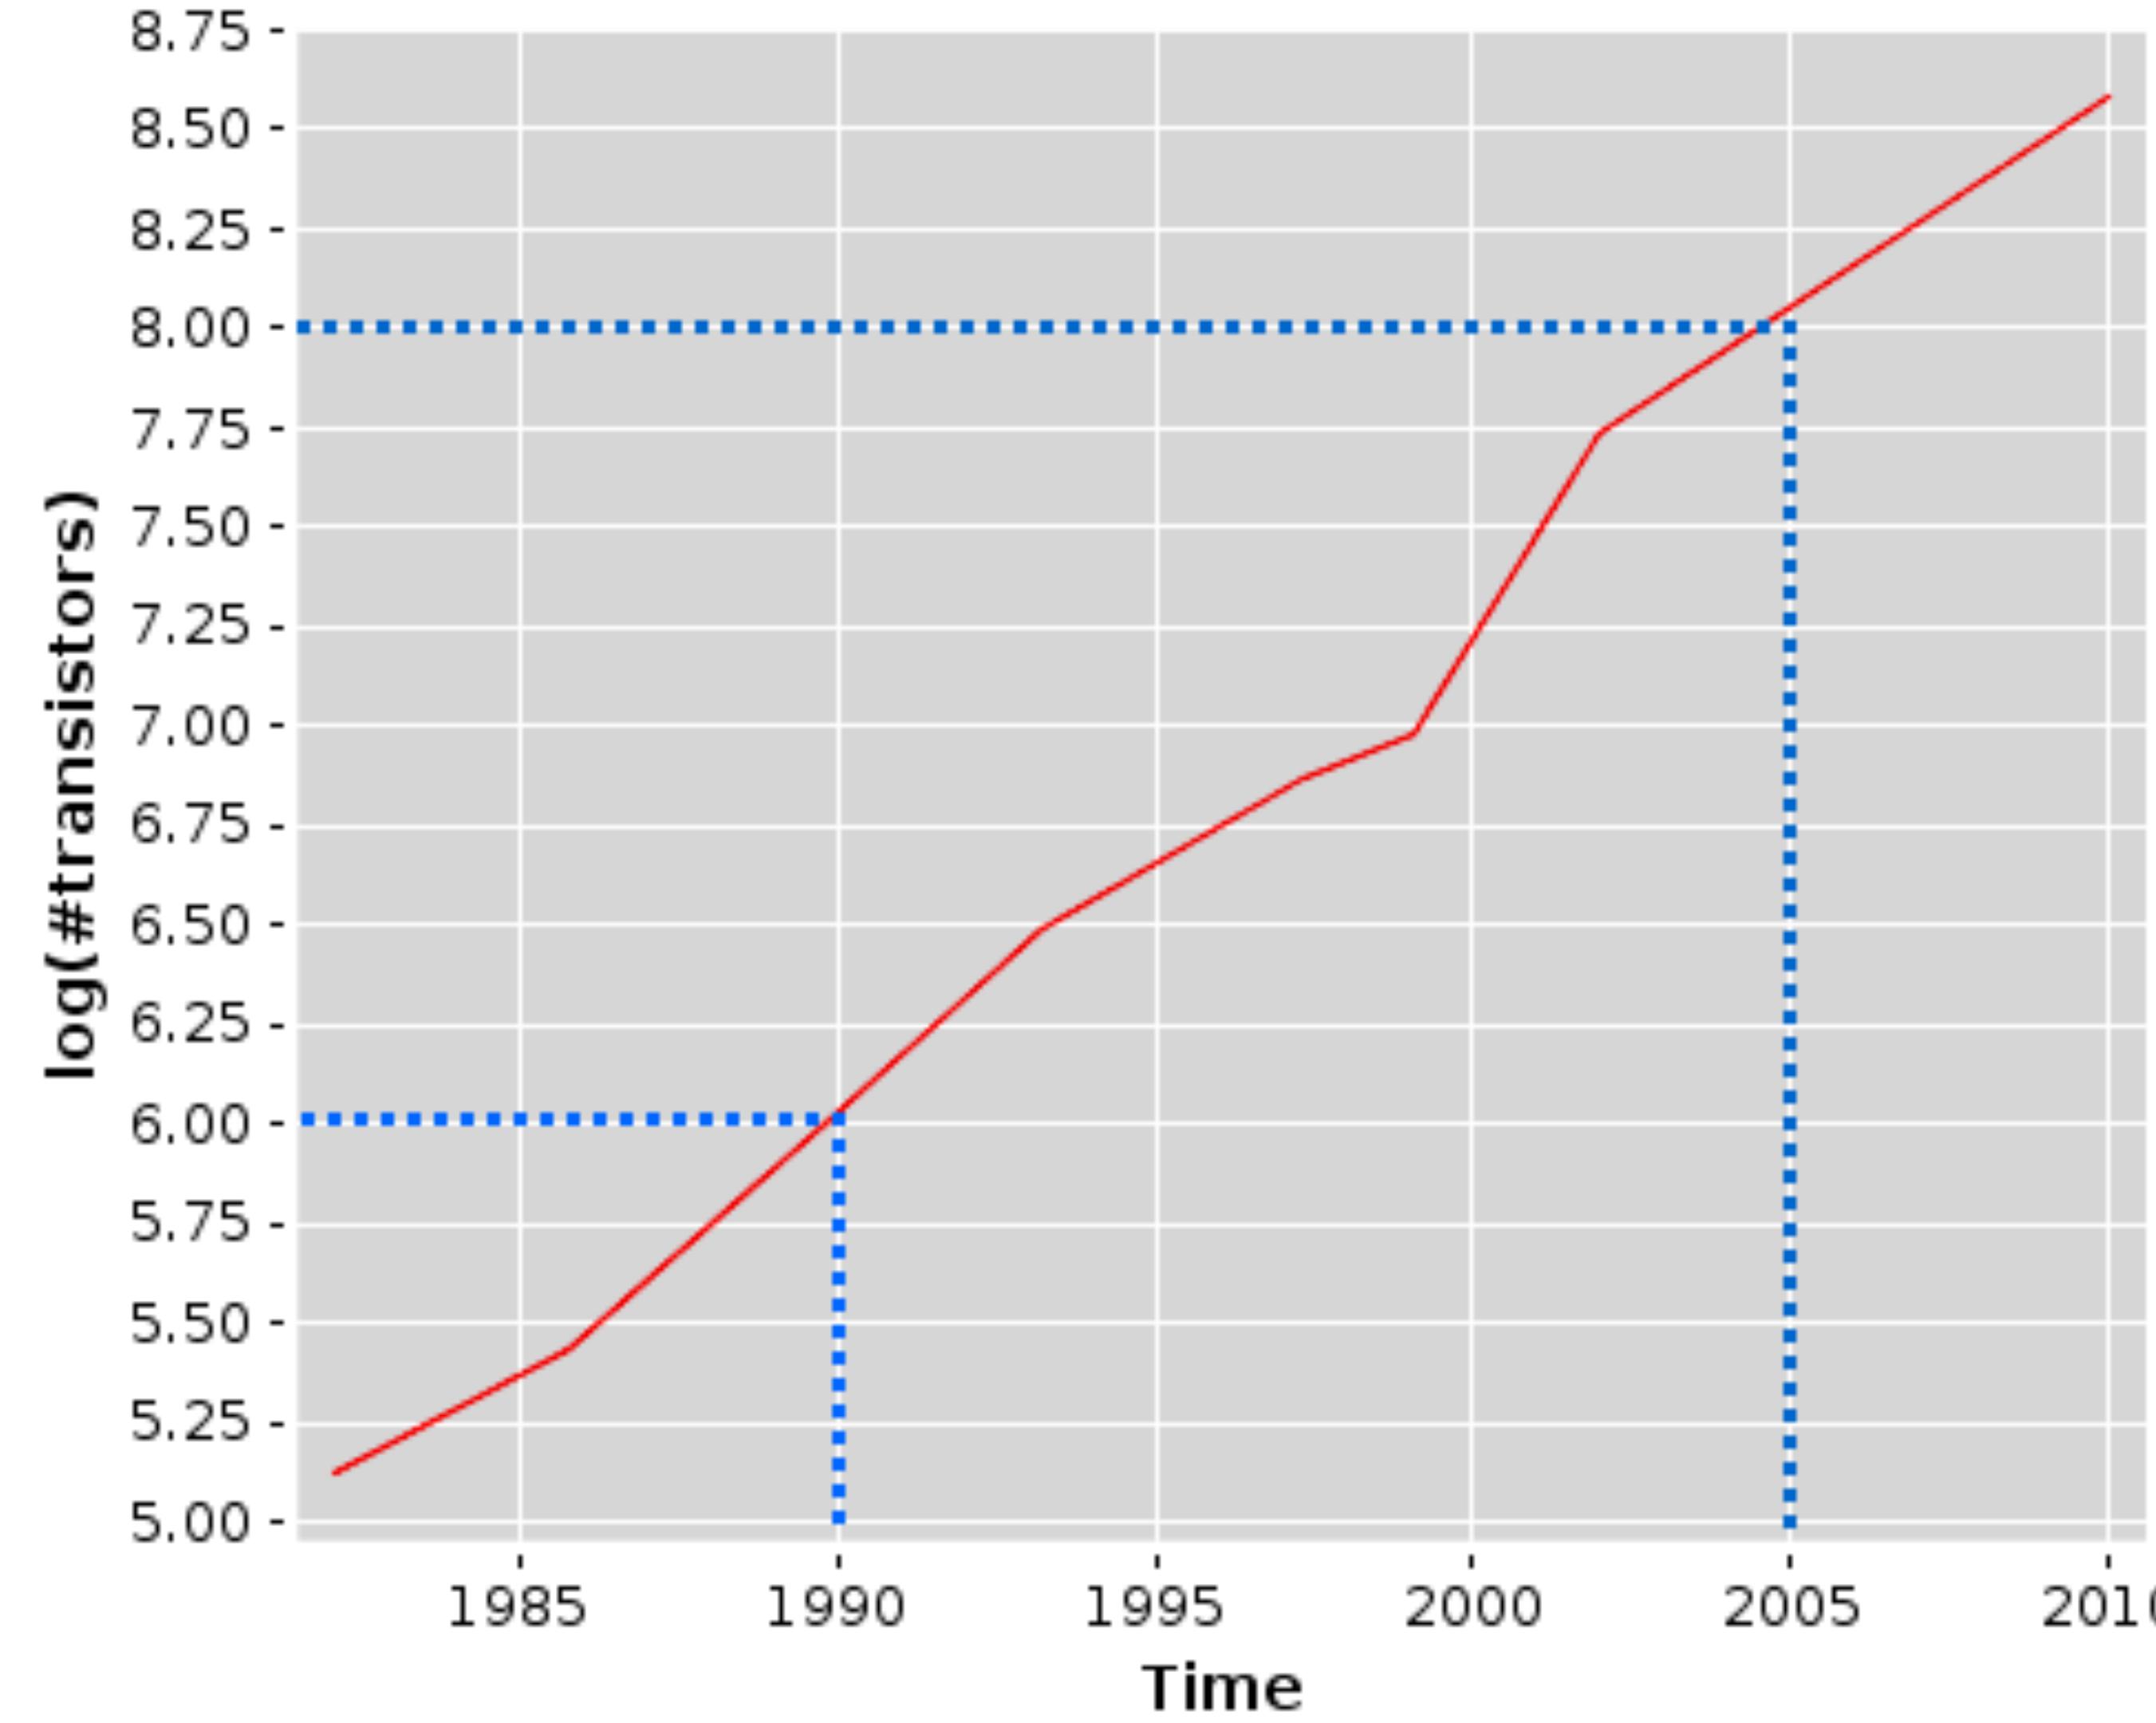


Mechanical Sympathy

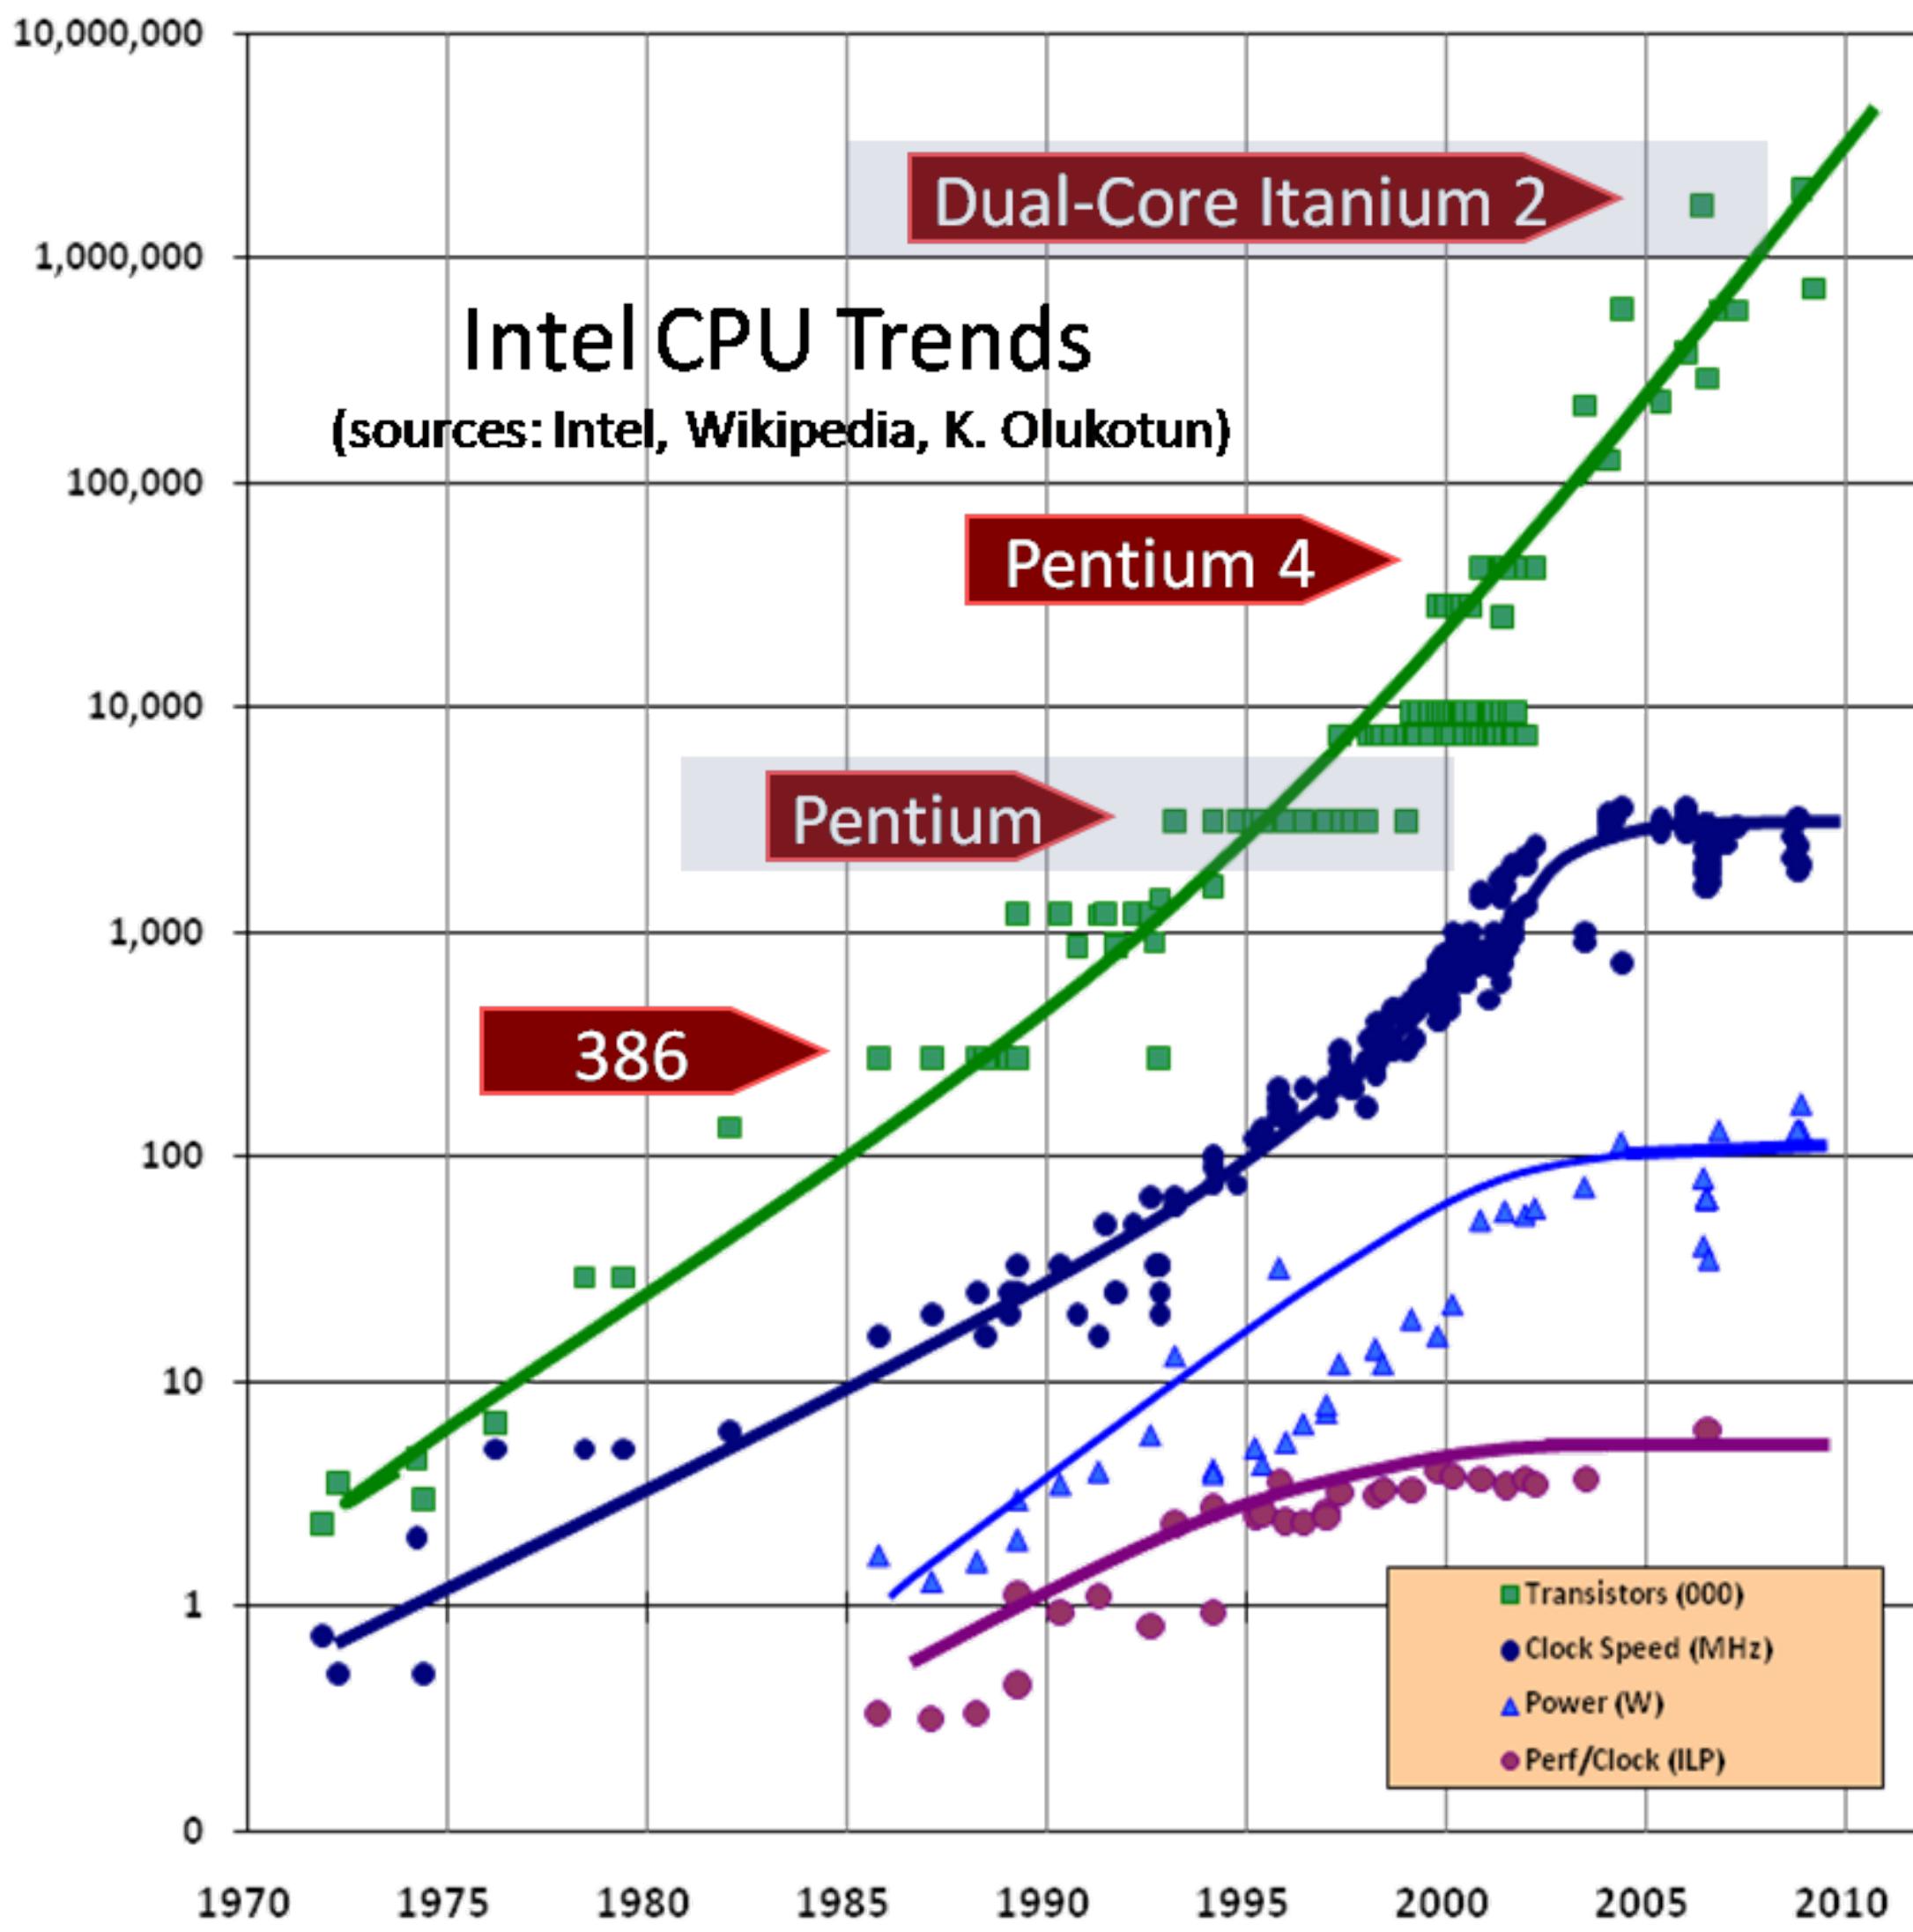
- Named for a quote by Jackie Stewart (F1 driver)
- Understanding the underlying hardware helps
- Pioneered at LMAX by Martin Thompson
 - See also: Disruptor Pattern, Trisha Gee, Mike Barker
- Performance engineers need to understand hardware
 - At least, more than most developers do



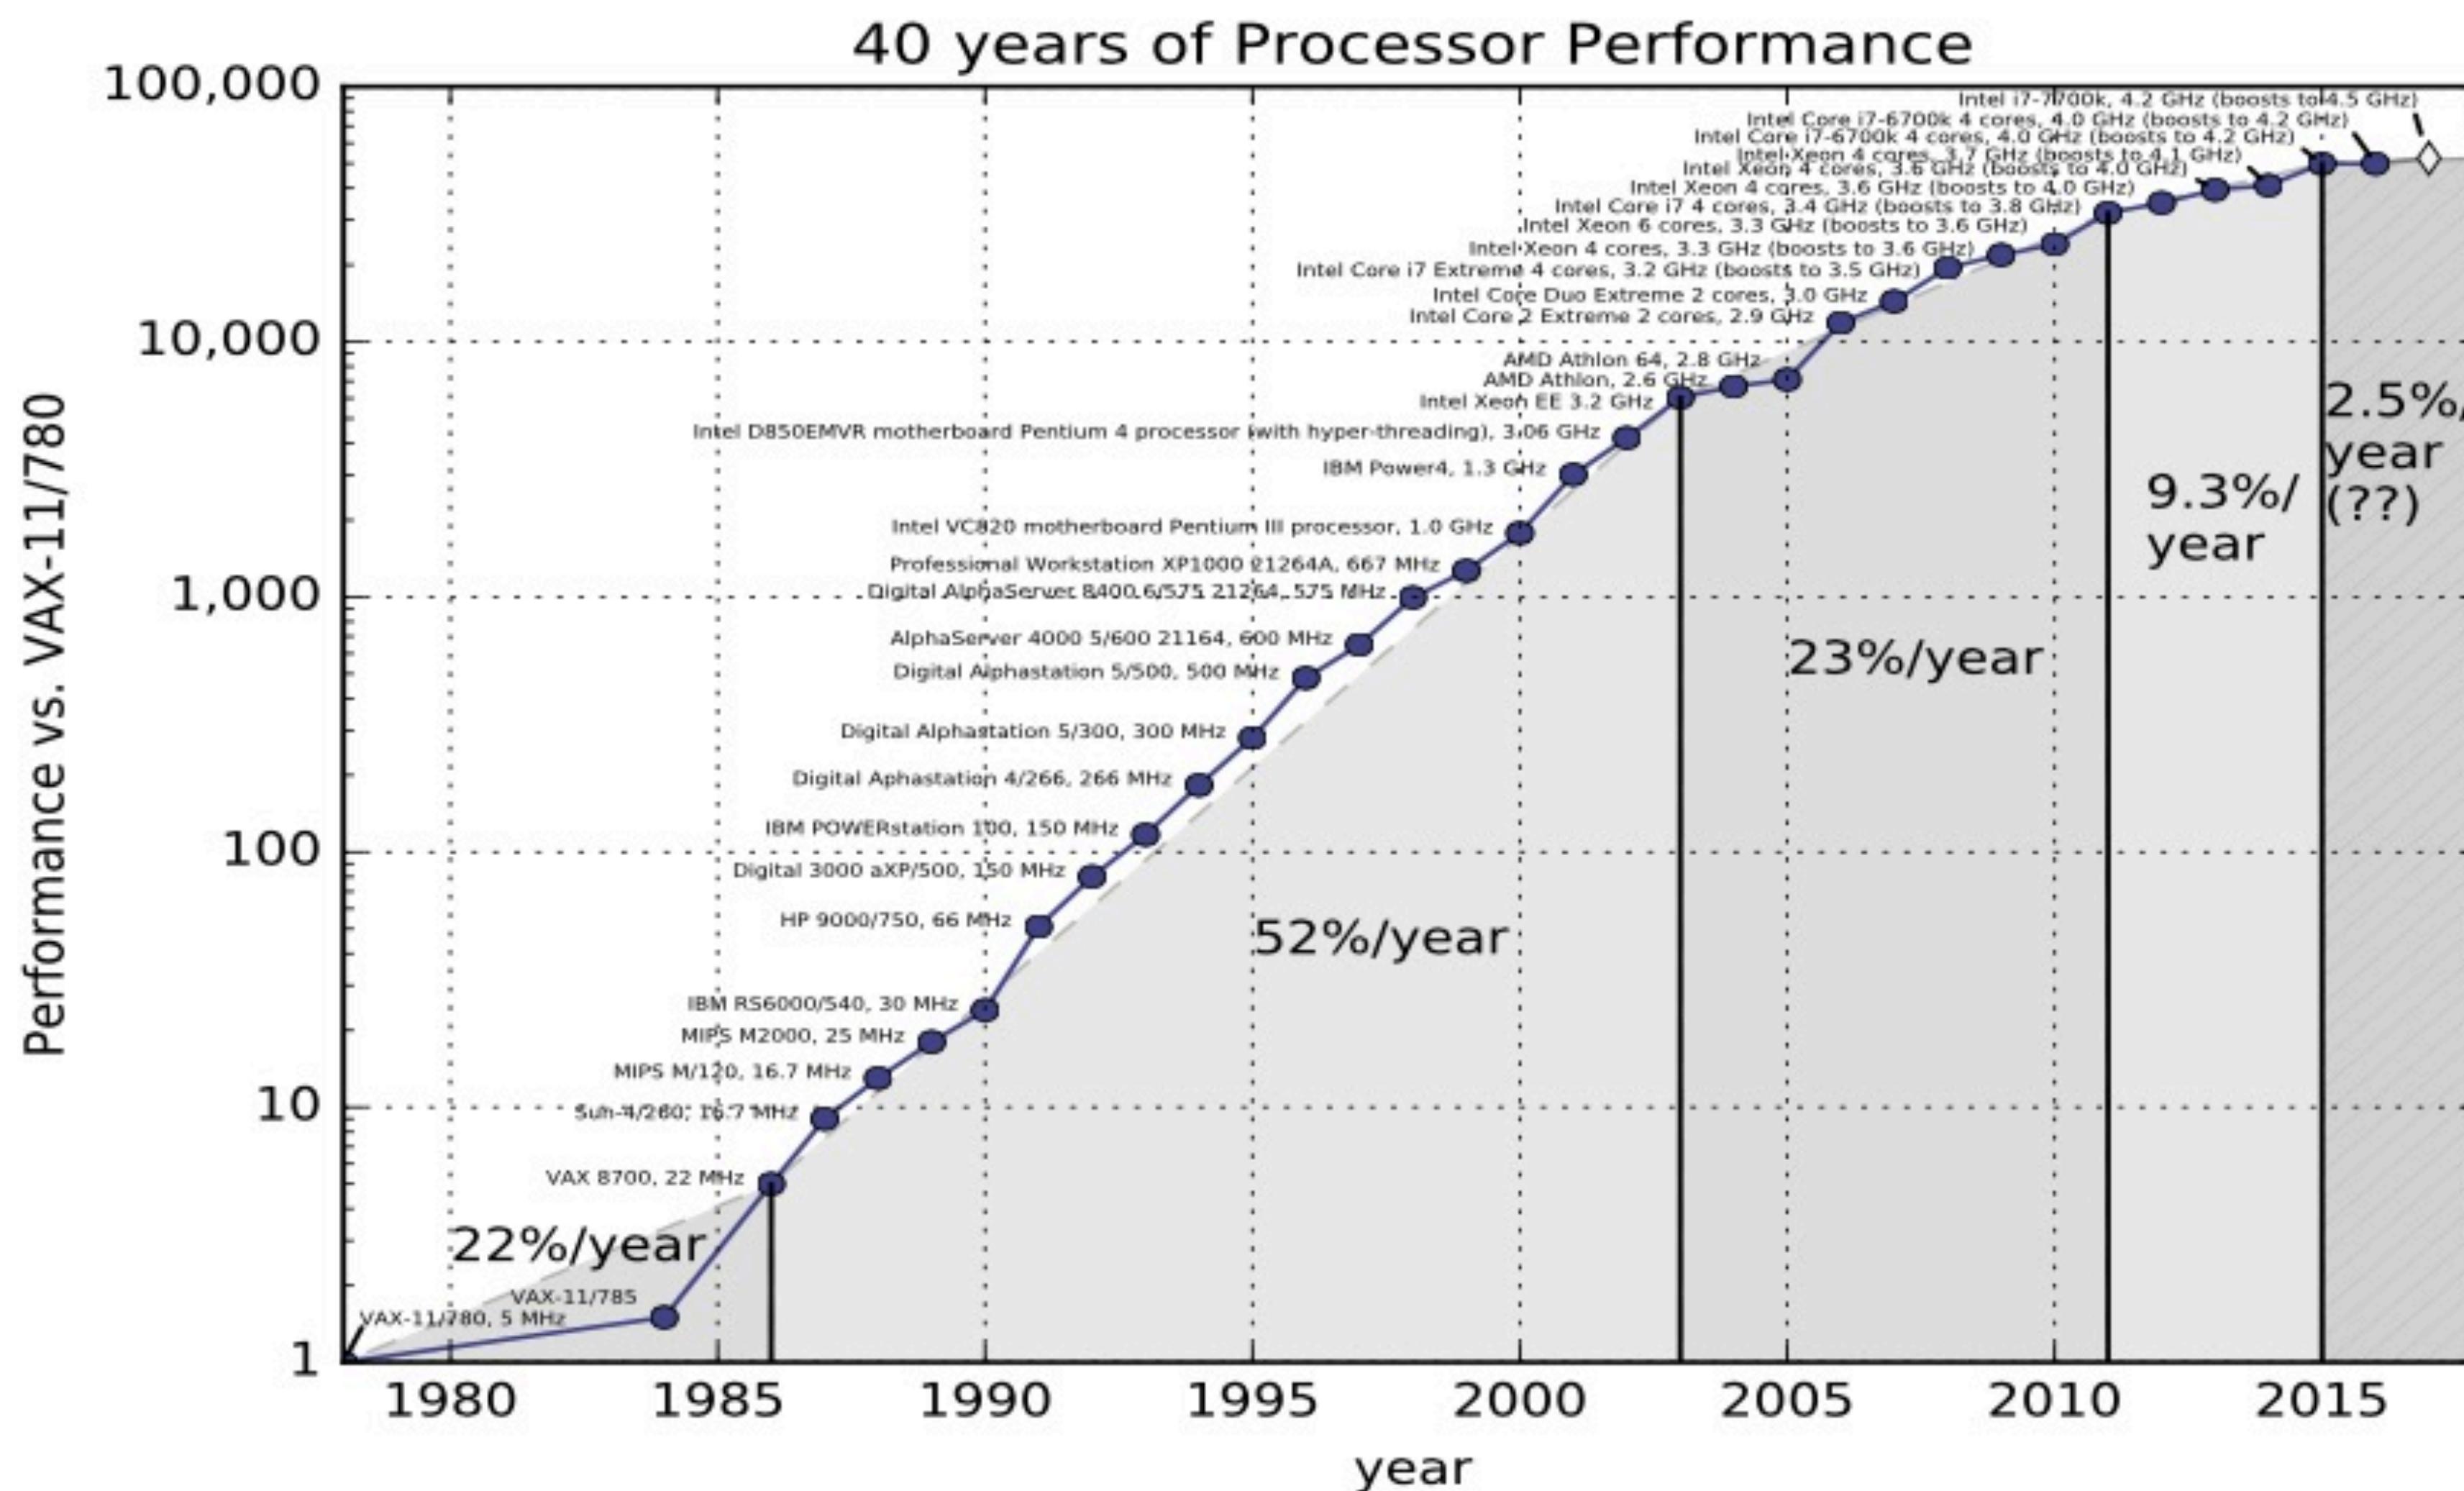
An Interesting Graph



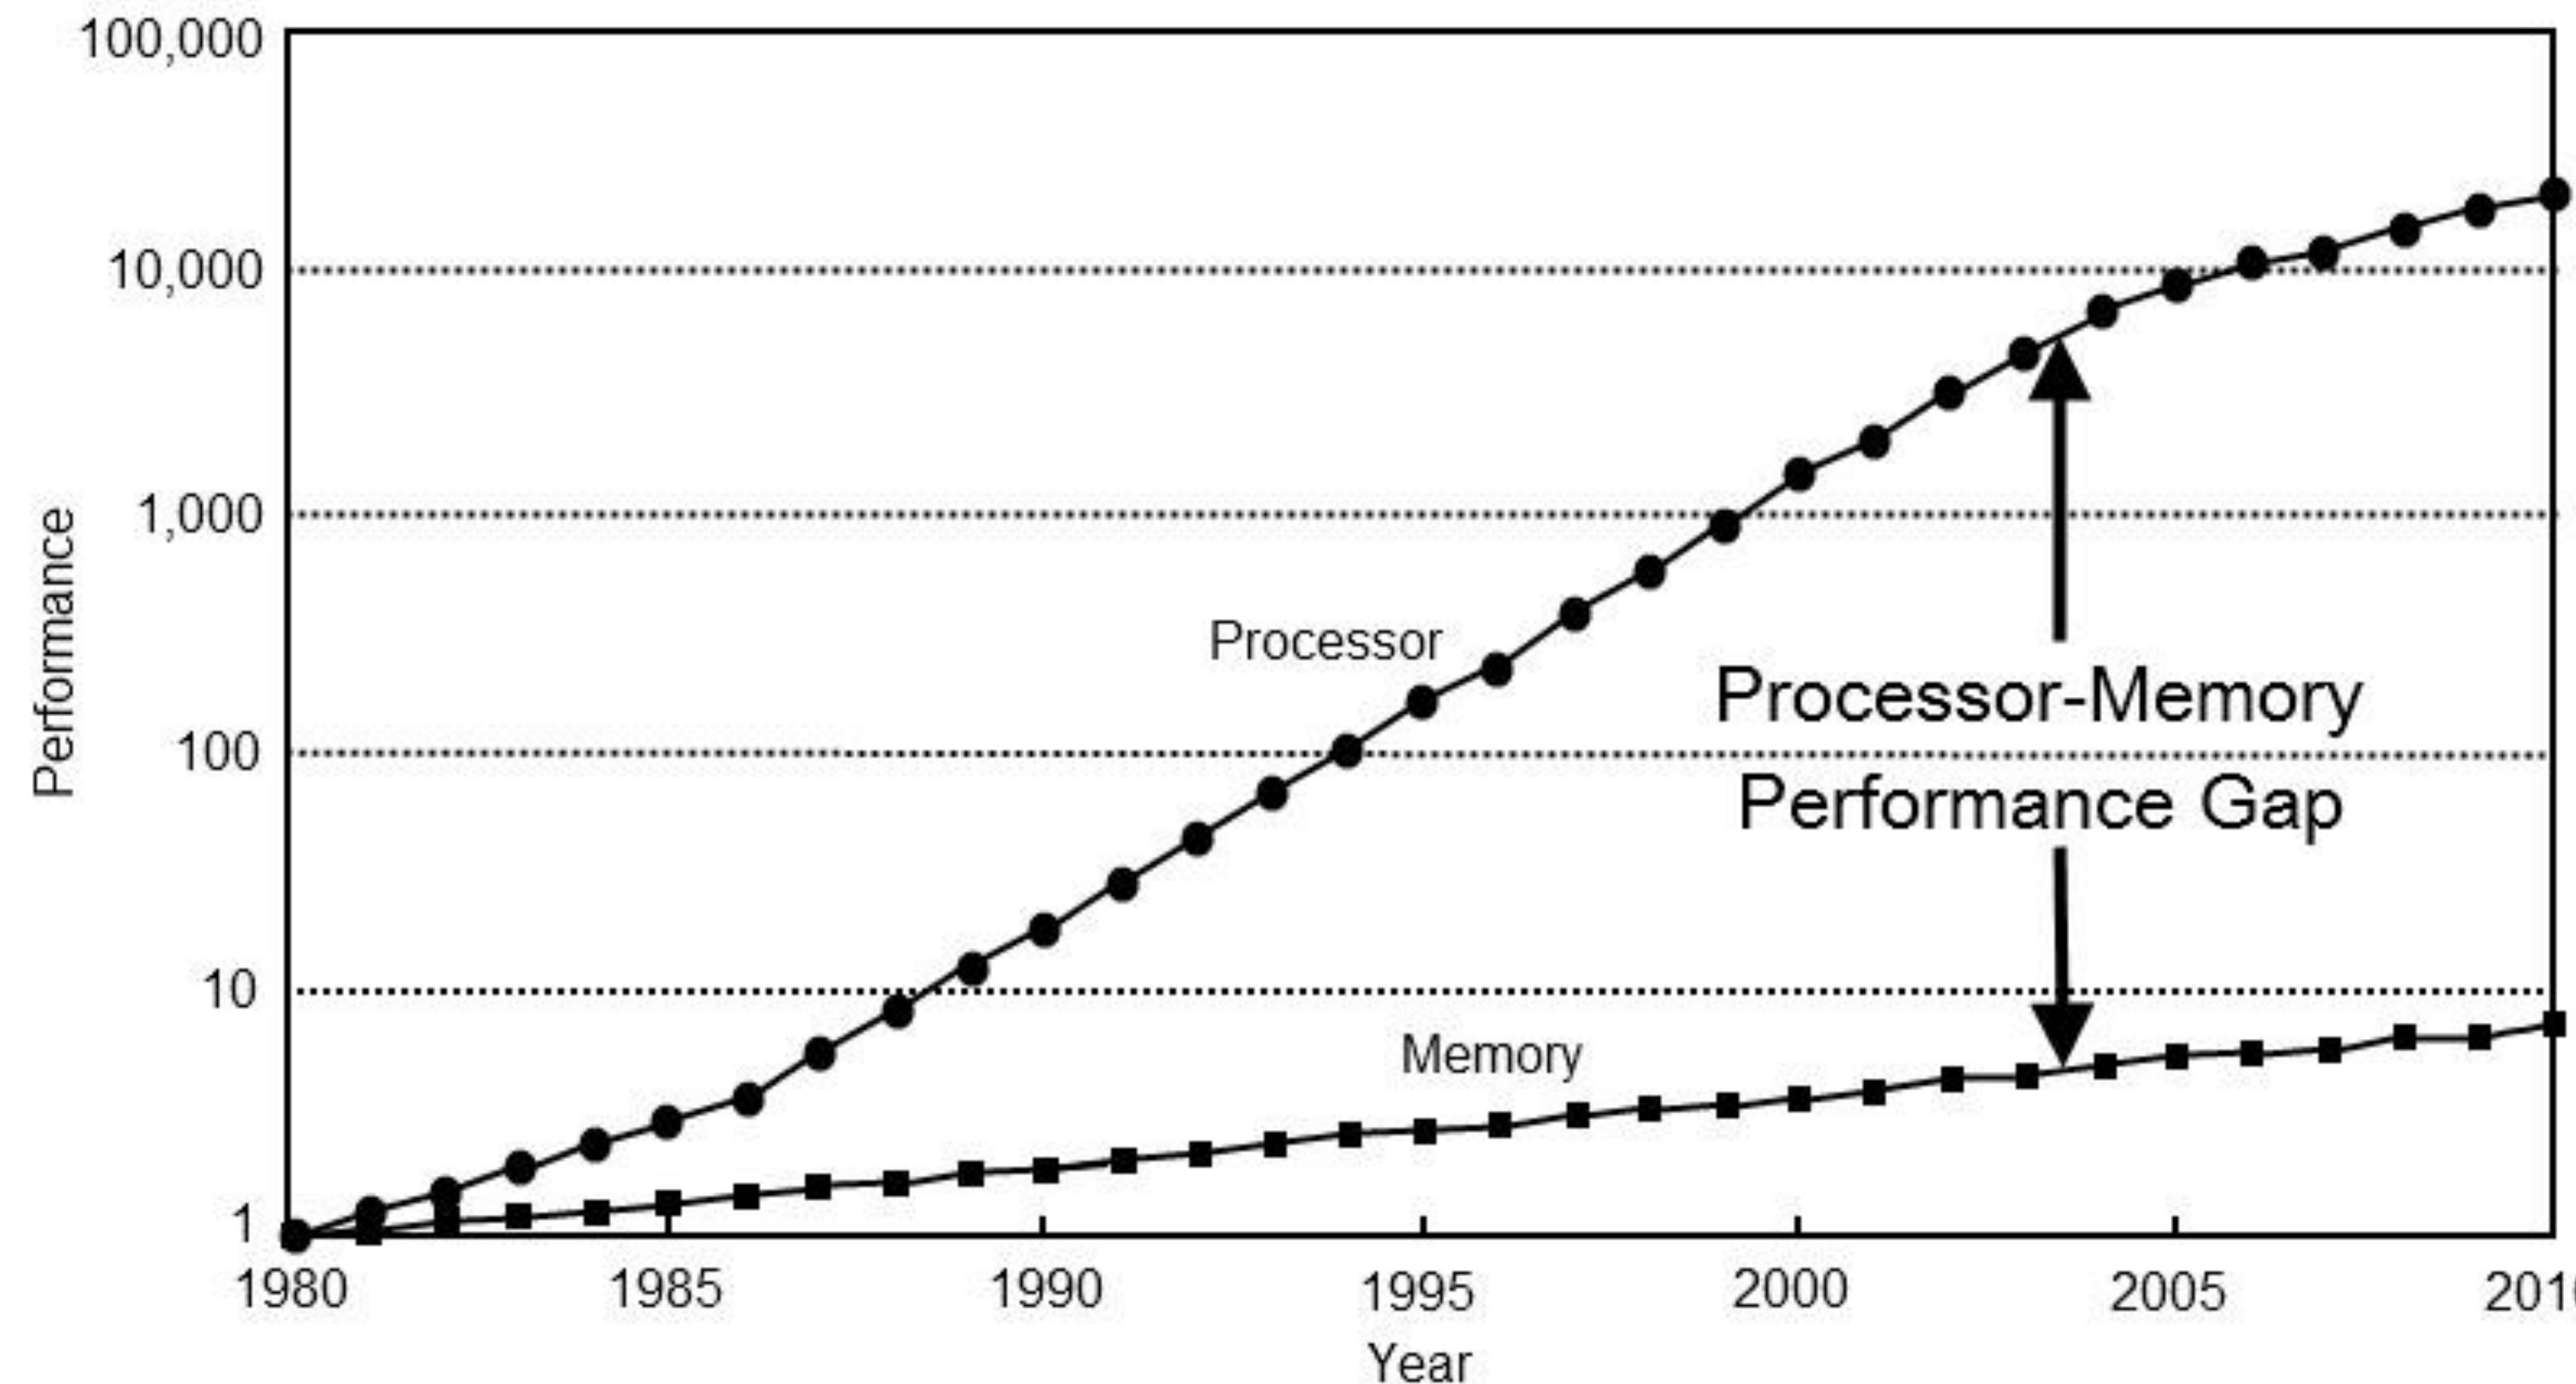
A Scary Graph



Another Interesting Graph



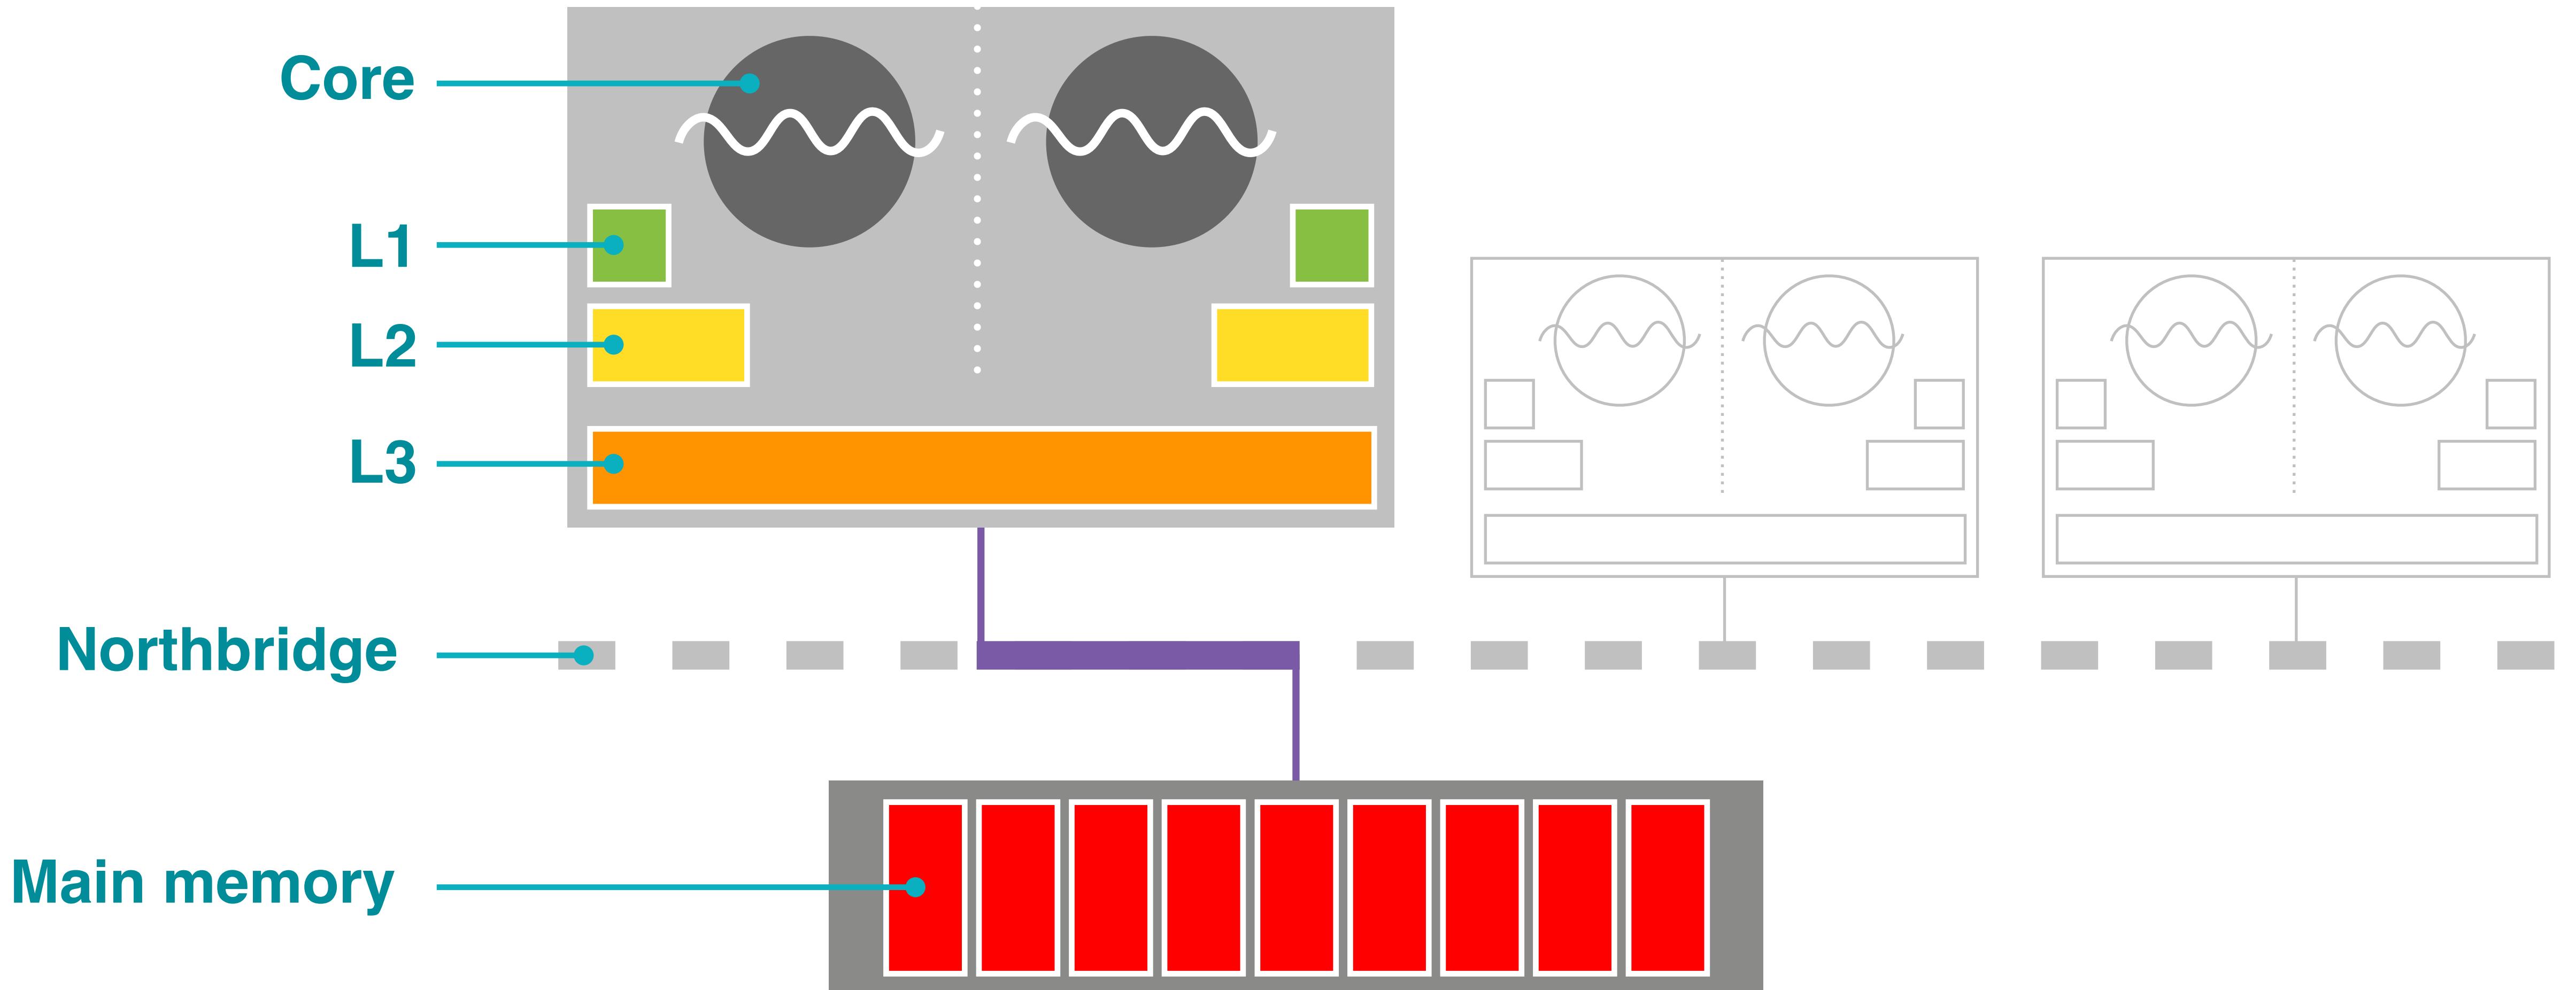
Another Interesting Graph



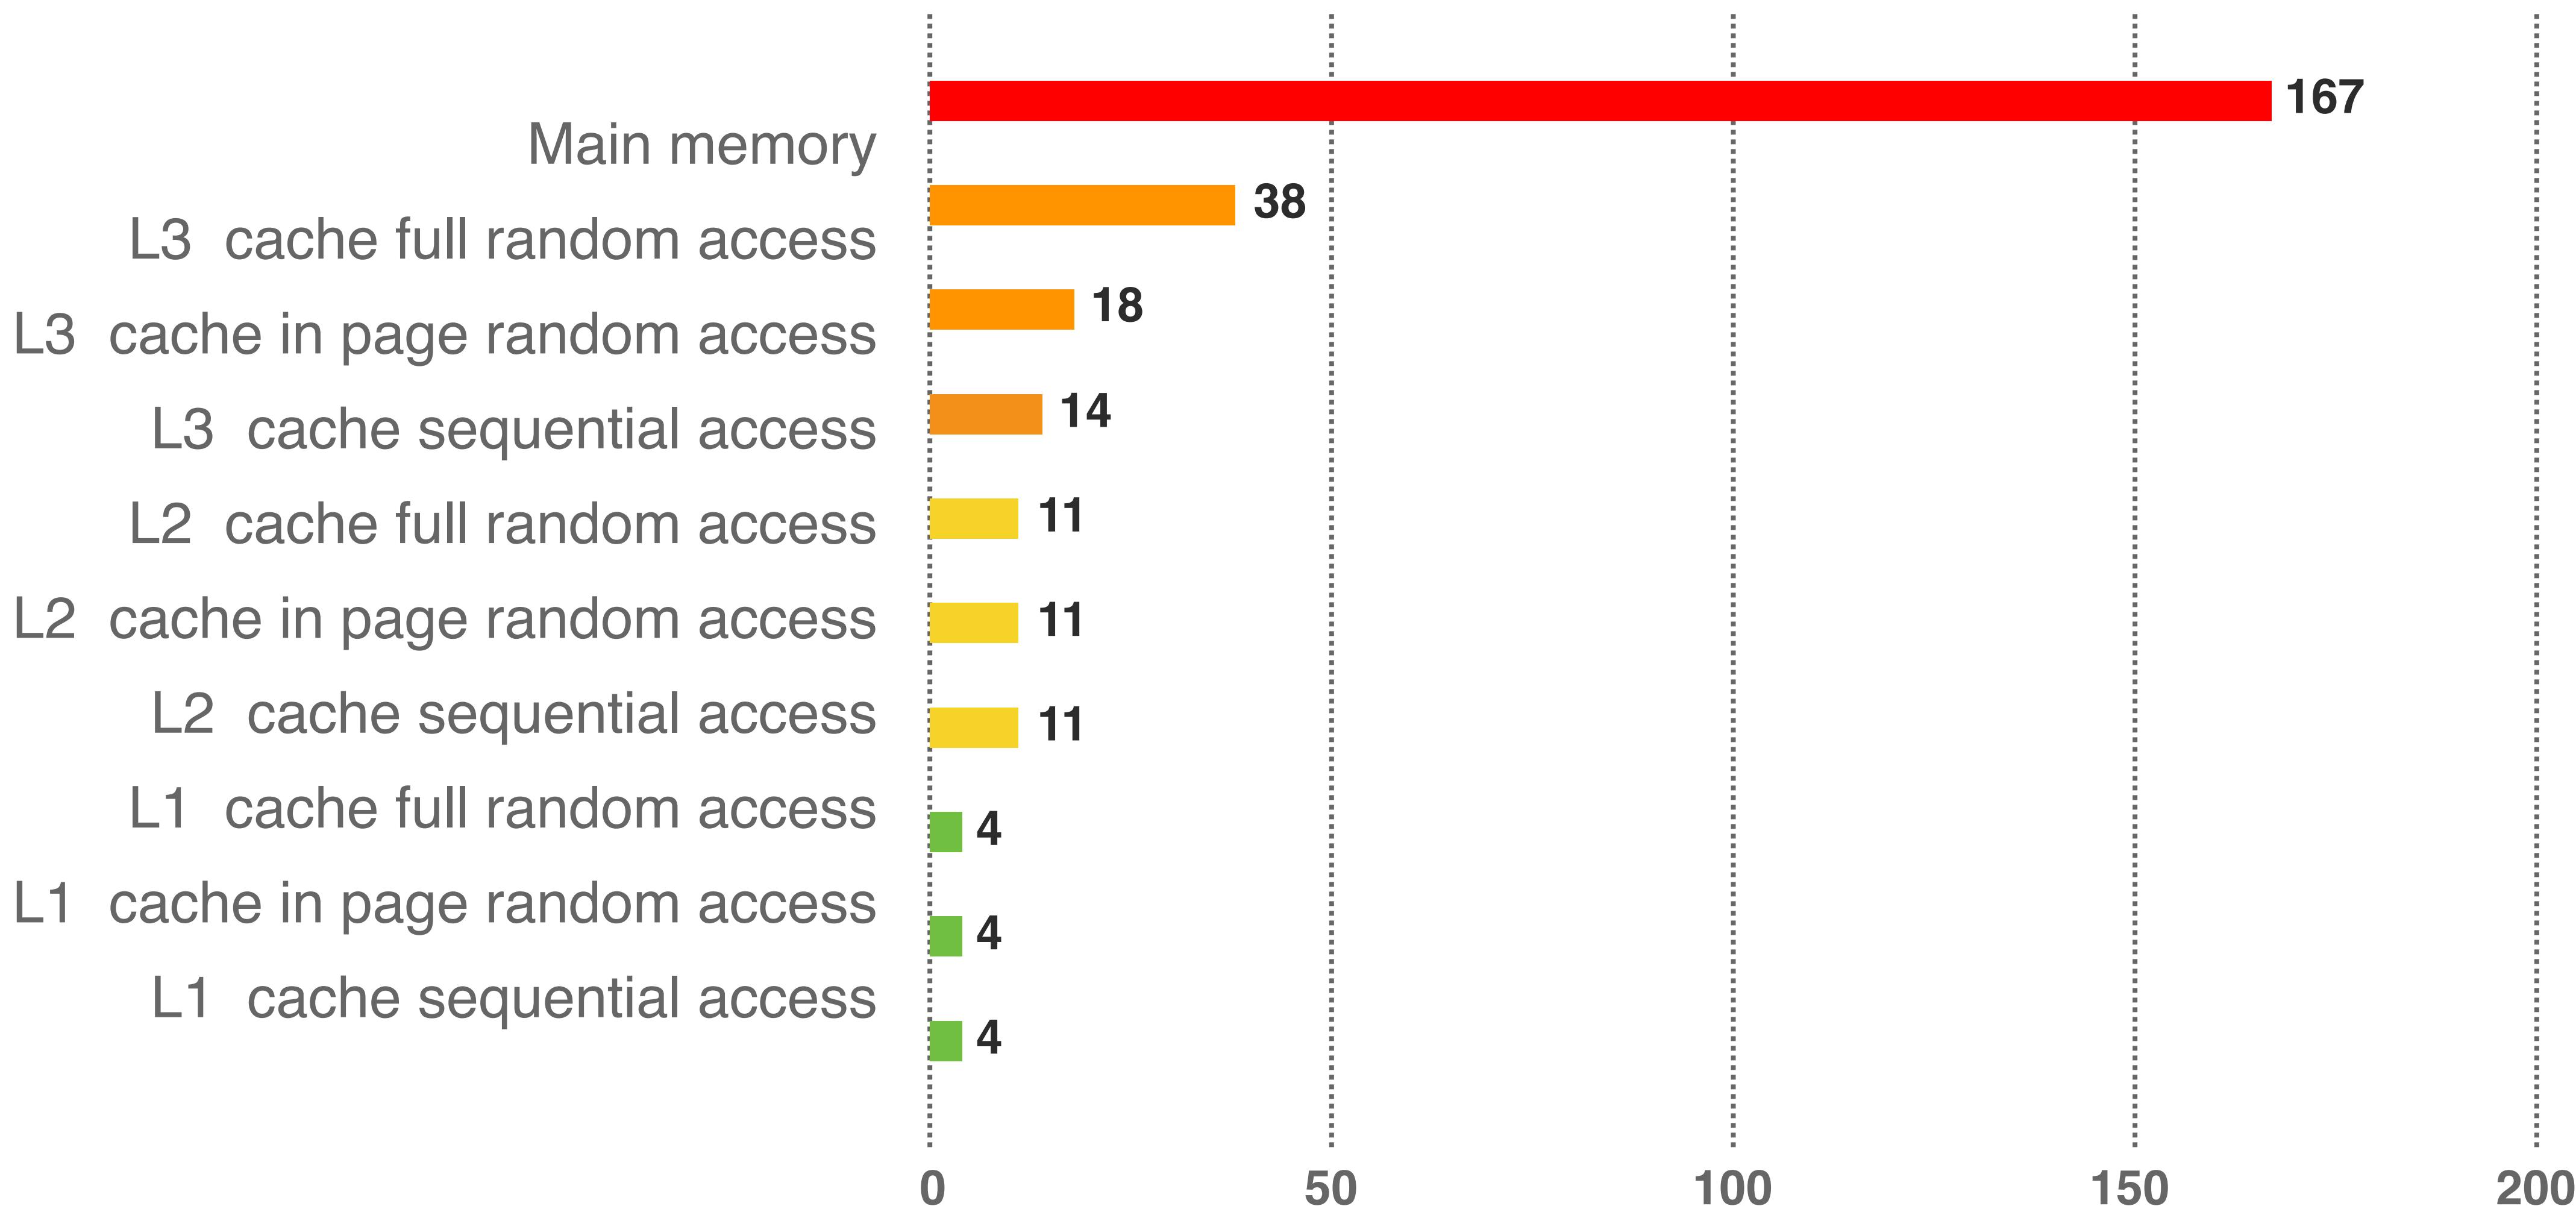
From: Computer Architecture: A Quantitative Approach by Hennessy, et al.



The CPU Landscape



The Speed of Memory



Exercise - Caching

EXERCISE

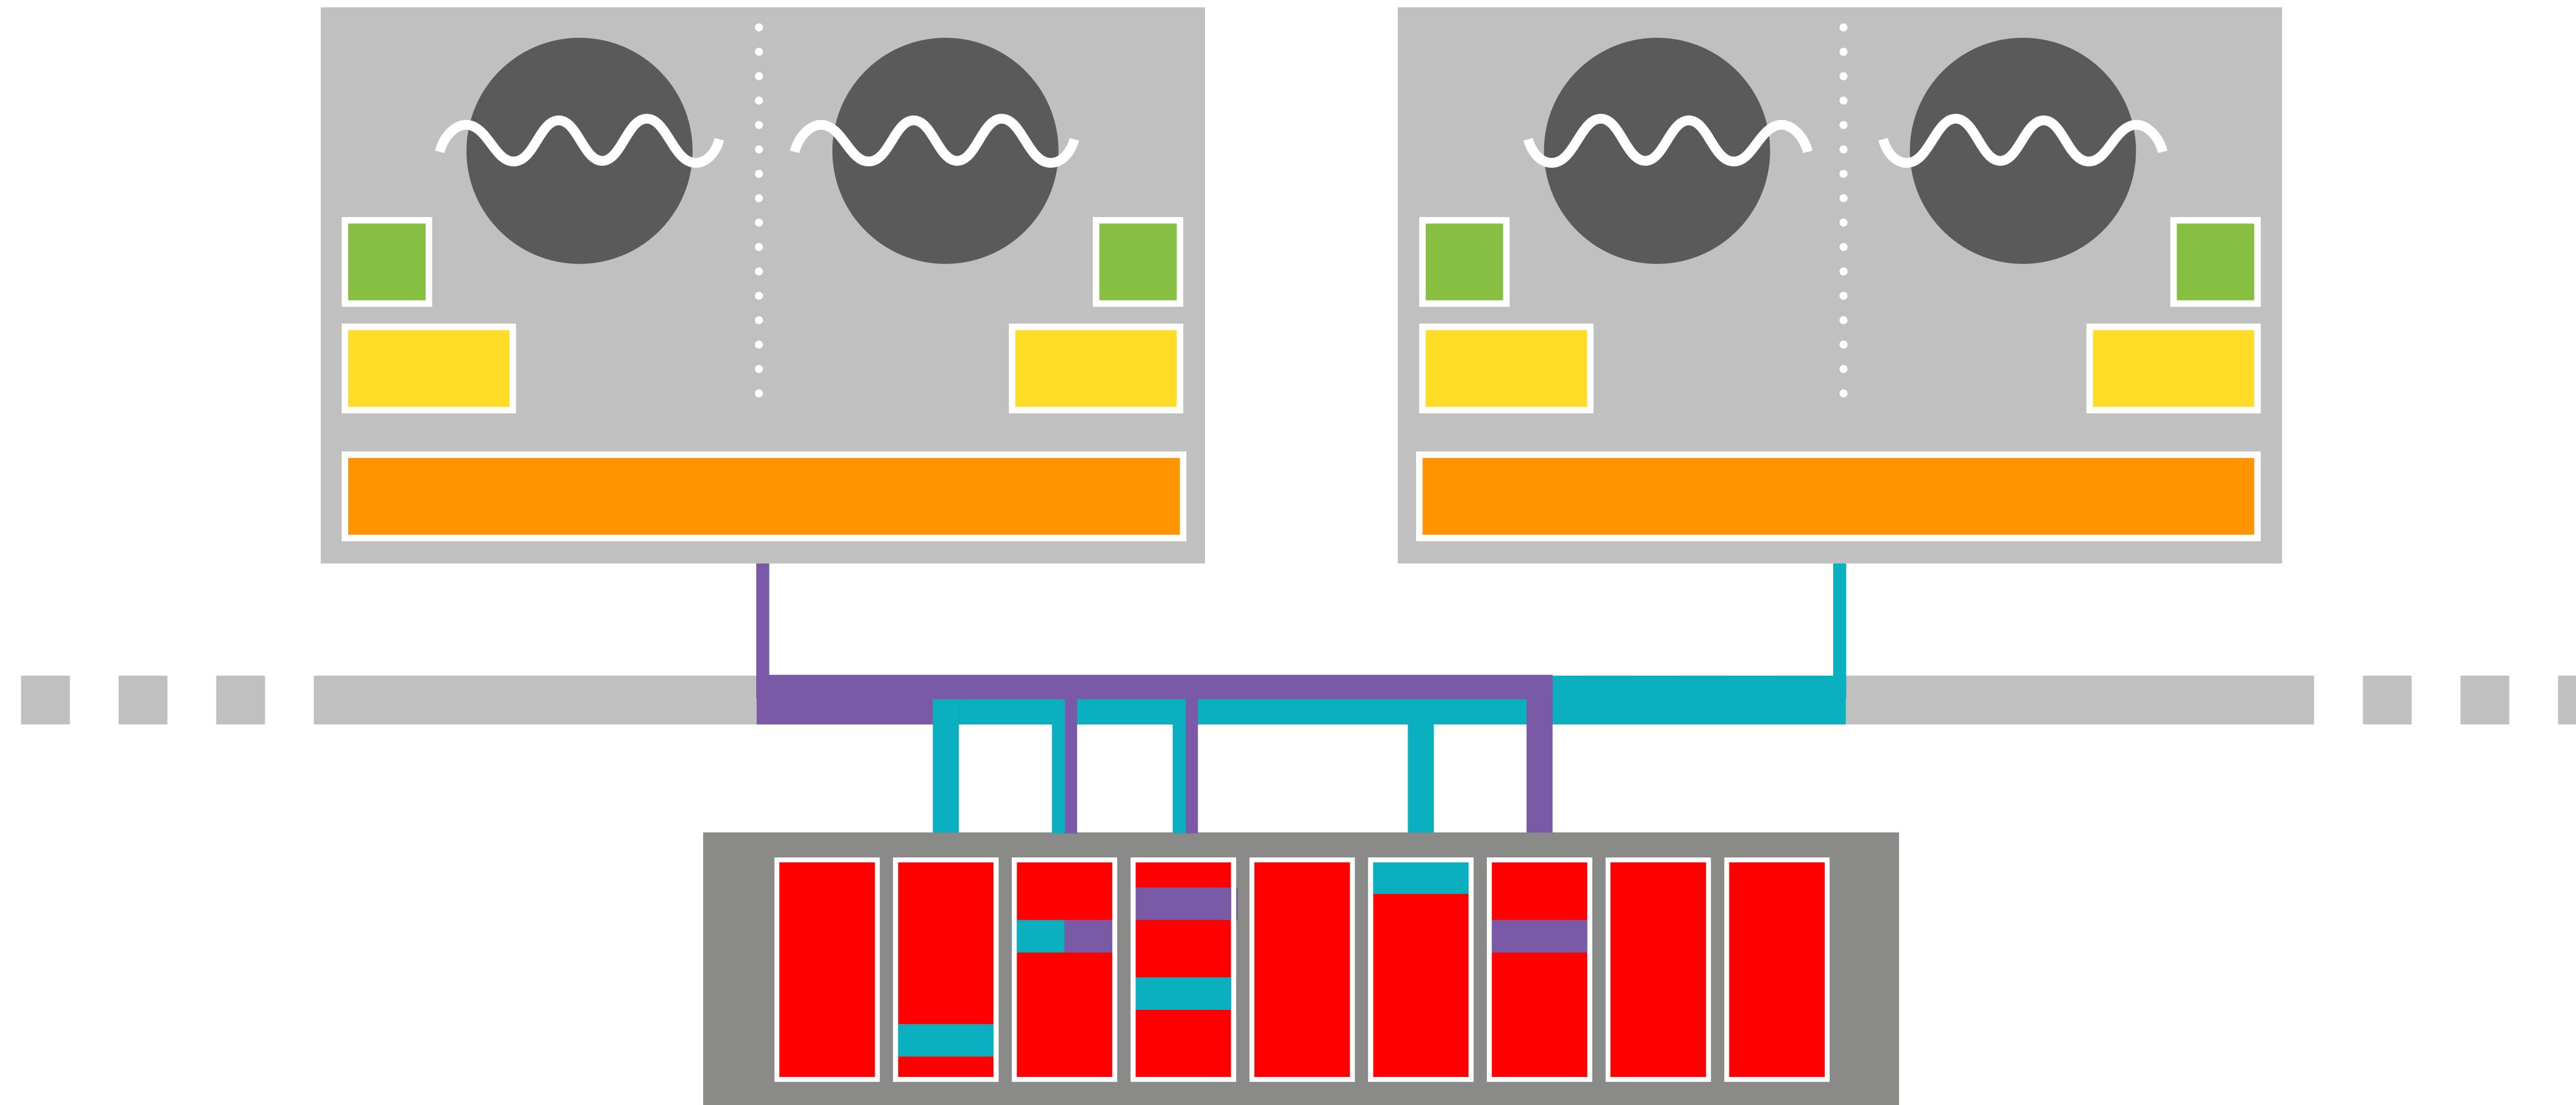


The Need for Synchronization

- Java's concurrency model
- Three fundamental concepts
- **synchronized** and **volatile**

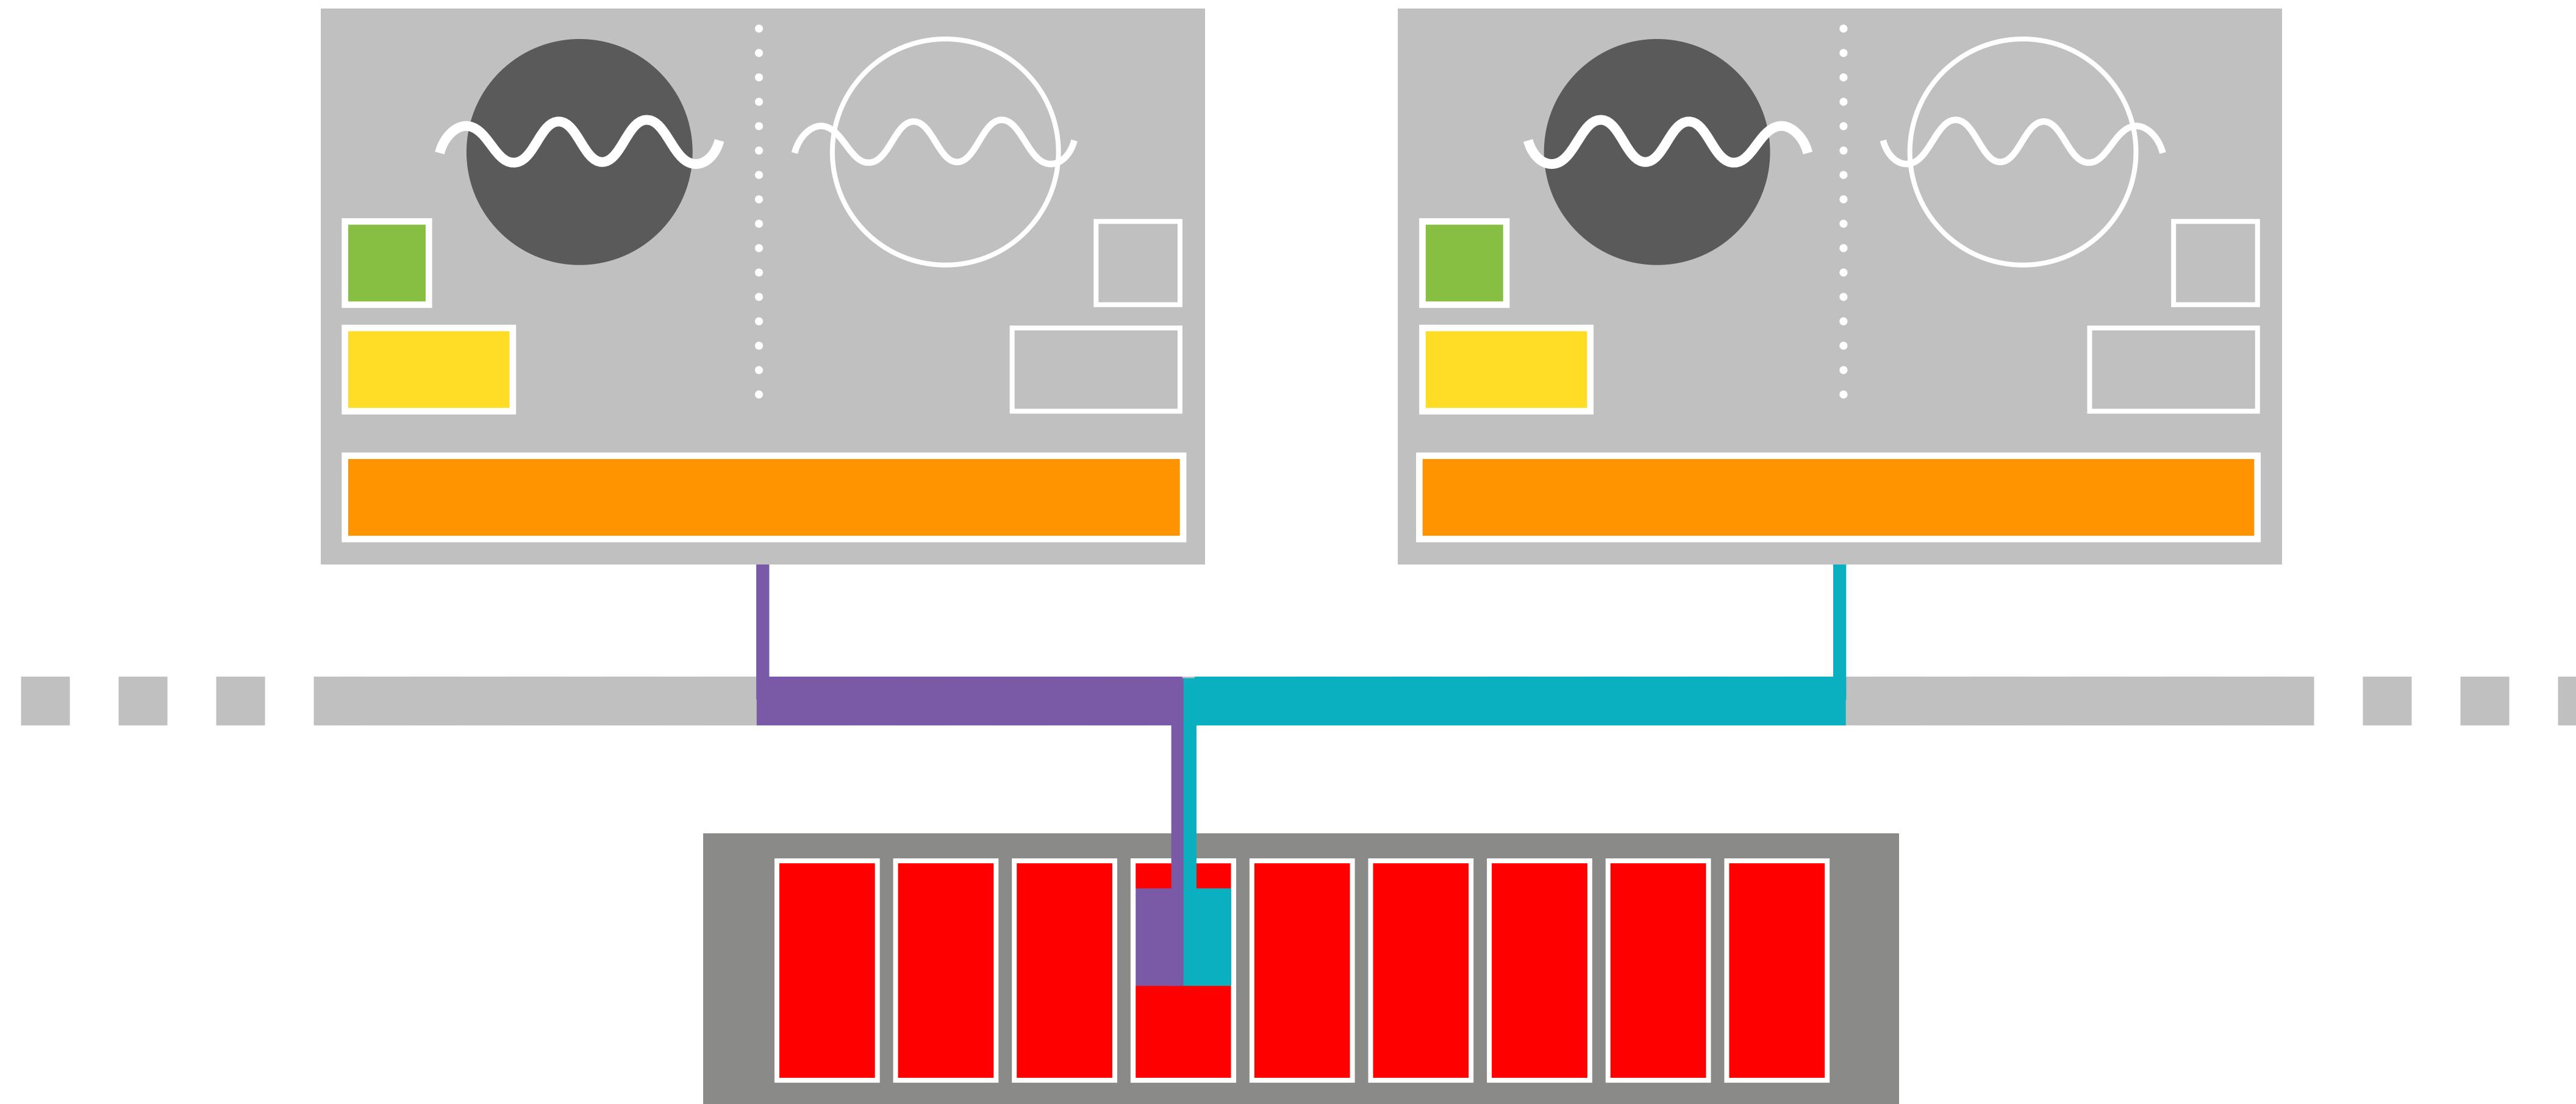


Accessing the NorthBridge

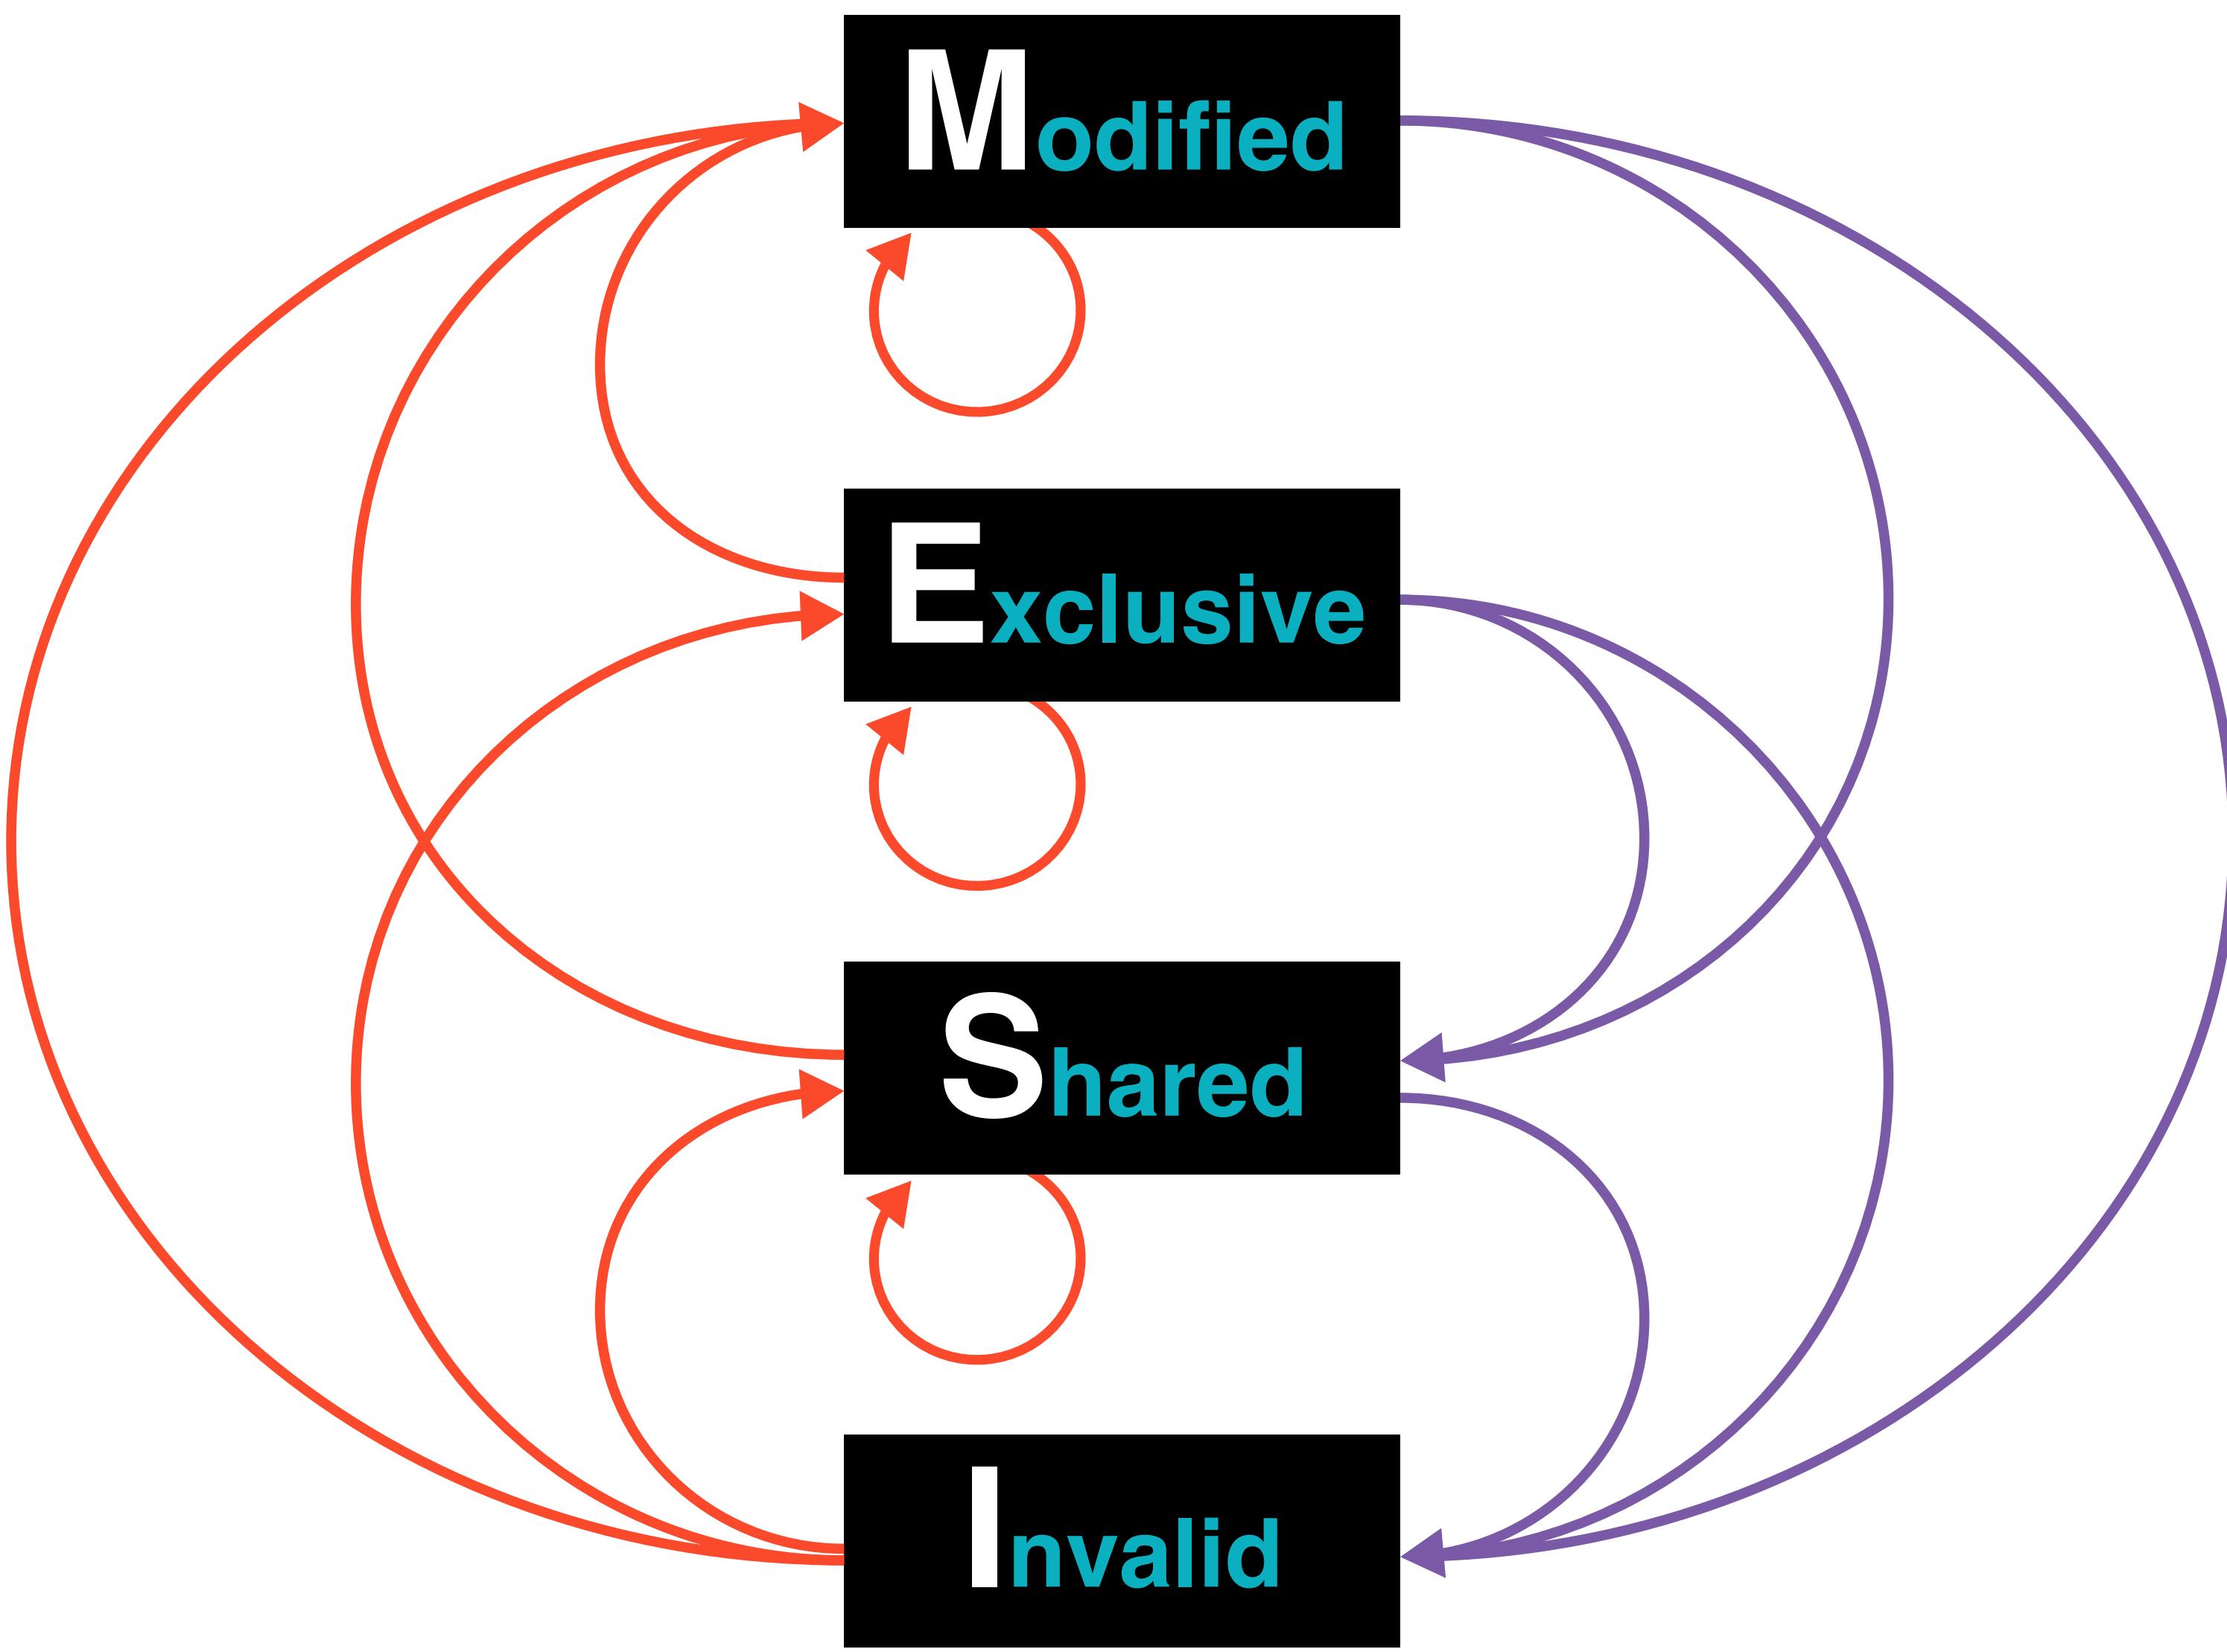


Cache Coherency

- If 2 cores access the same data, what is “true” view?



MESI



Fundamental Concepts of Java Concurrency

- Shared, visible-by-default mutable state
- Preemptive thread scheduling
- Objects must be locked to protect vulnerable state



Shared, Visible-by-Default Mutable State

- Objects are easily shared between all threads in a process
- Objects can be changed (“mutated”) by any thread
 - Provided the thread has a reference to the object



Preemptive Thread Scheduling

- What happens when a thread is swapped off a core?



Preemptive Thread Scheduling

- What happens when a thread is swapped off a core?
- Thread is suspended immediately
- It may be in an inconsistent state



Java's Concurrency model

- To protect against inconsistency, need to use exclusion
 - Only available mechanism
- All code using shared mutable state inconsistent must cooperate
 - **Modifies or reads**
- Code which does this is **Concurrently Safe**
- Remember: Scheduler can swap threads off cores at any time



Objects Must be Locked to Protect Vulnerable State

- Basis of “intrinsic concurrency” approach
- A Java thread is an OS thread
- `Thread.start()` calls to the OS to `clone()` a new thread
- A `java.lang.Thread` is just a metadata object
- OS monitors used to awake threads



Java's Concurrency Model

- What does `synchronized` do?



Java's Concurrency Model

- What does **synchronized** do?
- Why do we used the keyword **synchronized** for this?
- What is being synchronized?
- Why not **locked** or **critical** or similar?



Java's Concurrency Model

- What is being **synchronized**?
 - Main memory representation of object being locked



Java's Concurrency Model

- What is being **synchronized**?
 - Main memory representation of object being locked
- Two separate operations
 - Read from memory at start of synchronized block
 - Write-back to memory at end of block



Sequence of Events for Synchronized Block

- Thread needs to modify an object
 - Indicates it requires temporary exclusive access to the object



Sequence of Events for Synchronized Block

- Thread needs to modify an object
 - Indicates it requires temporary exclusive access to the object
- Thread acquires object monitor
 - May have had to block to acquire



Sequence of Events for Synchronized Block

- Thread needs to modify an object
 - Indicates it requires temporary exclusive access to the object
- Thread acquires object monitor
 - May have had to block to acquire
- Thread modifies (or reads) the object
 - While locked, object may become briefly inconsistent
 - Leaves it in a consistent, legal state when done



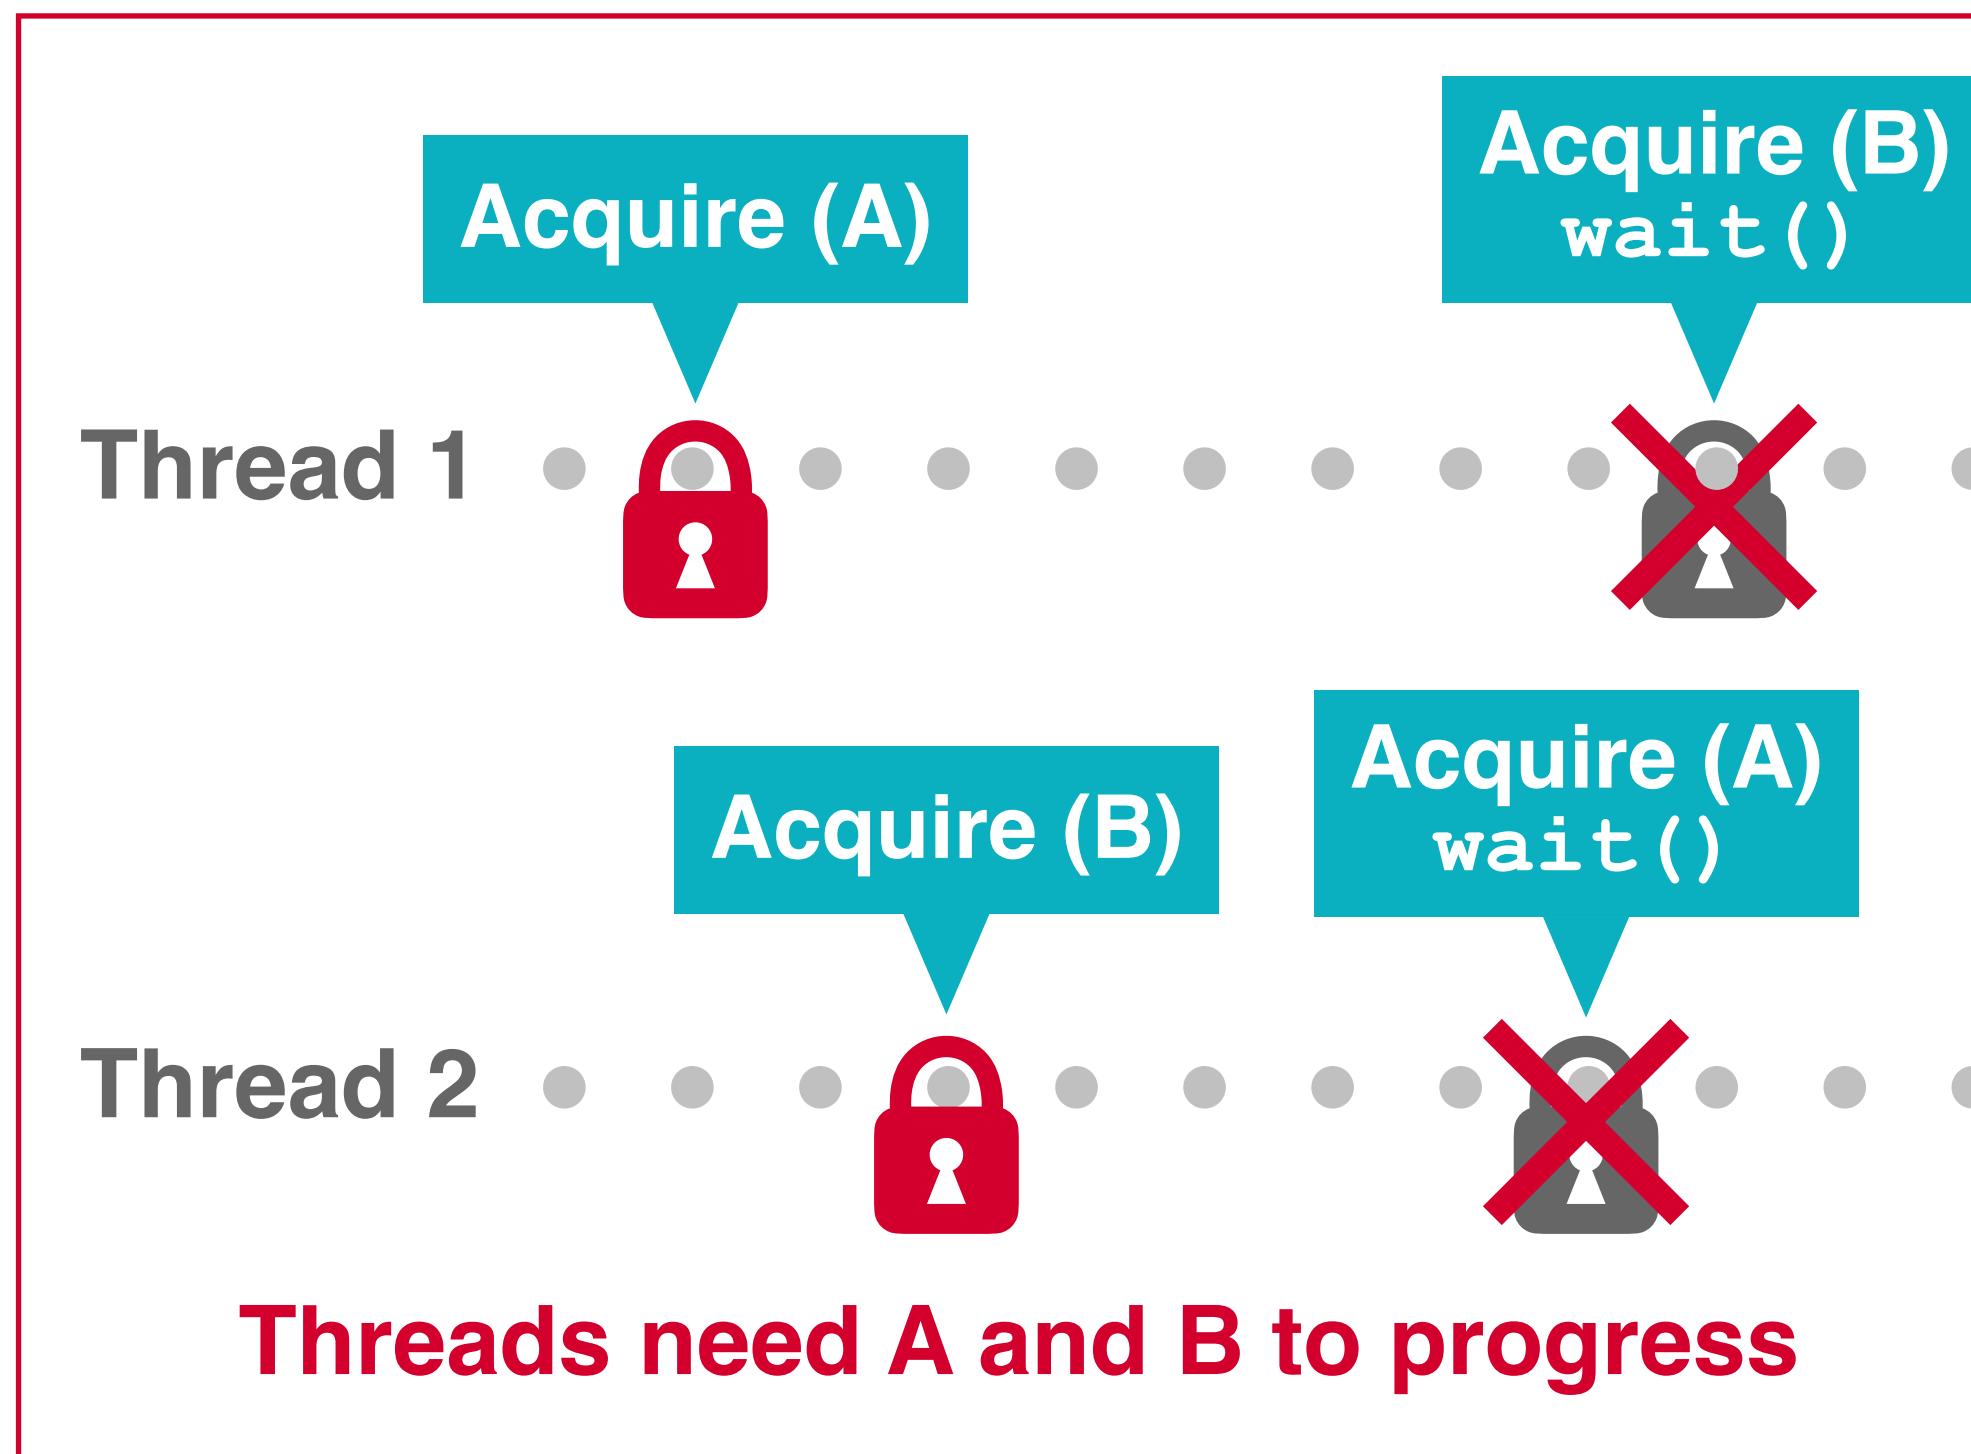
Sequence of Events for Synchronized Block

- Thread needs to modify an object
 - Indicates it requires temporary exclusive access to the object
- Thread acquires object monitor
 - May have had to block to acquire
- Thread modifies (or reads) the object
 - While locked, object may become briefly inconsistent
 - Leaves it in a consistent, legal state when done
- Thread releases the monitor
 - Other threads can now acquire



Deadlocks

- When several threads are competing over locks
 - A Liveness hazard



```
Object a = new Object();
Object b = new Object();

Thread t1 = new Thread(new Runnable() {
    public void run() {
        synchronized(a) {
            try {
                Thread.sleep(100);
            } catch (InterruptedException e) { }
            synchronized(b) { compute(); }
        }
    }
});

Thread t2 = new Thread(new Runnable() {
    public void run() {
        synchronized(b) {
            synchronized(a) { compute(); }
        }
    }
});

t1.start();
t2.start();
```



Avoiding Deadlocks

- Fully-Synchronized Objects are immune
- Immutable Objects are immune
- For other object types
 - Always acquire locks in same order in all threads
 - Use the Lock construct in j.u.c
 - Use strategies with a timeout / backoff
- Be careful of partial solutions
 - Extra complexity can cause schedulers to make different choices
 - Some “fixes” are really this in disguise



volatile

- What does **volatile** do?



volatile

- What does **volatile** do?
 - Provides **1** operation to or from main memory
 - No locking is involved
 - Caches are flushed



volatile

- What does **volatile** do?
 - Provides **1** operation to or from main memory
 - No locking is involved
 - Caches are flushed
- The value is:
 - Re-read from main memory before use
 - **Or** written back to main memory



volatile

- Volatile variables cannot be used for state-dependent updates
 - e.g. can't do `v++` on a volatile int
 - For these, must use a synchronized block
 - Example of `Counter` – works the same way for volatile as non-volatile
- Volatiles cannot cause deadlocks
 - There are no locks



Exercise - Counters

EXERCISE



Fun with Counters

```
public final class UnprotectedCounter
    implements Counter {

    private int i = 0;

    public int increment() {
        return i = i + 1;
    }

    public int get() {
        return i;
    }
}
```

```
public int increment();
Code:
0:  aload_0
1:  aload_0
2:  getfield #2 // Field i:I
5:  iconst_1
6:  iadd
7:  dup_x1
8:  putfield #2 // Field i:I
11: ireturn
```



Fun with Counters

```
Counter c = new UnprotectedCounter(); // init to 0

Runnable r = () -> {
    for (int i = 0; i < REPS; i++)
        c.increment();
};

Thread tA = new Thread(r);
Thread tB = new Thread(r);
long start = System.currentTimeMillis();
tA.start();
tB.start();
tA.join();
tB.join();
long fin = System.currentTimeMillis();
int diff = c.increment() - (2 * REPS + 1);
System.out.println("Diff: " + diff);
System.out.println("Elapsed: " + (fin - start));
```

```
A0: aload_0
A1: aload_0
A2: getfield #2 // Field i:I
A5: iconst_1
A6: iadd
A7: dup_x1
B0: aload_0
B1: aload_0
B2: getfield #2 // Field i:I
B5: iconst_1
B6: iadd
B7: dup_x1
A8: putfield #2 // Field i:I
A11: ireturn
B8: putfield #2 // Field i:I
B11: ireturn
```

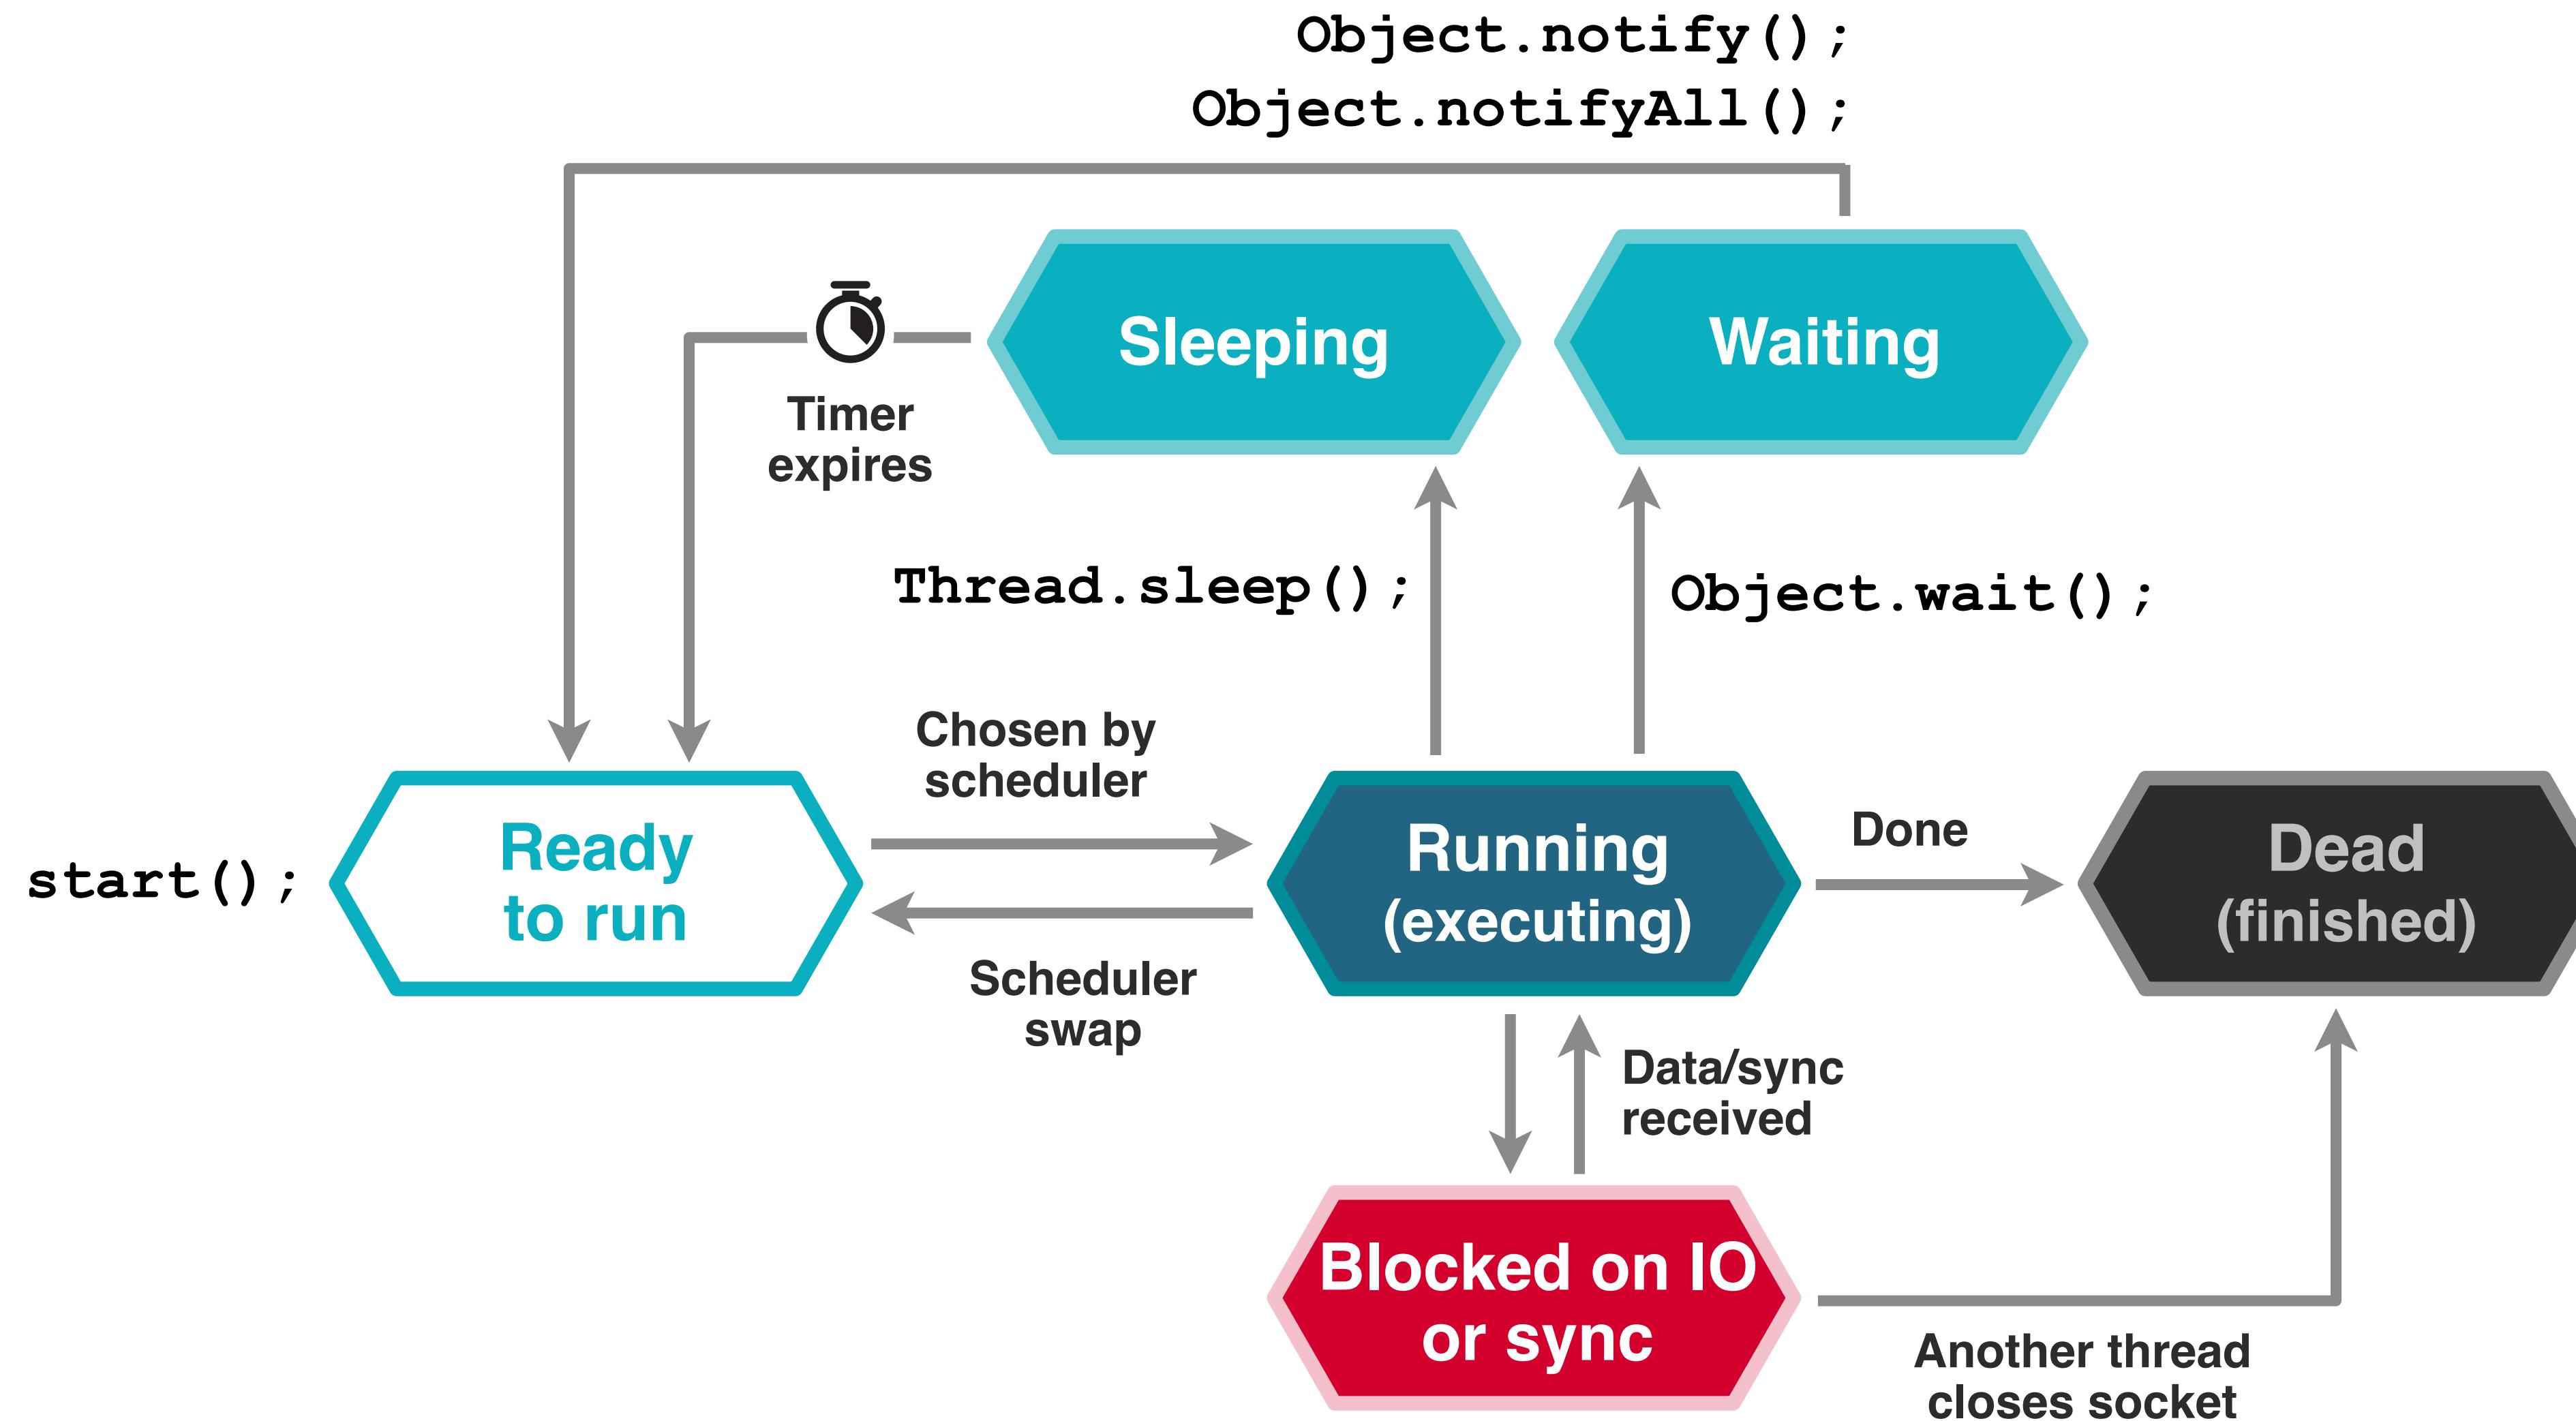


Java's Original Thread & Synchronization API

- Scheduling Basics
- The Java Thread class
- Thread States
- Thread Methods



Thread Lifecycle



The Java Thread Class

- Thread is a lightweight unit of execution
 - Smaller than a process
 - Still capable of executing arbitrary Java code
- Each thread is a full unit of execution to OS
 - Own stack and program counter
 - Main address space (heap) shared between all threads in process
- Java has supported multithreaded programming since 1.0



Thread States

- NEW
 - Thread has been created but `start()` method has not been called
 - All threads start in this state.
- RUNNABLE
 - The thread is running or is available to run
 - Operating system is responsible for scheduling it



Thread States

- **BLOCKED**
 - The thread is not running
 - It is waiting to acquire a lock or is in a system call
- **WAITING**
 - The thread is not running
 - It has called `Object.wait()` or `Thread.join()`



Thread States

- **TIMED_WAITING**
 - The thread is not running
 - It has called `Thread.sleep()` or `Object.wait()`
- **TERMINATED**
 - Thread has completed execution.
 - Its `run()` method has exited normally or by throwing an exception.



Java Thread Methods

- `getId()`
- `setName()` and `getName()`
- `getState()`
- `isAlive()`
- `start()`
- `interrupt()`
- `join()`



Java Thread Methods

- `getPriority()` and `setPriority()`
- `setDaemon()`
- `setUncaughtExceptionHandler()`



setUncaughtExceptionHandler

```
// This thread just throws an exception
Thread handledThread = new Thread(() -> {
    throw new UnsupportedOperationException(); });

// Giving threads a name helps with debugging
handledThread.setName("My Broken Thread");

// Here's a handler for the error
handledThread.setUncaughtExceptionHandler((t, e) -> {
    System.err.printf("Exception in thread %d '%s':"
        "%s at line %d of %s%n",
        t.getId(), // Thread id
        t.getName(), // Thread name
        e.toString(), // Exception name and message
        e.getStackTrace()[0].getLineNumber(),
        e.getStackTrace()[0].getFileName()); });

handledThread.start();
```



Deprecated Java Thread Methods

- **stop()**
 - Almost impossible to use without violating safety
- **suspend()**, **resume()** and **countStackFrames()**
 - The **suspend()** mechanism does not release any monitors
- **destroy()**
 - Never implemented
 - Would have suffered from the same races as **suspend()**



Native Methods of Object

```
static JNINativeMethod methods[ ] = {  
    {"hashCode",      "()I",                      (void *)&JVM_IHashCode},  
    {"wait",          "(J)V",                      (void *)&JVM_MonitorWait},  
    {"notify",         "()V",                      (void *)&JVM_MonitorNotify},  
    {"notifyAll",      "()V",                      (void *)&JVM_MonitorNotifyAll},  
    {"clone",          "()Ljava/lang/Object;",     (void *)&JVM_Clone},  
};
```



Native Code for an Object Wait

```
JVM_ENTRY(void, JVM_MonitorWait(JNIEnv* env, jobject handle, jlong ms))
    JVMWrapper("JVM_MonitorWait");
Handle obj(THREAD, JNIHandles::resolve_non_null(handle));
JavaThreadInObjectWaitState jtiows(thread, ms != 0);
if (JvmtiExport::should_post_monitor_wait()) {
    JvmtiExport::post_monitor_wait((JavaThread *)THREAD, (oop)obj(), ms);
}
ObjectSynchronizer::wait(obj, ms, CHECK);
JVM_END
```

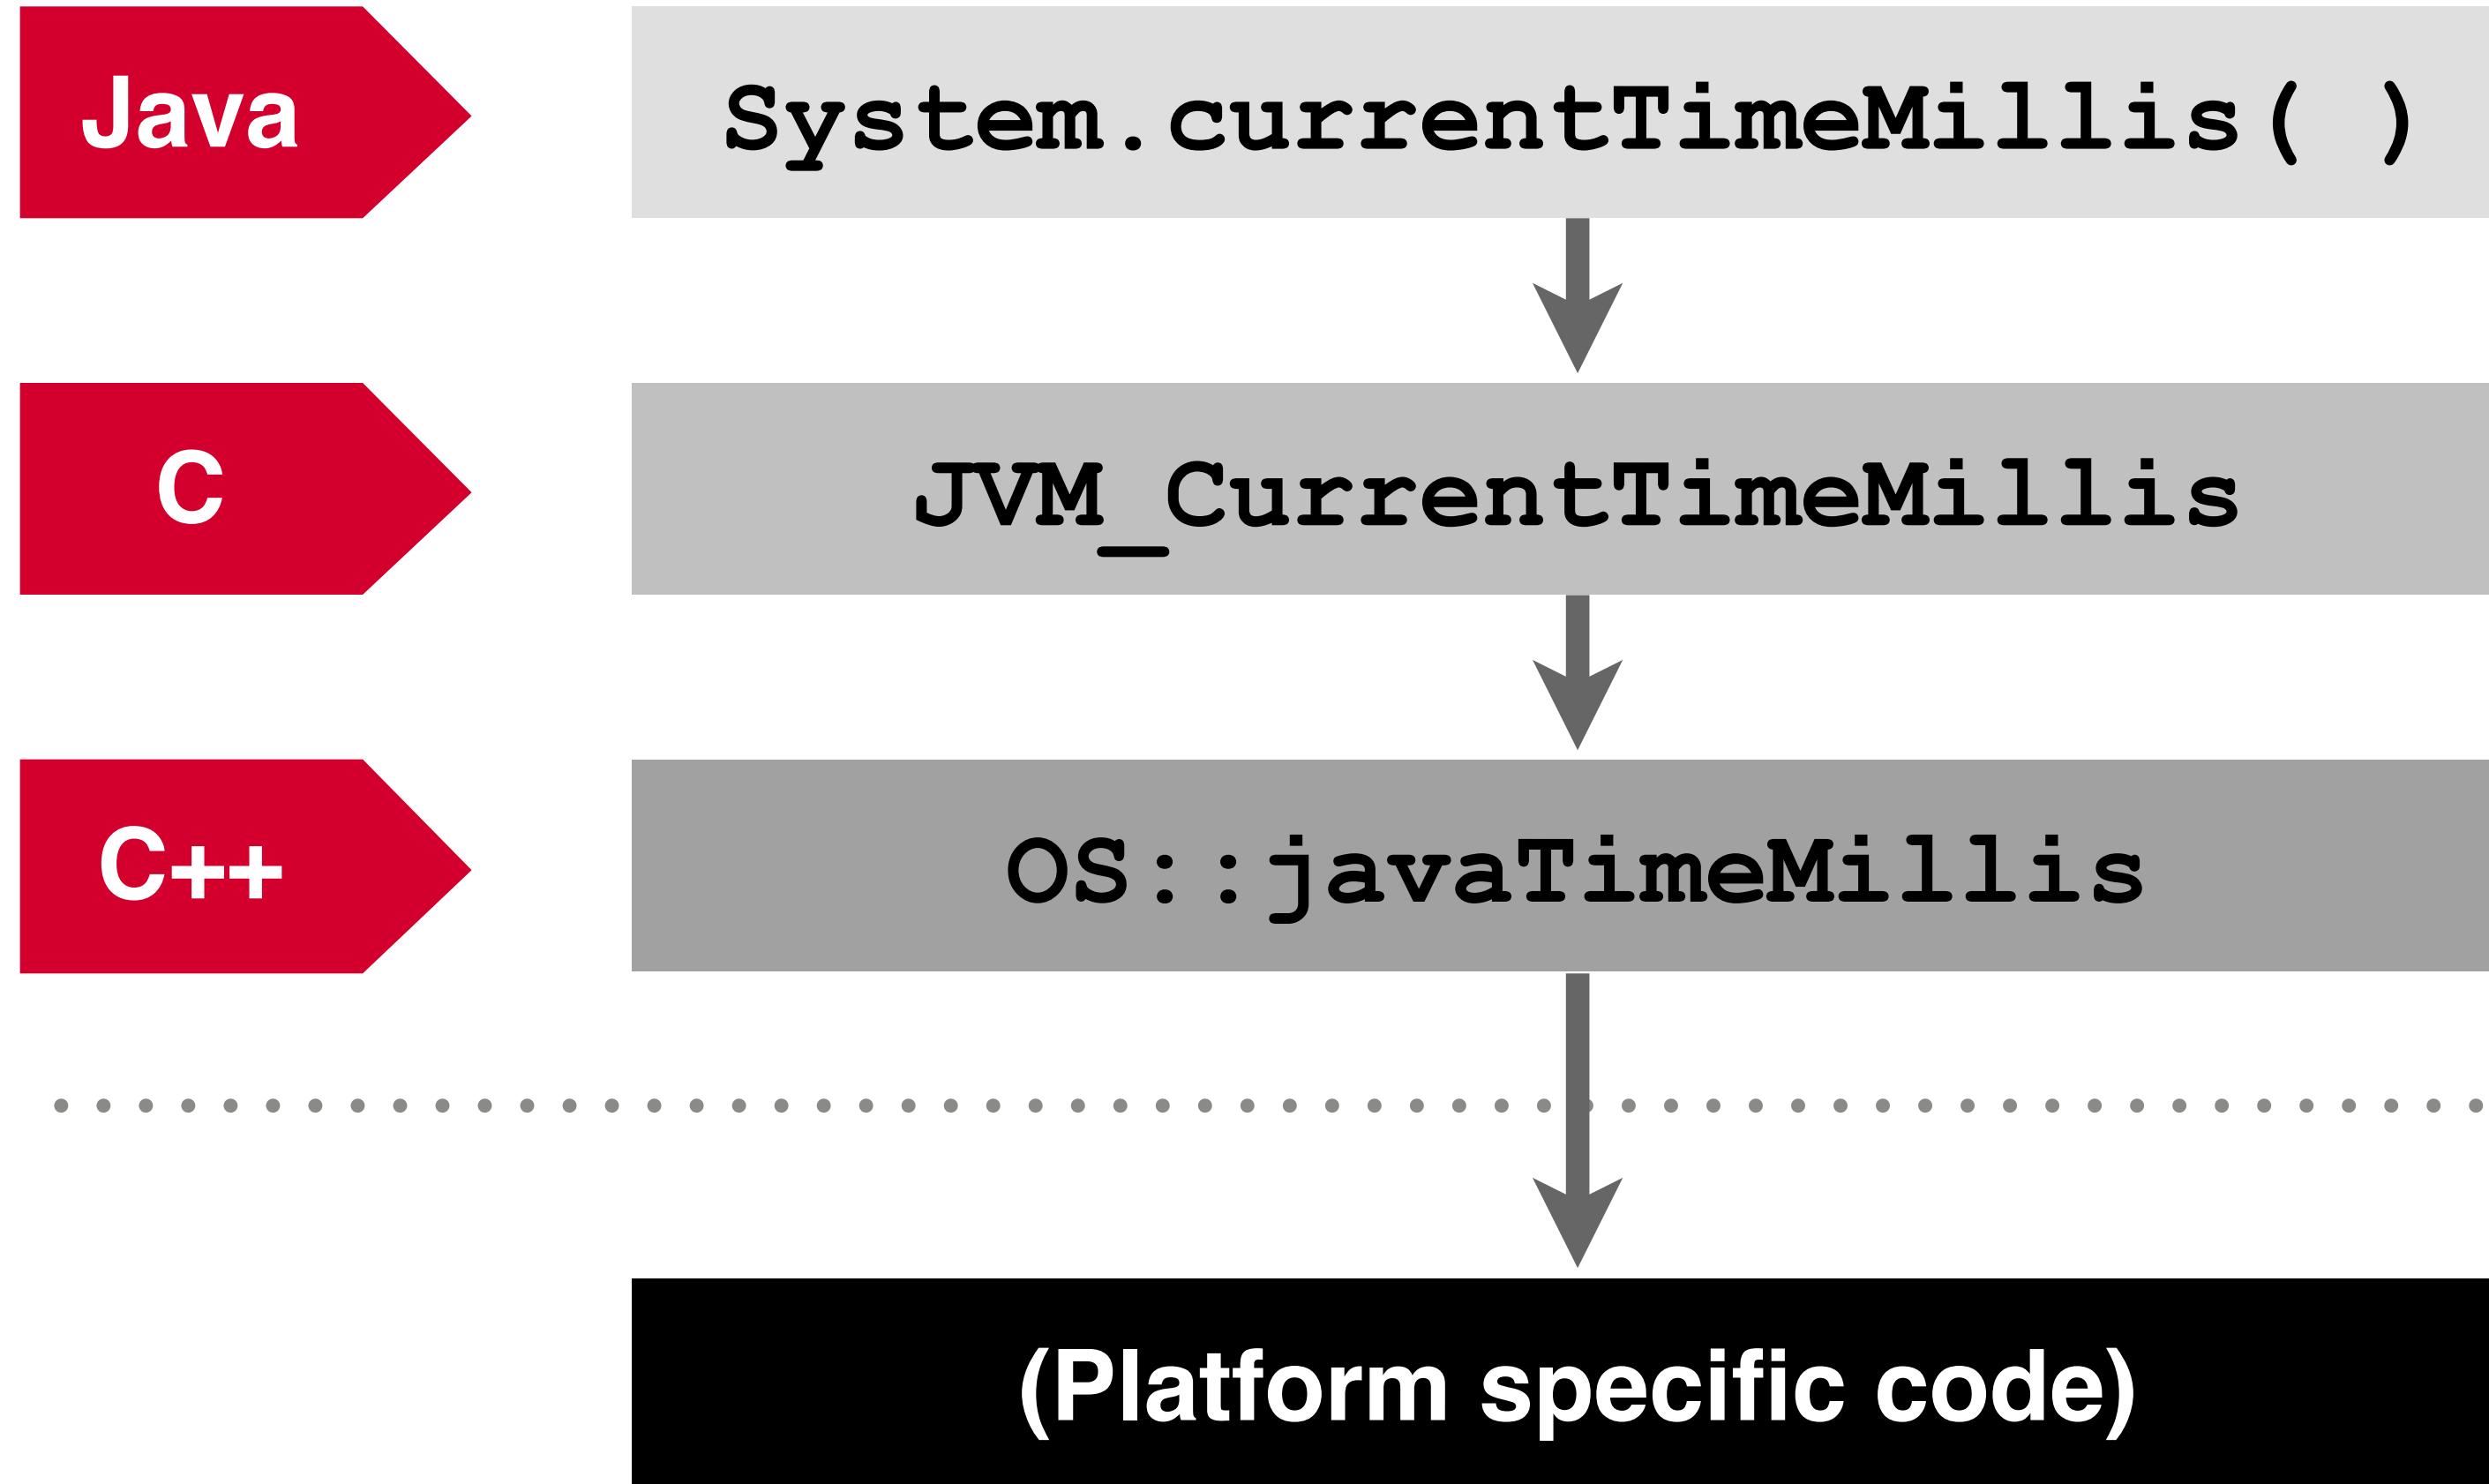


A Question of Time

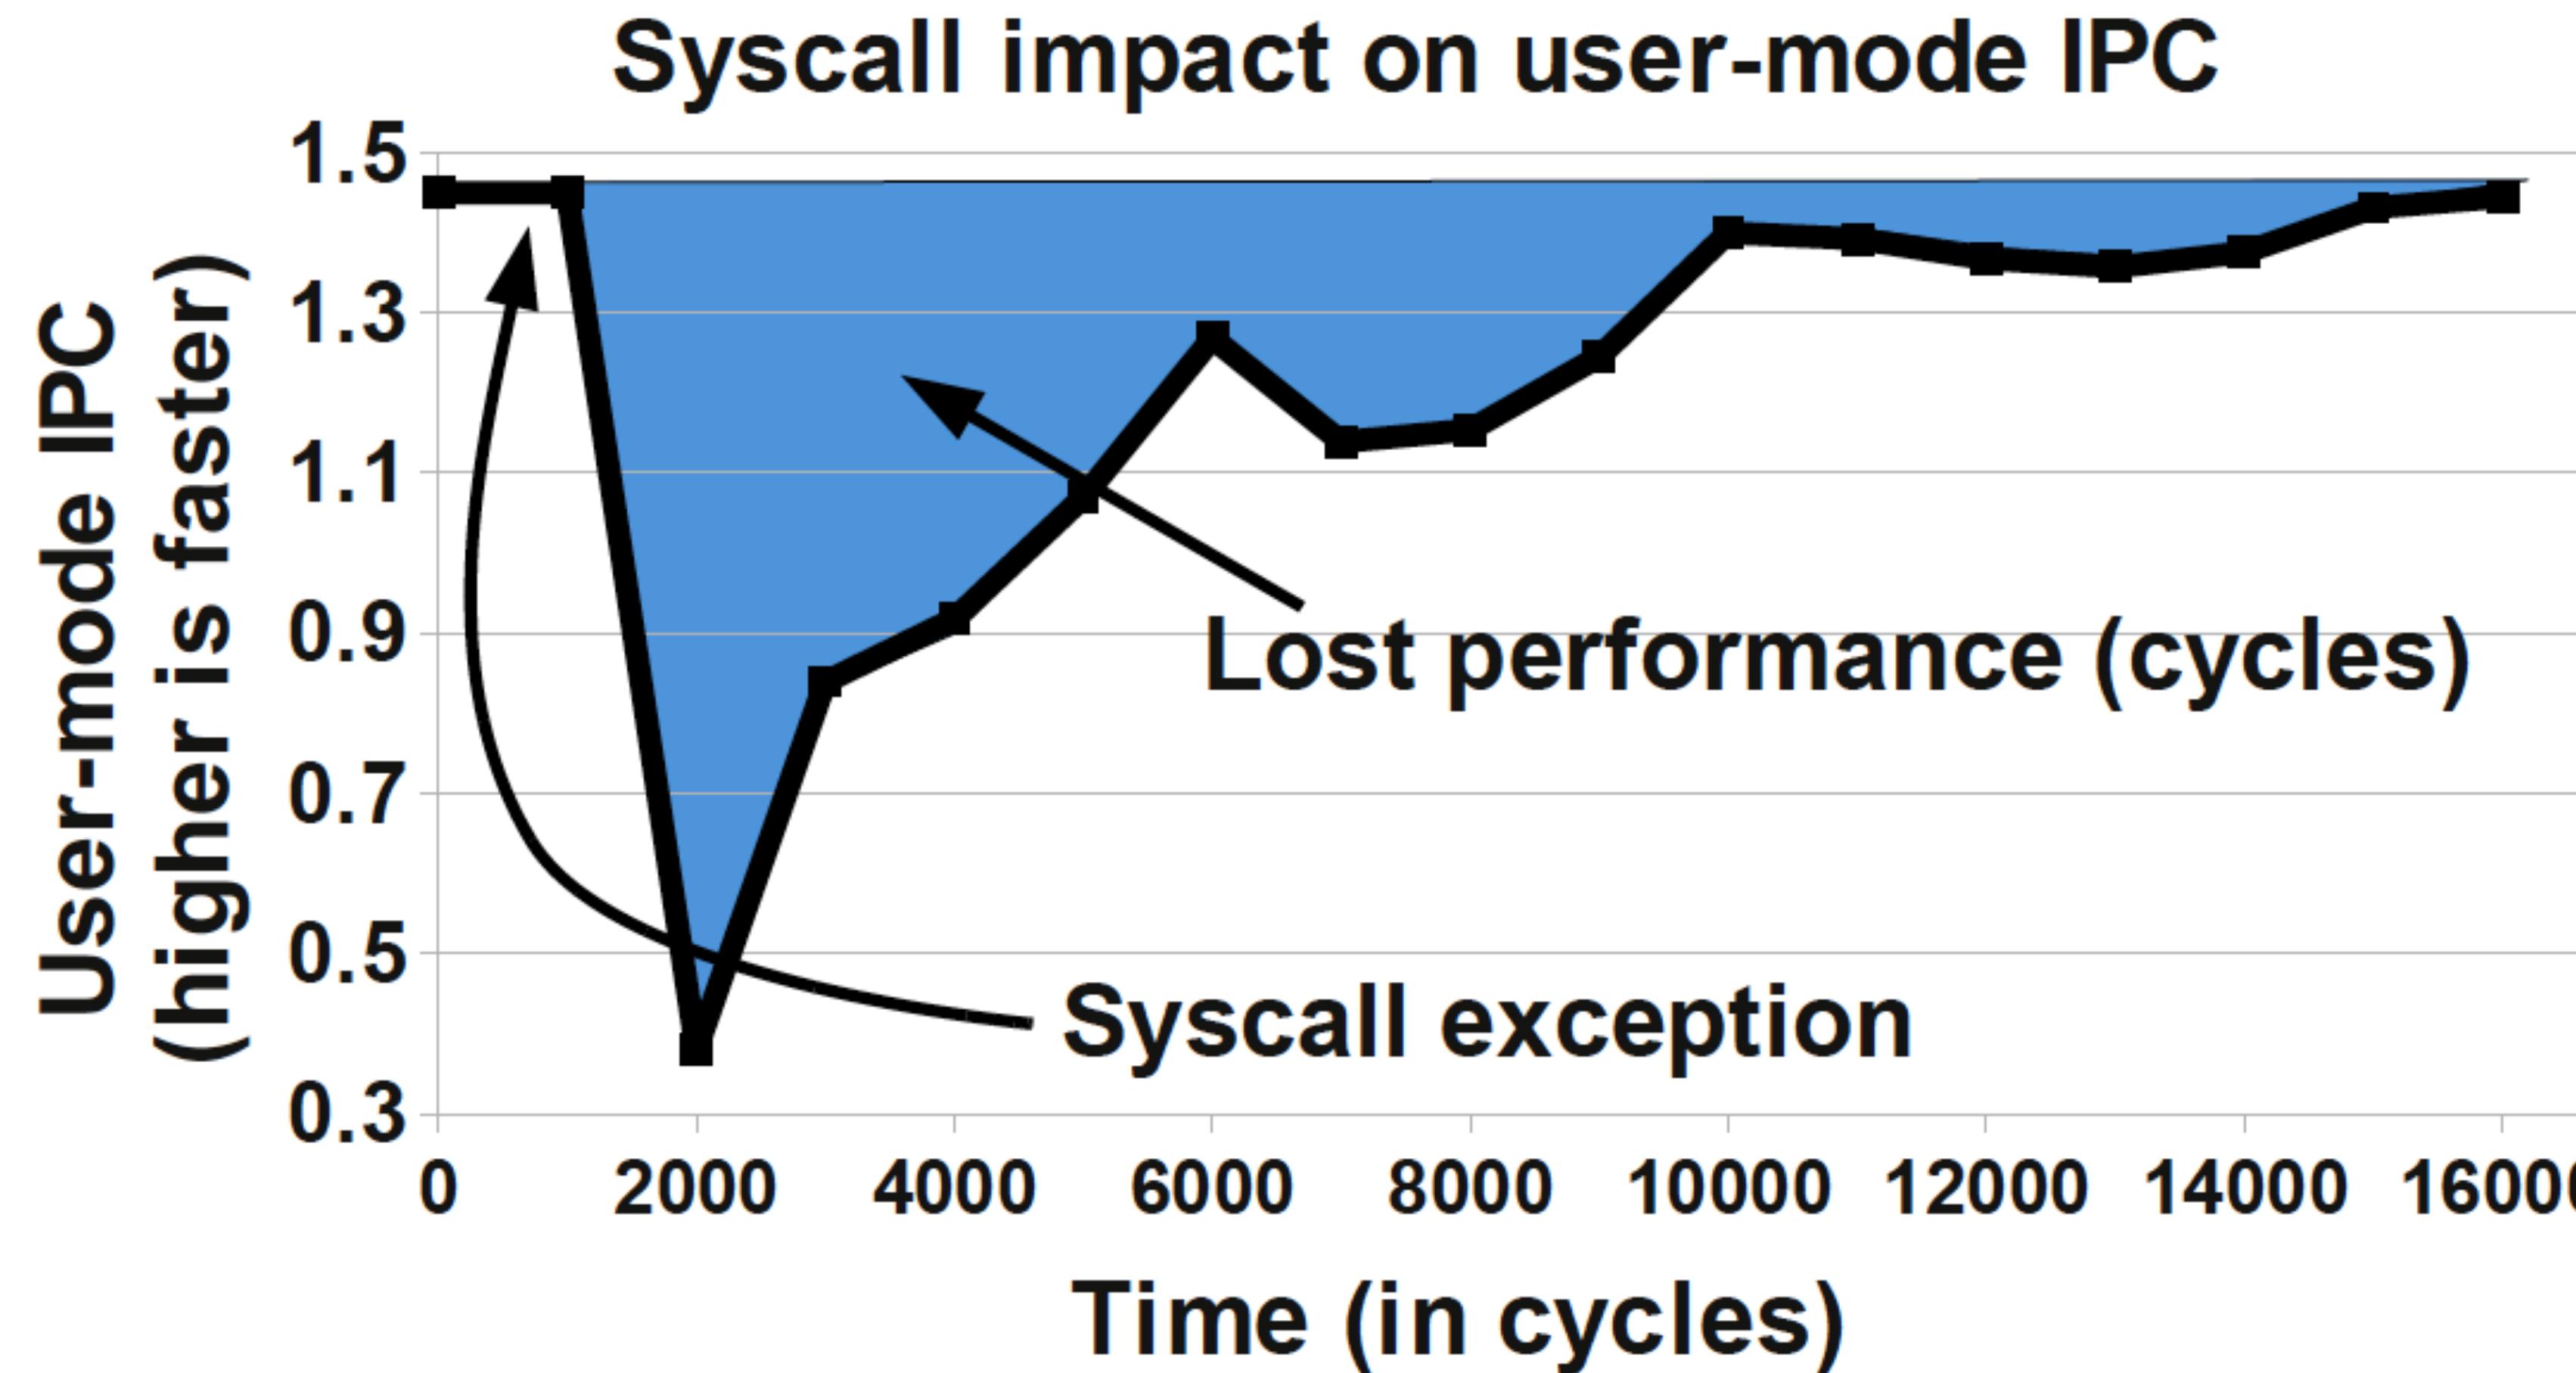
- Schedulers depend on timing subsystems
- Modern hardware may have up to 4 timers
 - RTC, 8254, TSC and possibly HPET
- Also NTP
- `System.nanoTime()` vs `currentTimeMillis()`



A Question of Time – Accessing the OS



Impact of System Calls



From: FlexSC: Flexible System Call Scheduling with Exception-Less System Calls,
by Livio Stoares & Michael Stumm



The Java Memory Model (JMM)

- Java has had a formal model of memory since 1.0
 - Heavily revised & fixed for Java 5 (JSR 133)
 - Continuing to evolve (new updates in Java 9)
- Increasingly important
- What happens when two cores access same data?
 - When are they guaranteed to see the same value?
 - How do memory caches affect this?
 - Cache coherency protocols (MESI)



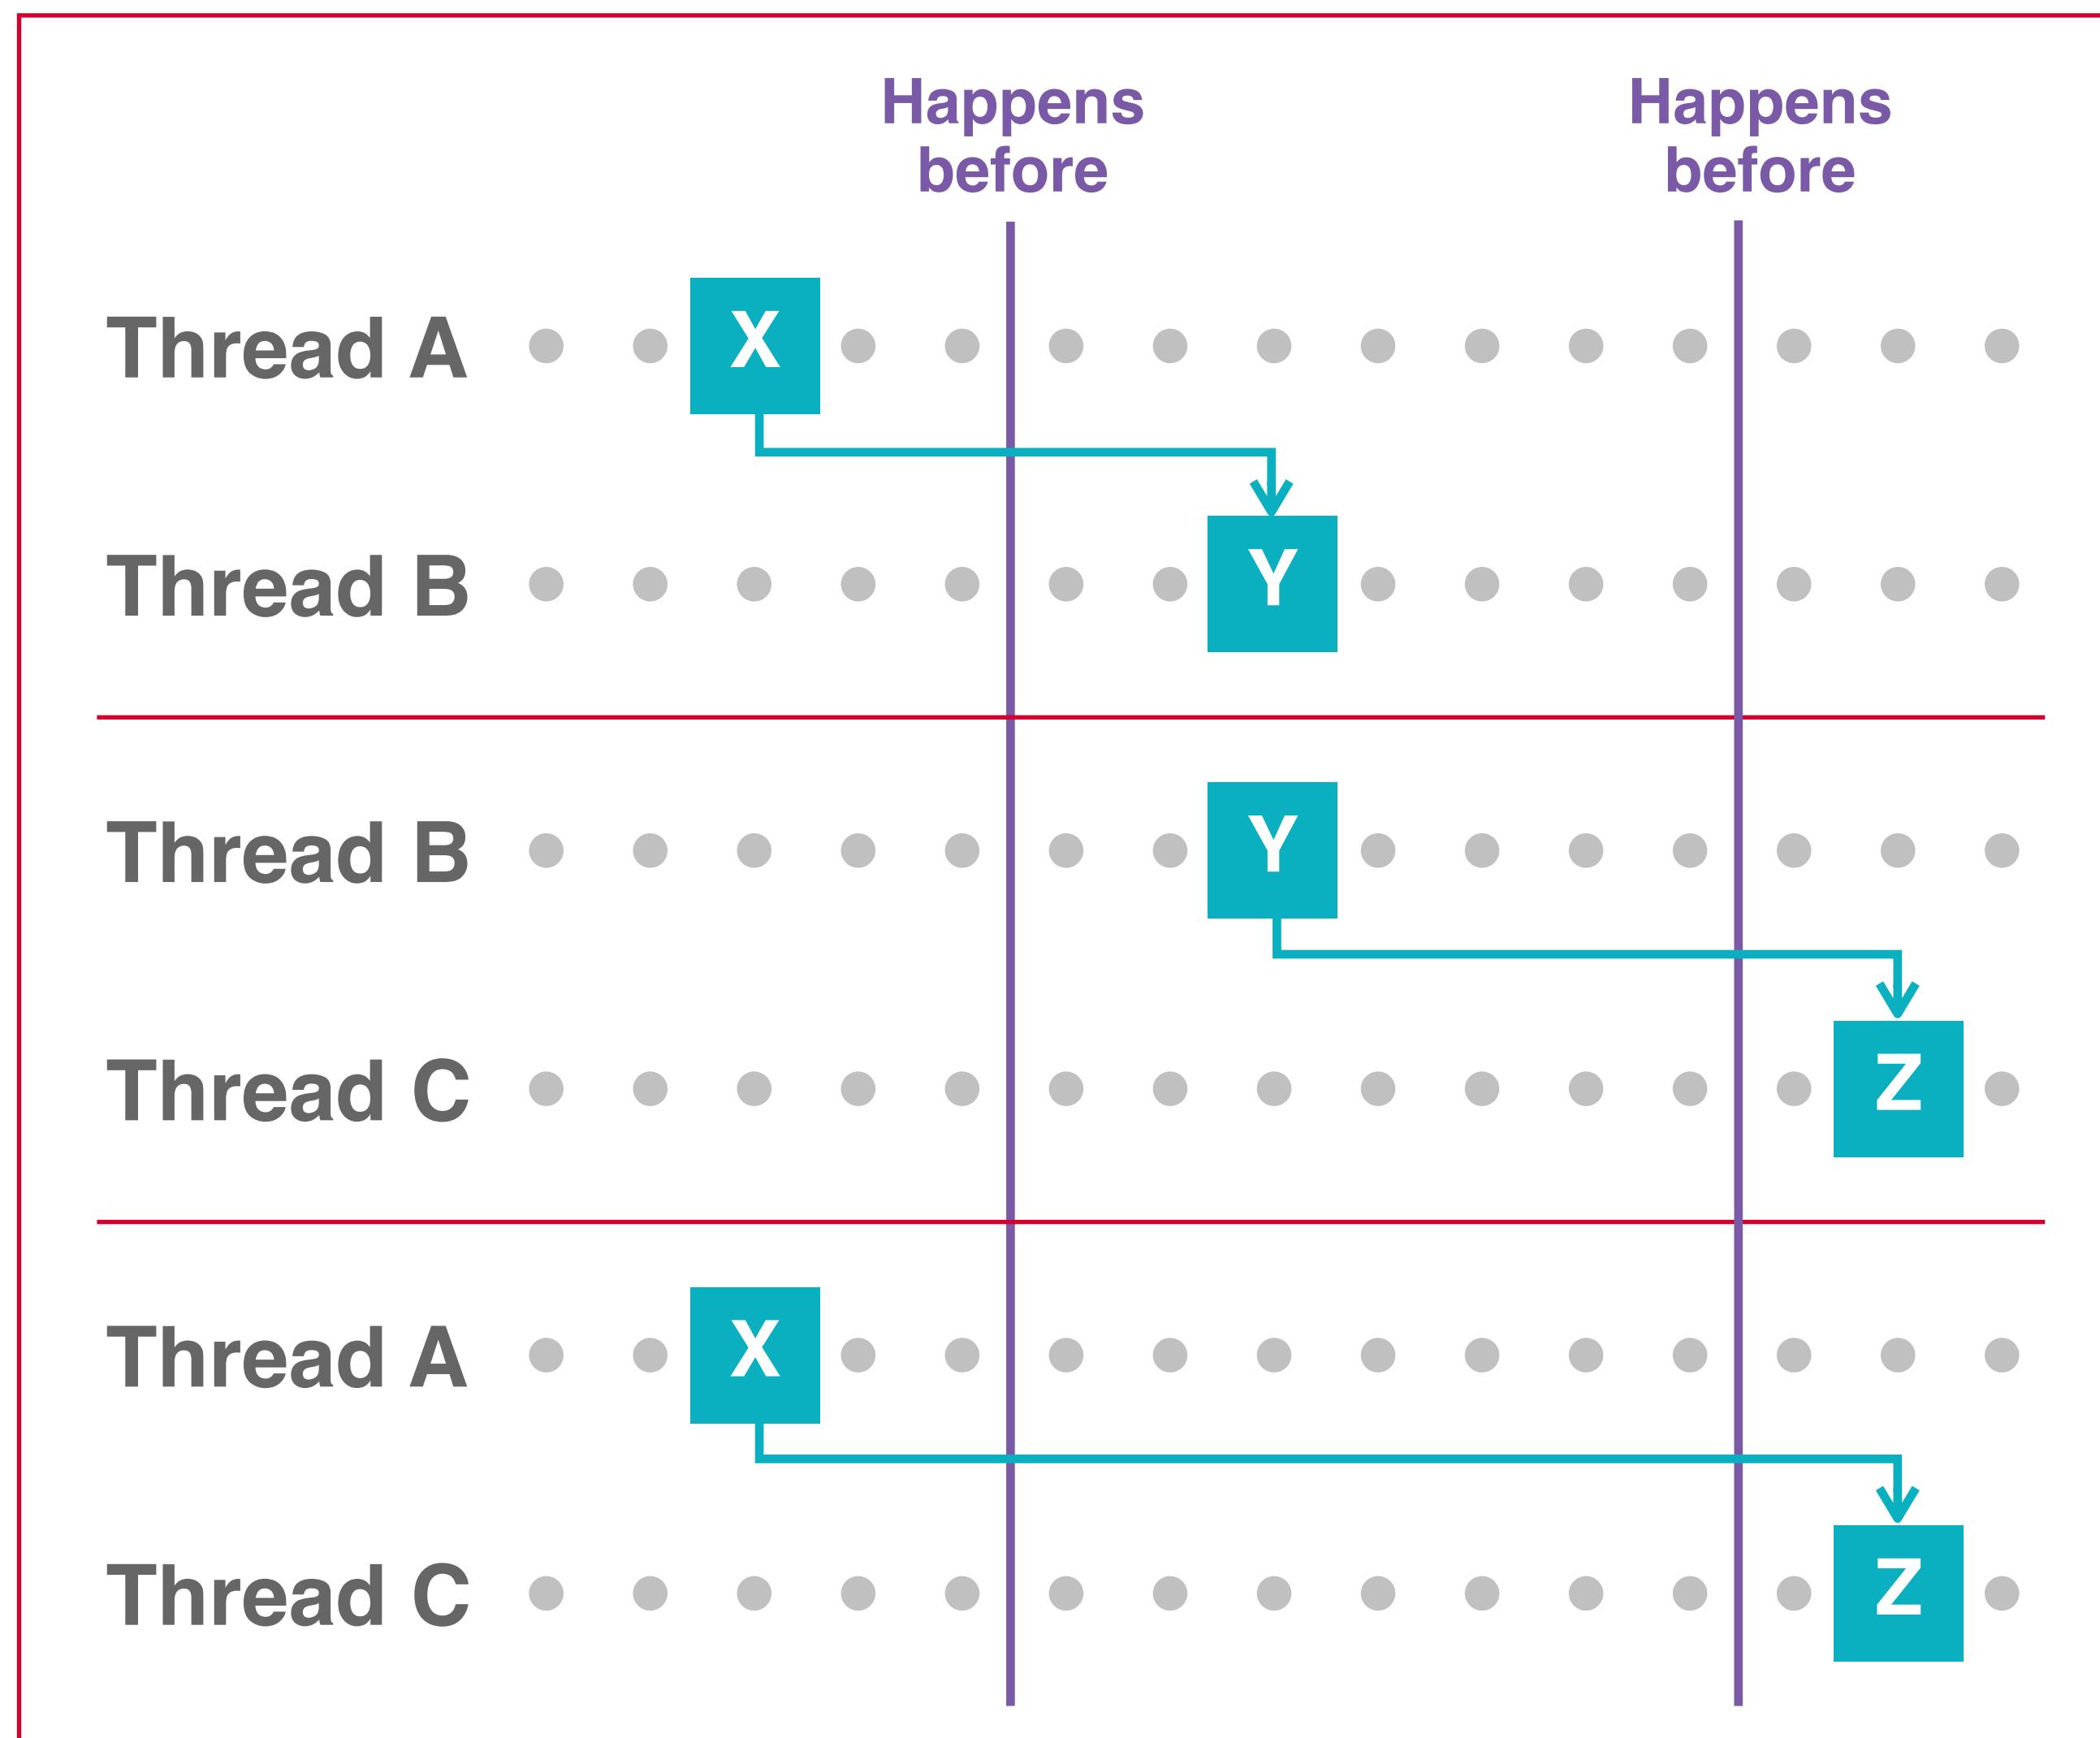
The Java Memory Model (JMM)

- Two possible approaches
 - Strong memory model (all cores always see same values)
 - Weaker memory model (special cache instructions)
- The JMM has a weak memory model
 - Fits better with trend in processors
 - Makes porting easier
- A strong memory model
 - requires more implementation code on top of weak hardware
 - scales worse as number of cores increases
 - Cache invalidation swamps the memory bus



The JMM

- Mathematical description of memory
- Most impenetrable part of JLS
- JMM makes minimum guarantees
 - Real JVMs (and CPUs) may do more



JMM – Key concepts

- Happens-Before
 - One block of code fully completes before other can start
- Synchronizes-With
 - Action will synchronize its view of an object with main memory
- As-If-Serial
 - Instructions appear to execute in-order
- Release-Before-Acquire



Synchronizes-with

- Threads have their own description of an object's state
 - Must be flushed to main memory (& other threads)
- **synchronized** means that local view has been **synchronized with** with other threads
- Defines touch-points where threads must perform synching



Synchronization Summary

- Unsynchronized methods
 - Don't look at or care about the state of any locks
 - Can progress while synchronized methods are running



Synchronization Summary

- Unsynchronized methods
 - Don't look at or care about the state of any locks
 - Can progress while synchronized methods are running
- Only objects can be locked – not primitives.
- Locking an `Object[]` doesn't lock the individual objects



Synchronization Summary

- Unsynchronized methods
 - Don't look at or care about the state of any locks
 - Can progress while synchronized methods are running
- Only objects can be locked – not primitives.
- Locking an `Object[]` doesn't lock the individual objects
- A synchronized method is equivalent to a `synchronized (this) { ... }` block that covers the entire method
- A static synchronized method locks the `Class` object, because there's no instance object to lock.



Synchronization Summary

- Synchronization in an inner class is independent of the outer class
- When locking class object, two options



Synchronization Summary

- Synchronization in an inner class is independent of the outer class
- When locking class object, two options
 - Explicit lock on the class literal
 - Using `getClass()`
 - The approaches have different behaviour in a subclass



Synchronization Summary

- Synchronization in an inner class is independent of the outer class
- When locking class object, two options
 - Explicit lock on the class literal
 - Using `getClass()`
- The approaches have different behaviour in a subclass
- **synchronized** isn't part of the method signature
 - It can't appear on a method declaration in an interface



Synchronization Summary

- Synchronization in an inner class is independent of the outer class
- When locking class object, two options
 - Explicit lock on the class literal
 - Using `getClass()`
- The approaches have different behaviour in a subclass
- **synchronized** isn't part of the method signature
 - It can't appear on a method declaration in an interface
- Java's locks are reentrant
 - A recursive call to a synchronized method won't self-deadlock



Modern JDK Concurrency Classes

- Executors
- Concurrent Data Structures
- Exercise: Getting to Know ConcurrentHashMap
- Queues, Latches and Barriers
- Locks in `java.util.concurrent`
- Q&A



Execution & Executors

- Deferred Execution
- Executors & Threadpools



Deferred Execution

- A class is the smallest unit of functionality
 - Functionality is loaded & linked as classes
 - Must go through classloading
 - There's no other way to do it
- Java has no “free functions”



Deferred Execution

- Representing Execution means we're deferring execution
 - Have a reference to some code which will be executed later
 - Callable: "Here's some code. Run it when we call `call()`"
 - Runnable: "Start this code up in a different thread (via a Thread ctor)"
 - Callbacks: "Here's some code - call it back when something happens"
- Deferring Execution comes up in other contexts
 - Closures / Lambda Expressions / Function pointers
 - Libraries (e.g. map / filter idioms)



Executors Factory

- `java.util.concurrent` has factories for making threadpools
 - In the Executors class
 - Another plurals class!
- Like this:

```
ScheduledExecutorService stpe =  
    Executors.newScheduledThreadPool(2);
```



Don't Roll Your Own!

- Greenspun's Tenth Rule:

“Any sufficiently complicated C or FORTRAN program contains an ad hoc, informally-specified, bug-ridden, slow implementation of half of Common Lisp.”



Don't Roll Your Own!

- Greenspun's Tenth Rule:

“Any sufficiently complicated C or FORTRAN program contains an ad hoc, informally-specified, bug-ridden, slow implementation of half of Common Lisp.”

- There's a similar law about multithreaded programs
 - If you aren't using `java.util.concurrent`
 - Then you'll end up reimplementing half of the classes in it
 - Badly, usually.



Ready-Made ThreadPools

- `Executors` has factory methods
 - `newFixedThreadPool(int nThreads)`
 - `newCachedThreadPool()`
 - `newSingleThreadExecutor()`
 - `newScheduledThreadPool(int corePoolSize)`



FixedSize ThreadPool

- Threads will be reused to run multiple tasks
- Prevents having to pay the cost of thread creation
- If threads are in use new tasks are stored in a queue
- Useful if task flow is stable and known



CachedThreadPool

- Will create new threads as required
- Will reuse threads where possible
- Created threads are kept for 60 seconds
 - After which they will be removed from the cache
- Can give better performance with small asynchronous tasks
 - Or bursty workloads



Single Thread Executor

- Backed by a single thread
- Newly submitted tasks are queued until the thread is available
- Can be useful to control number of tasks concurrently executed
- Also useful for testing



ScheduledThreadPoolExecutor

- Common choice of threadpool for many applications

```
ScheduledExecutorService stpe =  
    Executors.newScheduledThreadPool(poolsize);
```

- Schedule events to execute at a fixed rate

```
stpe.scheduleAtFixedRate(msgReader, initialDelay,  
    period, TimeUnit.MILLISECONDS);
```

- Returns a ScheduledFuture for pending execution
- Good for controlled concurrency work



ScheduledThreadPoolExecutor

- Schedule all threads to run
 - `stpe.invokeAll()`
 - Executes a bunch of tasks at once
 - Returns a collection of Future objects
 - Can be beneficial when you have sufficient hardware
- `shutdown()`, `shutdownNow()`
 - Used to shutdown threadpools



Concurrent Data Structures

- Atomic Data Structures
- Concurrent Maps
- Concurrent Lists and Sets



AtomicInteger & AtomicLong

- AtomicInteger & AtomicLong can be used as counters
- Can be used as a replacement for volatile variables
 - More flexible
 - Can do “state-dependent” updates (e.g. increment)
 - Still lock-free & can’t deadlock



AtomicInteger & AtomicLong

- AtomicInteger & AtomicLong can be used as counters
- Can be used as a replacement for volatile variables
 - More flexible
 - Can do “state-dependent” updates (e.g. increment)
 - Still lock-free & can’t deadlock
- Have composite operations
 - add or increment / decrement combined with get
 - Compare And Set (CAS) operations
 - These will execute atomically
 - Results are immediately seen by other threads



Atomic Classes

- Atomic classes don't inherit from similarly named classes
 - `AtomicBoolean` can't be used in place of a `Boolean`
 - `AtomicInteger` isn't an `Integer`
 - Does extend `Number`
- Why Not?



Adders

- Java 8 introduces scalable accumulators & adders
- Performance enhancement for contented shared state
 - As number of threads adding to an `AtomicLong` rises performance worsens
- `LongAdder` is designed to solve this problem
 - More scalable in the number of threads that are adding

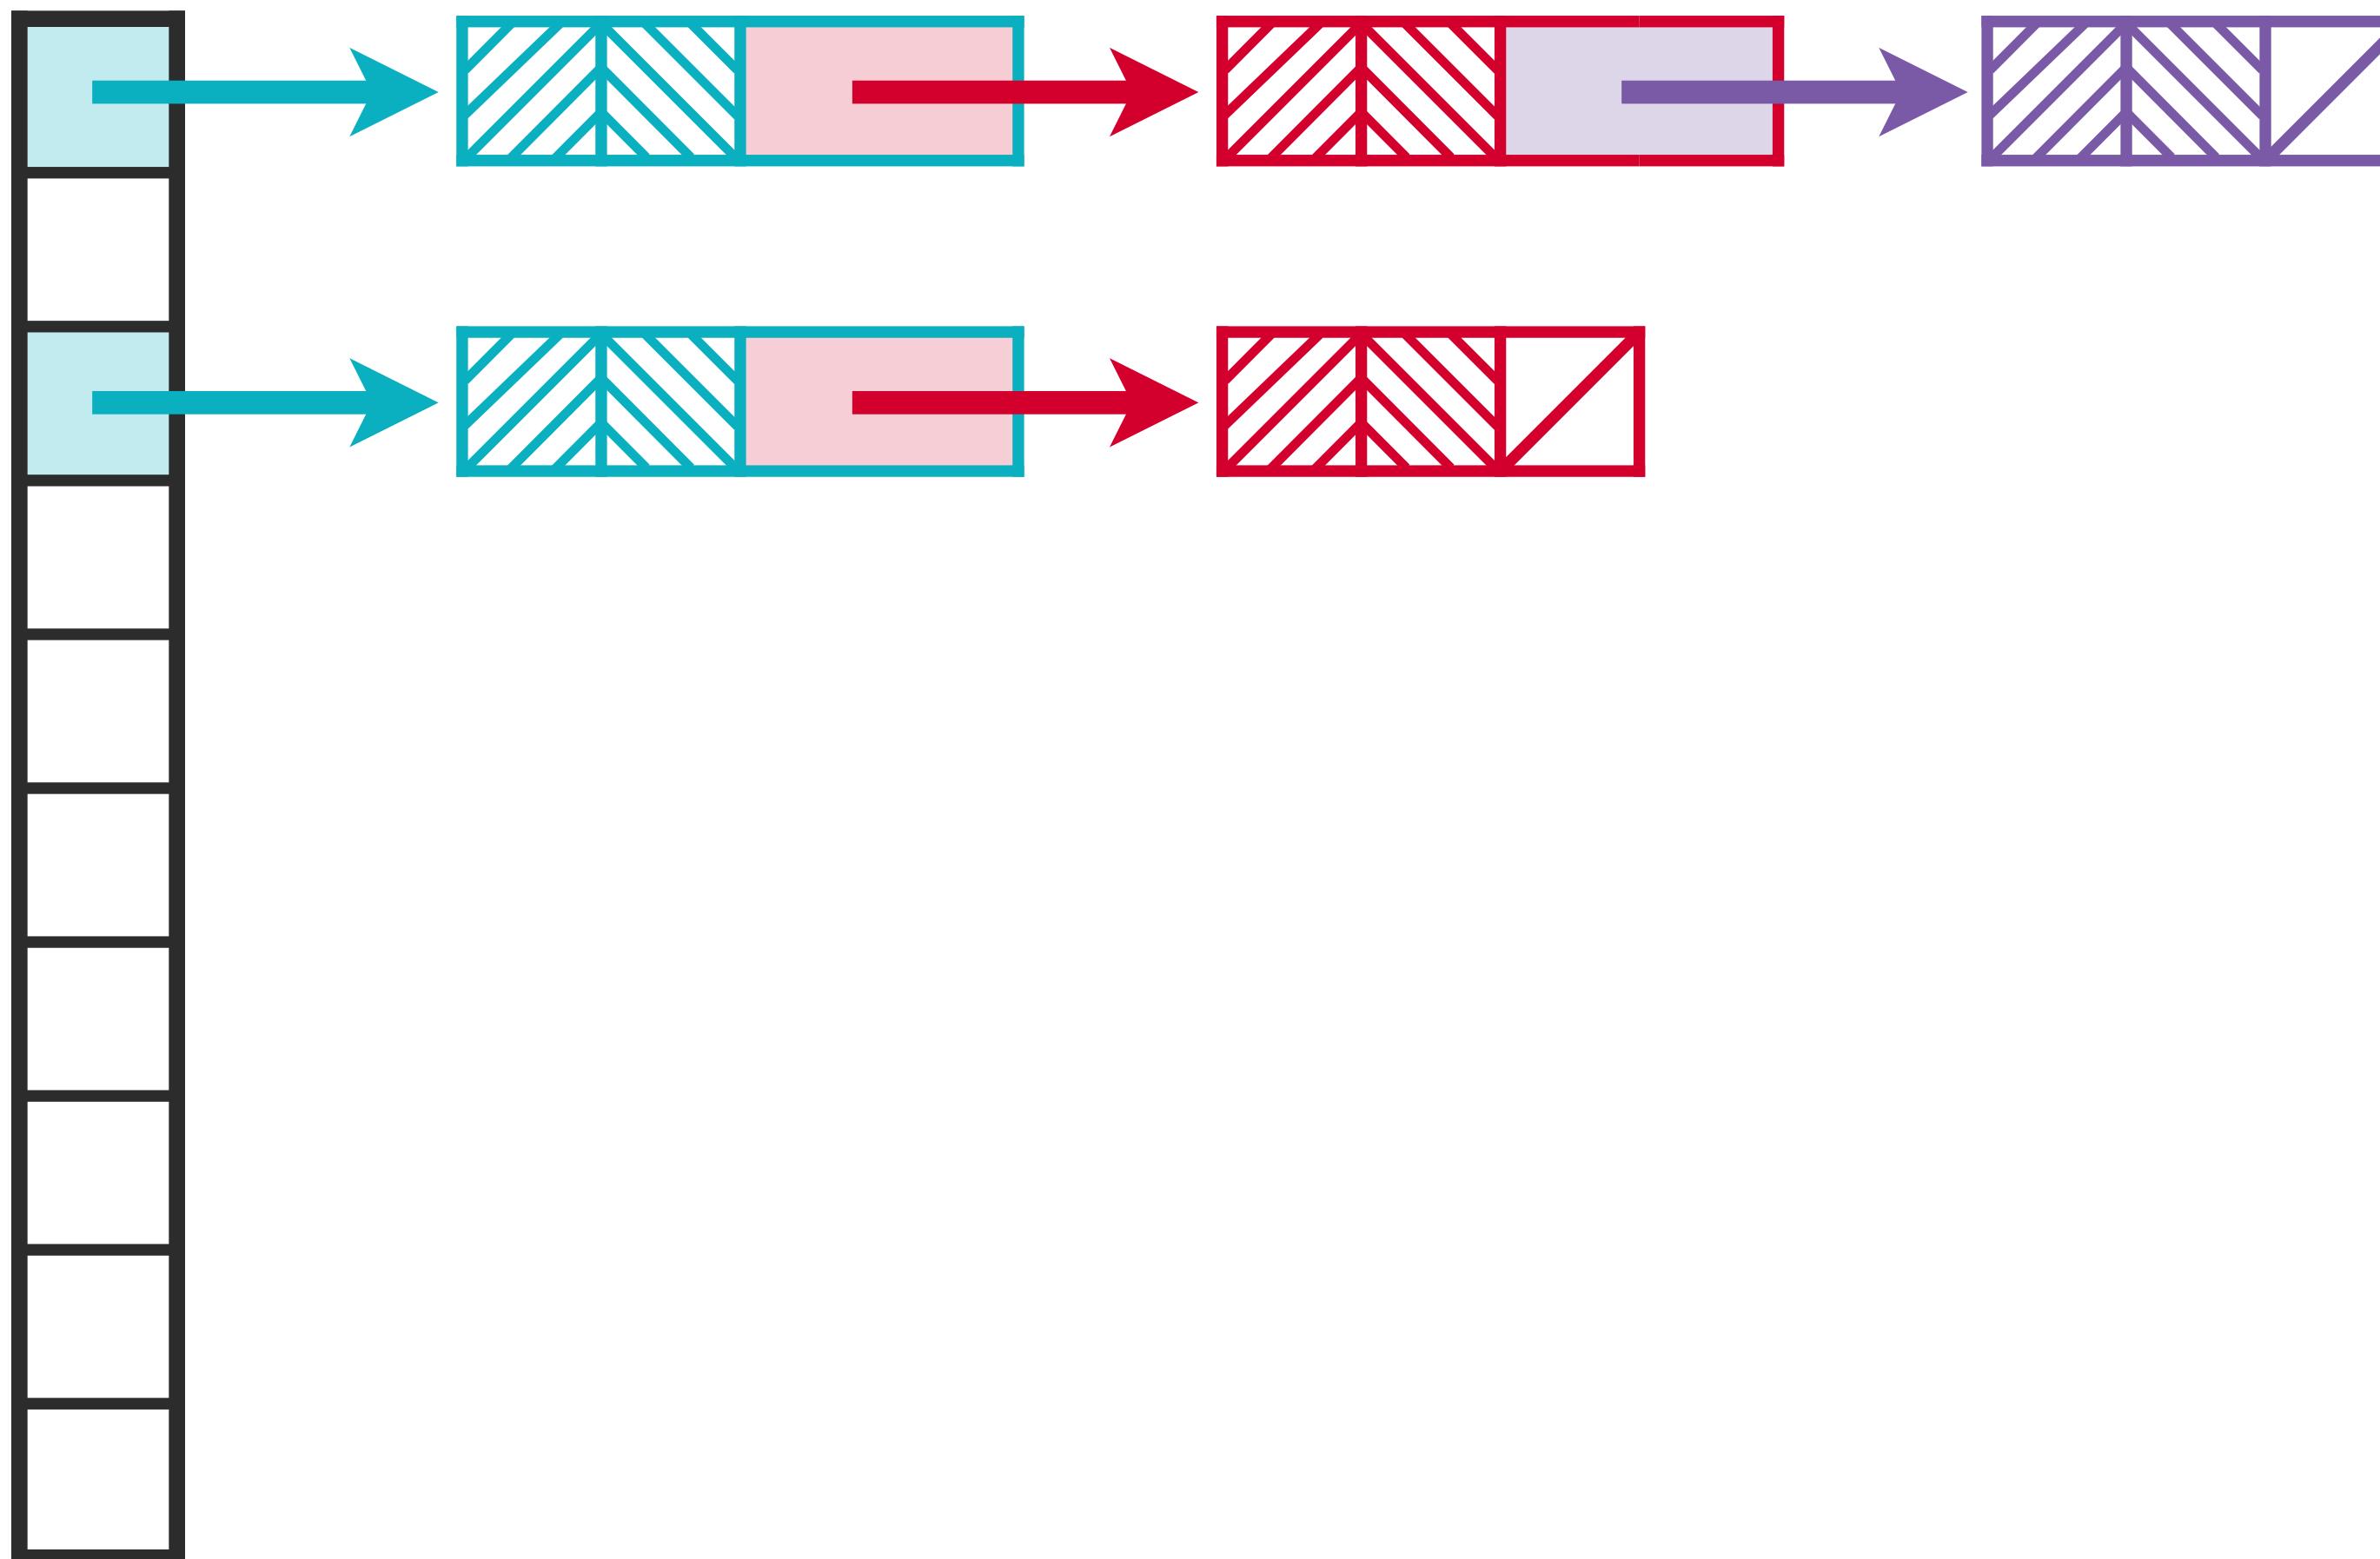


Accumulators & Adders

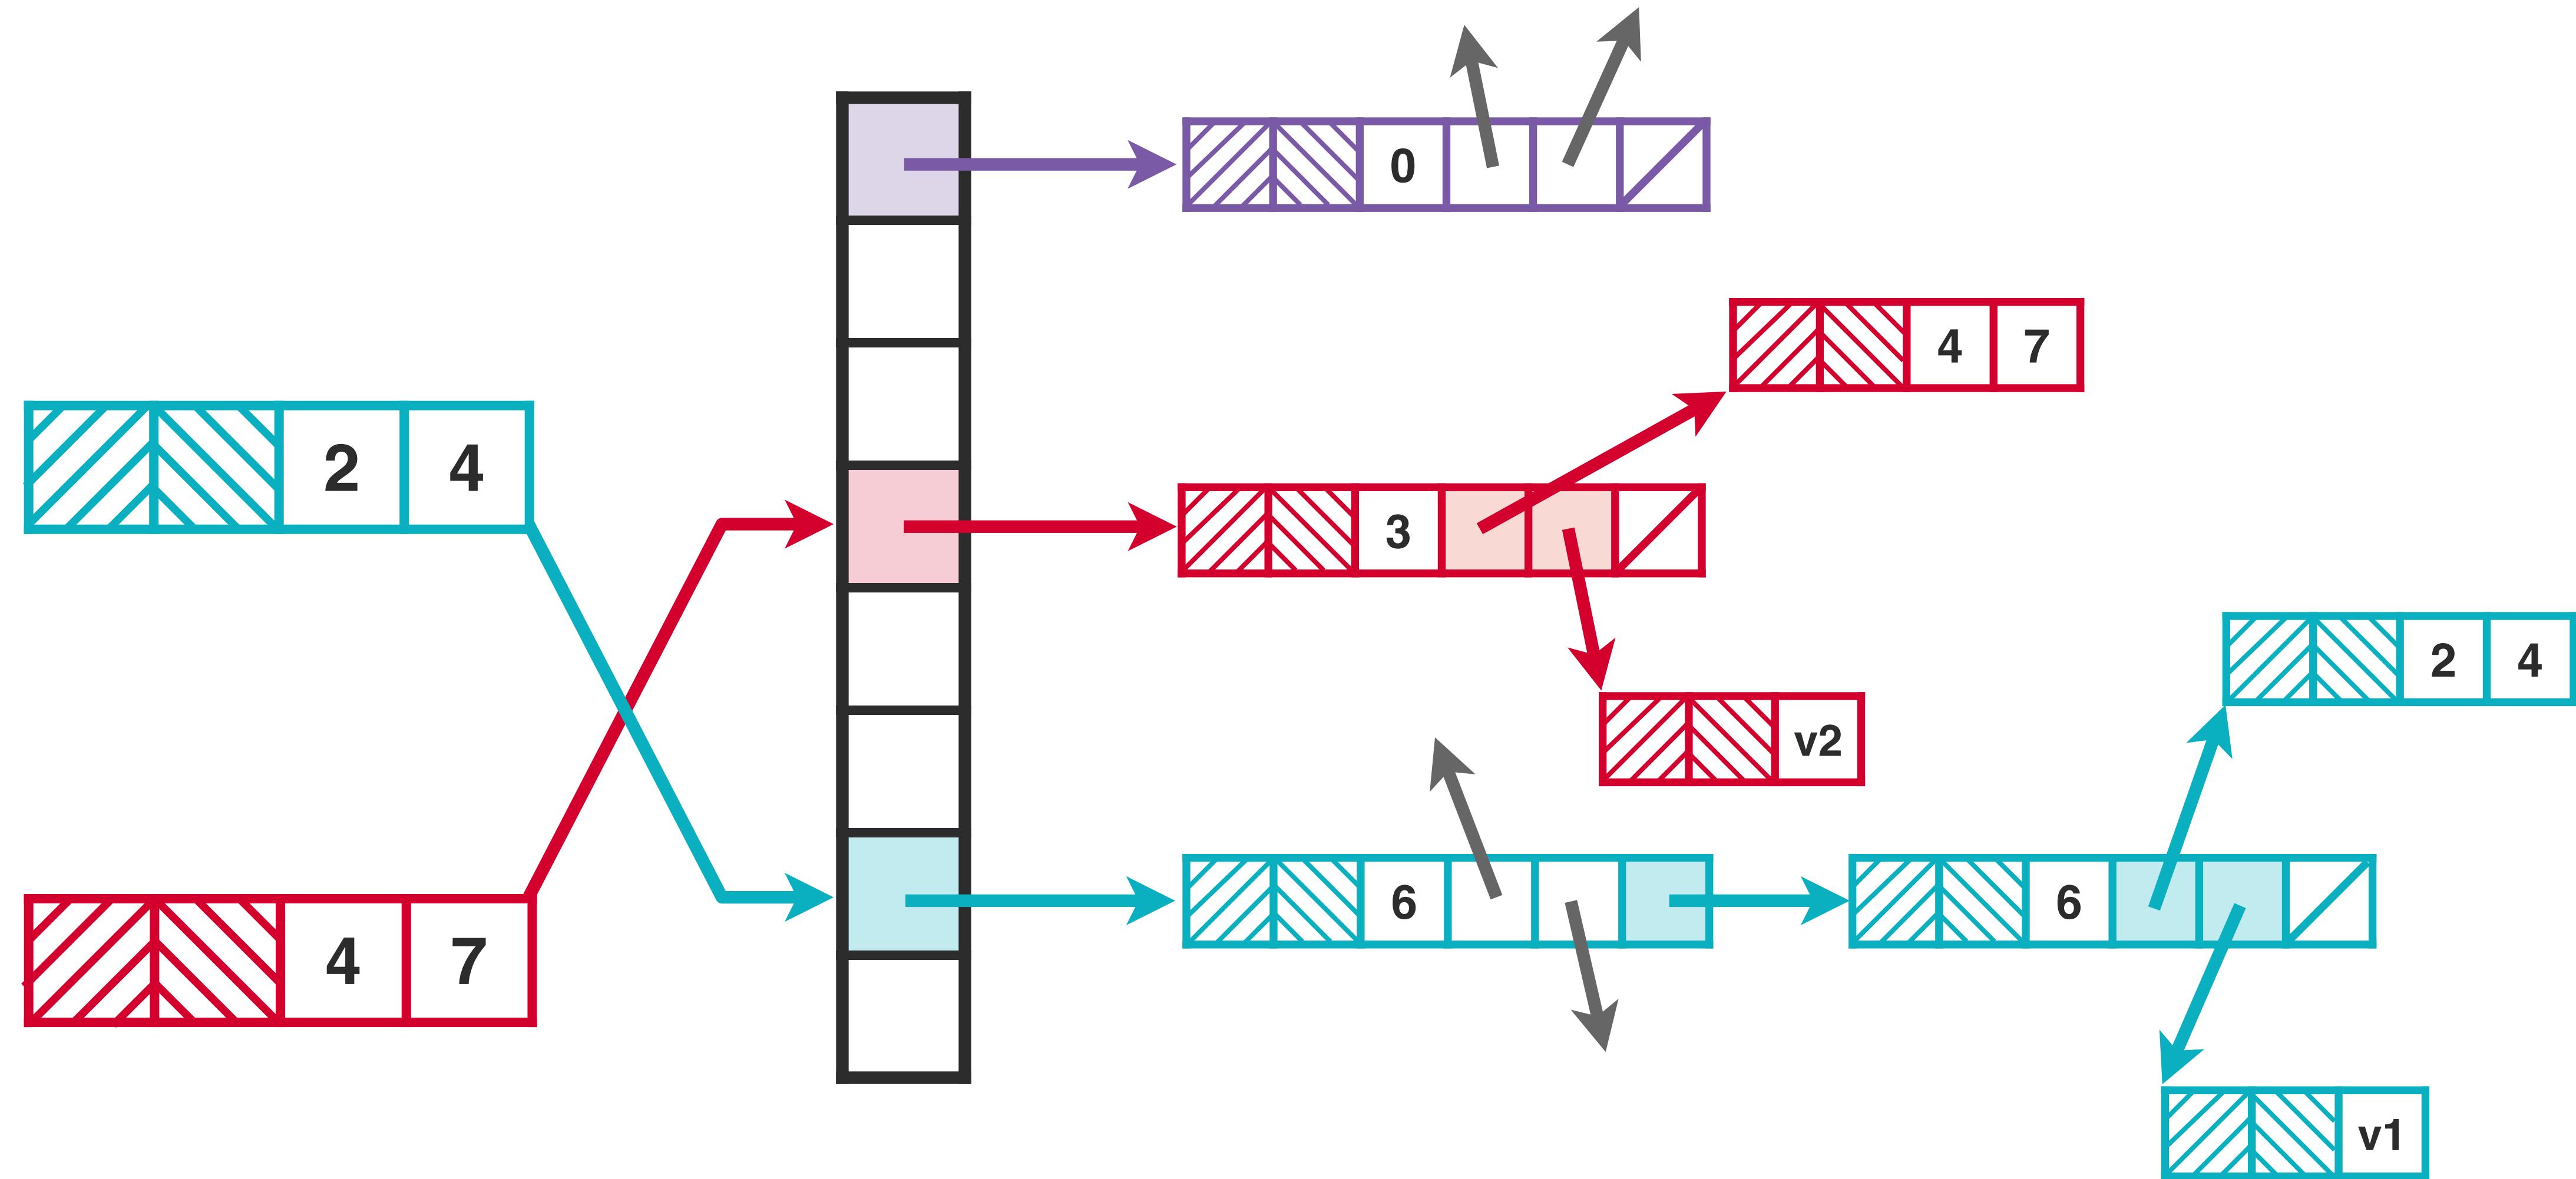
- Accumulators are a generalisation of adders
- Related to the reduce operations from functional code
 - `new DoubleAdder()` is just
`new DoubleAccumulator((x, y) -> x + y, 0.0)`
- Same contention behaviour as adders



Hashing – for Data Structures



HashMap at Runtime



SimpleDict Internals

```
public class SimpleDict implements Map<String, String> {  
    private final Node[ ] table; // Hash buckets (head of link-lists)  
    private int size; // num of entries in the dict
```



Aside: SimpleDict Internals

```
public String get(Object key) {  
    // Null keys are not supported  
    if (key == null)  
        return null;  
  
    int hash = key.hashCode();  
    int i = indexFor(hash, table.length);  
    for (Node e = table[i]; e != null; e = e.next) {  
        String k;  
        if (e.hash == hash  
            && ((k = e.key) == key || key.equals(k)))  
            return e.value;  
    }  
  
    return null;  
}  
  
private int indexFor(int h, int length) {  
    return h & (length - 1);  
}
```

```
private class Node implements  
Entry<String, String> {  
    final int hash;  
    final String key;  
    String value;  
    Node next;  
  
    Node(int h, String k,  
        String v, Node n) {  
        hash = h;  
        key = k;  
        value = v;  
        next = n;  
    }  
}
```



CHM v7

- Simple segmented model
- One **ReentrantLock** object per hash bucket
- Default concurrency level: 16
- So CHM v7 is memory intensive
 - $\sim \text{sizeof}(\text{HashMap}) + 16 * \text{sizeof}(\text{ReentrantLock})$

Class	Size
HashMap	216
ReentrantLock	72
ConcurrentHashMap	2272



CHM v8

- In Java 8 implementation of **ConcurrentHashMap** changes
- Once hash chain crosses threshold
- Replace linked list with balanced tree
- Improves worst-case perf from $O(n)$ to $O(\log n)$
 - The case of high hash collisions
- Increased code complexity



New Methods on Map

- Java 8 introduced new methods on Map
 - Using the default methods feature
 - Add “missing” features from CHM
 - lambda expressions



New Methods on Map

default	V	compute(K key, BiFunction<? super K, ? super V, ? extends V> remappingFunction)
default	V	computeIfAbsent(K key, Function<? super K, ? extends V>
default	V	computeIfPresent(K key, BiFunction<? super K, ? super V, ? extends V> remappingFunction)
default	void	forEach(BiConsumer<? super K, ? super V> action)
default	V	getOrDefault(Object key, V defaultValue)
default	V	merge(K key, V value, BiFunction<? super V, ? super V, ? extends V> remappingFunction)
default	V	putIfAbsent(K key, V value)
default	boolean	remove(Object key, Object value)
default	V	replace(K key, V value)
default	boolean	replace(K key, V oldValue, V newValue)
default	void	replaceAll(BiFunction<? super K, ? super V, ? extends V> function)



What about ArrayList?

- CHM intended to be a drop-in replacement for `HashMap`
- What about the sequential case (`ArrayList`)?
- Could just use `Collections.synchronizedList()`
 - Usually a lot poorer performance
- Instead, want a sequential data structure that is consistent
 - Will not throw a `ConcurrentModificationException`
- Uses a versioned data model
- Also `CopyOnWriteArrayList`



More Building Blocks

- Queues
- Latches
- Other Barriers

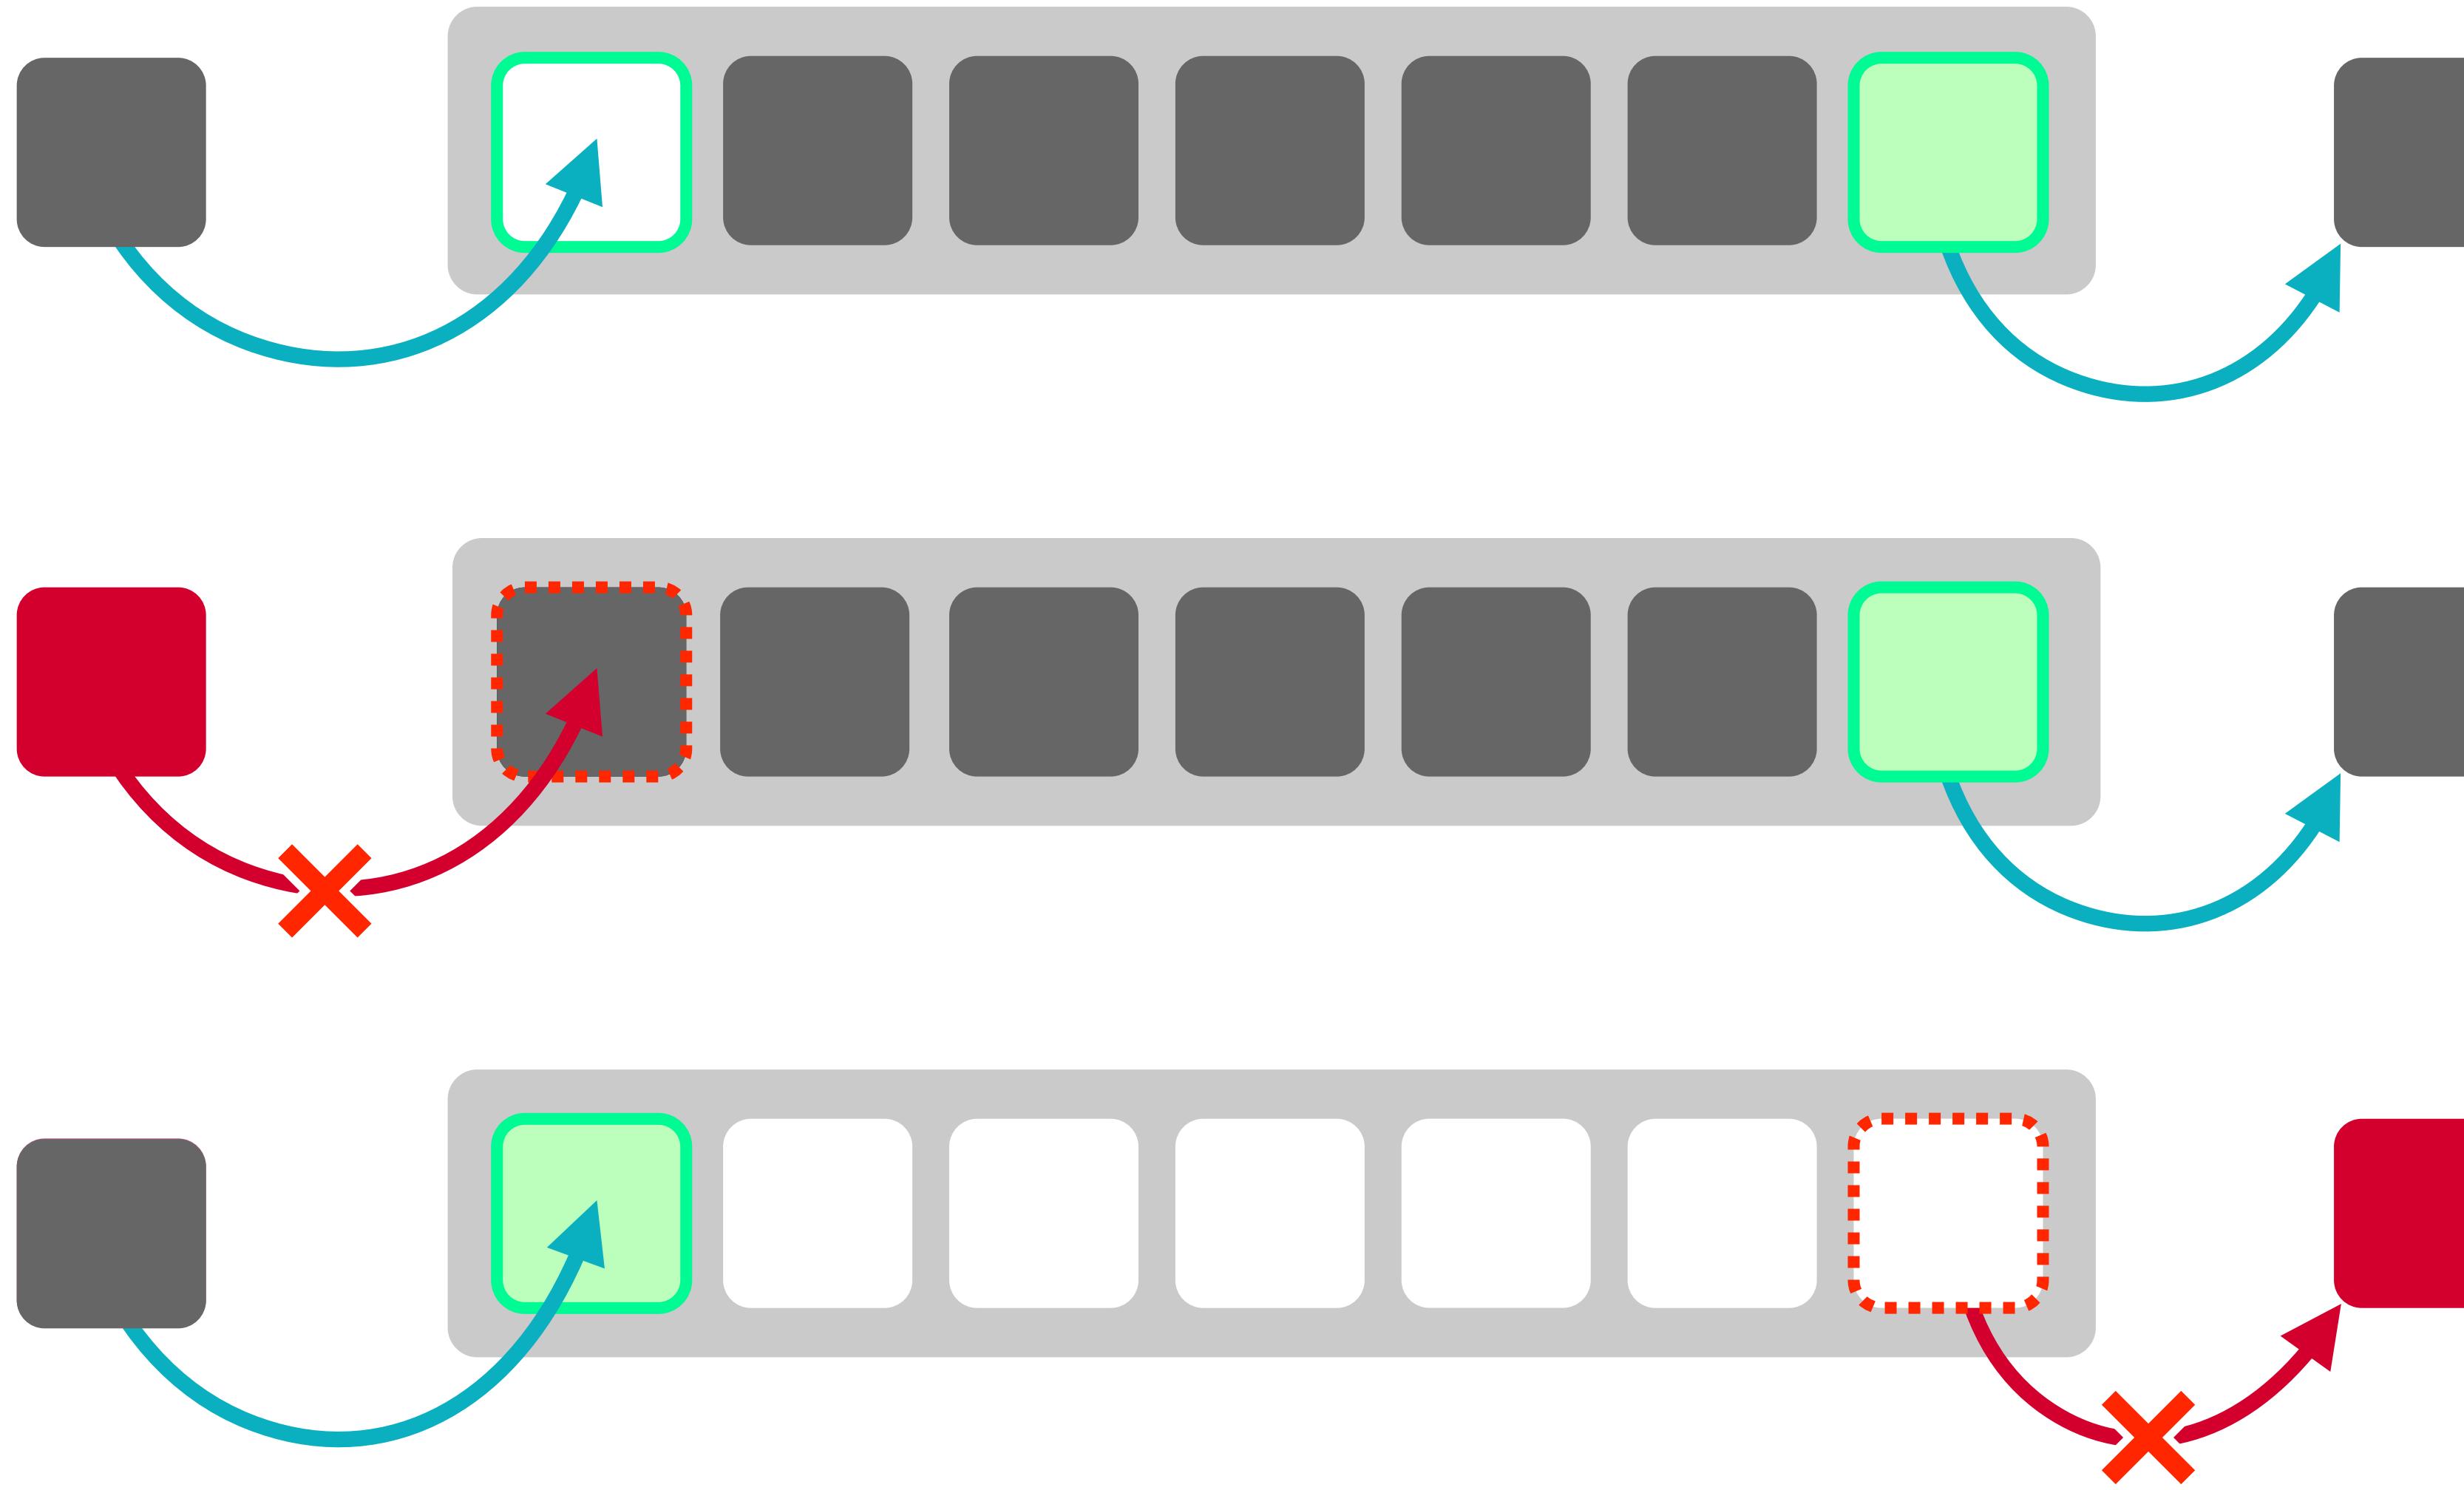


Blocking Queues

- Queues are often used to link execution together
- Efficient way to hand off work between threads
- **BlockingQueue** is the usual choice
 - Two basic impls – **ArrayList** and **LinkedList** backed
 - Java 7 introduces the new **TransferQueue**



Blocking Queue



BlockingQueue

- Offers several different ways to interact with it
 - `add()`, `remove()`, `element()` throw exceptions on failure
 - `offer()`, `poll()`, `peek()` return a special value on failure
 - `put()`, `take()` block until they succeed (note - no “examine” method)
- See Javadoc for full details



BlockingQueue methods

```
public interface BlockingQueue<E> extends Queue<E>
```

	Throws exception	Special value	Blocks	Times out
Insert	<code>add(e)</code>	<code>offer(e)</code>	<code>put(e)</code>	<code>offer(e, time, unit)</code>
Remove	<code>remove()</code>	<code>poll()</code>	<code>take()</code>	<code>poll(time, unit)</code>
Examine	<code>element()</code>	<code>peek()</code>	not applicable	not applicable



BlockingQueue

- It's best to pick one set of operations and stick to it
 - Mixing and matching can be confusing
- Blocking operations with timeouts can be very useful
 - e.g. For backoff / retry
 - `put()`, `take()` have overloads which allow for backoff



LinkedBlockingQueue

- Implemented using `LinkedList`
- Can specify a capacity
 - By default, is `Integer.MAX_VALUE`
 - Usually thought of as an unbounded queue
- Slightly better perf characteristics for low-latency apps
- A common default choice for producer-consumer apps



ArrayBlockingQueue

- Implemented using `ArrayList`
- Like `LinkedBlockingQueue`, has interesting special methods
 - e.g. `drainTo(Collection T)`
 - Drains all available elements in the queue into the collection
- Faster than `LinkedBlockingQueue` in single thread case
- Like `LinkedBlockingQueue`, is a FIFO implementation



CountDownLatch

- A group consensus construct
- Constructor takes an `int` (the count)
- `countDown()` decrements the count
- `await()` blocks until `count == 0`
 - i.e. consensus
 - Releases all waiting threads at once



CountDownLatch



CountDownLatch

- No way to increase the count - only decrease
- Use cases
 - Cache population & startup
 - Multithreaded testing
- countdown-and-await pattern common
 - Simulates Happens-Before edge
- Not the only way to use this



CountDownLatch Example

```
public class LatchExample implements Runnable {
    private final CountDownLatch latch;

    public LatchExample(CountDownLatch latch) {
        this.latch = latch;
    }

    @Override
    public void run() {
        //Call an API
        System.out.println(Thread.currentThread().getName() + " Done API Call");
        try {
            latch.countDown();
            latch.await();
        } catch (InterruptedException e) {
            e.printStackTrace();
        }
        System.out.println(Thread.currentThread().getName() + " Continue processing");
    }
}
```



CyclicBarrier

- **CountDownLatch** is 1-shot
- What if we want to reset or reuse the barrier?
- **CyclicBarrier**
 - Clean reset if the latch has been released by threads reaching it
 - Otherwise, throws a **BrokenBarrierException**



CyclicBarrier

```
public class ResynchingWorker implements Runnable {
    private final CyclicBarrier barrier;
    public ResynchingWorker(CyclicBarrier b) {
        barrier = b;
    }
    public void run() {
        try {
            // Wait until all Worker threads are ready
            barrier.await();
            // Simulate some work...
            Thread.sleep((int) (Math.random() * 1000));
            // Wait until all Worker threads are finished
            barrier.await();
        } catch (InterruptedException ignore) {
            // Ignore
        } catch (BrokenBarrierException bbe) {
            System.err.println("Barrier broken!");
        }
    }
}
```



Phaser

- Synchronization barrier similar to **CyclicBarrier**
 - More flexible
- Registered threads can change over time
 - Methods `register()` and `bulkRegister(int)`
- Can repeatedly await (like **CyclicBarrier**)



Phaser

- Methods `arrive()` & `arriveAndDeregister()` record arrival
- Non-blocking, return an associated arrival phase number
- When the final thread for a given phase arrives
 - An optional action is performed
 - The phase advances
- Wait for given phase to advance – `awaitAdvance(int phase)`
- Very flexible, but not often needed



Intrinsic Locks – Drawbacks

- There is only one type of lock
- Applies equally to all synchronized operations on locked object



Intrinsic Locks – Drawbacks

- There is only one type of lock
- Applies equally to all synchronized operations on locked object
- Lock acquired at the start of synchronized block or method
- Lock released at the end of block or method



Intrinsic Locks – Drawbacks

- There is only one type of lock
- Applies equally to all synchronized operations on locked object
- Lock acquired at the start of synchronized block or method
- Lock released at the end of block or method
- Lock is either acquired or the thread blocks
 - No other outcomes are possible
- We can imagine more general classes of locks



Reimagined Locks

- Different types of locks
 - e.g. reader and writer locks



Reimagined Locks

- Different types of locks
 - e.g. reader and writer locks
- Not restrict locks to blocks
 - Allow a lock in one method and unlock in another



Reimagined Locks

- Different types of locks
 - e.g. reader and writer locks
- Not restrict locks to blocks
 - Allow a lock in one method and unlock in another
- If a thread cannot acquire a lock allow back out – `tryLock()`
- Allow a thread attempted acquire with timeout



Lock

- The main interface for this reimagining is **Lock**
 - `java.util.concurrent.locks.Lock`
- Key methods
 - `lock()` - acquires the lock
 - `newCondition()` - creates a condition for wait / notify
 - `tryLock()` - acquire the lock (with optional timeout)
 - `unlock()` - release the lock



Lock

- Two main implementations
 - **ReentrantLock** - replicates the “synchronized” lock
 - **ReentrantReadWriteLock** - used for reader / writer locks
- Also the non-reentrant implementation
 - Very rarely seen in practice

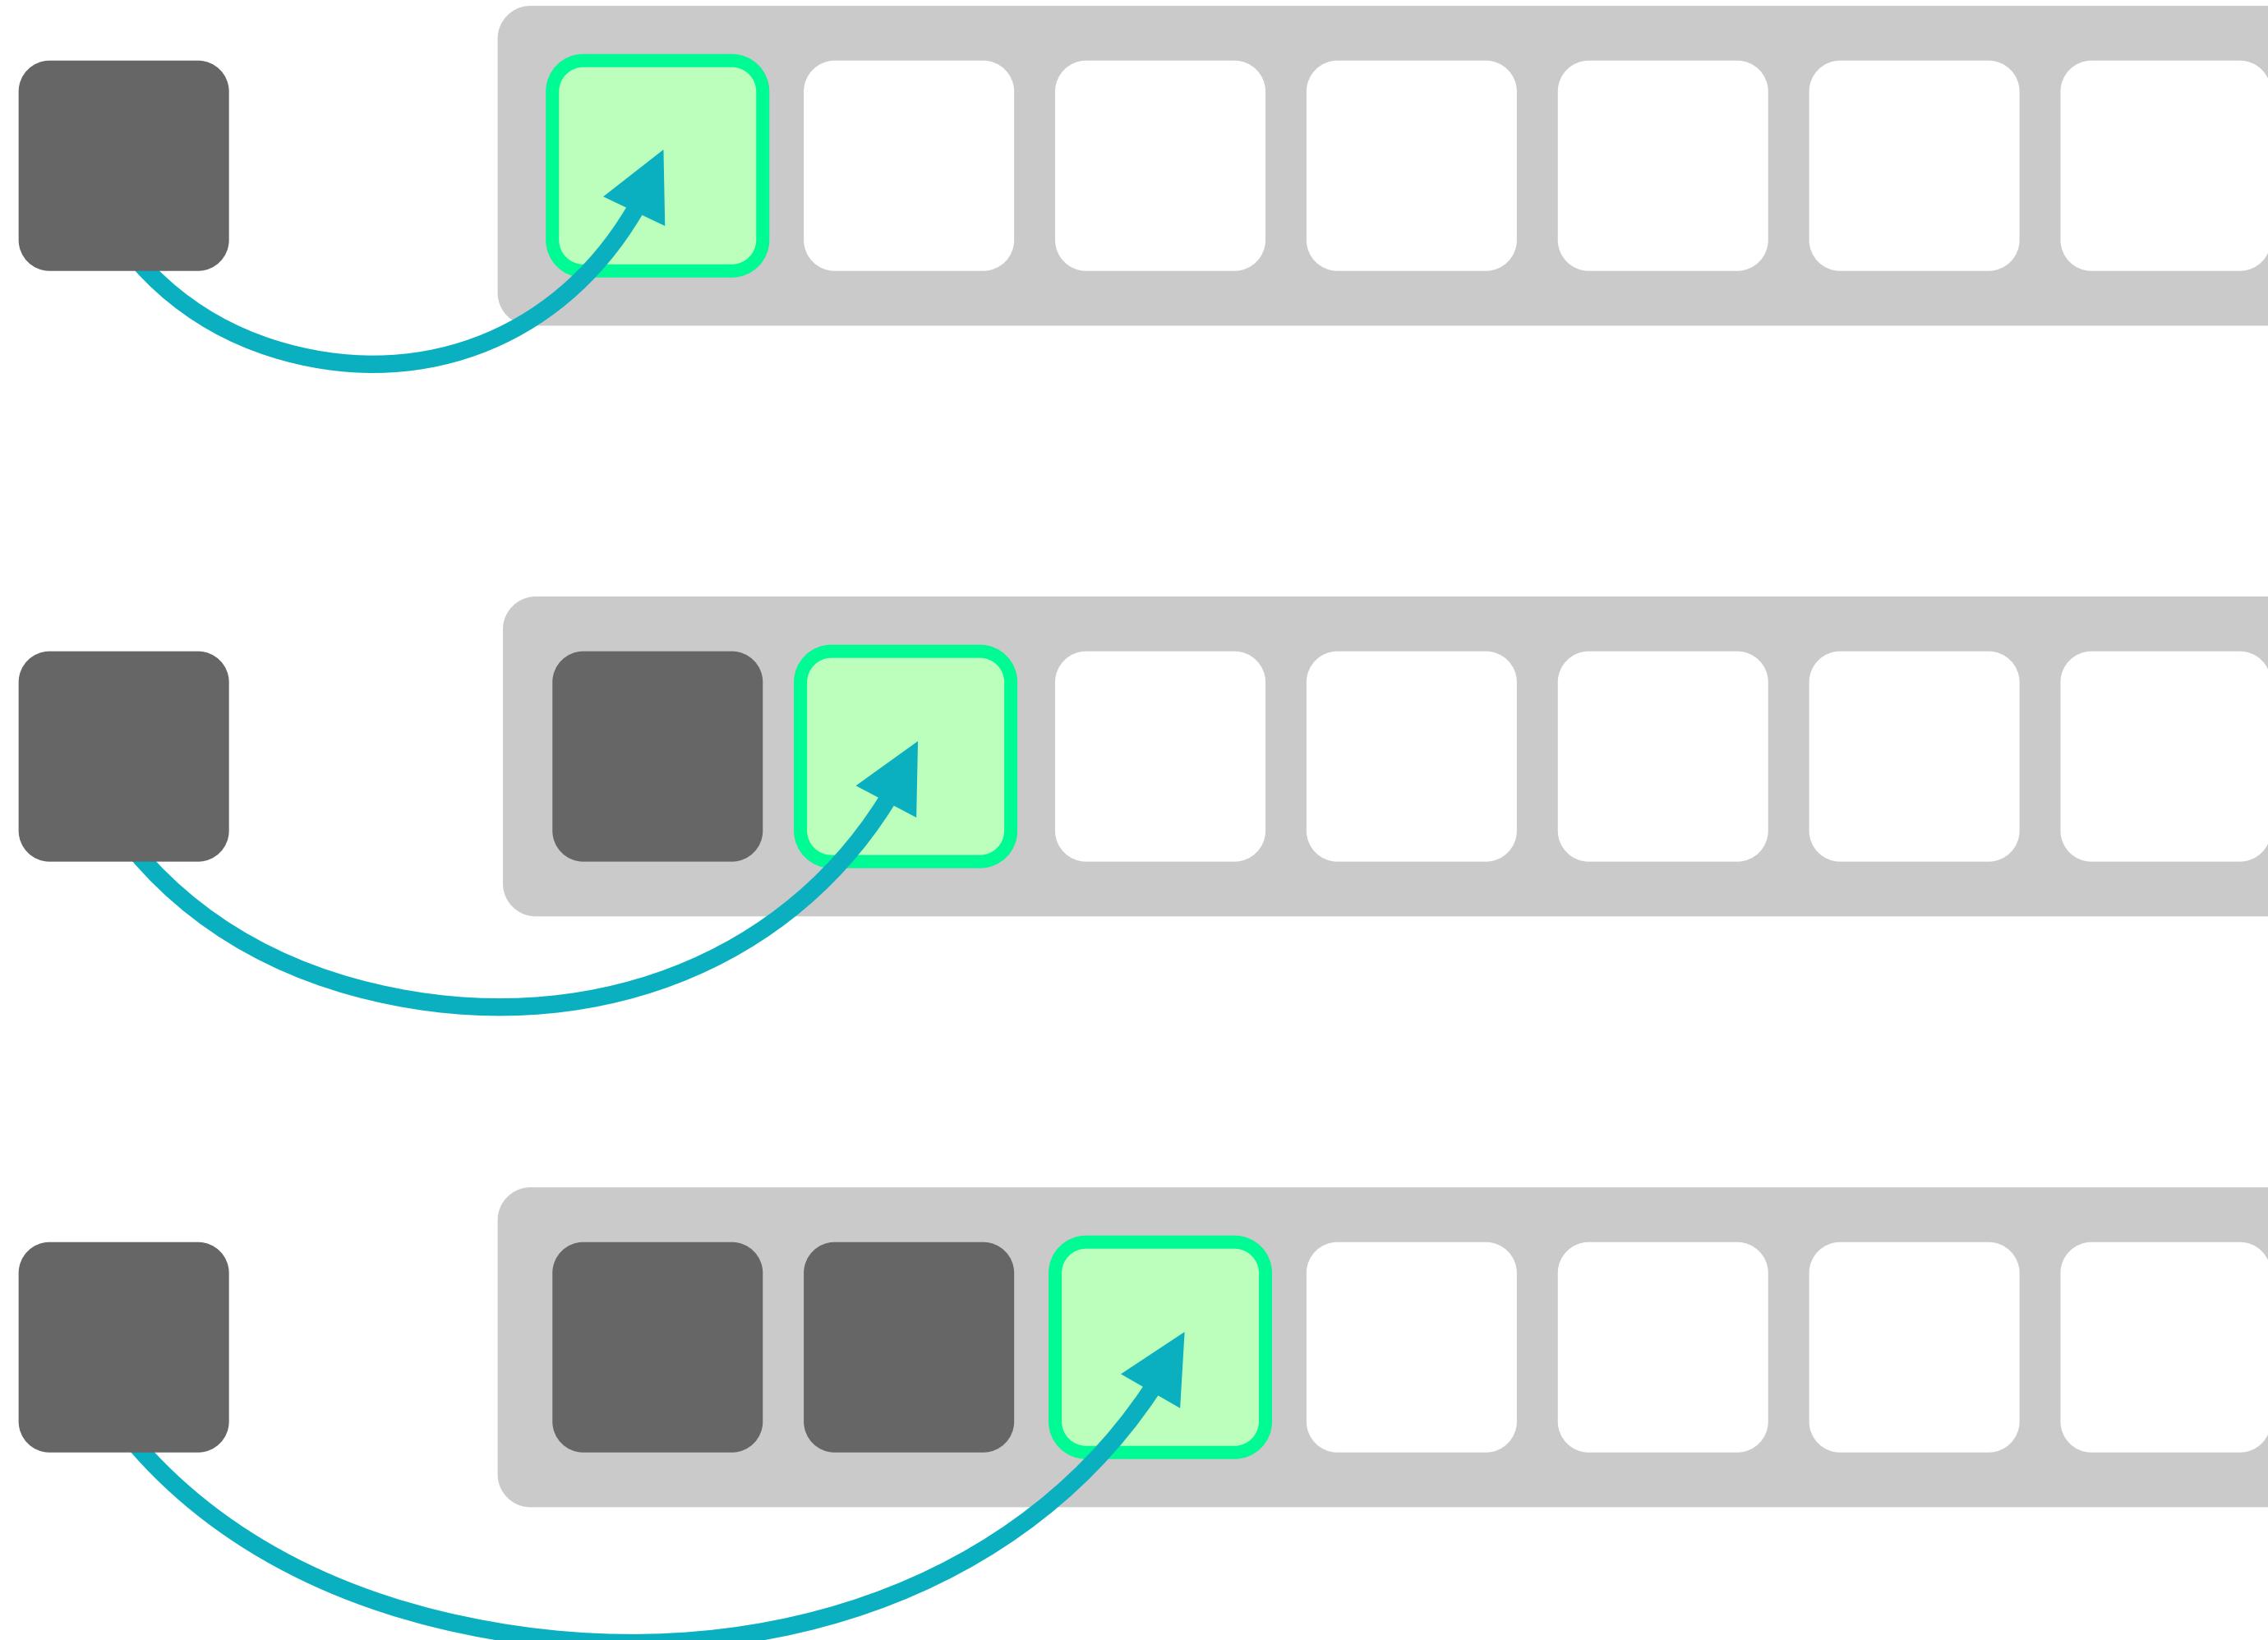


Condition

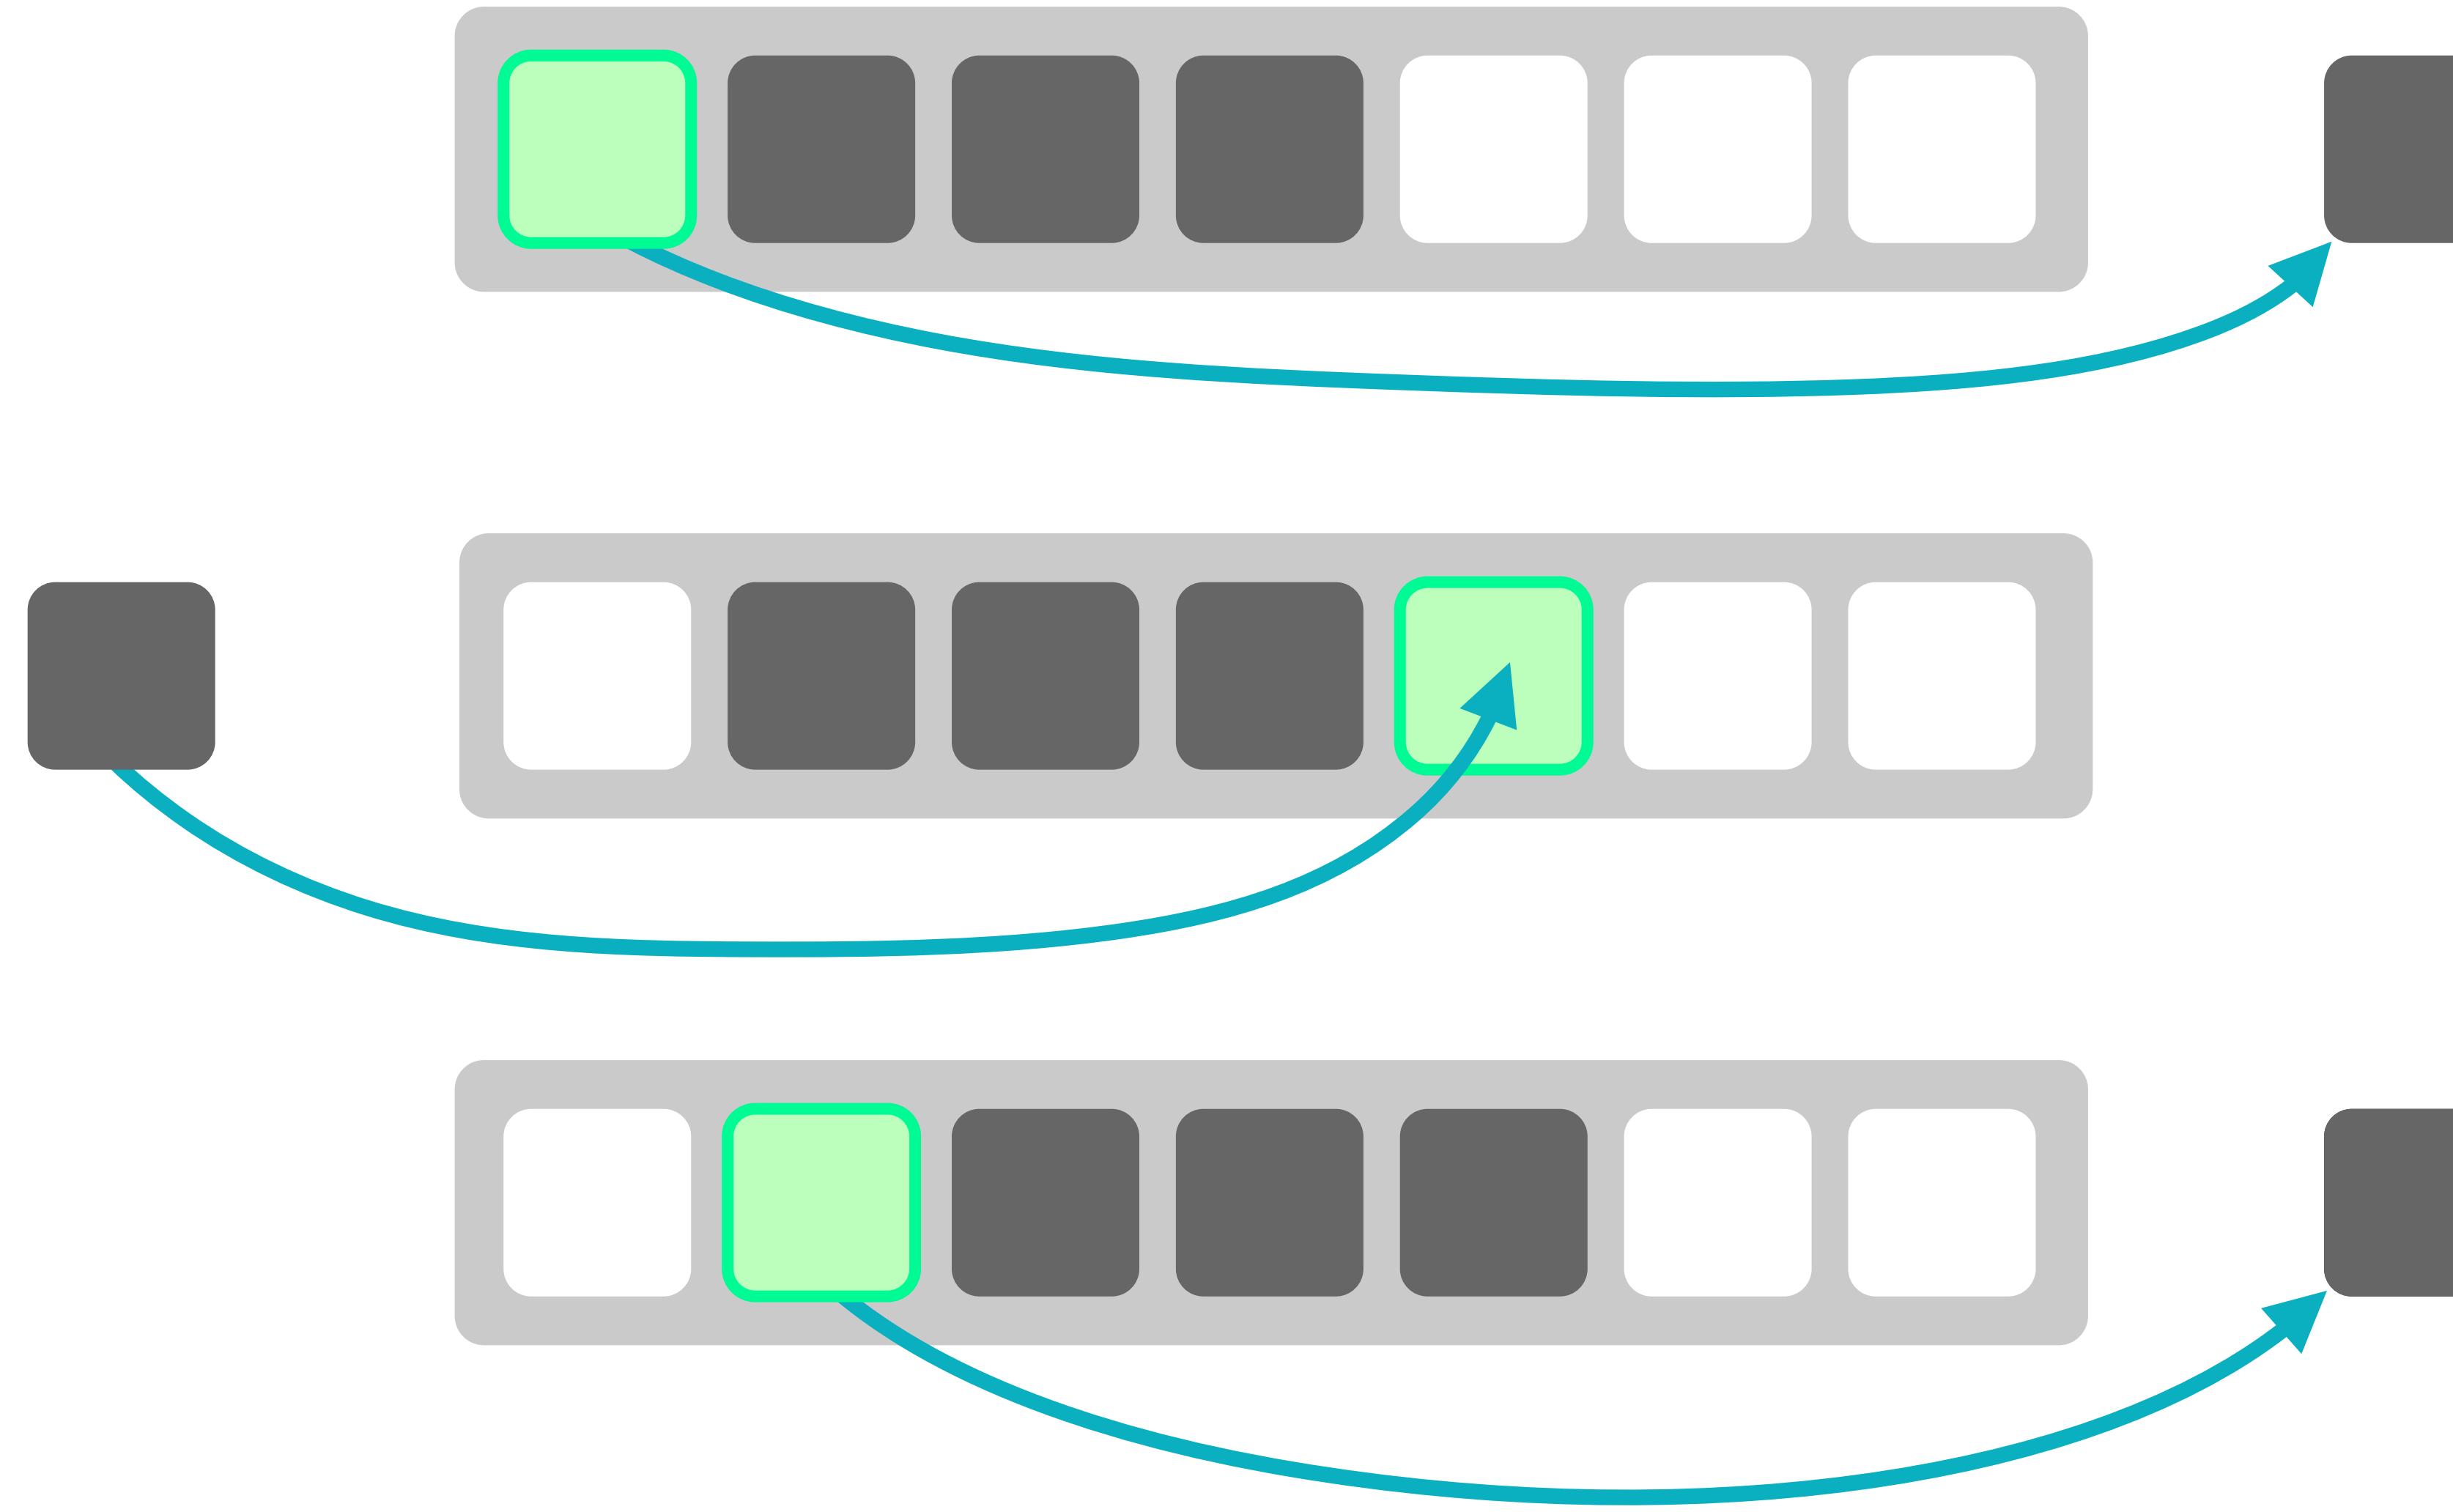
- Intrinsic concurrency only has 1 condition to wait on
- With j.u.c we can be more flexible
- `Lock.newCondition()` creates a `Condition` variable
 - Can represent any condition
 - Multiple conditions can be created on a lock
 - No “semantic” meaning as in intrinsic concurrency



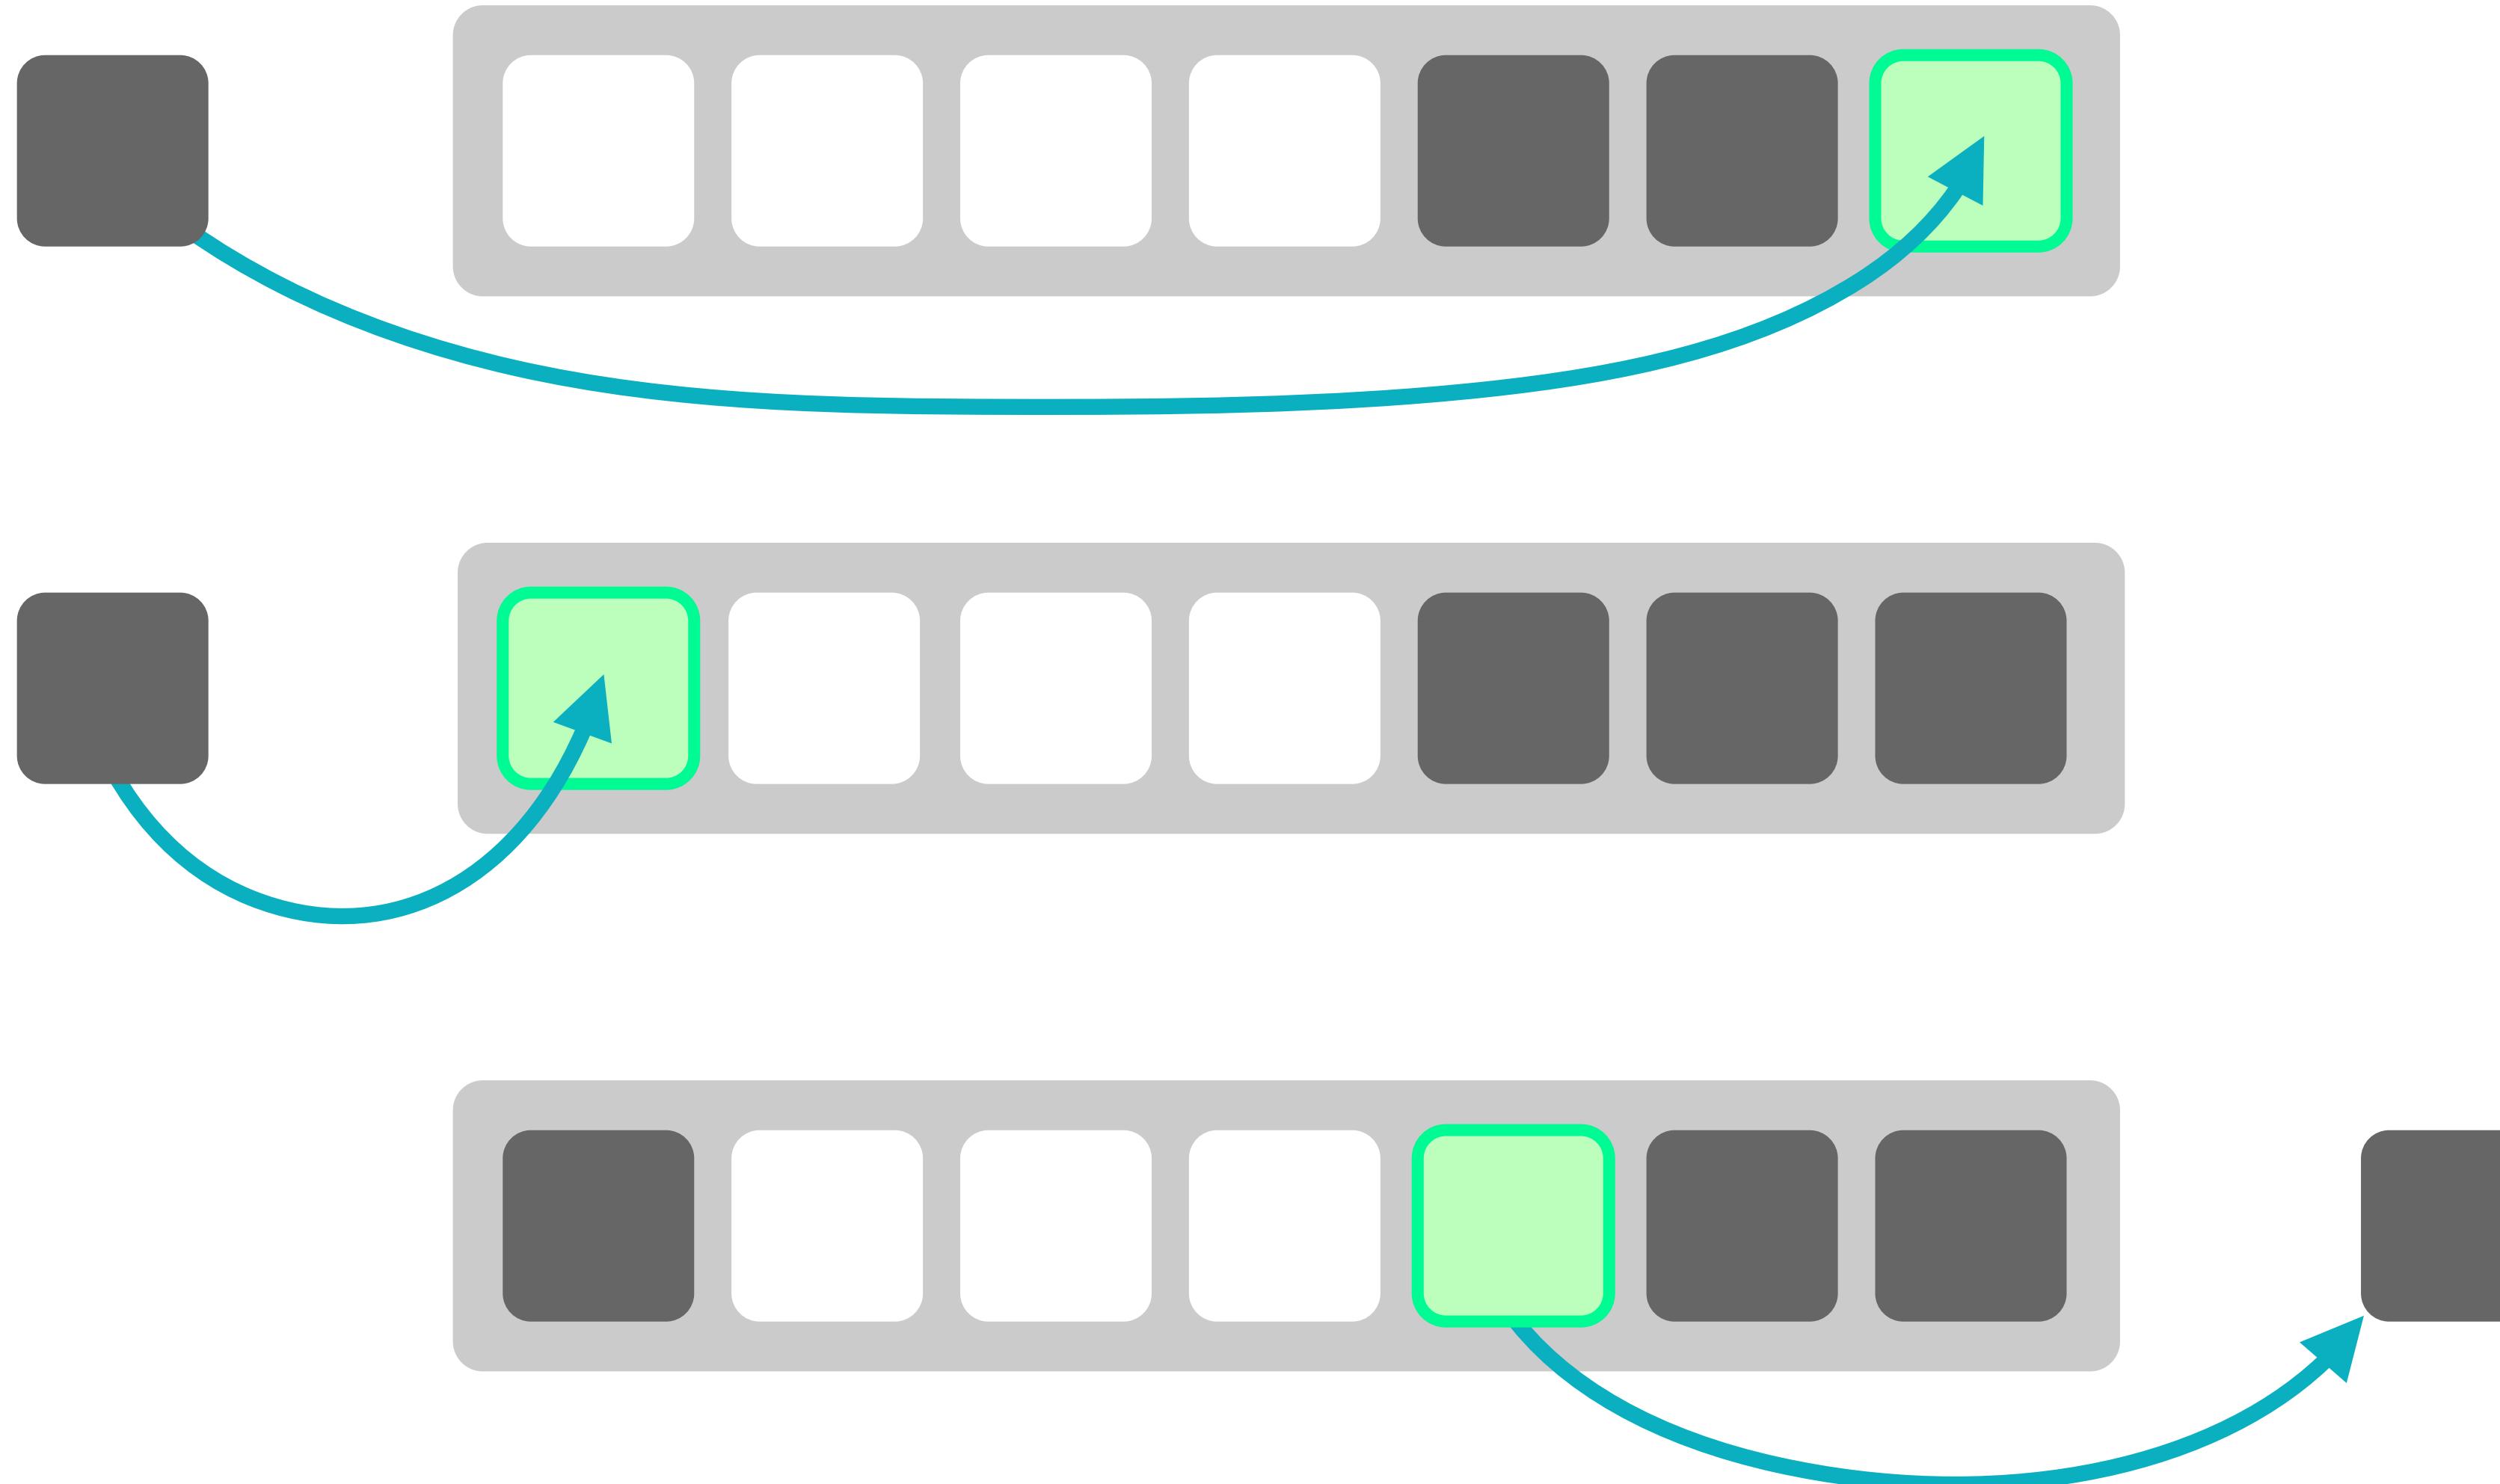
Example: Bounded Queue



Example: Bounded Queue



Example: Bounded Queue



Lock & Condition

```
Lock lock = new ReentrantLock();
Condition notFull    = lock.newCondition();
Condition notEmpty   = lock.newCondition();
Object[ ] items = new Object[100];
int putptr, takeptr, count;
```



Lock & Condition

```
// Thread 1
public void put(Object x) throws InterruptedException {
    lock.lock();
    try {
        // If full, release lock & wait until we get it back
        while (count == items.length) notFull.await();

        items[putptr] = x;
        if (++putptr == items.length) putptr = 0;
        ++count;

        // signal notEmpty that it can try for the lock again
        notEmpty.signal();
    } finally {
        lock.unlock();
    }
}
```



Lock & Condition

```
// Thread 2
public Object take() throws InterruptedException {
    lock.lock();
    try {
        // If empty, release the lock & wait
        while (count == 0) notEmpty.await();

        Object x = items[takeptr];
        if (++takeptr == items.length) takeptr = 0;
        --count;

        // Signal notFull that it can stop waiting & try for the lock
        notFull.signal();
        return x;
    } finally {
        lock.unlock();
    }
}
```



Exercise – Intrinsic Lock vs ReentrantLock

- Take the BoundedQueue example
- Write it using intrinsic concurrency as well as j.u.c
- Compare the throughput in single-reader, single-writer case
- What about 2 writers, 1 reader? 2 readers, 1 writer?
- What happens if Biased Locking is switched off?

<https://github.com/kittylst/optimizing-java>



Under the Hood

- Inside AtomicInteger (VarHandles and Unsafe)
- Exercise: Build your own Atomic
- Q&A



Implementing Atomics

- Lock-free
 - Use CAS hardware where available
 - On all modern platforms
 - Otherwise, fall back to software implementation (lock-based)
- Java Support provided via Unsafe
 - e.g. `Unsafe::compareAndSwapInt()`
 - See later...



CAS Operations

- Pass in 2 values
 - “Expected current value”
 - “Requested new value”



CAS Operations

- Pass in 2 values
 - “Expected current value”
 - “Requested new value”
- Compare expected current with actual current
- If they match, update to requested new value

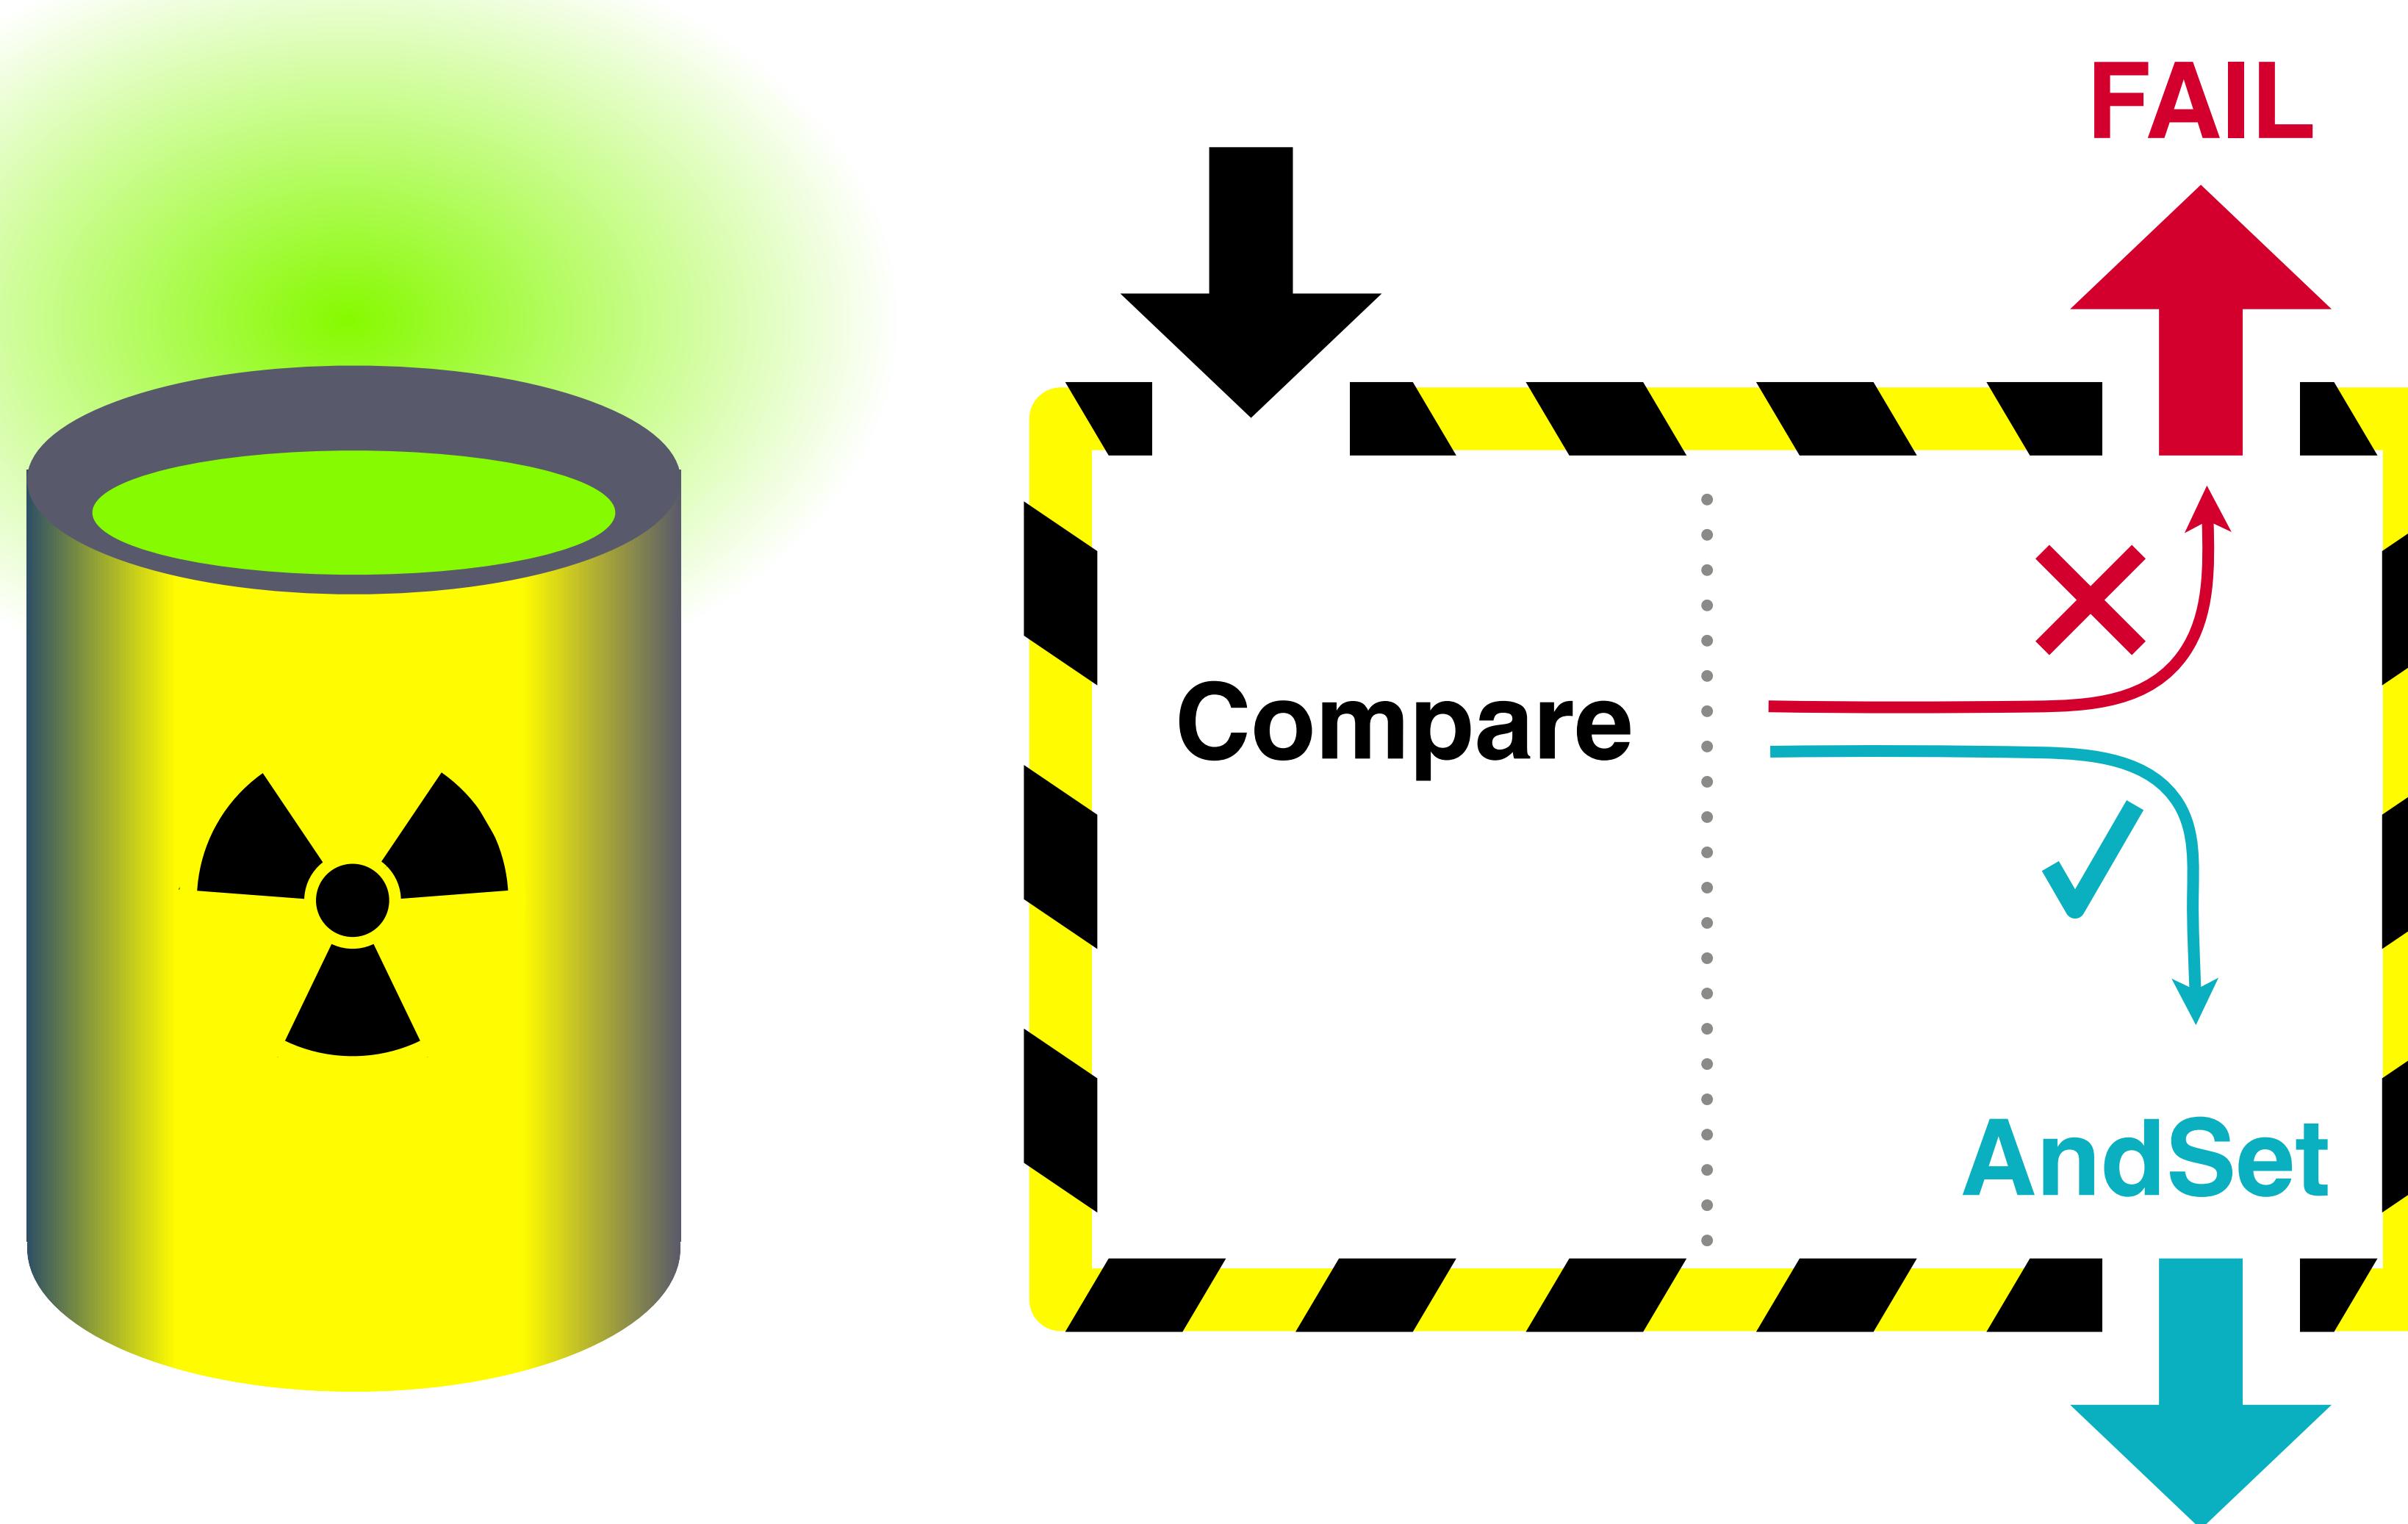


CAS Operations

- Pass in 2 values
 - “Expected current value”
 - “Requested new value”
- Compare expected current with actual current
- If they match, update to requested new value
- If the match fails
 - Another thread has modified it
 - The update fails
- In both cases return the current value after the attempt
- Performed as a single operation



CAS Operations



Implementing Atomics

```
public class AtomicInteger extends Number implements java.io.Serializable {  
    private static final long serialVersionUID = 6214790243416807050L;  
  
    // setup  
    // ...  
  
    static {  
    // ...  
    }  
  
    private volatile int value;
```



Implementing Atomics

```
/**  
 * Gets the current value.  
 *  
 * @return the current value  
 */  
public final int get() {  
    return value;  
}  
  
/**  
 * Sets to the given value.  
 *  
 * @param newValue the new value  
 */  
public final void set(int newValue) {  
    value = newValue;  
}
```



Implementing Atomics

```
public class AtomicInteger extends Number implements java.io.Serializable {  
    private static final long serialVersionUID = 6214790243416807050L;  
  
    // setup to use Unsafe.compareAndSwapInt for updates  
    private static final Unsafe unsafe = Unsafe.getUnsafe();  
    private static final long valueOffset;  
  
    static {  
        try {  
            valueOffset = unsafe.objectFieldOffset  
                (AtomicInteger.class.getDeclaredField("value"));  
        } catch (Exception ex) { throw new Error(ex); }  
    }  
  
    private volatile int value;
```



Implementing Atomics

```
/**  
 * Atomically sets to the given value and returns the old value  
 *  
 * @param newValue the new value  
 * @return the previous value  
 */  
public final int getAndSet(int newValue) {  
    return unsafe.getAndSetInt(this, valueOffset, newValue);  
}
```



Implementing Atomics – Unsafe

```
public final int getAndSetInt(Object o, long offset, int newValue) {  
    int v;  
    do {  
        v = getIntVolatile(o, offset);  
    } while (!compareAndSwapInt(o, offset, v, newValue));  
    return v;  
}
```



Implementing Atomics – Unsafe

```
public native int getIntVolatile(Object o, long offset);
/**
 * Atomically update Java variable to <tt>x</tt> if it is currently
 * holding <tt>expected</tt>.
 * @return <tt>true</tt> if successful
 */
public final native boolean compareAndSwapInt(Object o, long offset,
                                              int expected,
                                              int x);
```



CAS Operation / Spinlock

- Pass in 2 values
 - “Expected current value”
 - “Requested new value”

```
locked:  
dd 0  
  
spin_lock:  
mov eax, 1  
xchg eax, [locked]  
test eax, eax  
jnz spin_lock  
ret  
  
spin_unlock:  
mov eax, 0  
xchg eax, [locked]  
ret
```



Removal of Unsafe

- Many frameworks/libraries unable to move to Java > 8
 - Without a replacement for some 'safe' sun.misc.Unsafe features
 - Indirectly impacts everyone using a wide range of frameworks
 - Basically every application in the ecosystem
- What can we do?
 - Give Unsafe a pass for now
 - Create new supported APIs for Java 11 & inform that migration



Exercise

EXERCISE



Practical Considerations

- Concurrency vs Parallelism
- Amdahl's Law
- Q&A



Concurrency vs Parallelism

- Concurrency
 - At least two threads are executing simultaneously
- Parallelism
 - At least two threads are making progress
 - Can include time-slicing
 - Including traditional UNIX timesharing

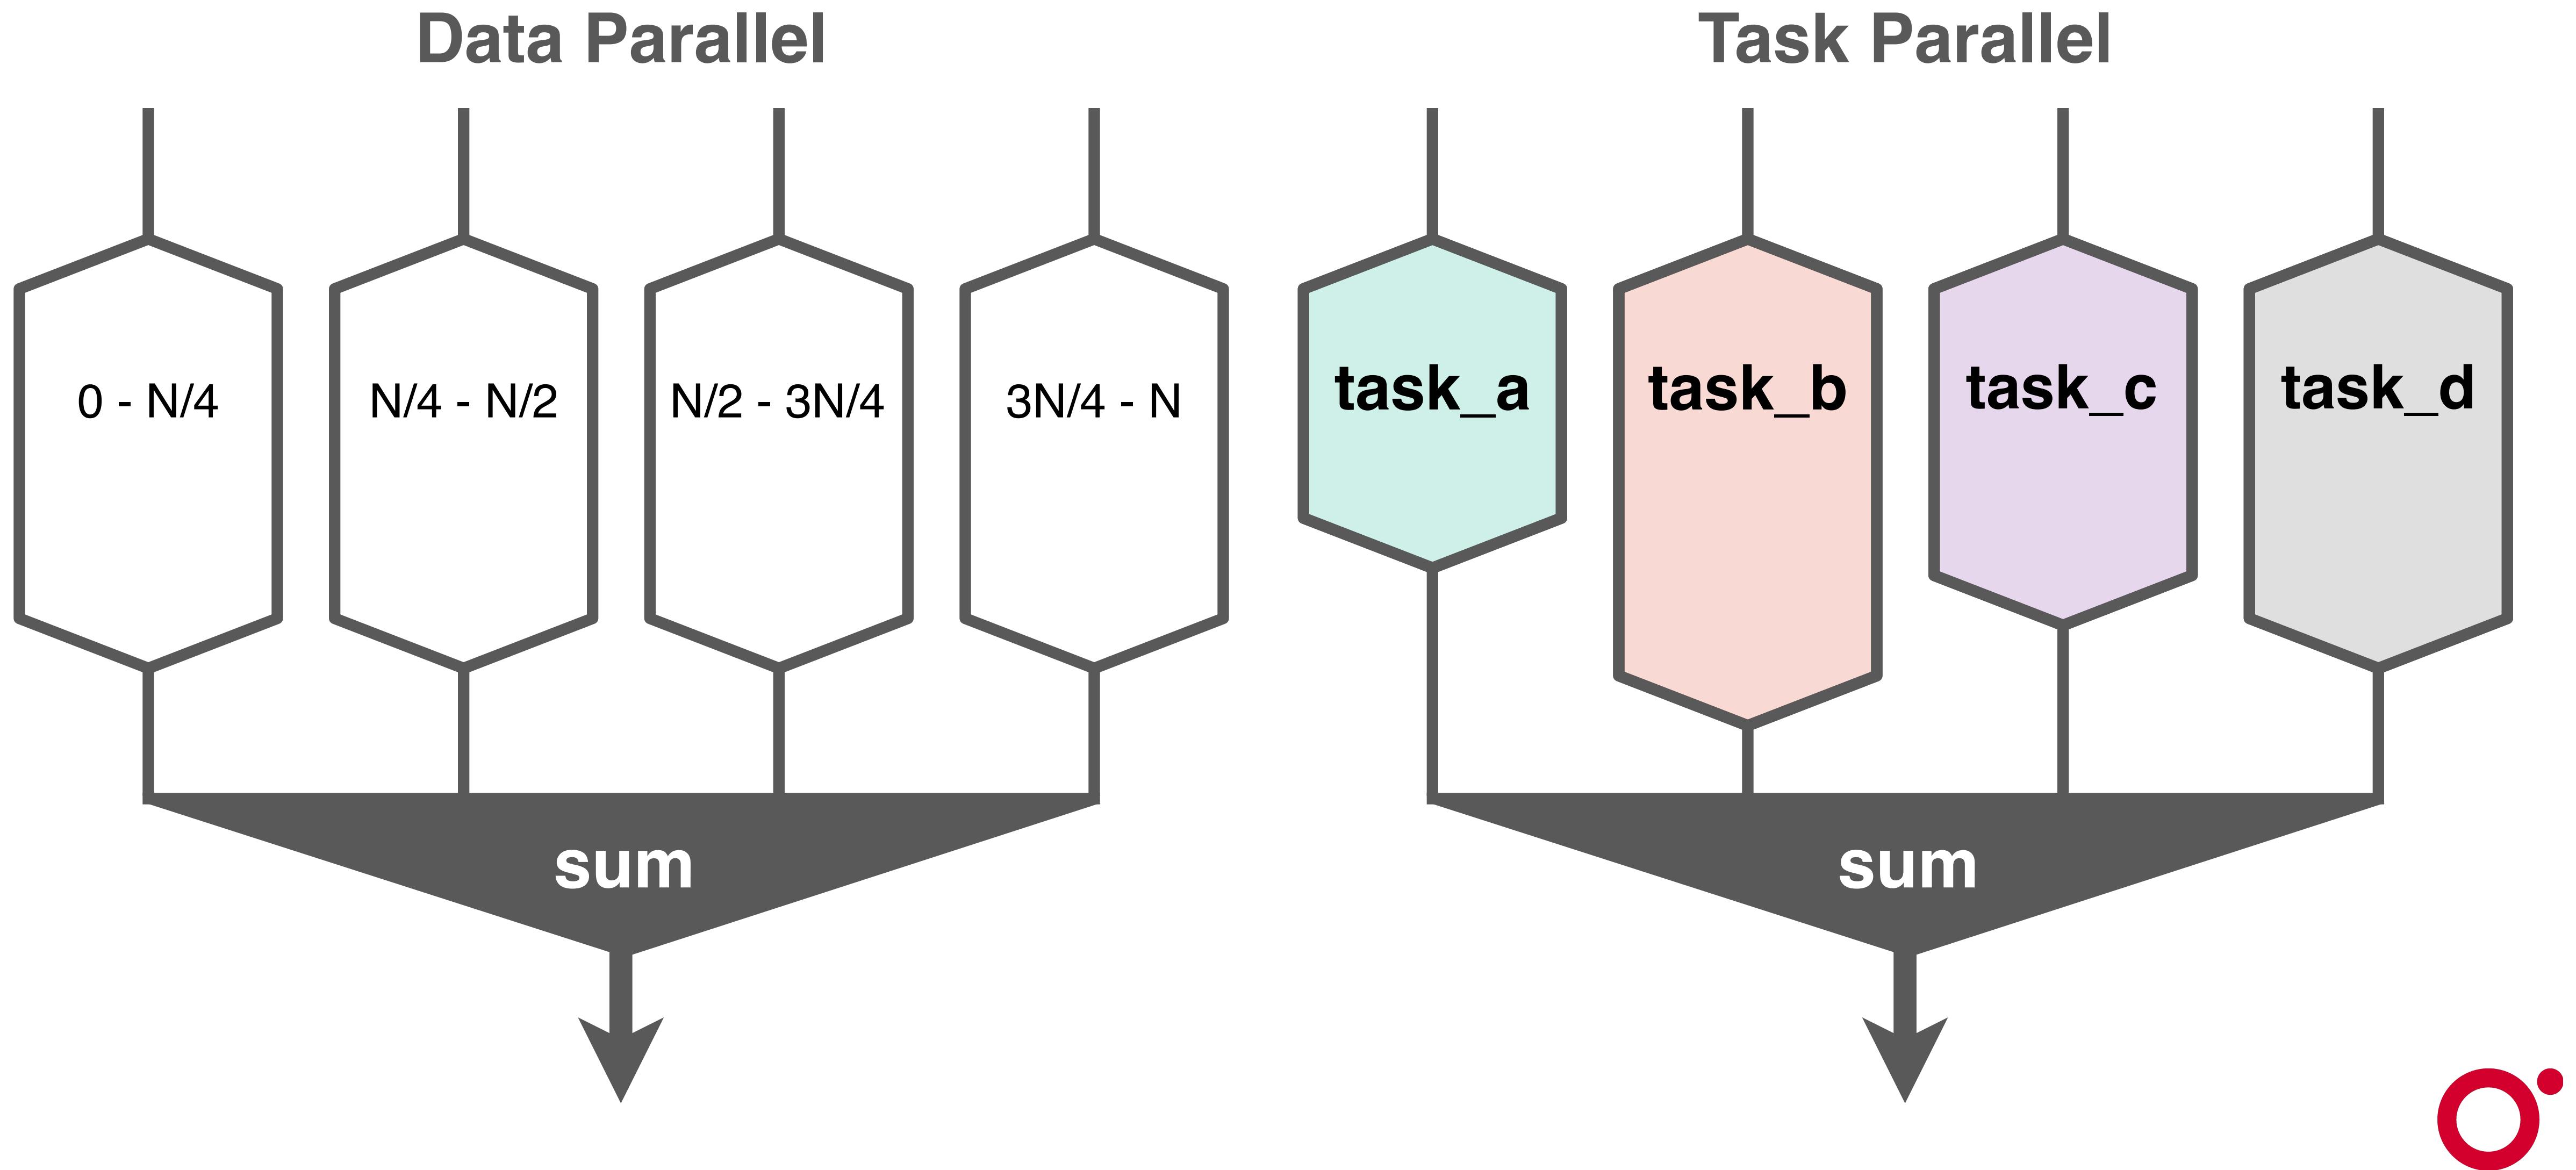


Parallelism

- Task
 - Distribute execution of operations over processes
 - Threads and Executors in Java
 - e.g. each thread services a user in Java EE App
- Data
 - Distribute data over different processes
 - Support built on top of Streams
 - e.g. process payroll and allocate each core 100 employees to process



Parallelism



Amdahl's Law

- Defines upper bound for parallel speedup

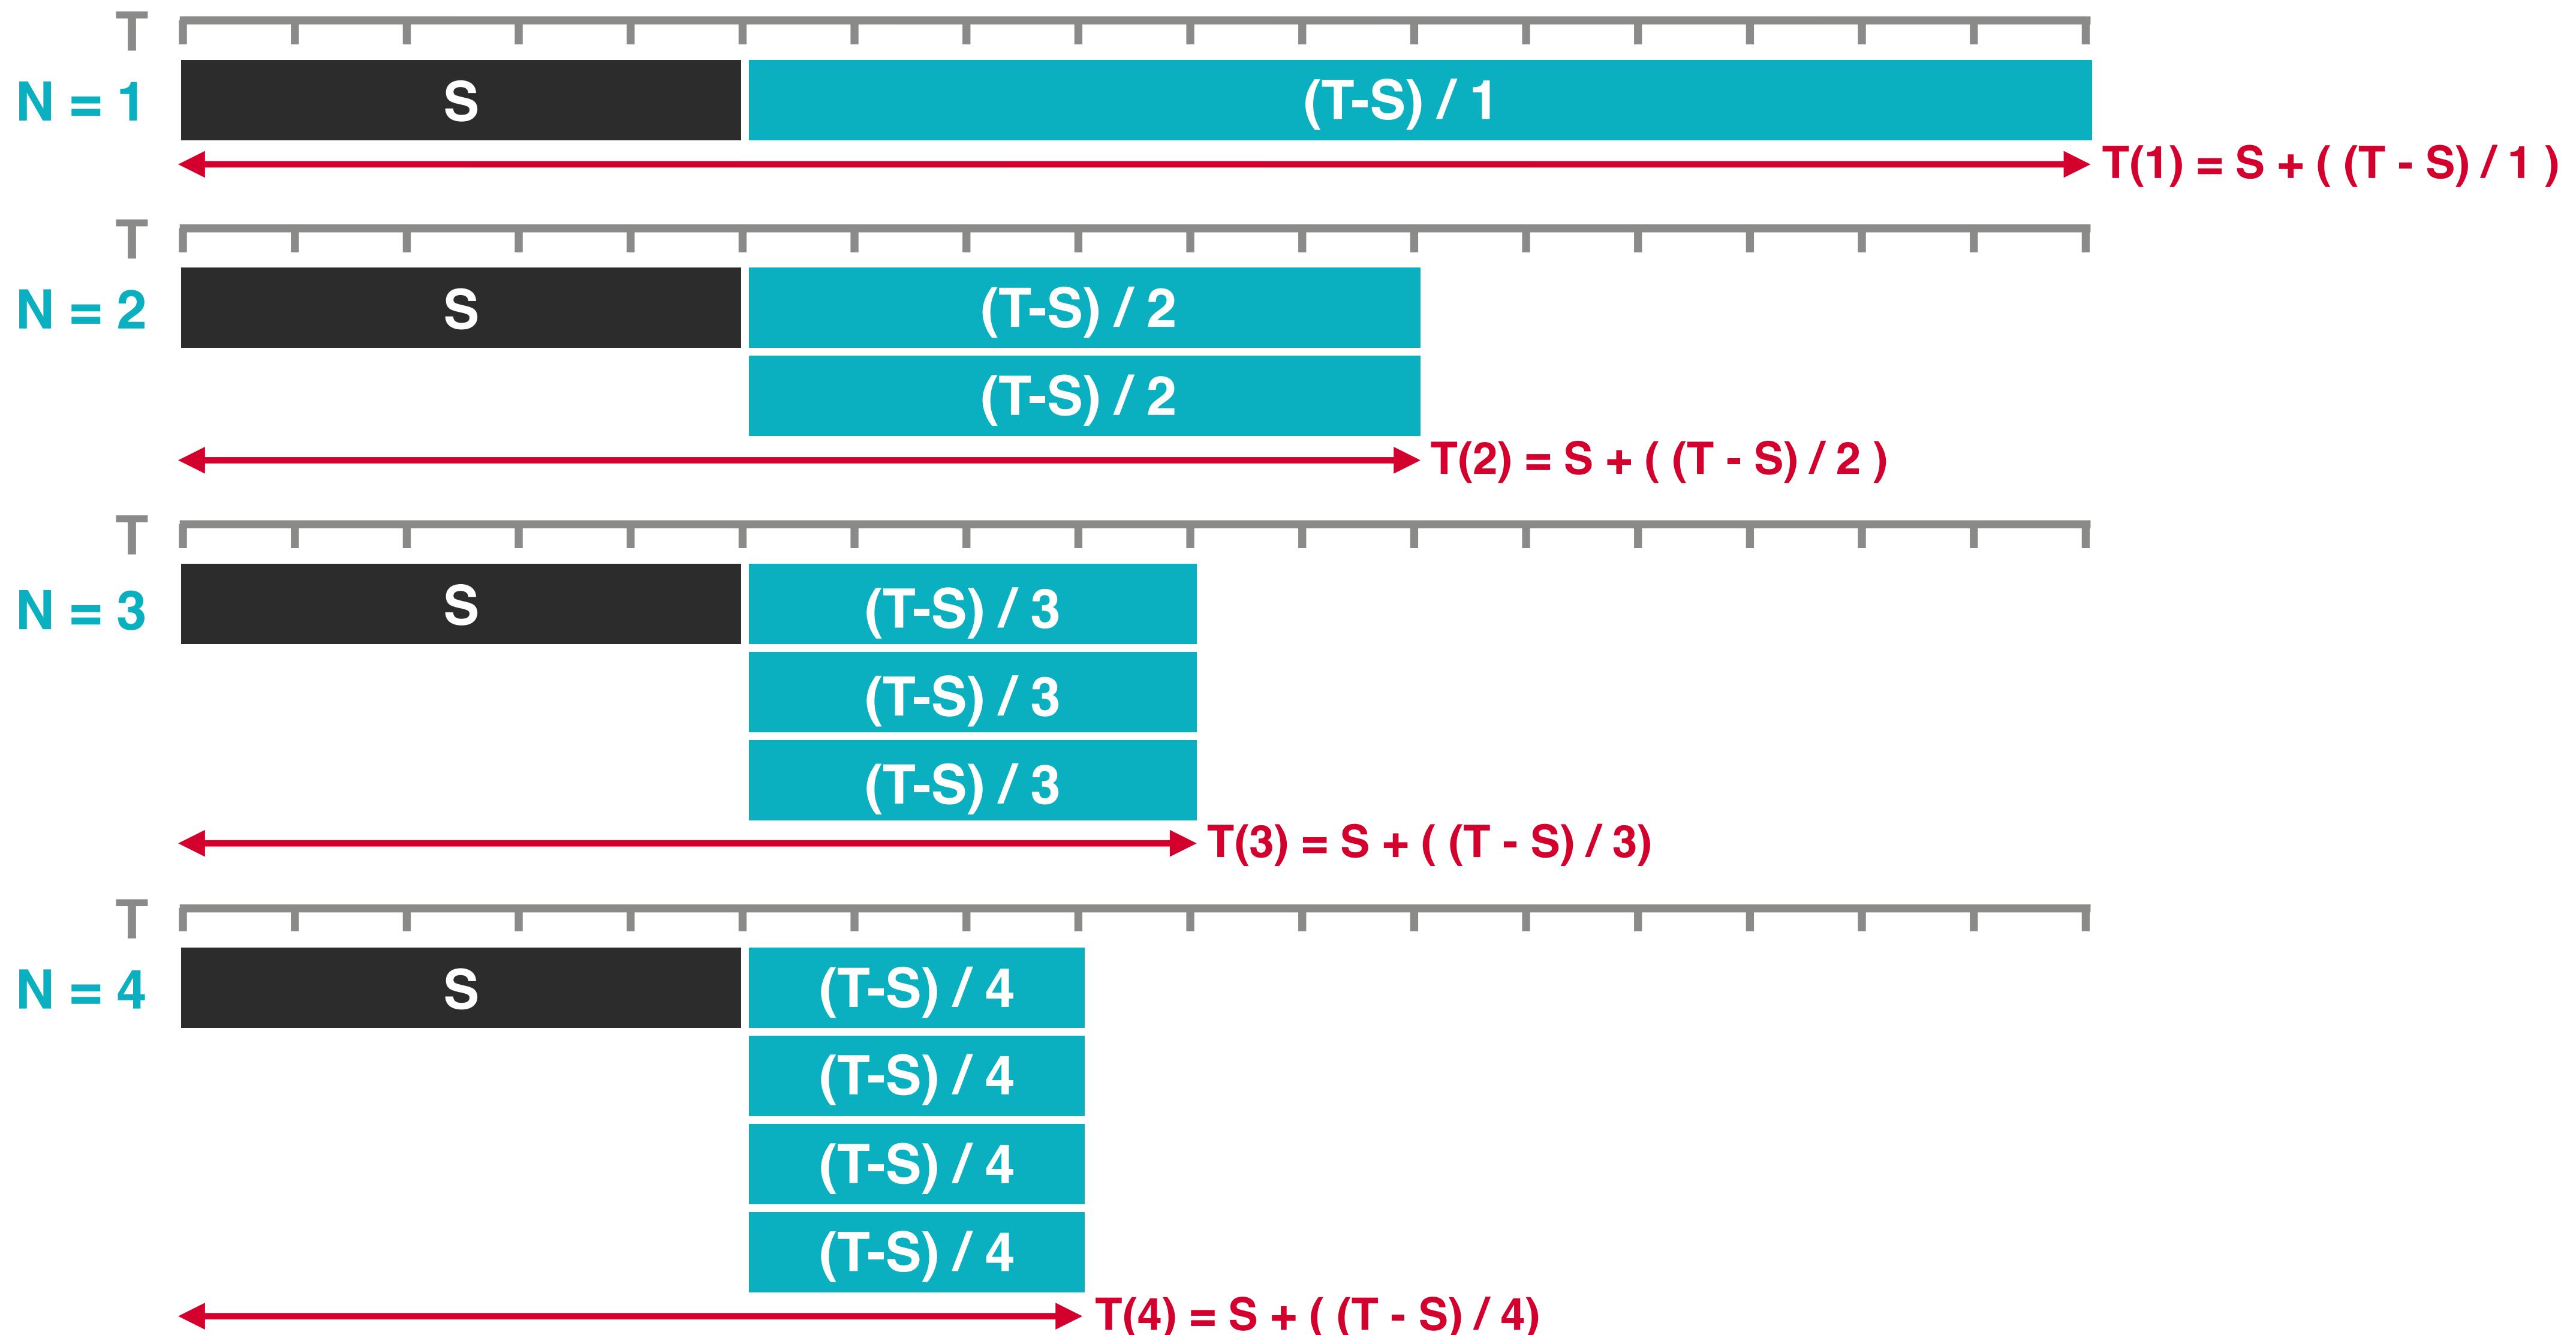
$$T(n) = T(1) * (s + (1 - s) /n)$$

$$u * n = 1 / (s + (1 - s) /n)$$

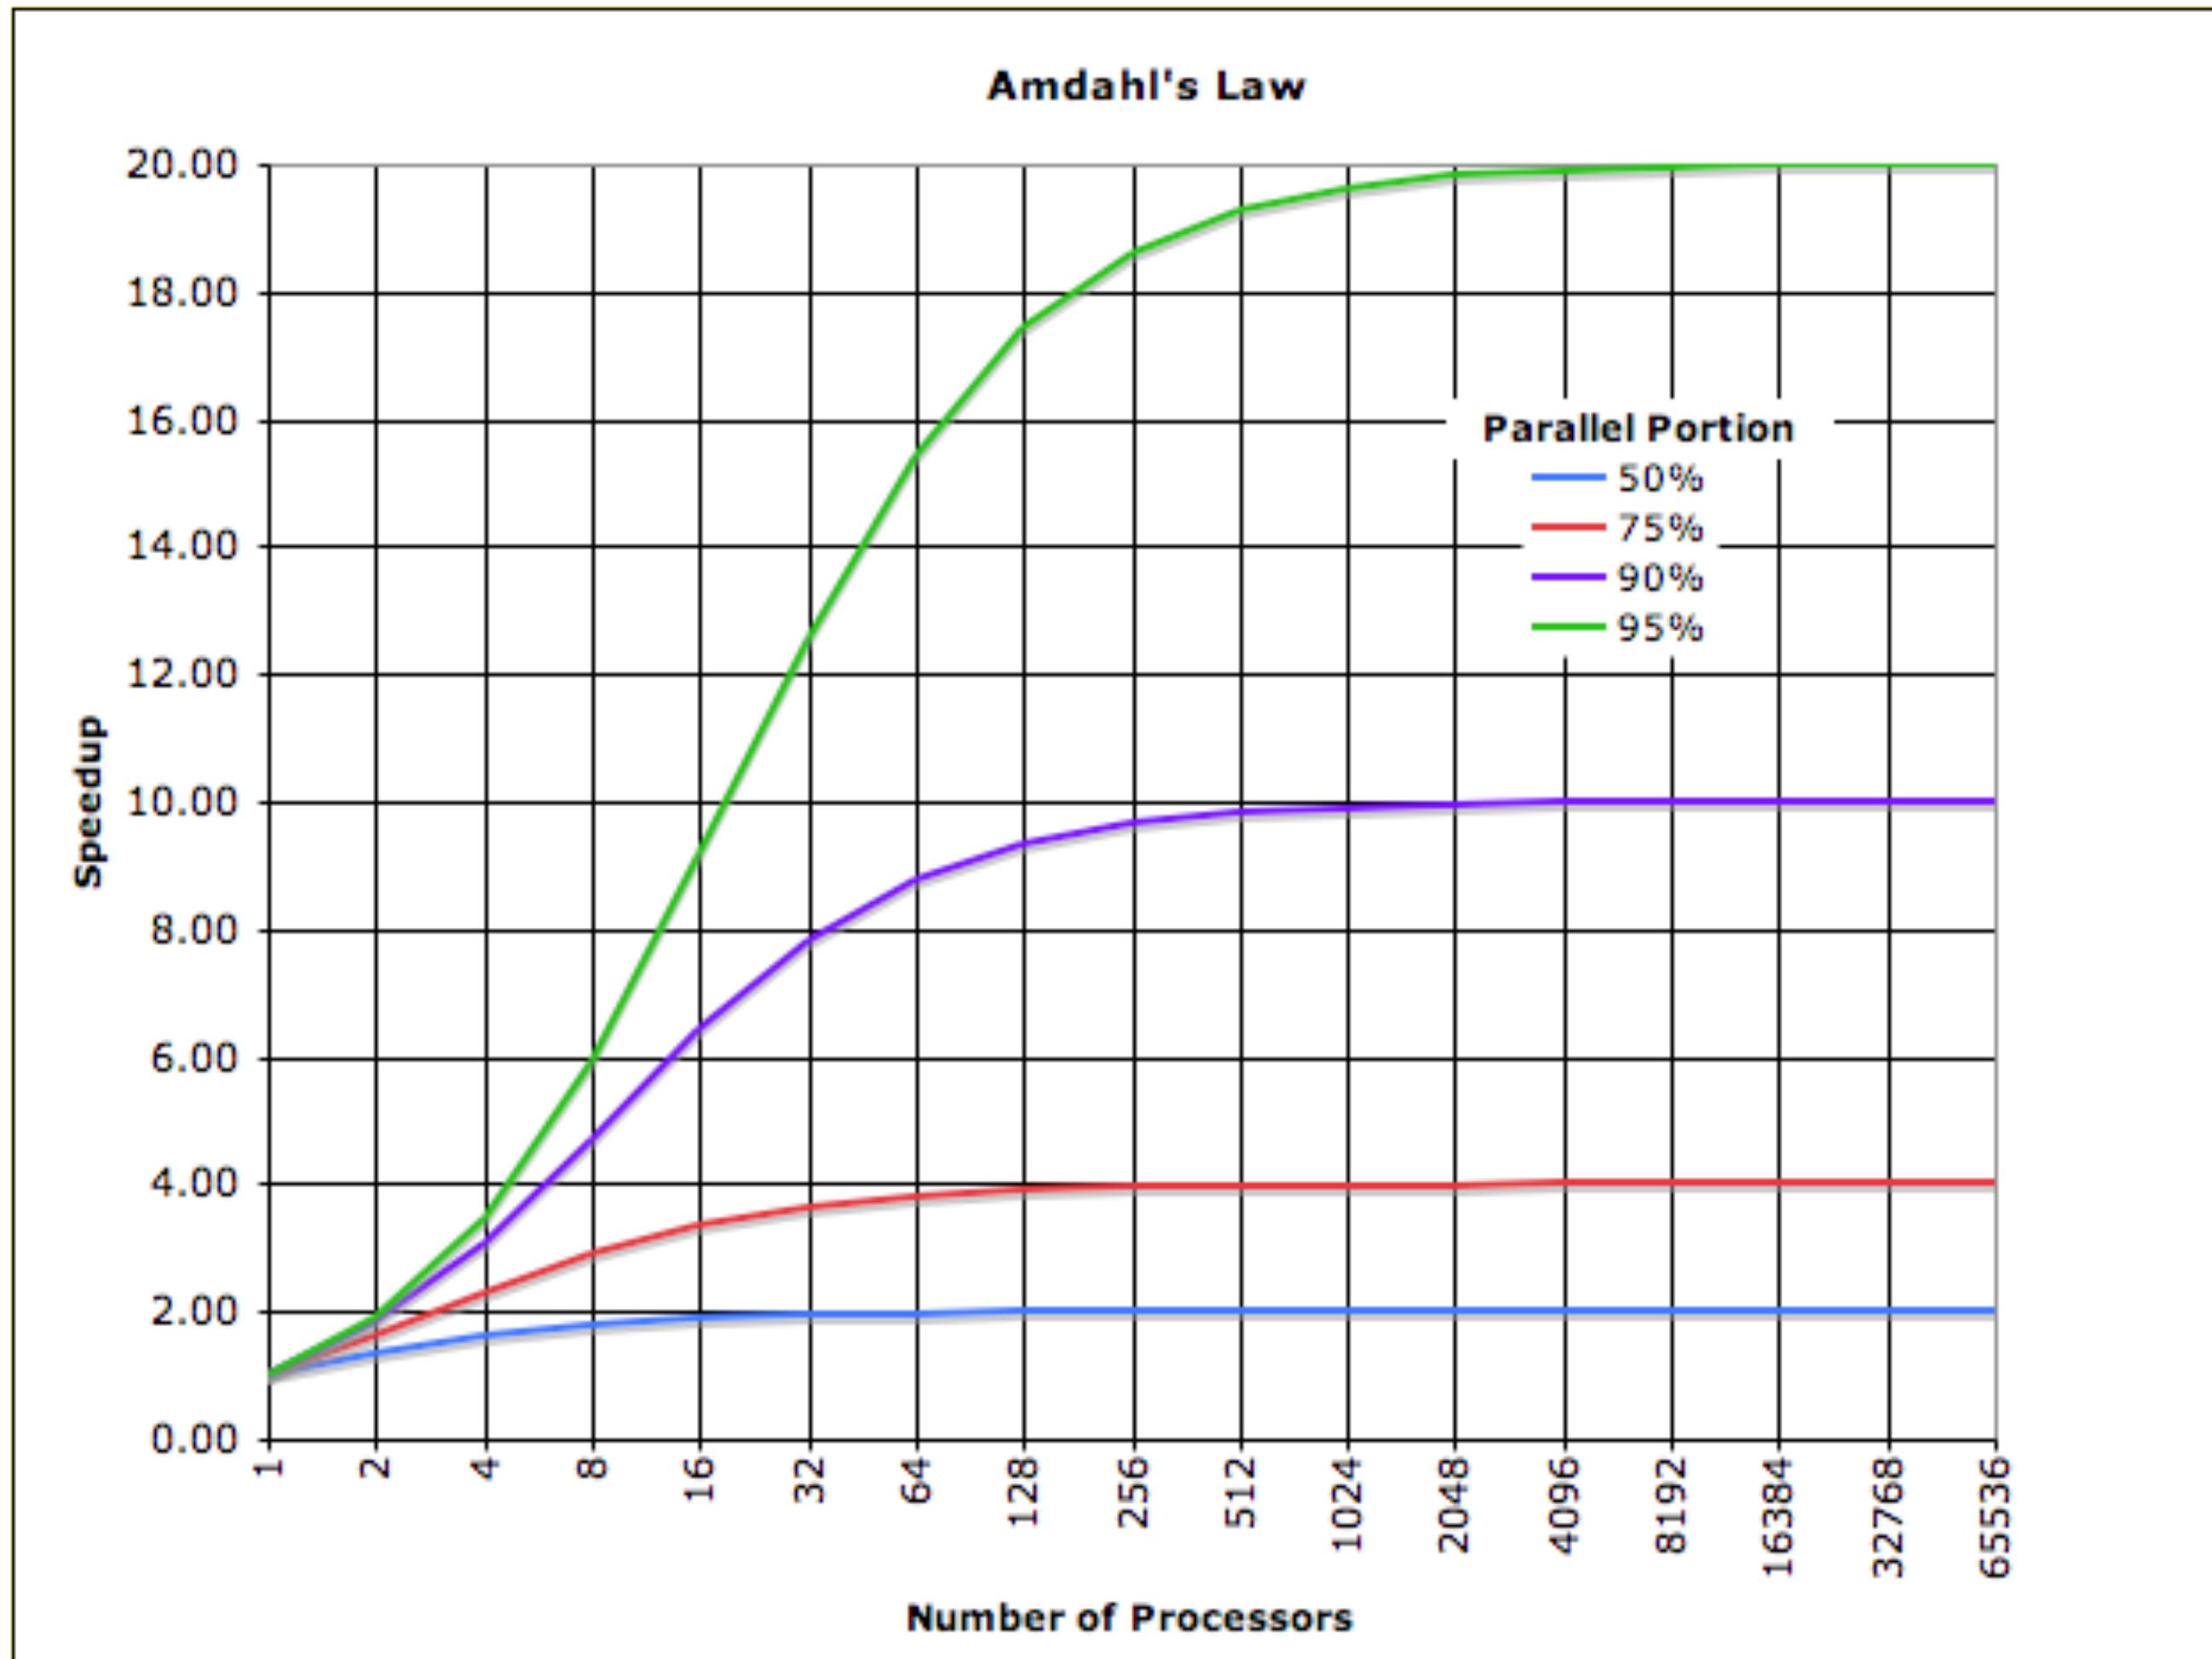
- n = number of cores
- s = proportion of code that must be serial
- u = % utilisation (across all cores)



Amdahl's Law



Parallelism



Amdahl's Law

- Invert Amadahl's Law
 - Calculate max s (serial overhead) for given utilisation
- E.g. To hit 70% utilisation across cores
 - $s = 3 / (7 * (n-1))$

#Cores (n)	Max Serial % (s)
1	100%
2	43%
4	14%
32	1.38%
48	<1%
400	< 0.1%



Concurrency Reading List

- **Java Concurrency in Practice** - B. Goetz et al.
- **Concurrent Programming in Java** - D. Lea
- **Optimizing Java** - B. Evans et al.
- **Mechanical Sympathy** (mailing list)
- **Concurrency Interest** (mailing list)

